

CSCE566-DATA MINING

WEEK 8

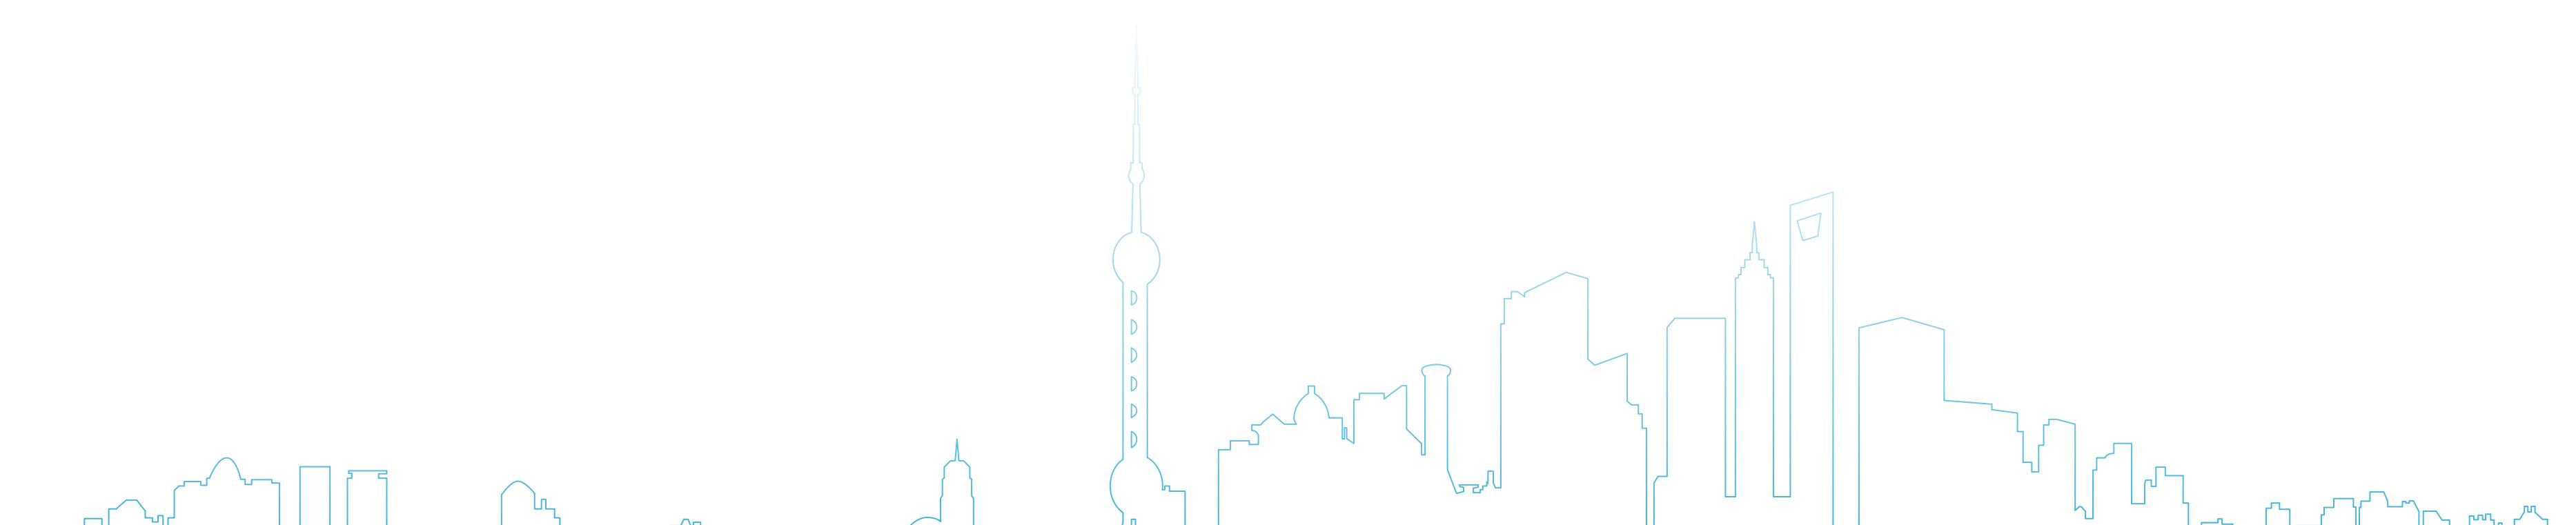
Image Data Mining

Min Shi
min.shi@louisiana.edu

Oct 14, 2024

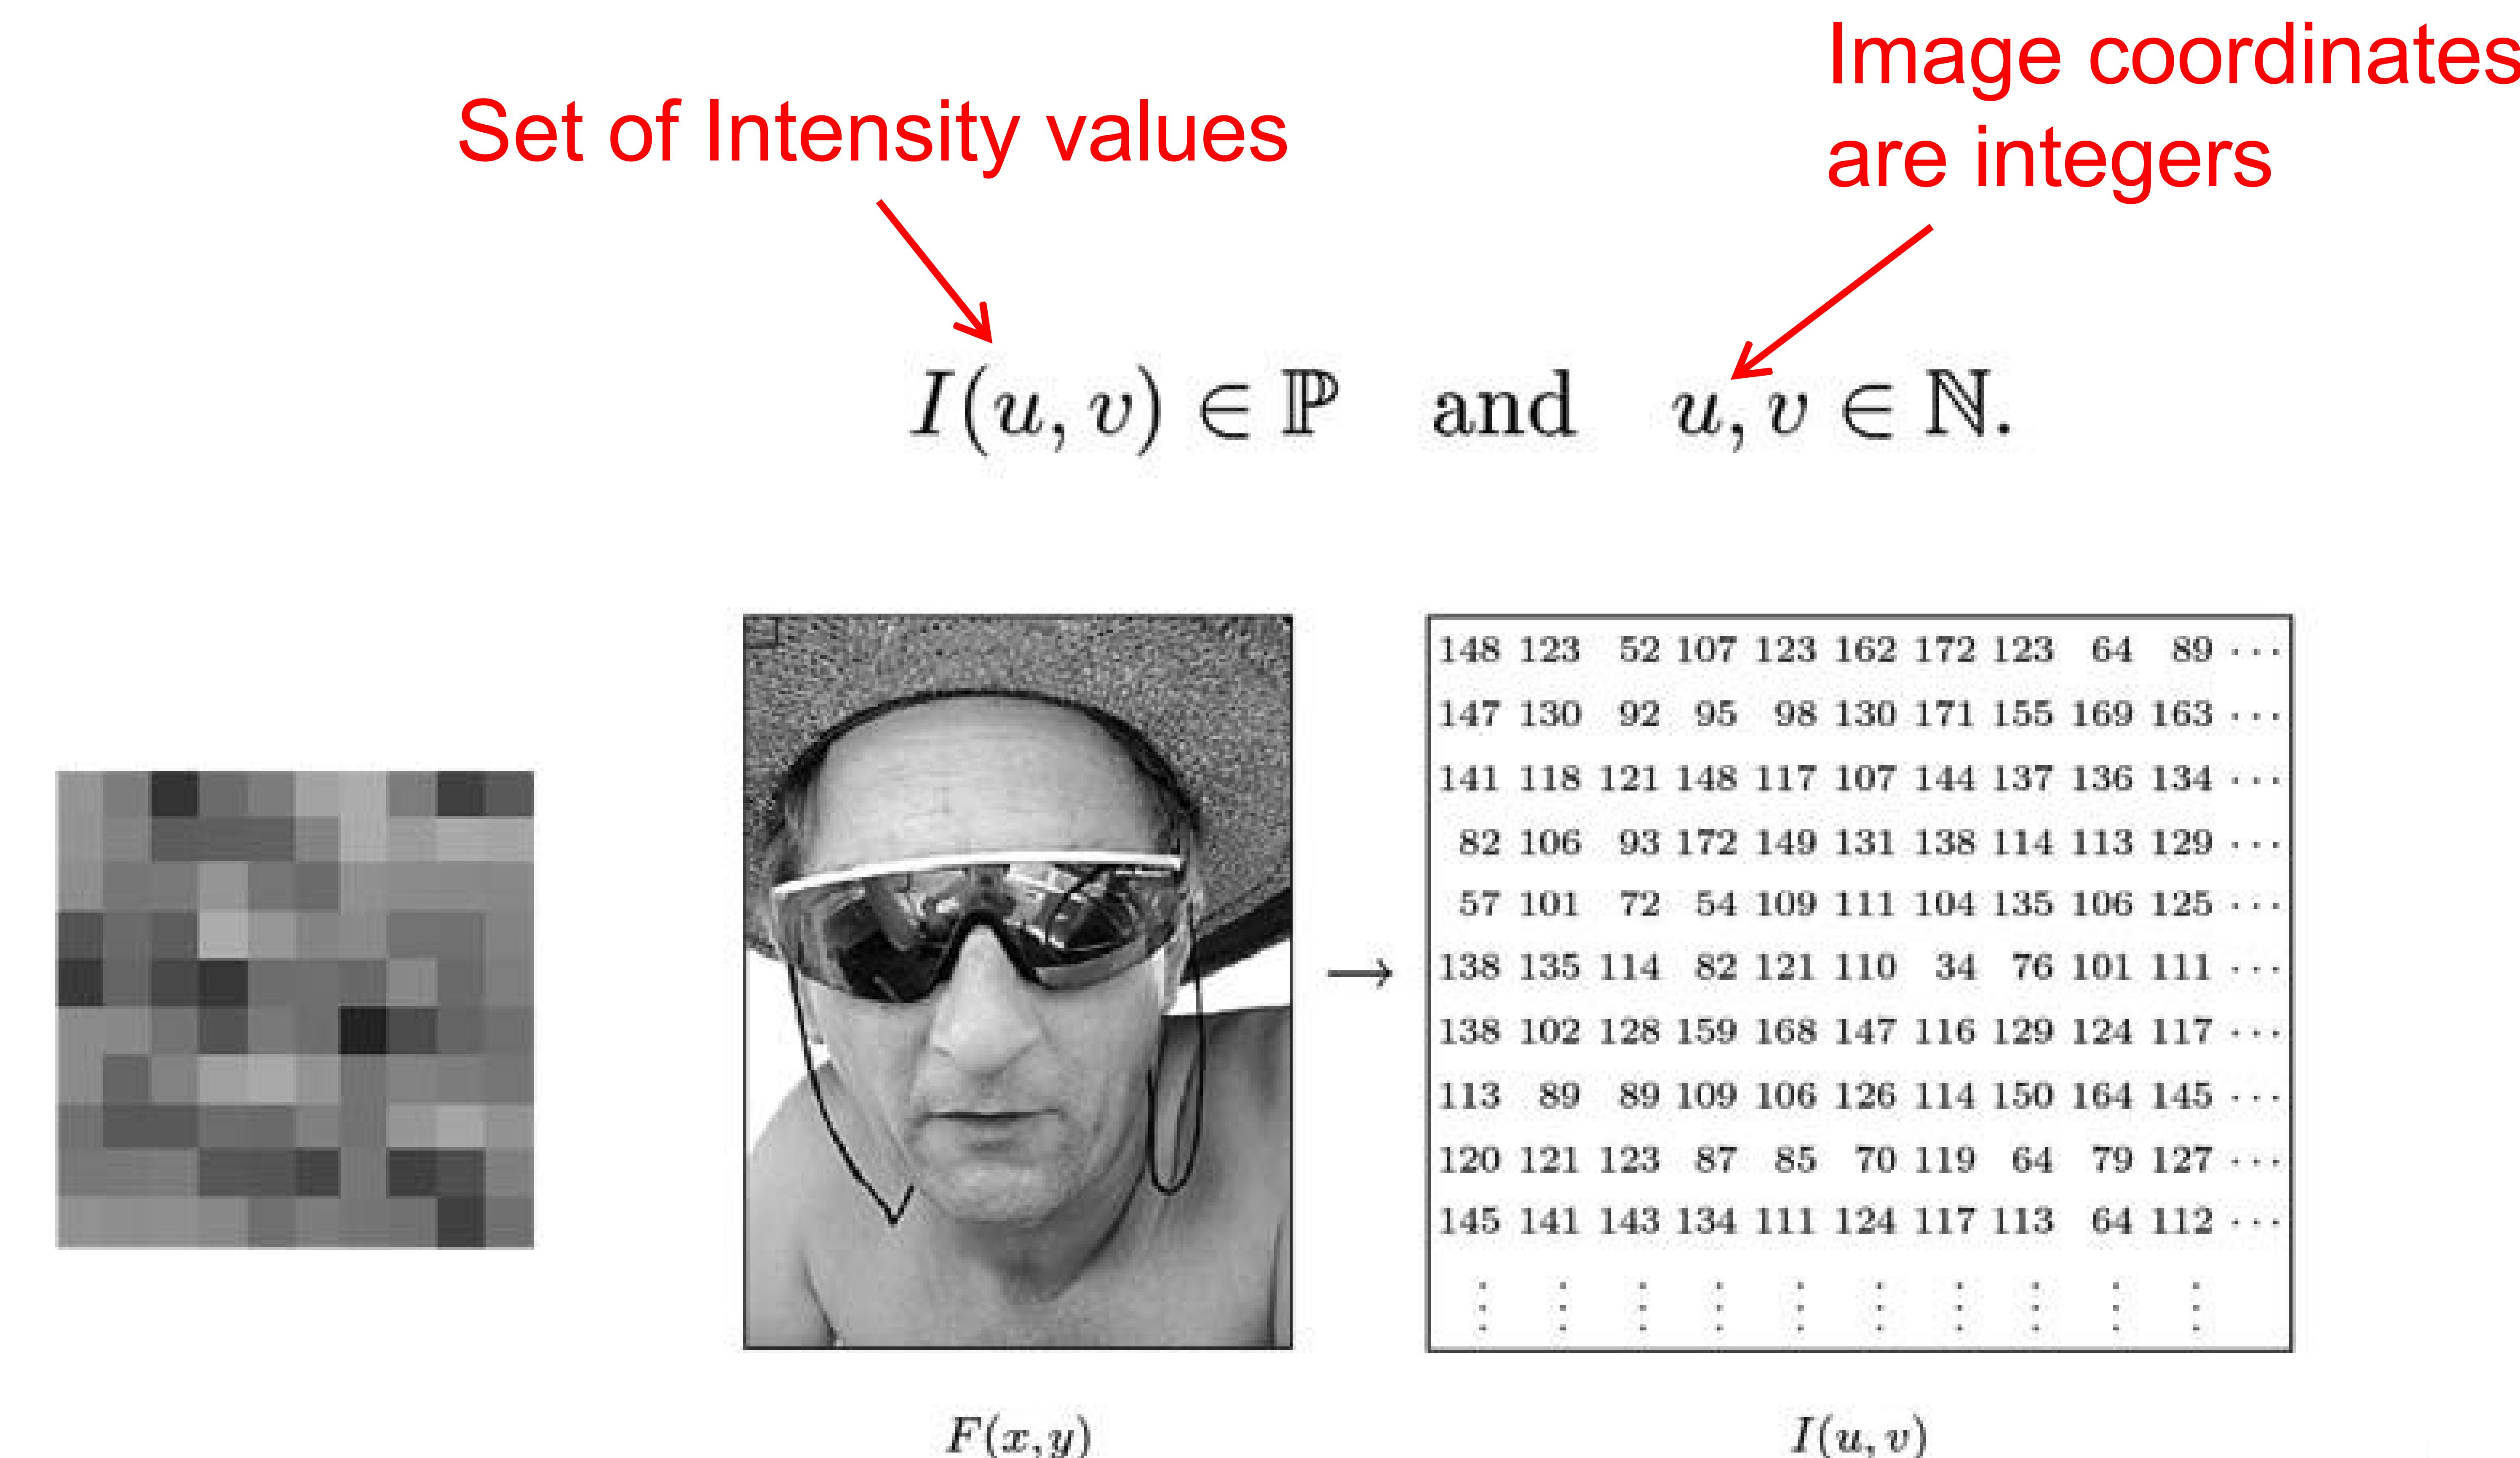
Outline: Image Data Mining

1. Introduction
2. Digital Image Processing
3. Convolutional Neural Network
4. Computer Vision Tasks



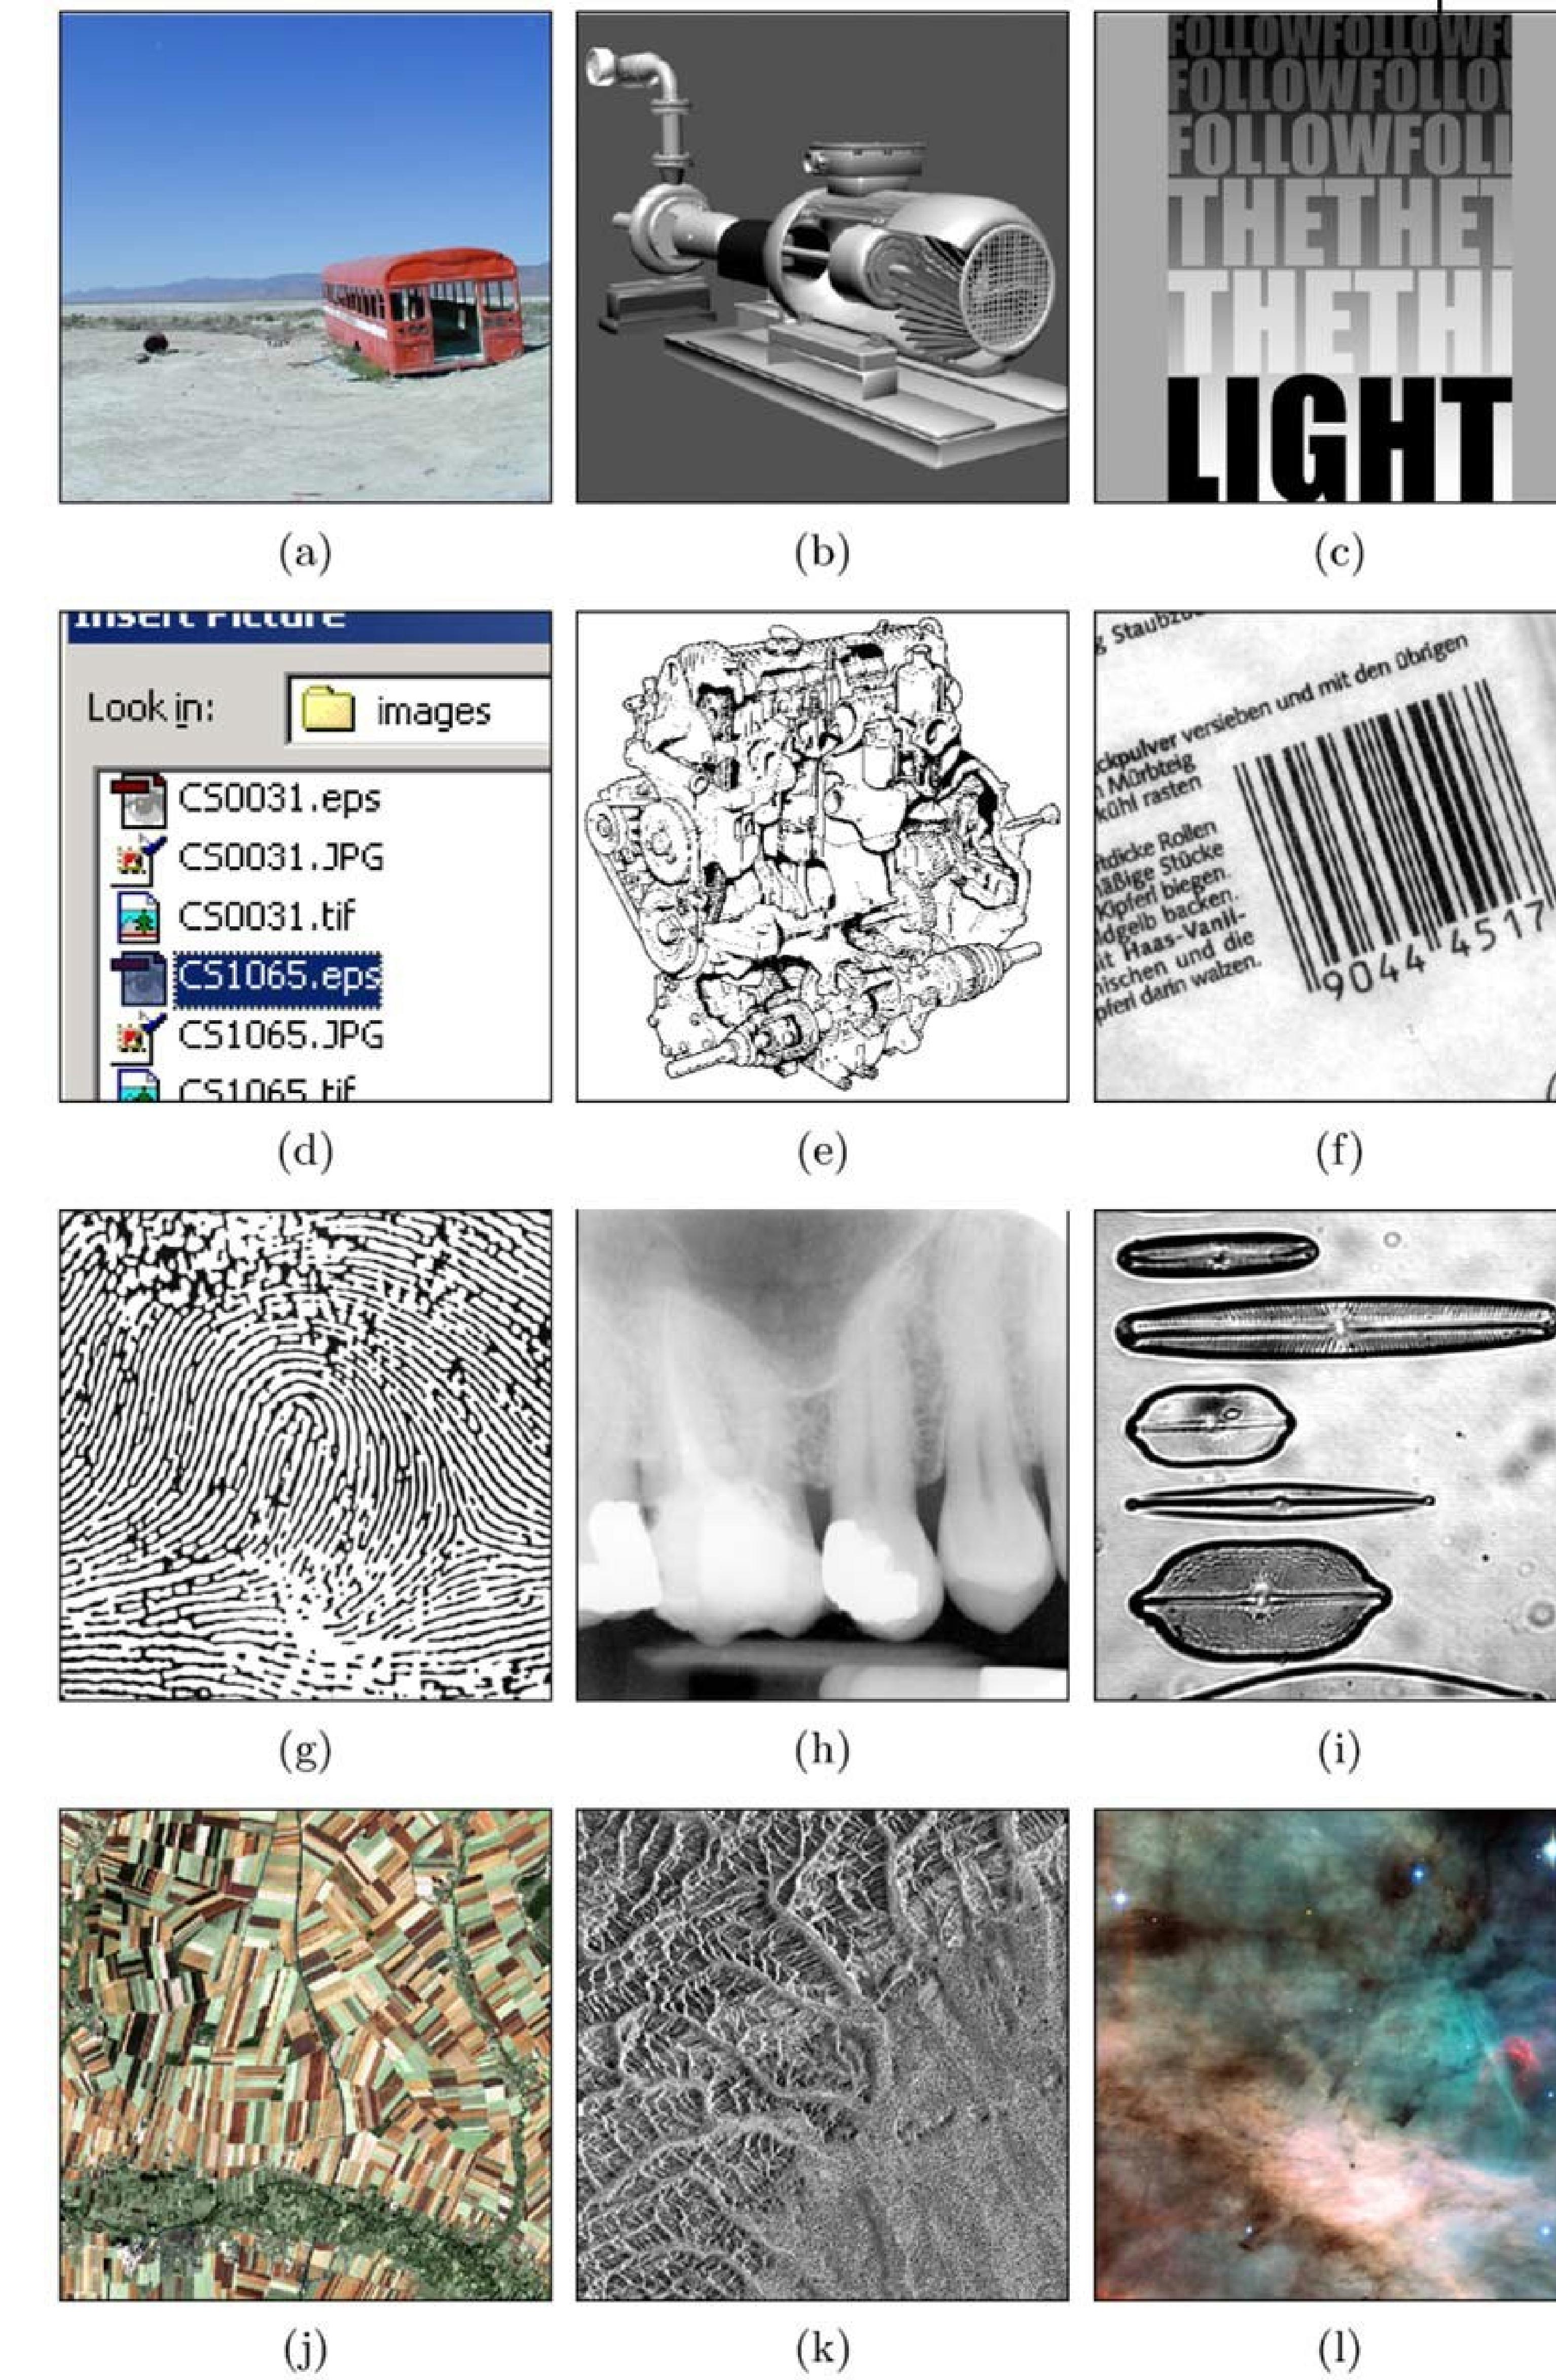
What is an image?

2-dimensional matrix of Intensity (gray or color) values

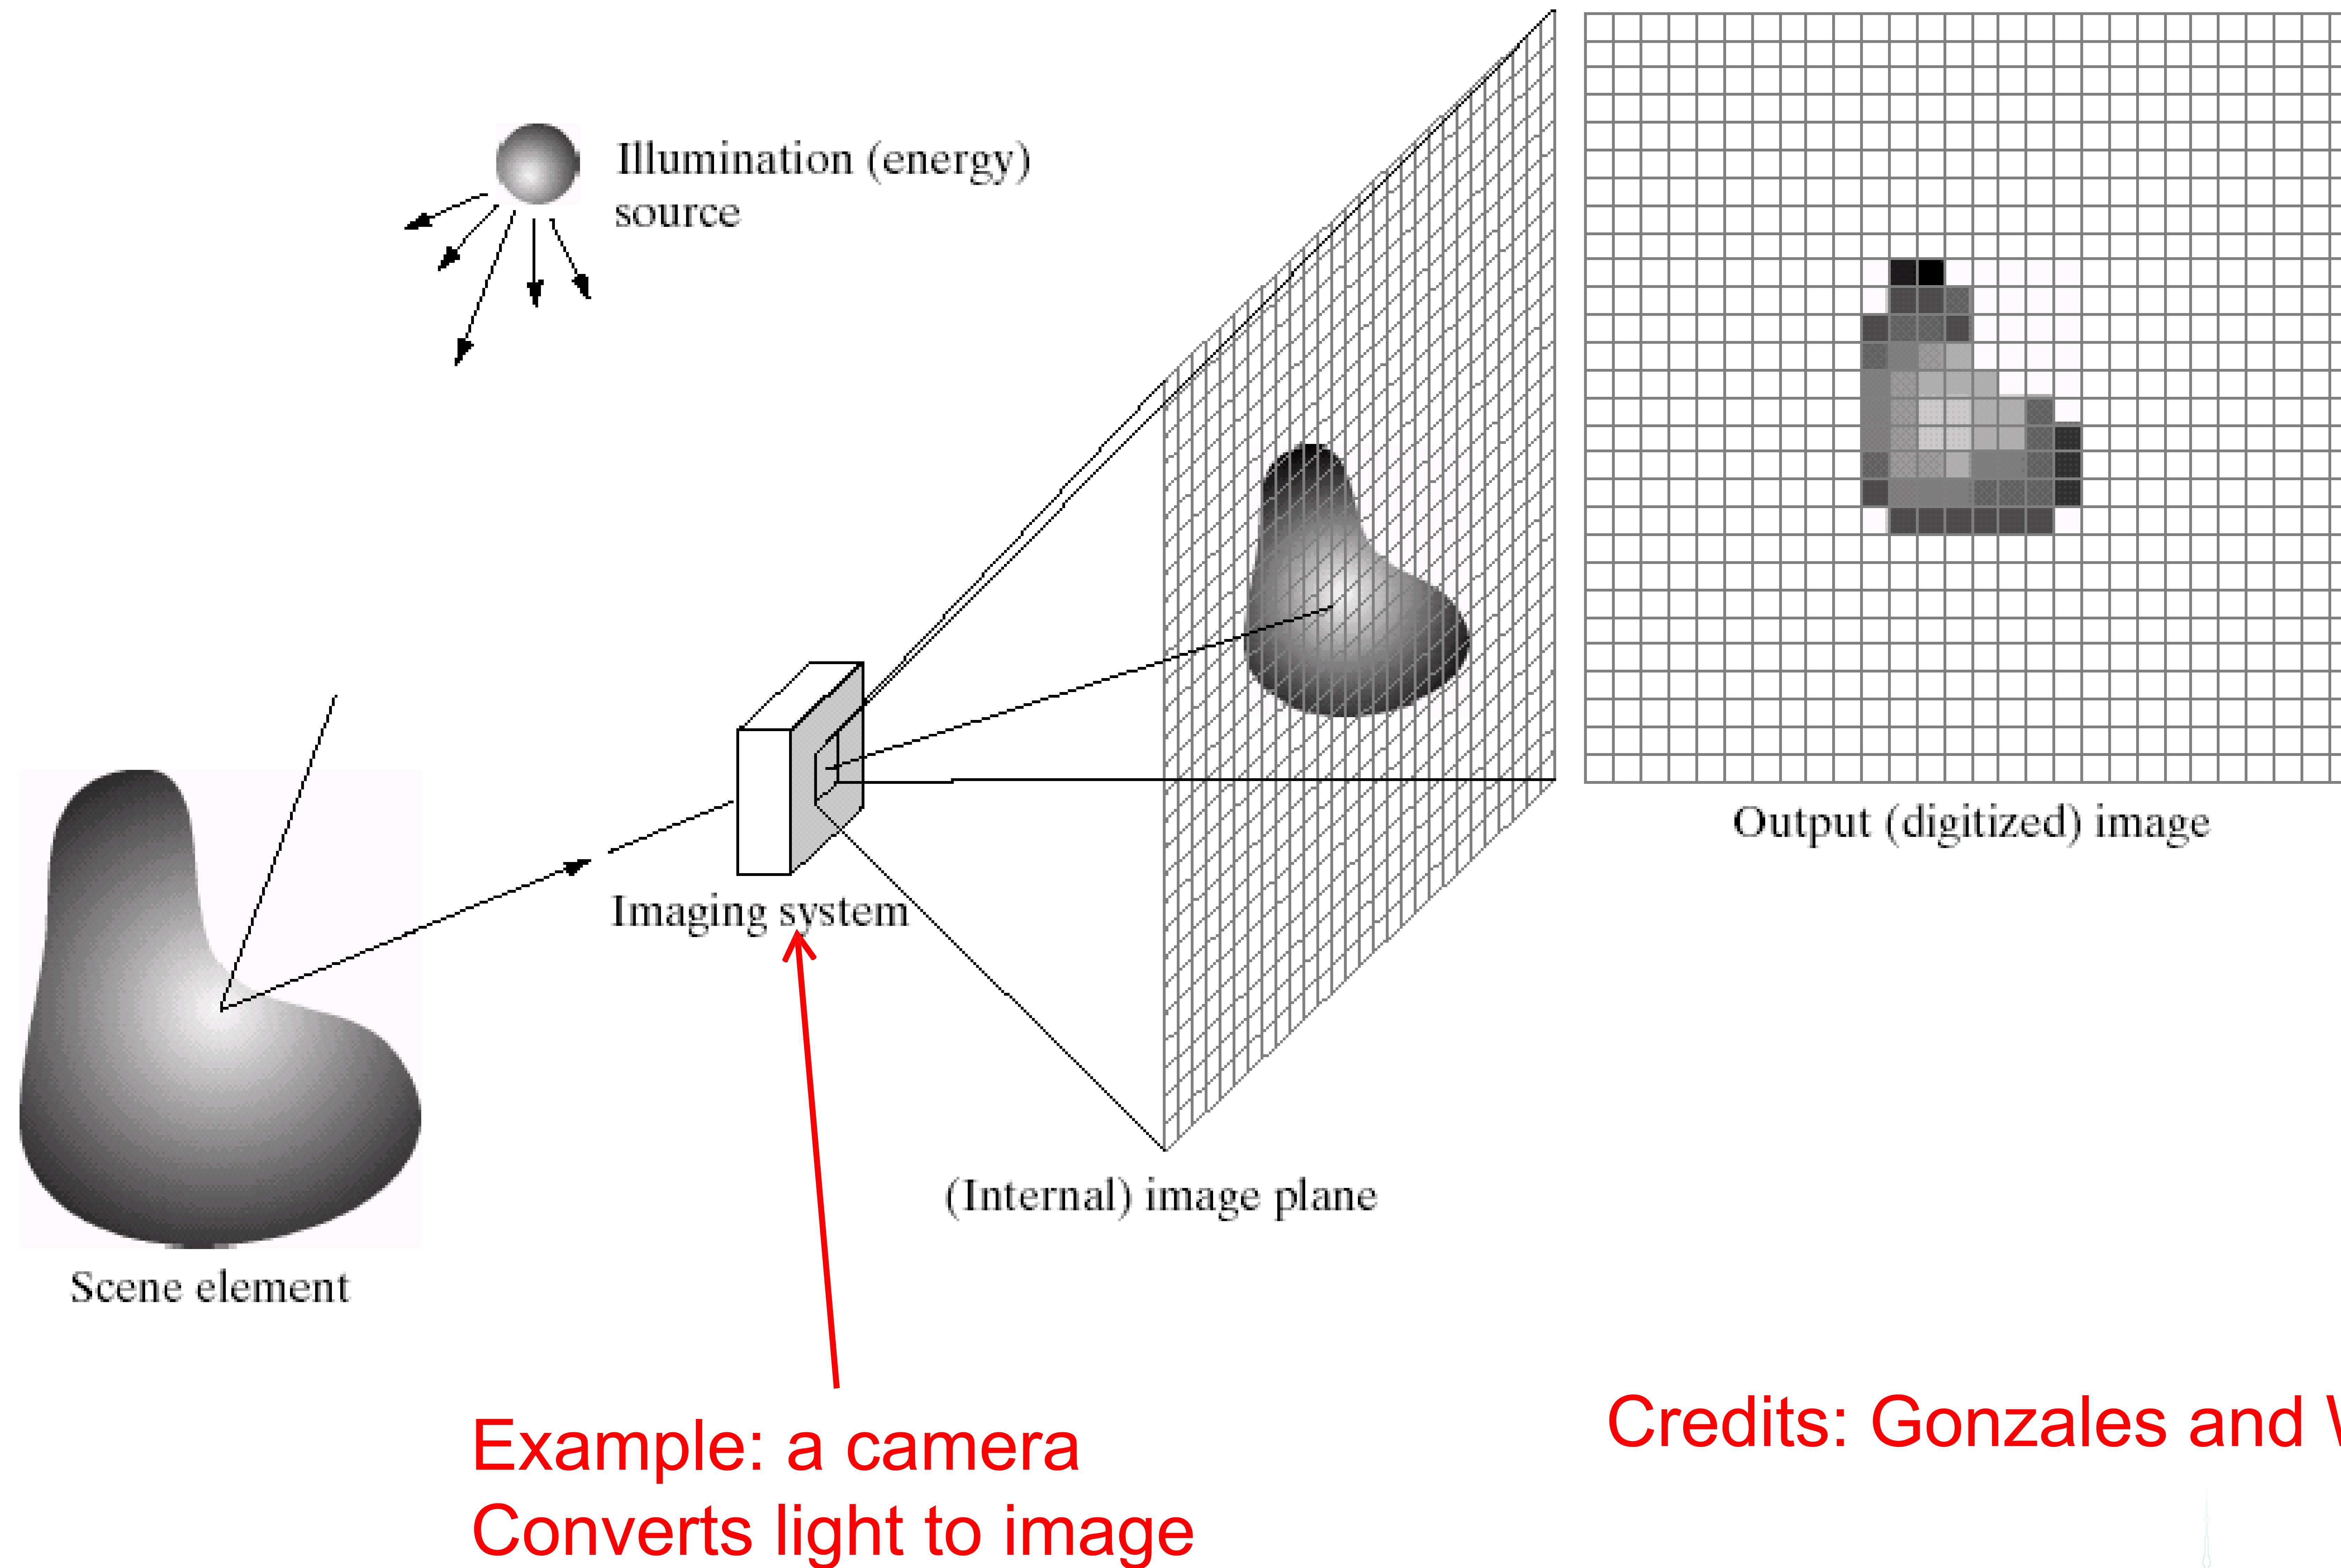


Example of digital images

- a) Natural landscape
- b) Synthetically generated scene
- c) Poster graphic
- d) Computer screenshot
- e) Black and white illustration
- f) Barcode
- g) Fingerprint
- h) X-ray
- i) Microscope slide
- j) Satellite Image
- k) Radar image
- l) Astronomical object

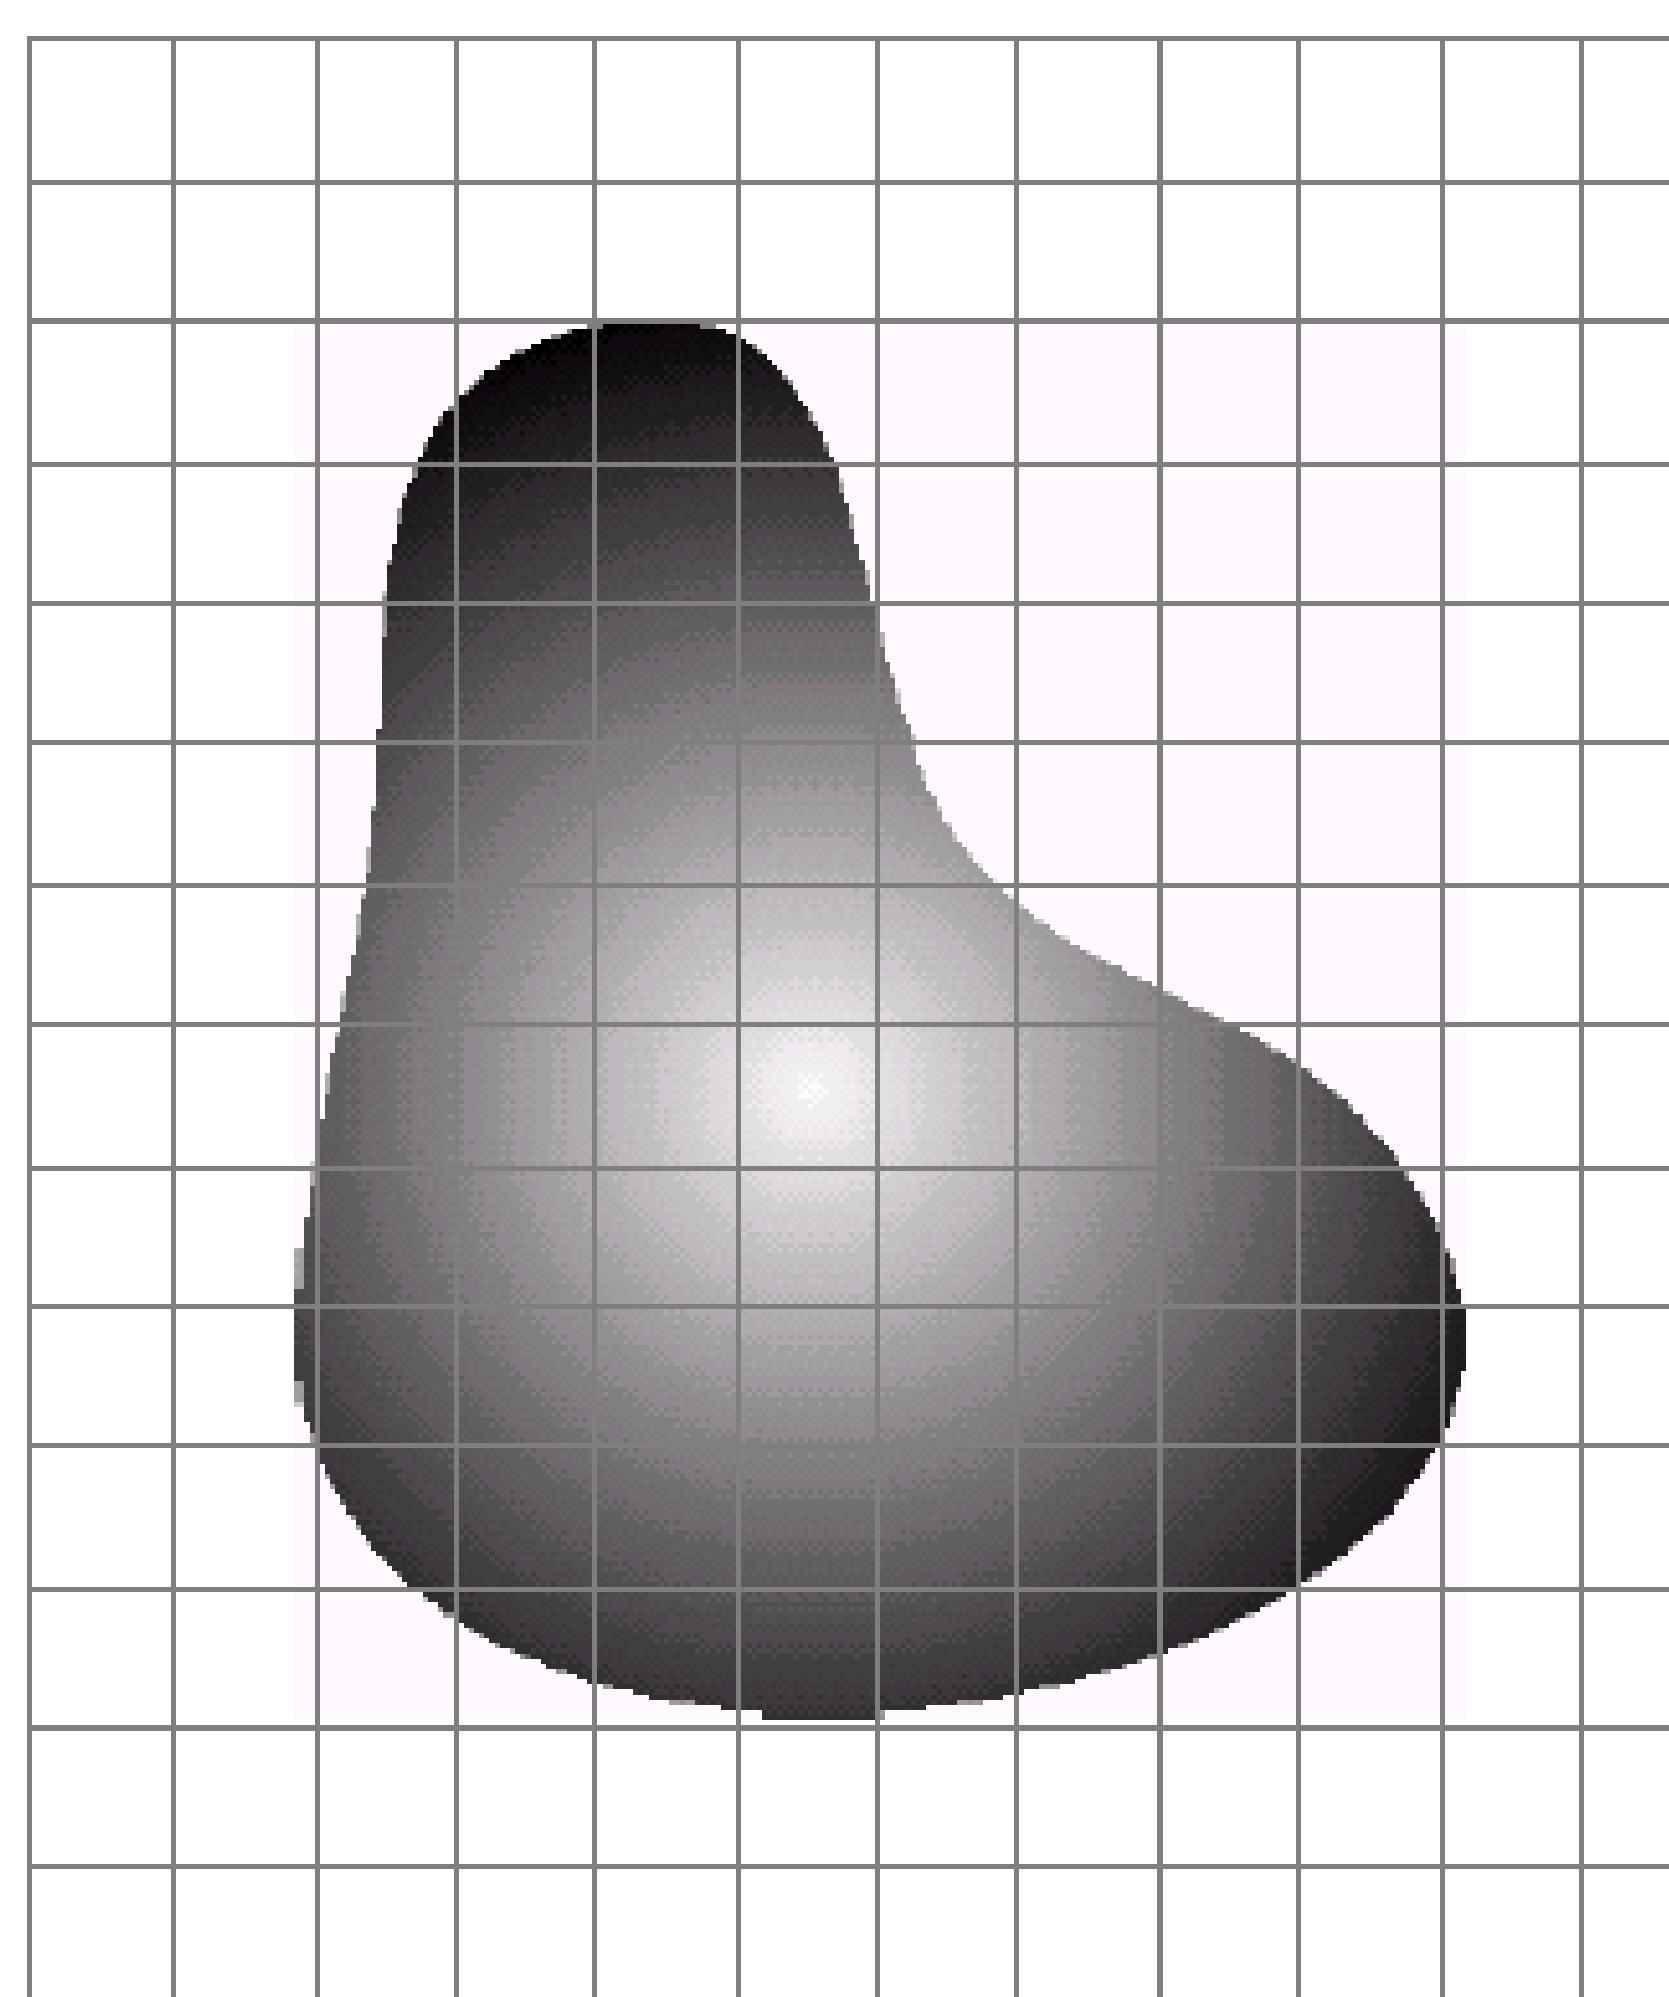


Imaging system

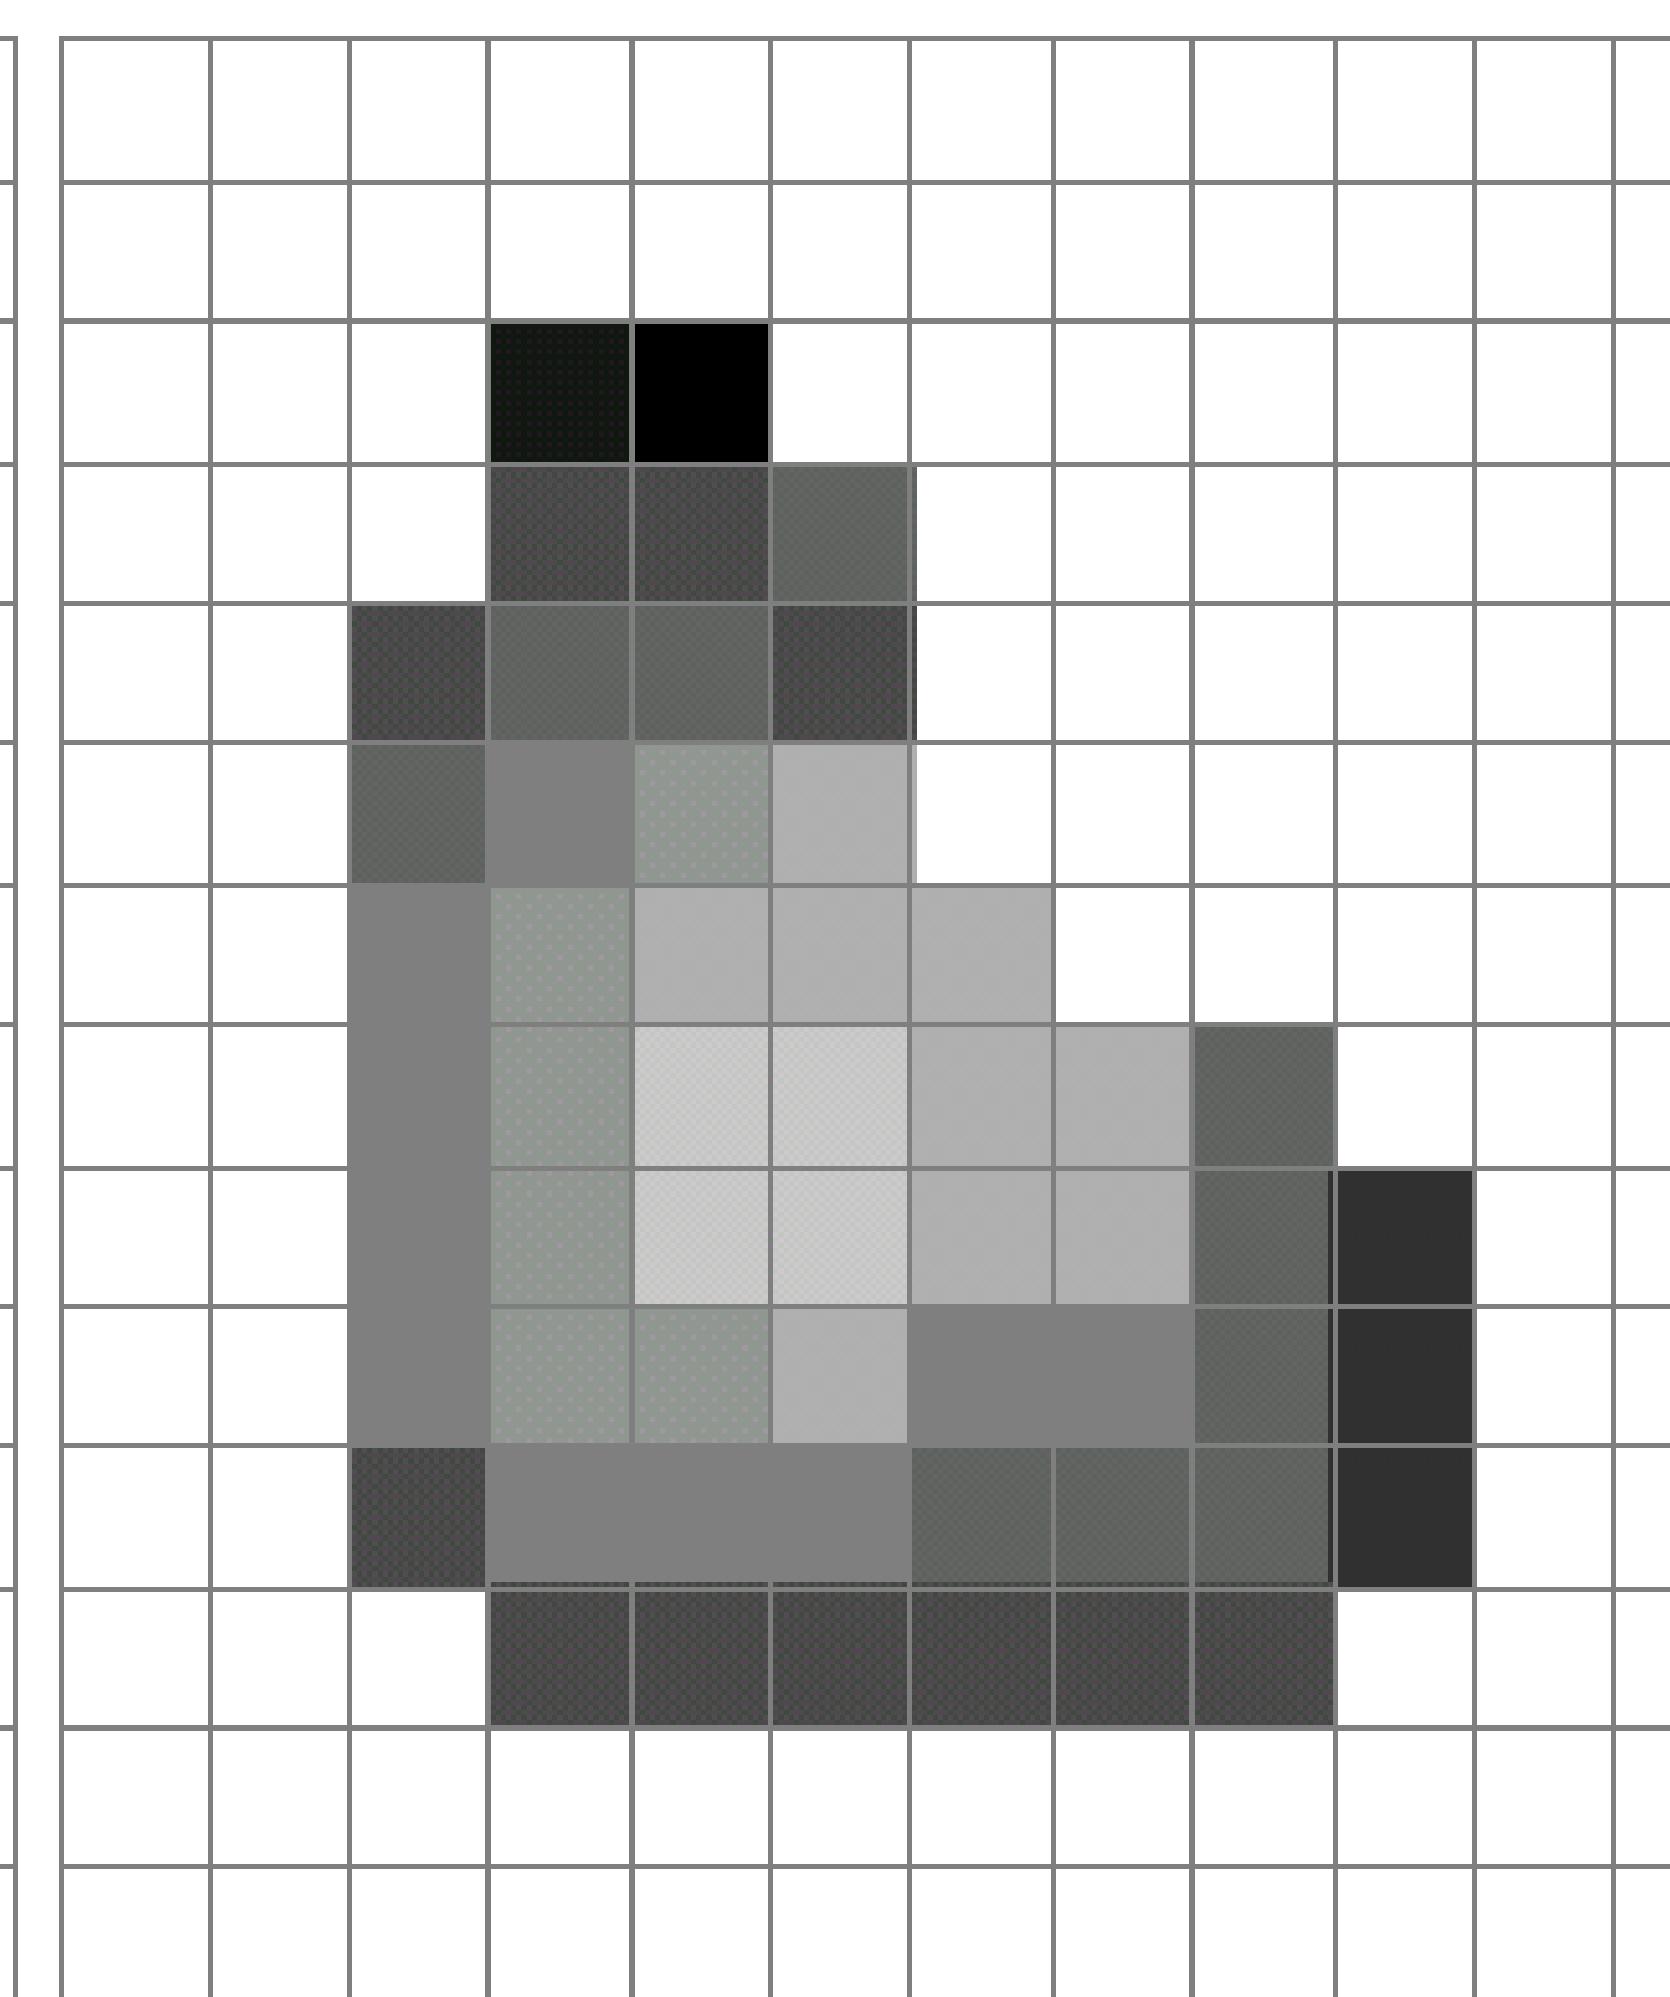


Digital image

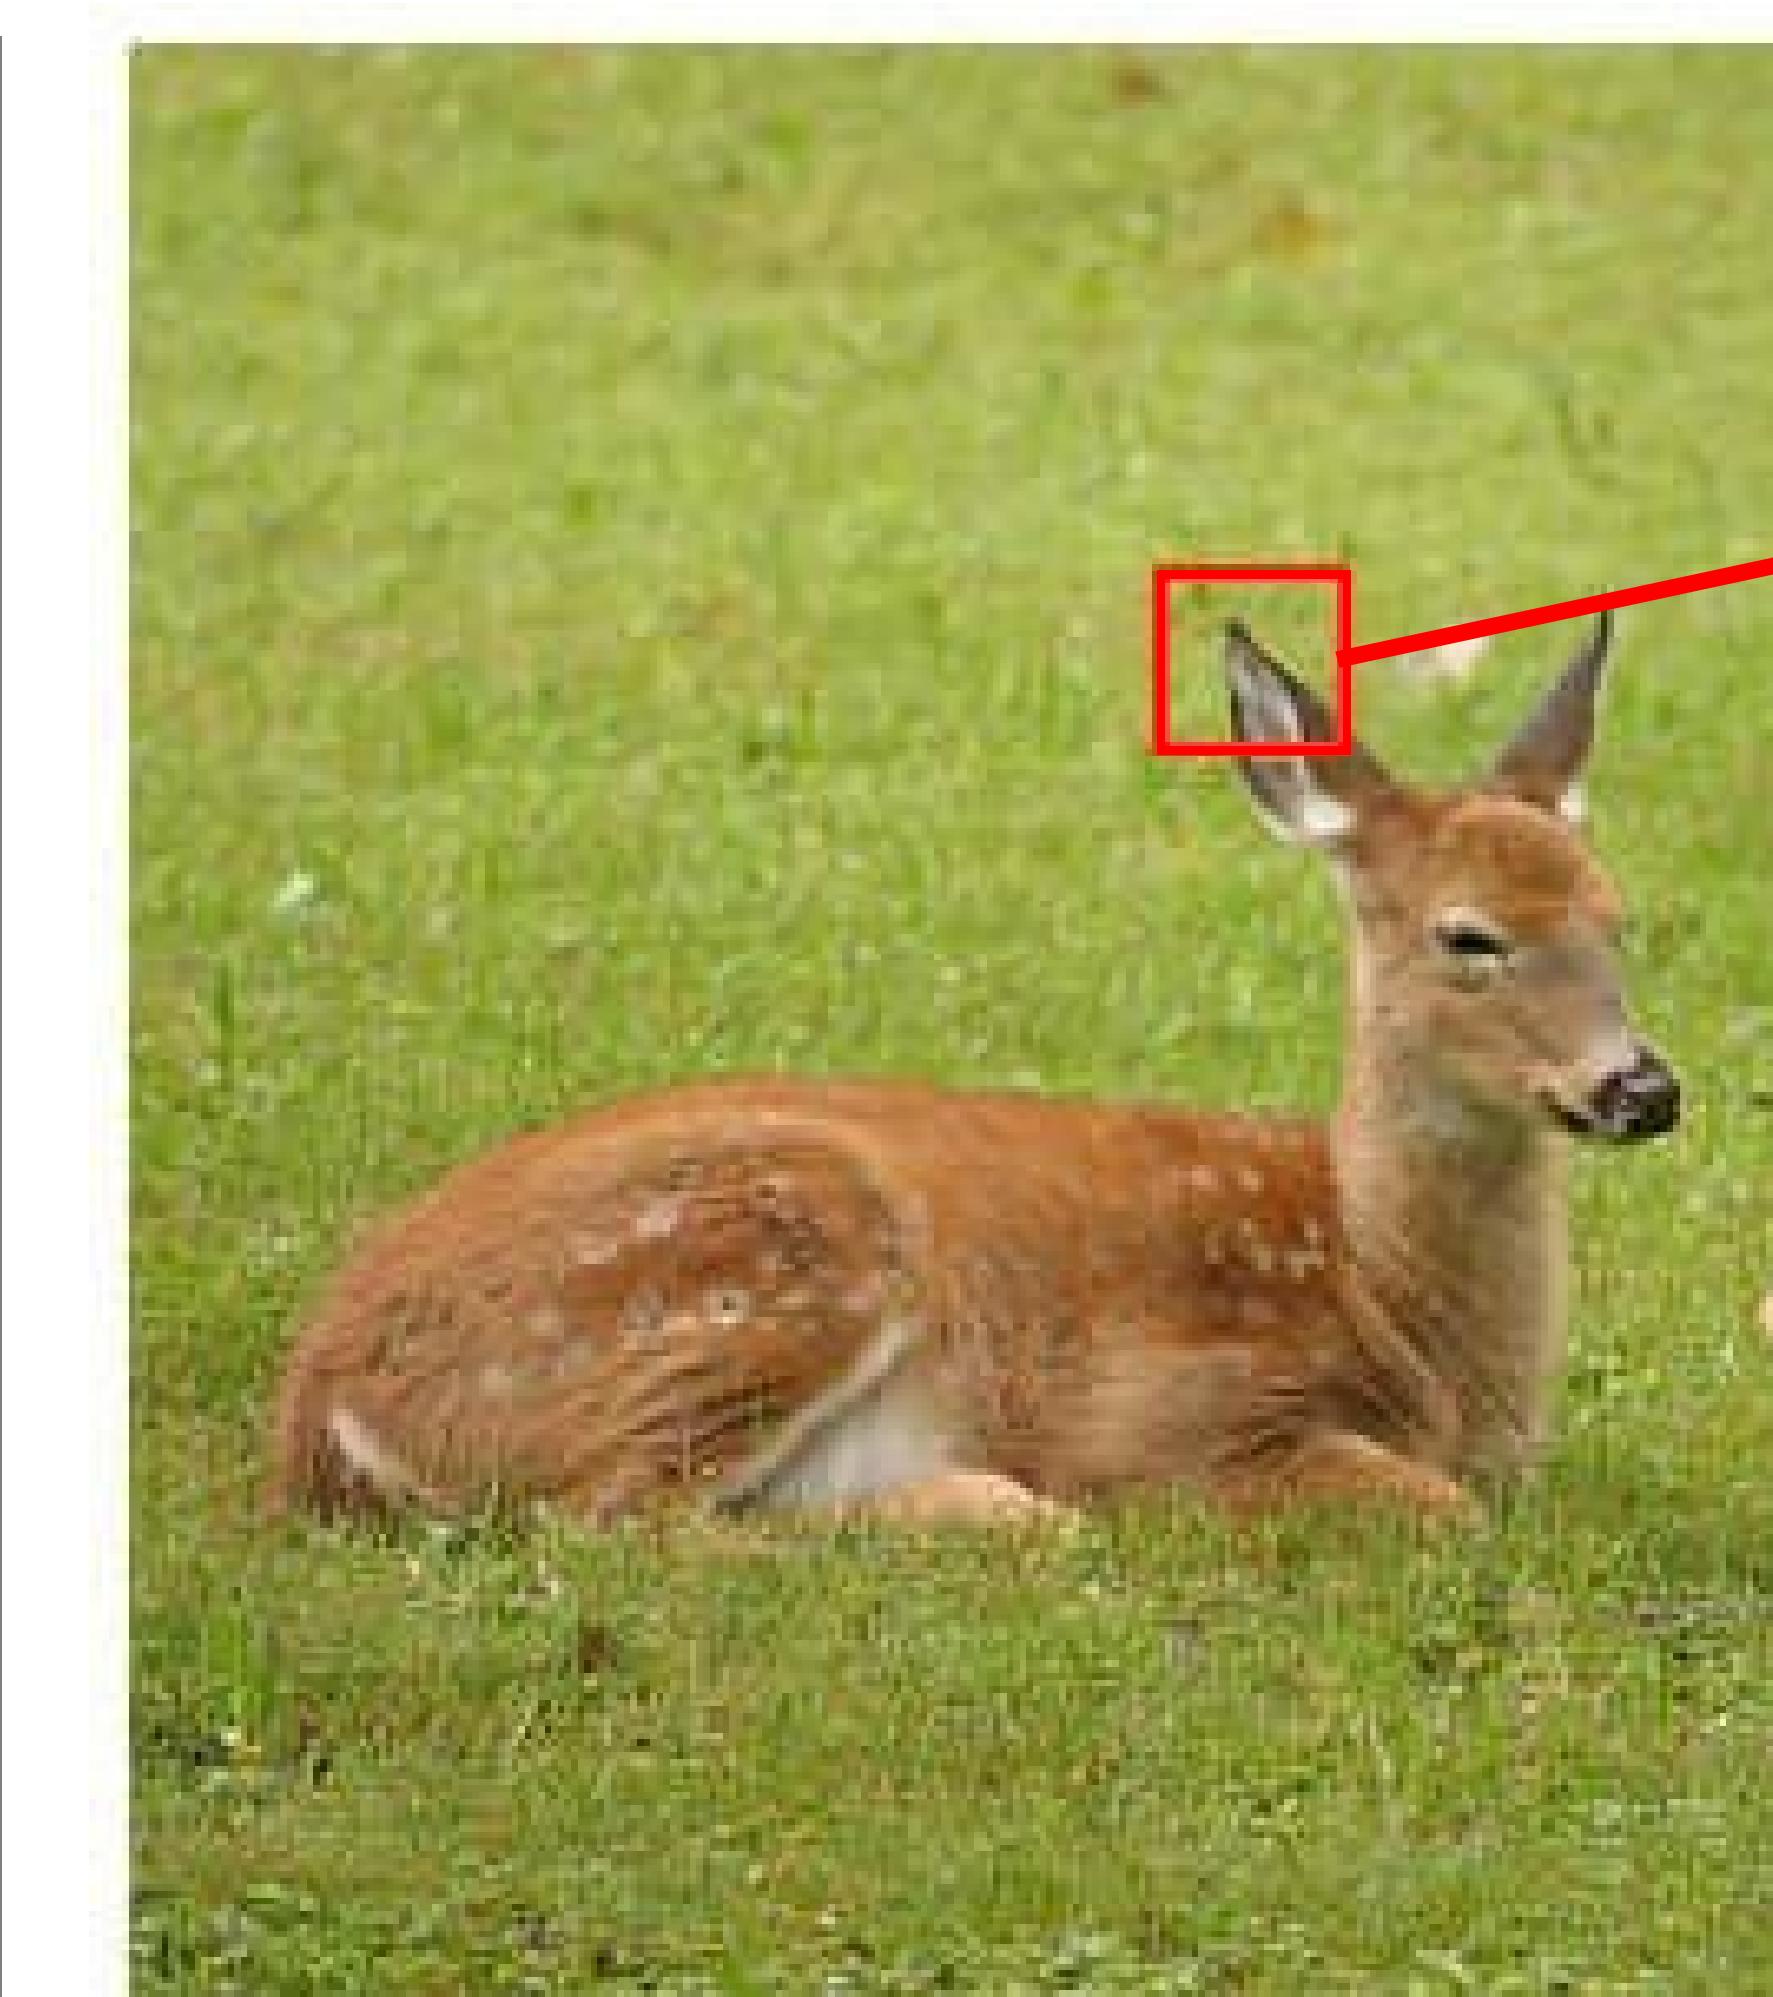
Digitization causes a digital image to become an *approximation* of a real scene



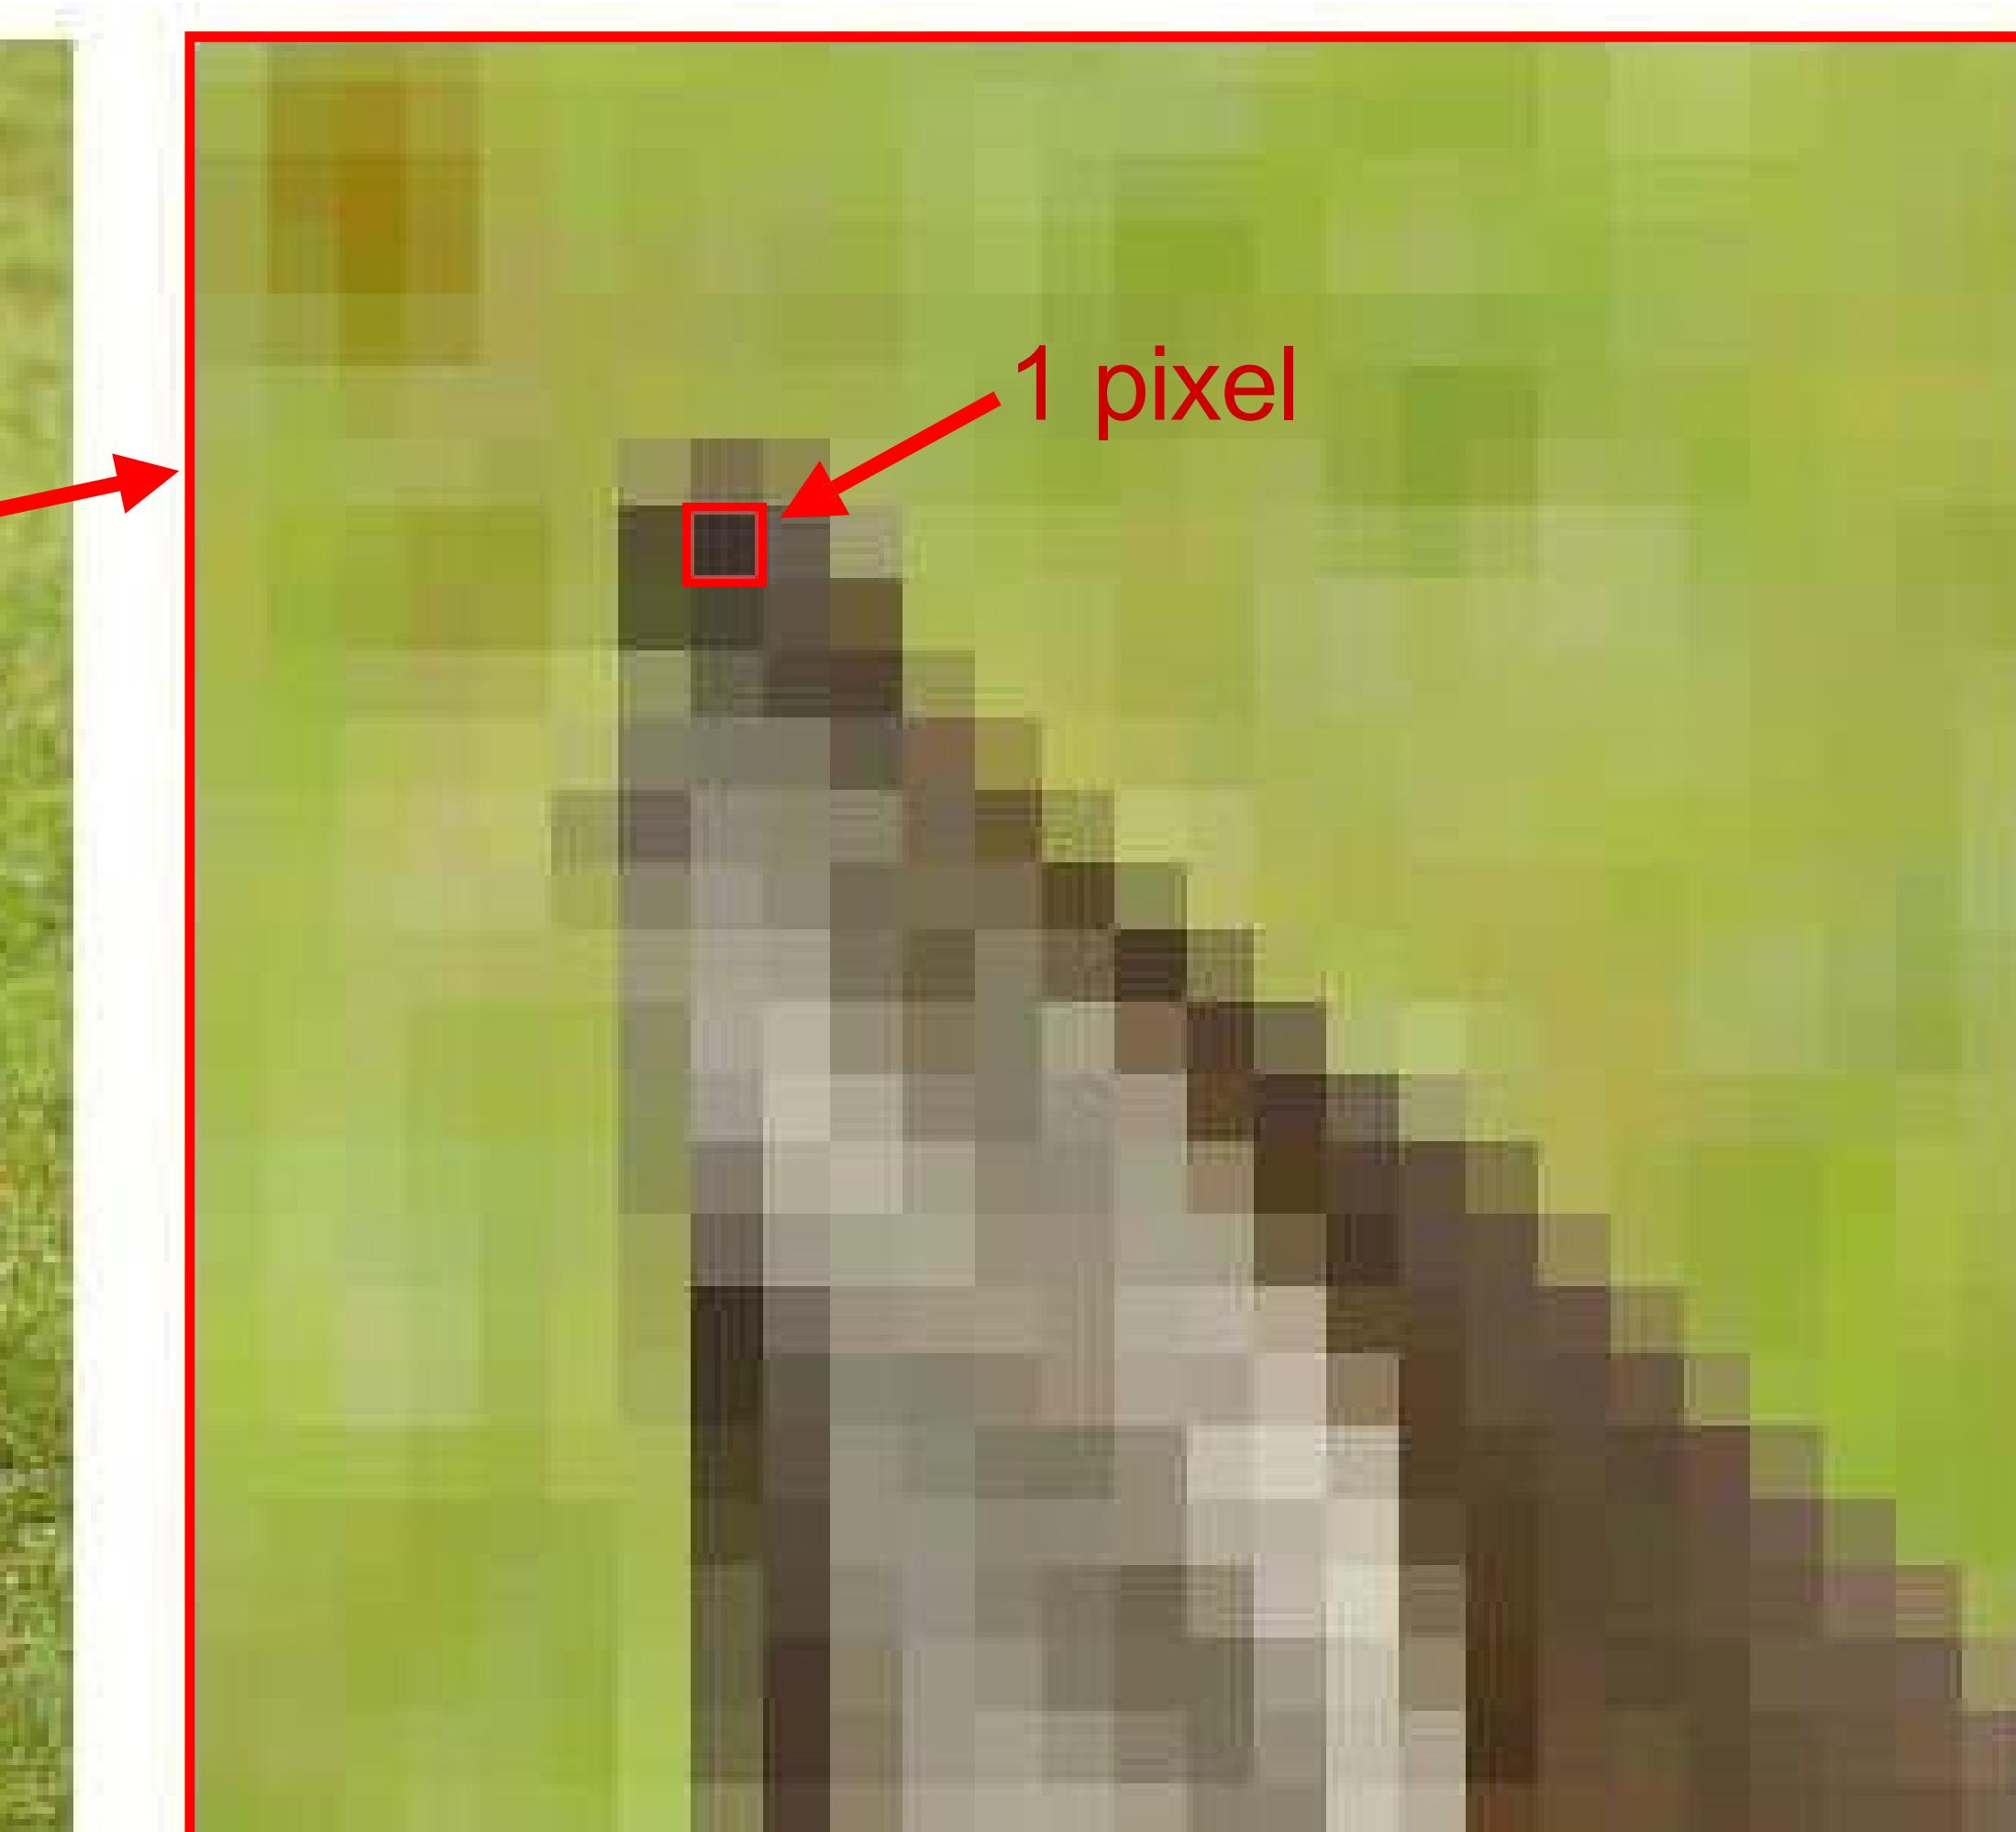
Real image



Digital Image
(an approximation)



Real image



Digital Image
(an approximation)

Digital image

- Common image formats include:
 - 1 values per point/pixel (B&W or Grayscale)
 - 3 values per point/pixel (Red, Green, and Blue)
 - 4 values per point/pixel (Red, Green, Blue, + “Alpha” or Opacity)



Grayscale



RGB

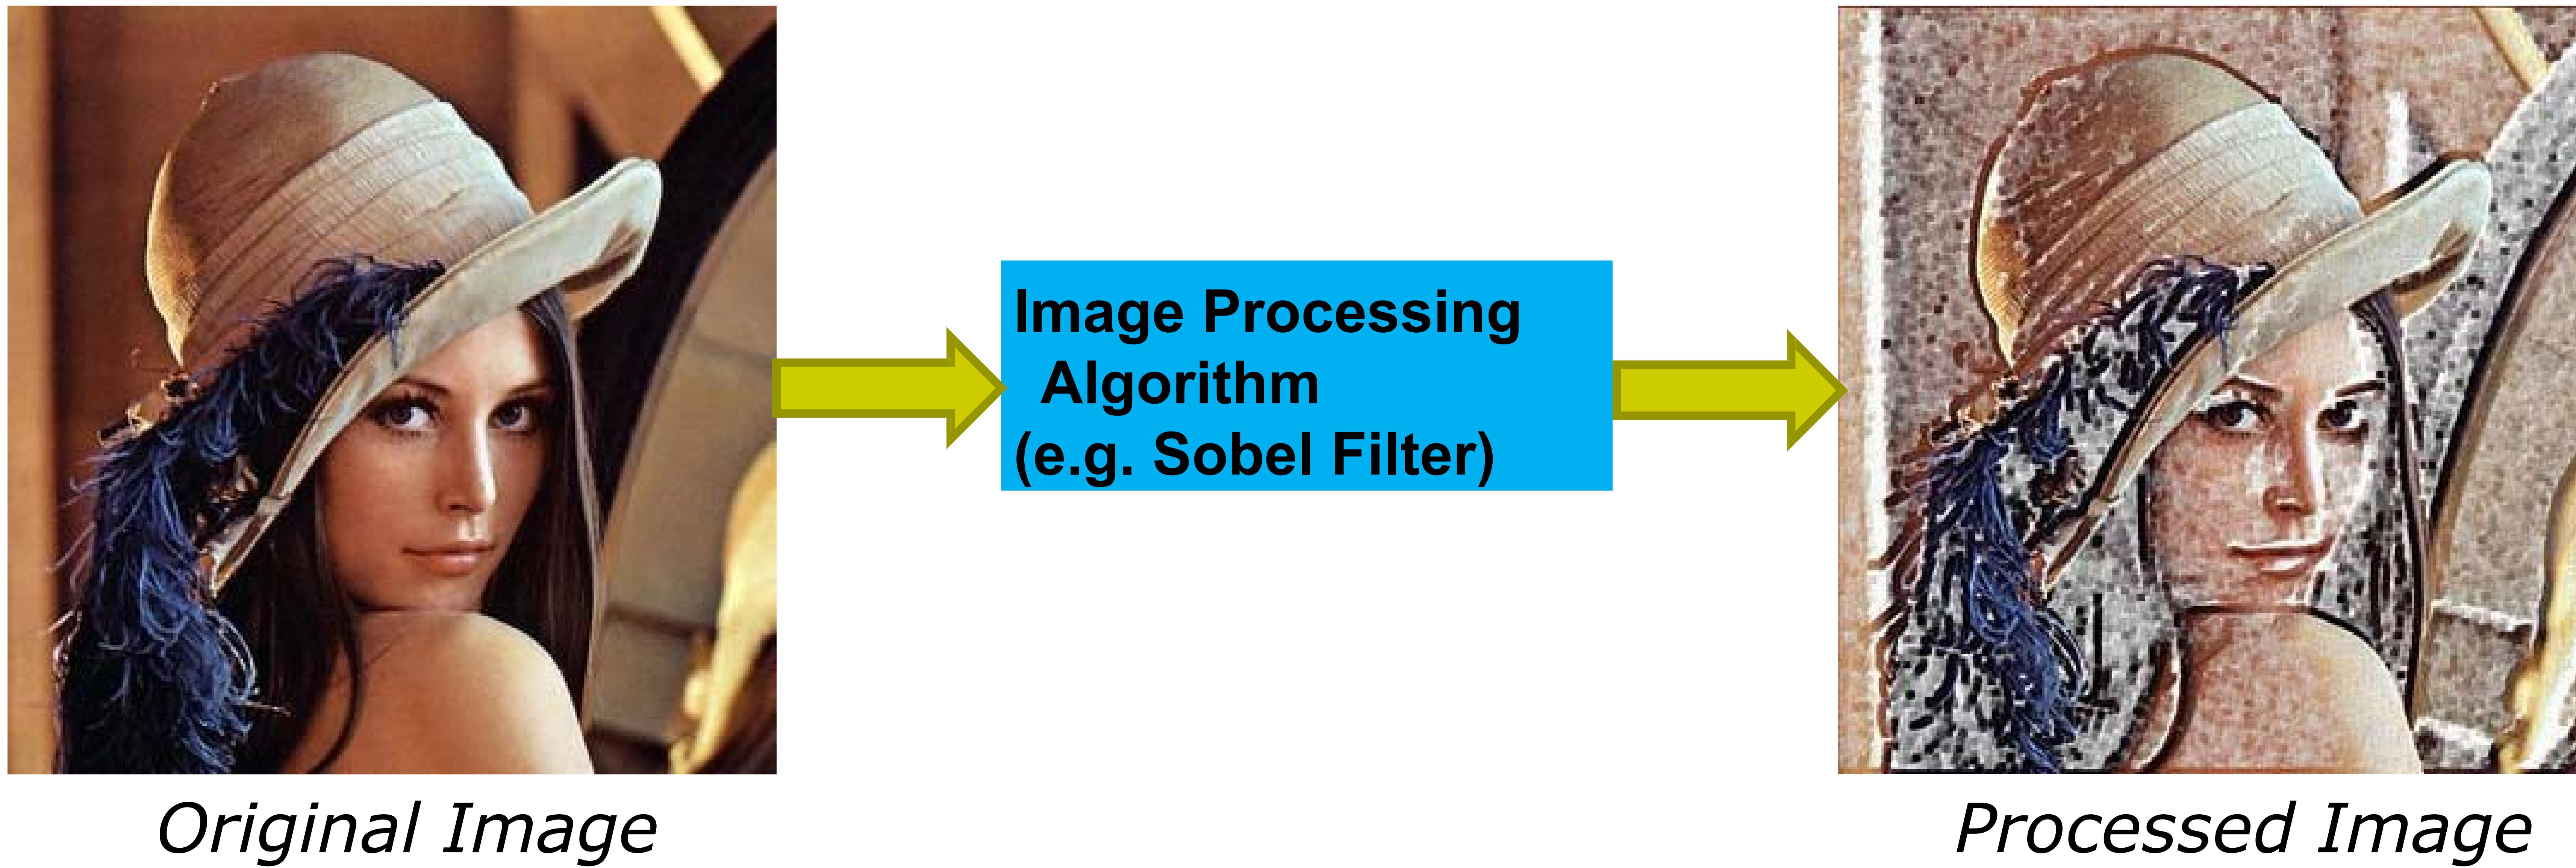


RGBA

- We will start with gray-scale images, extend to color later

What is image processing?

- Algorithms that alter an input image to create new image
- Input is image, output is image



- Improves an image for human interpretation in ways including:
 - Image display and printing
 - Image editing
 - Image enhancement
 - Image compression

Example operation: noise removal

Noisy Image

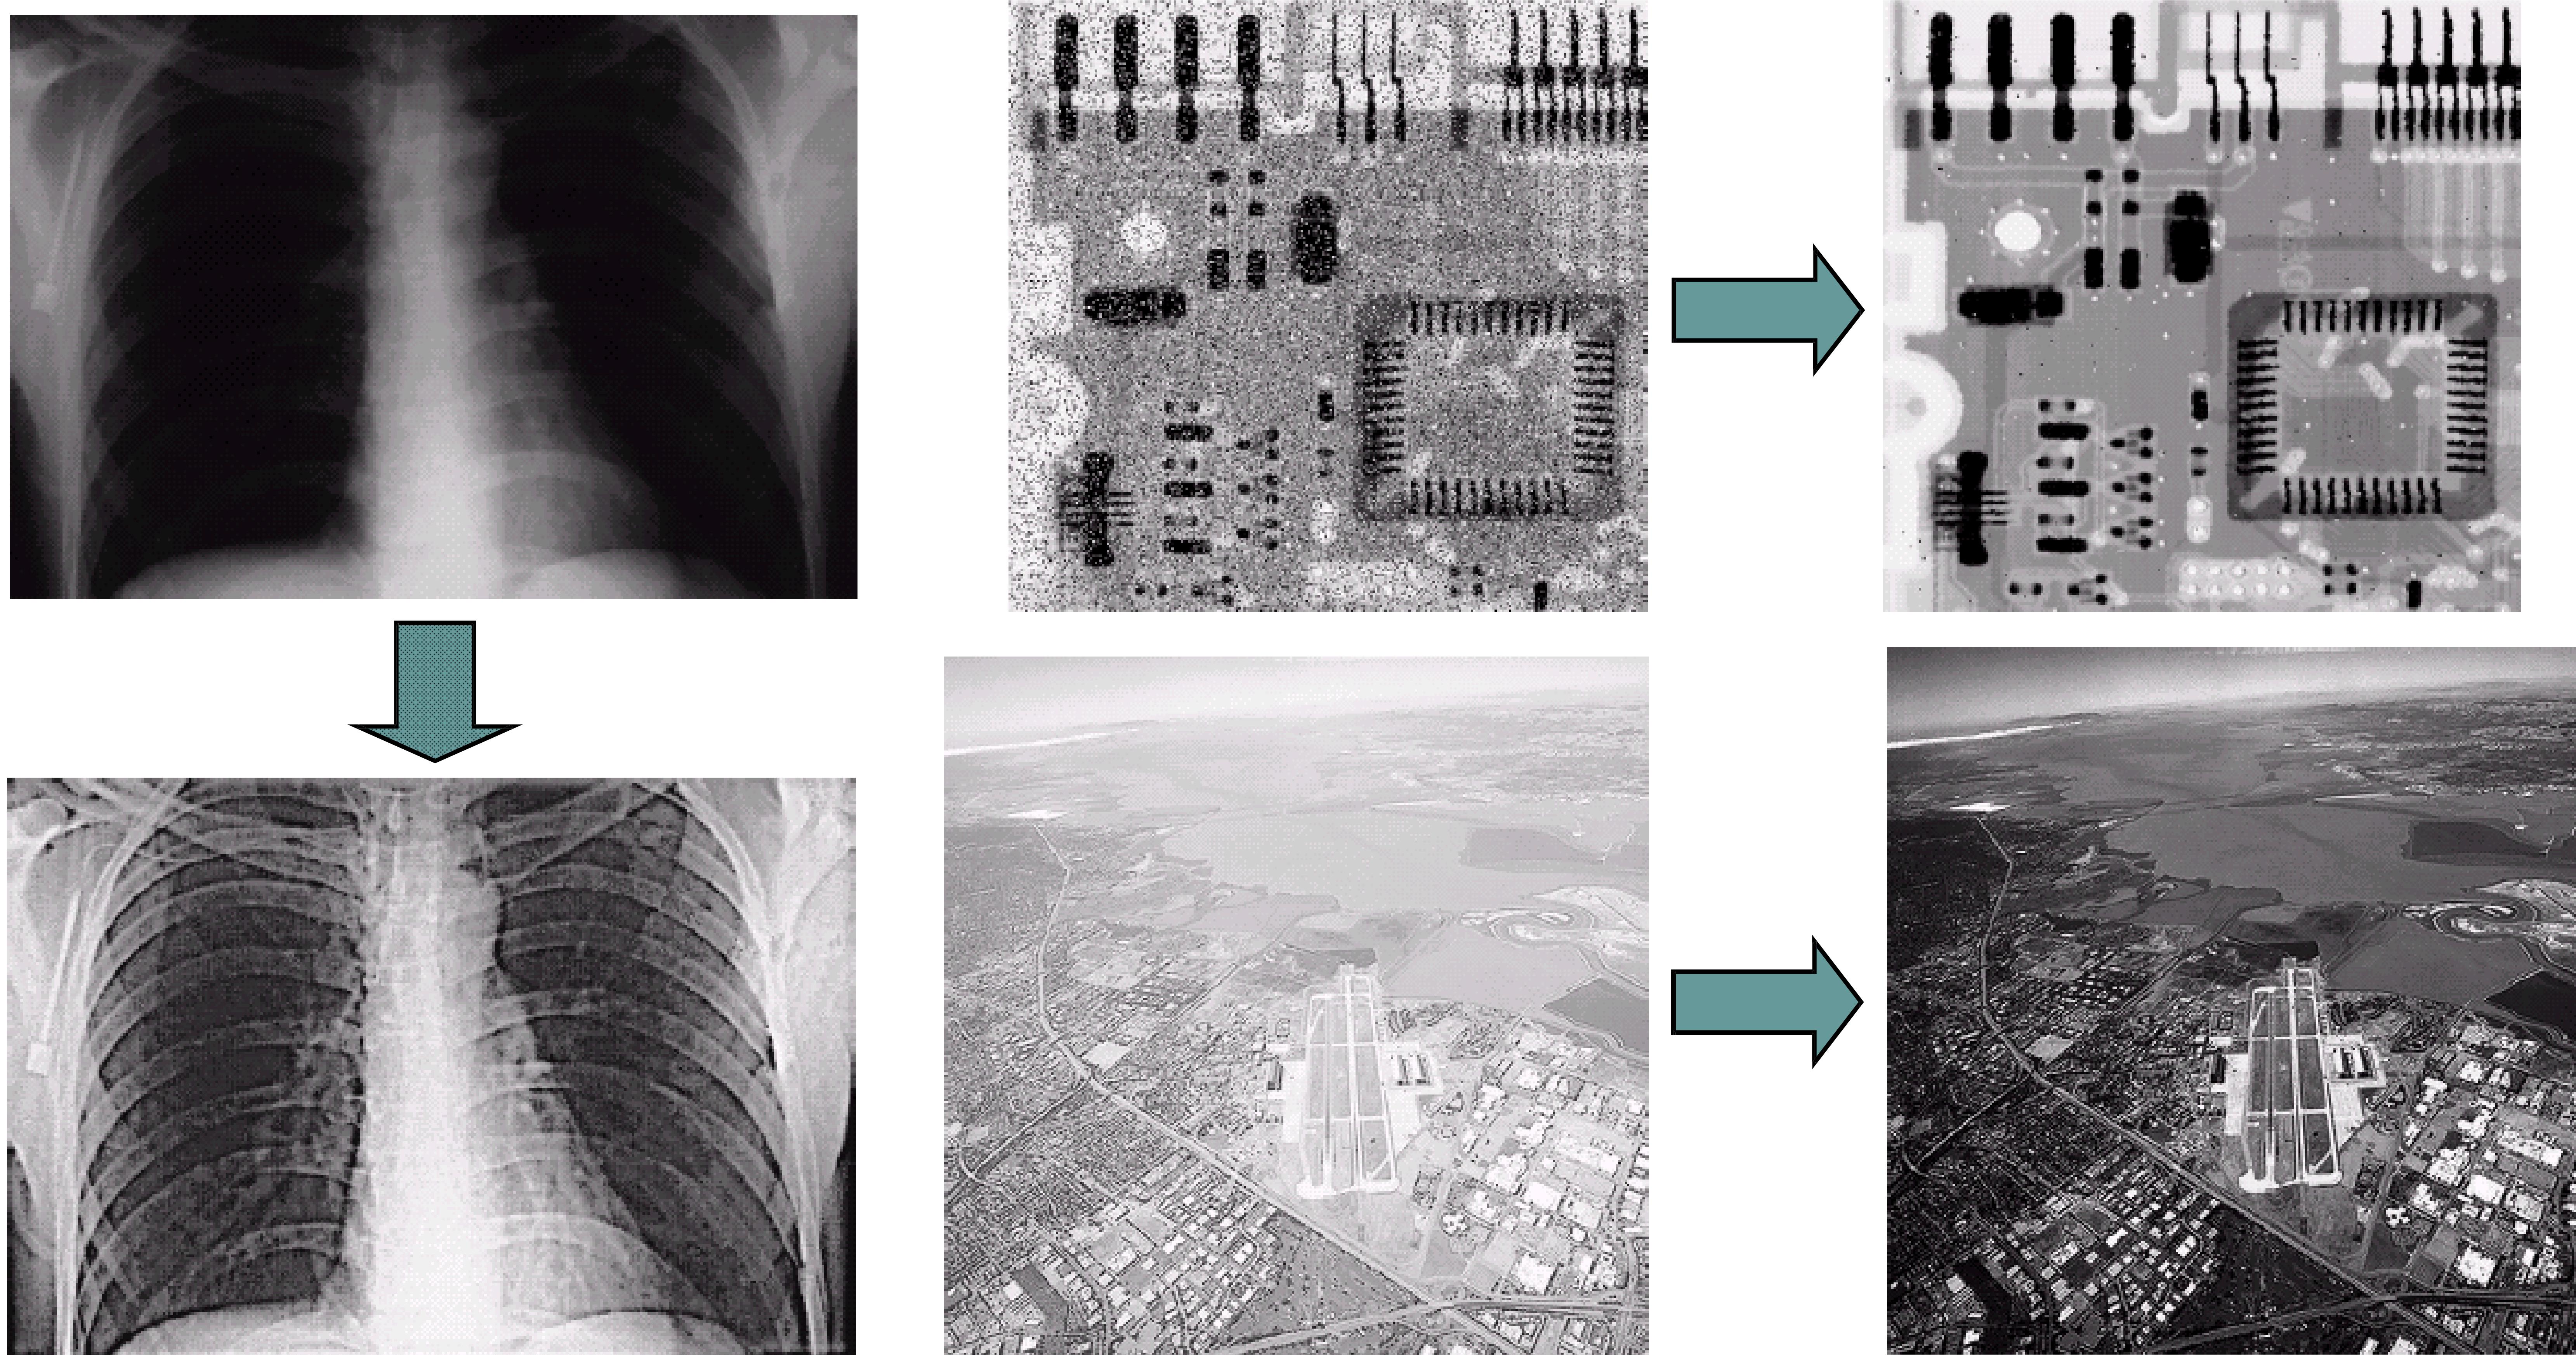


Denoised Image



Think of noise as white specks on a picture (random or non-random)

Example operation: noise removal



Example: contrast adjustment



Low Contrast

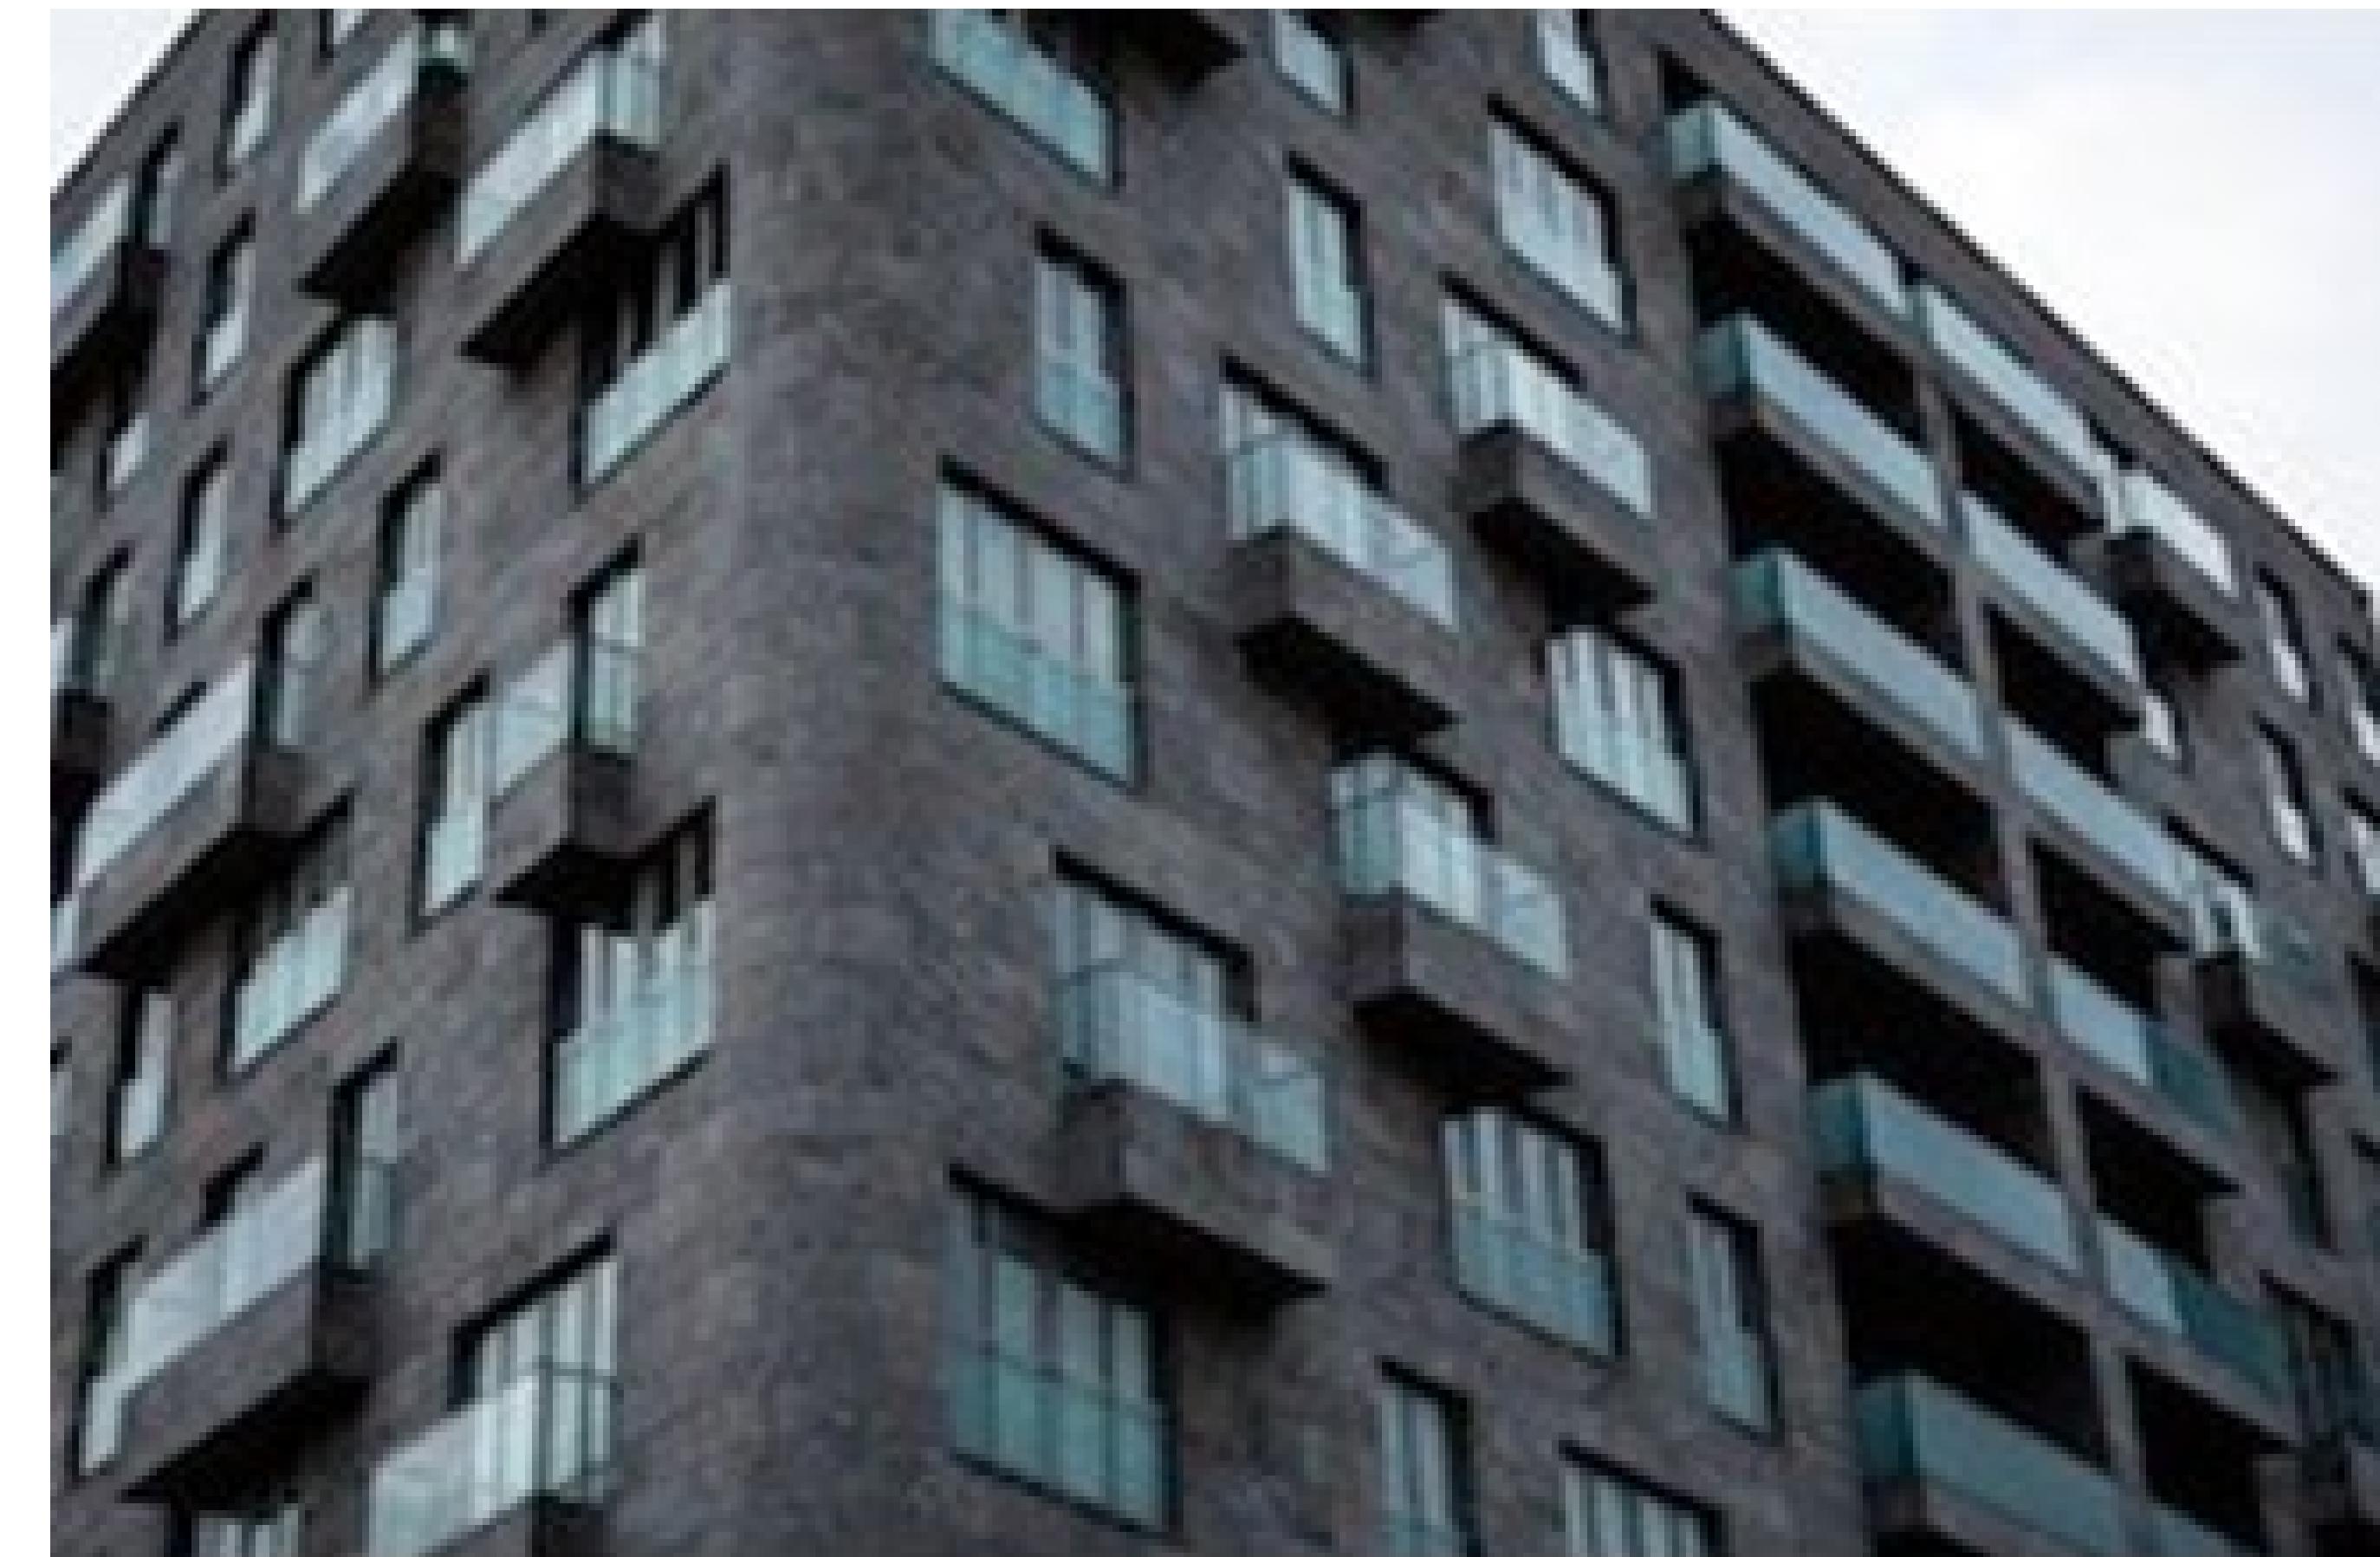


Original Contrast



High Contrast

Example: image super-resolution



input



4x original

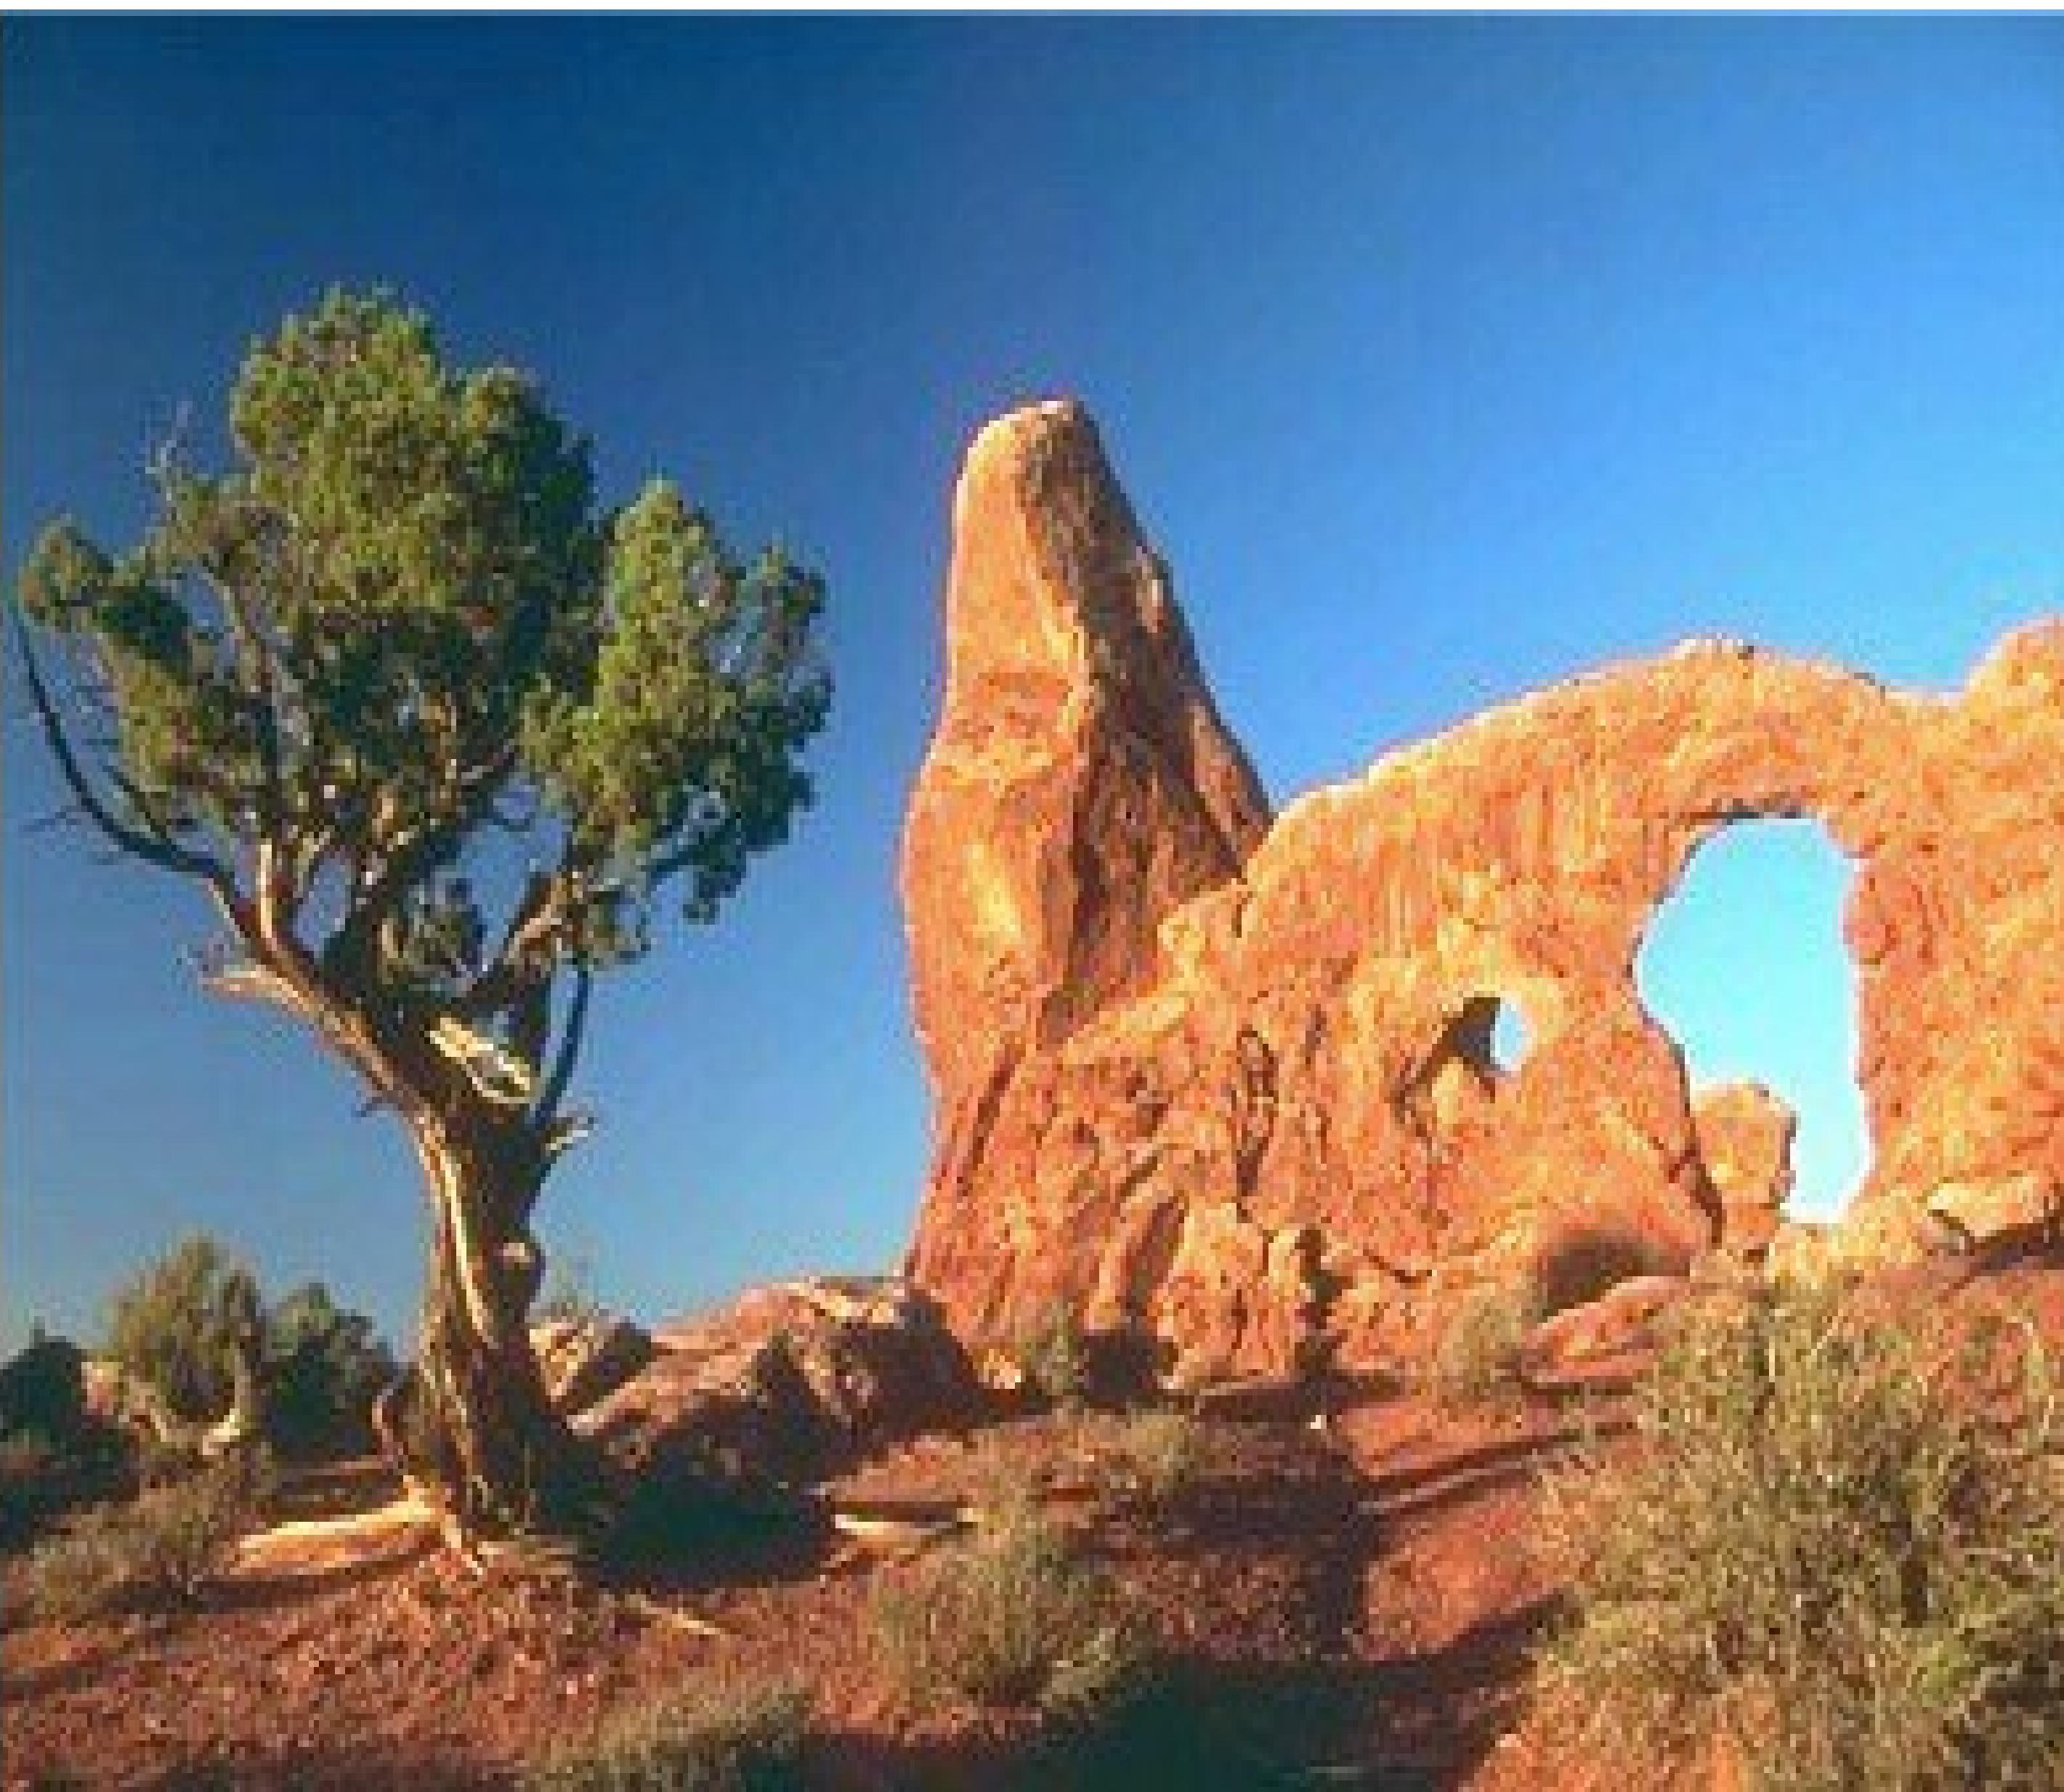


input

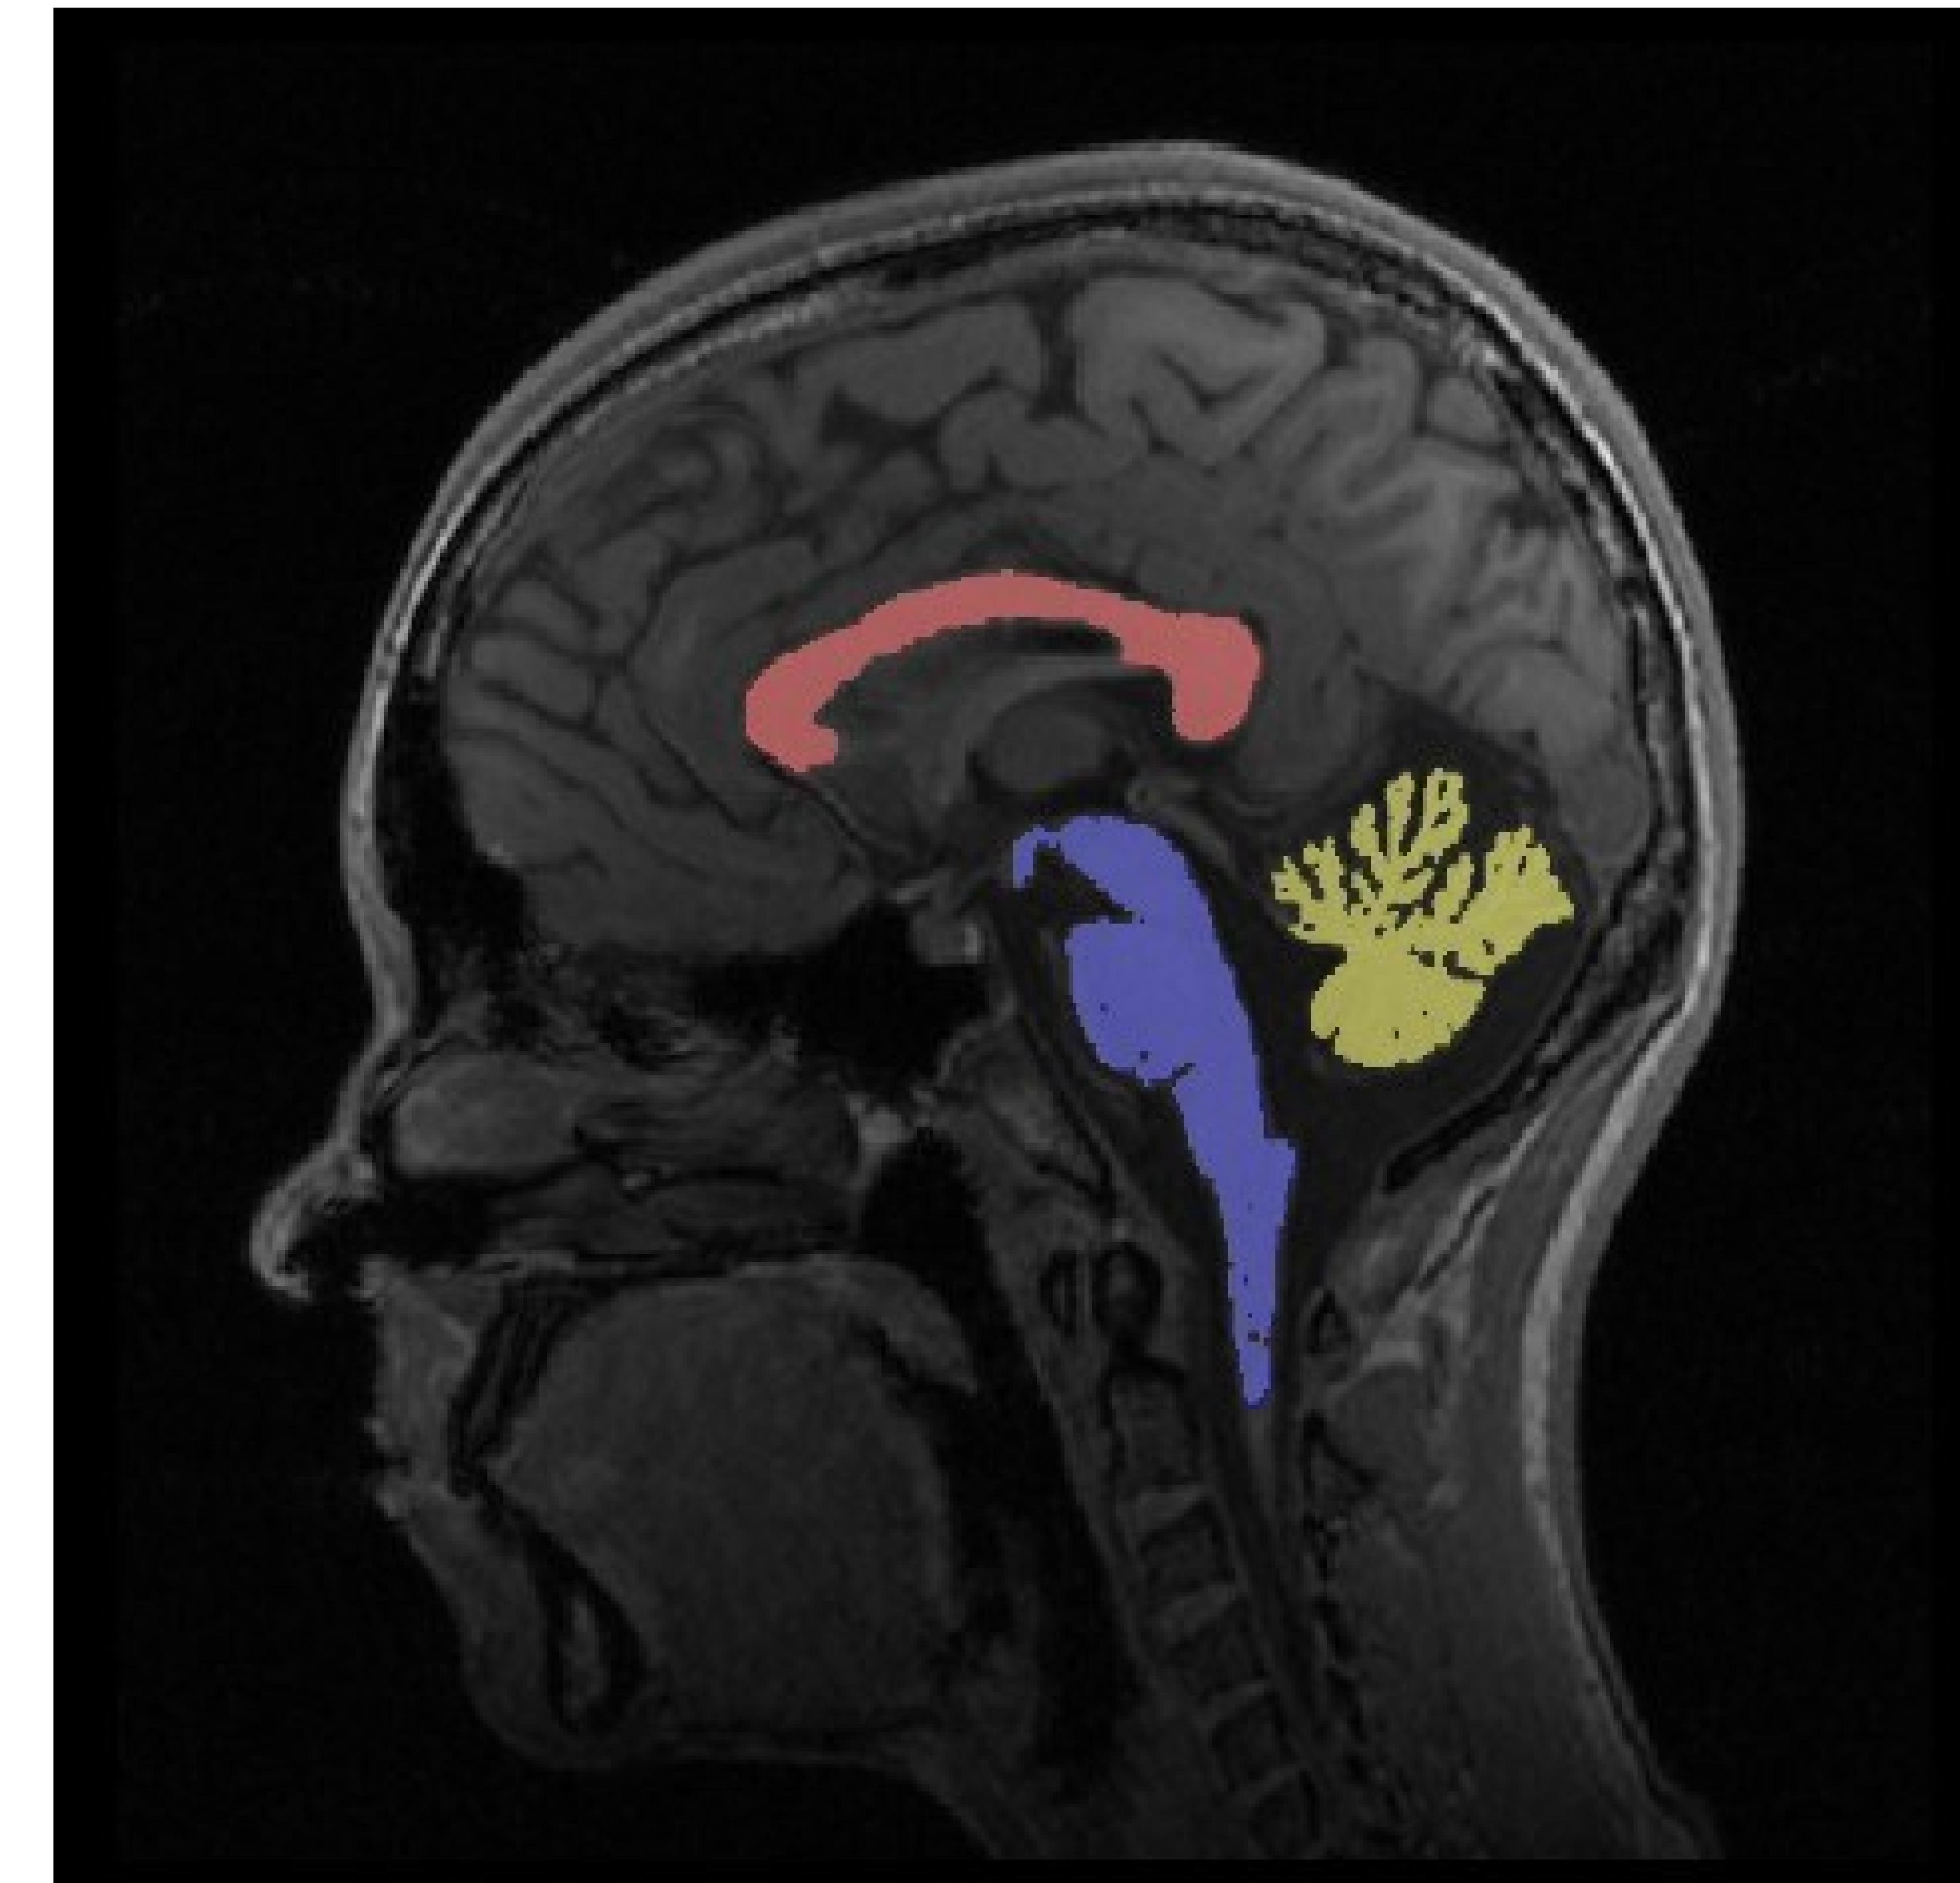
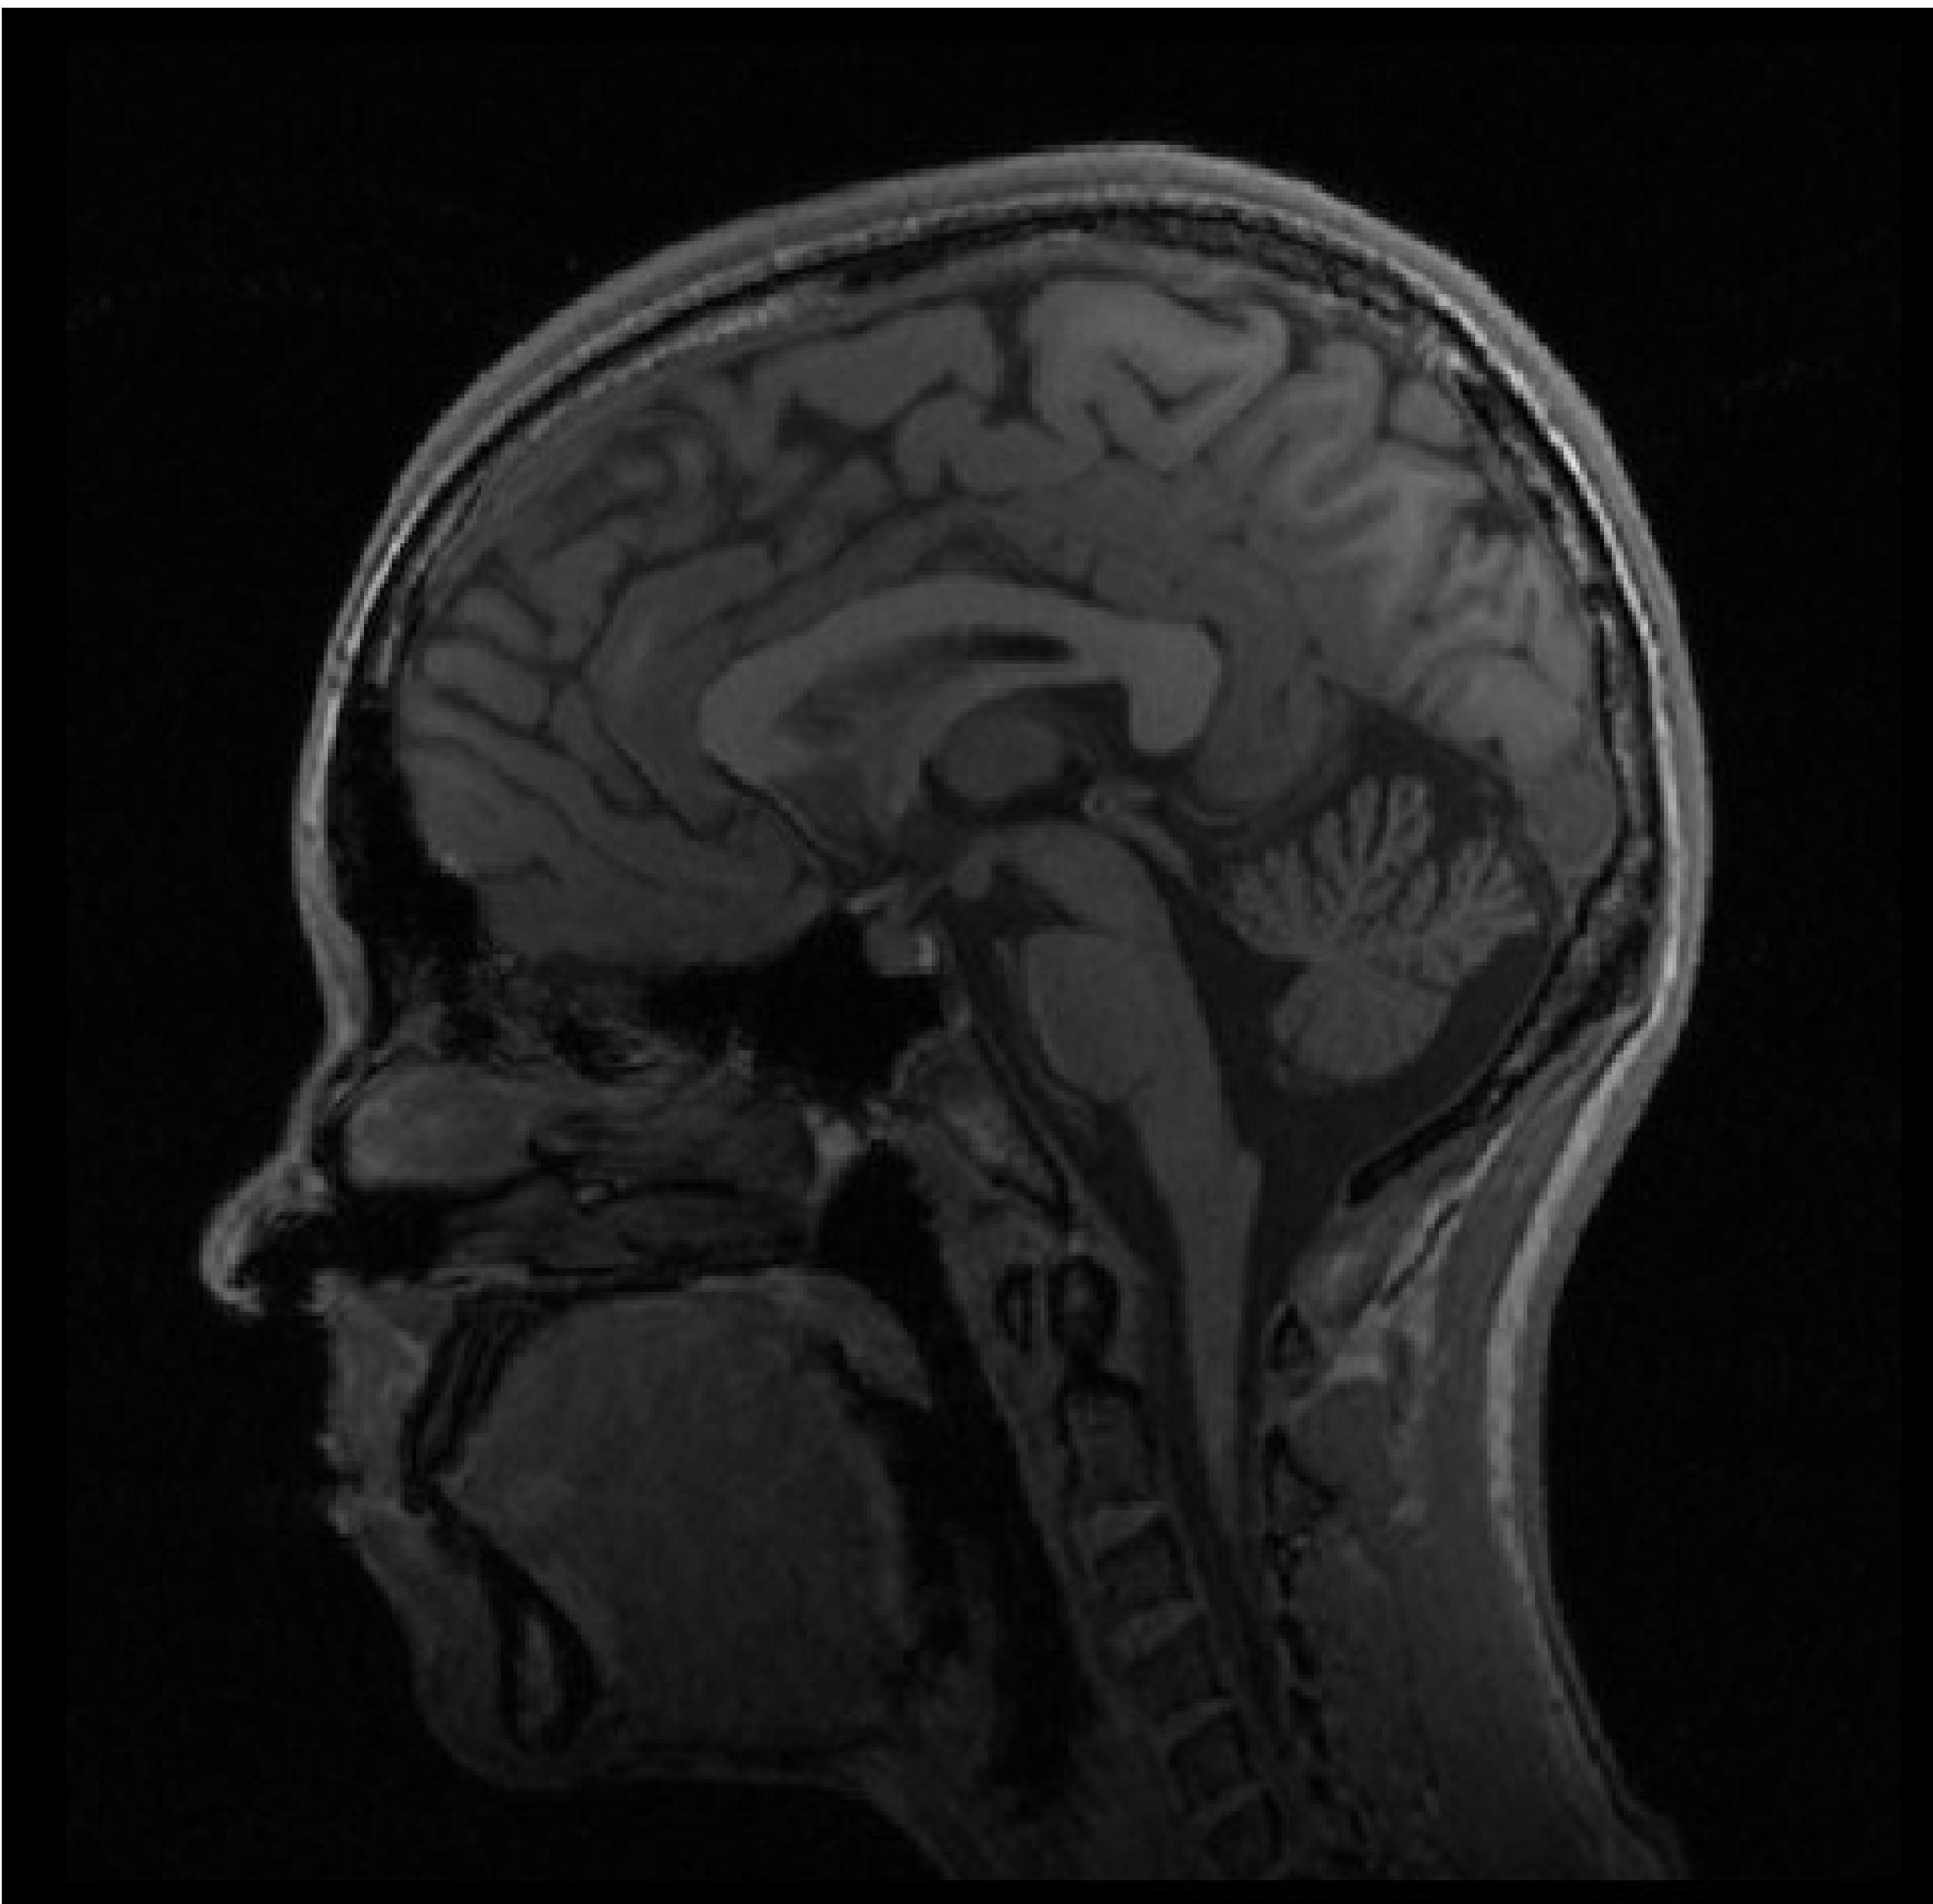


4x output

Example: edge detection



Example: region detection / segmentation



Example: image compression



Original, 2.1MB



JPEG Compression, 308KB (15%)

Example: image inpainting

Damaged Image



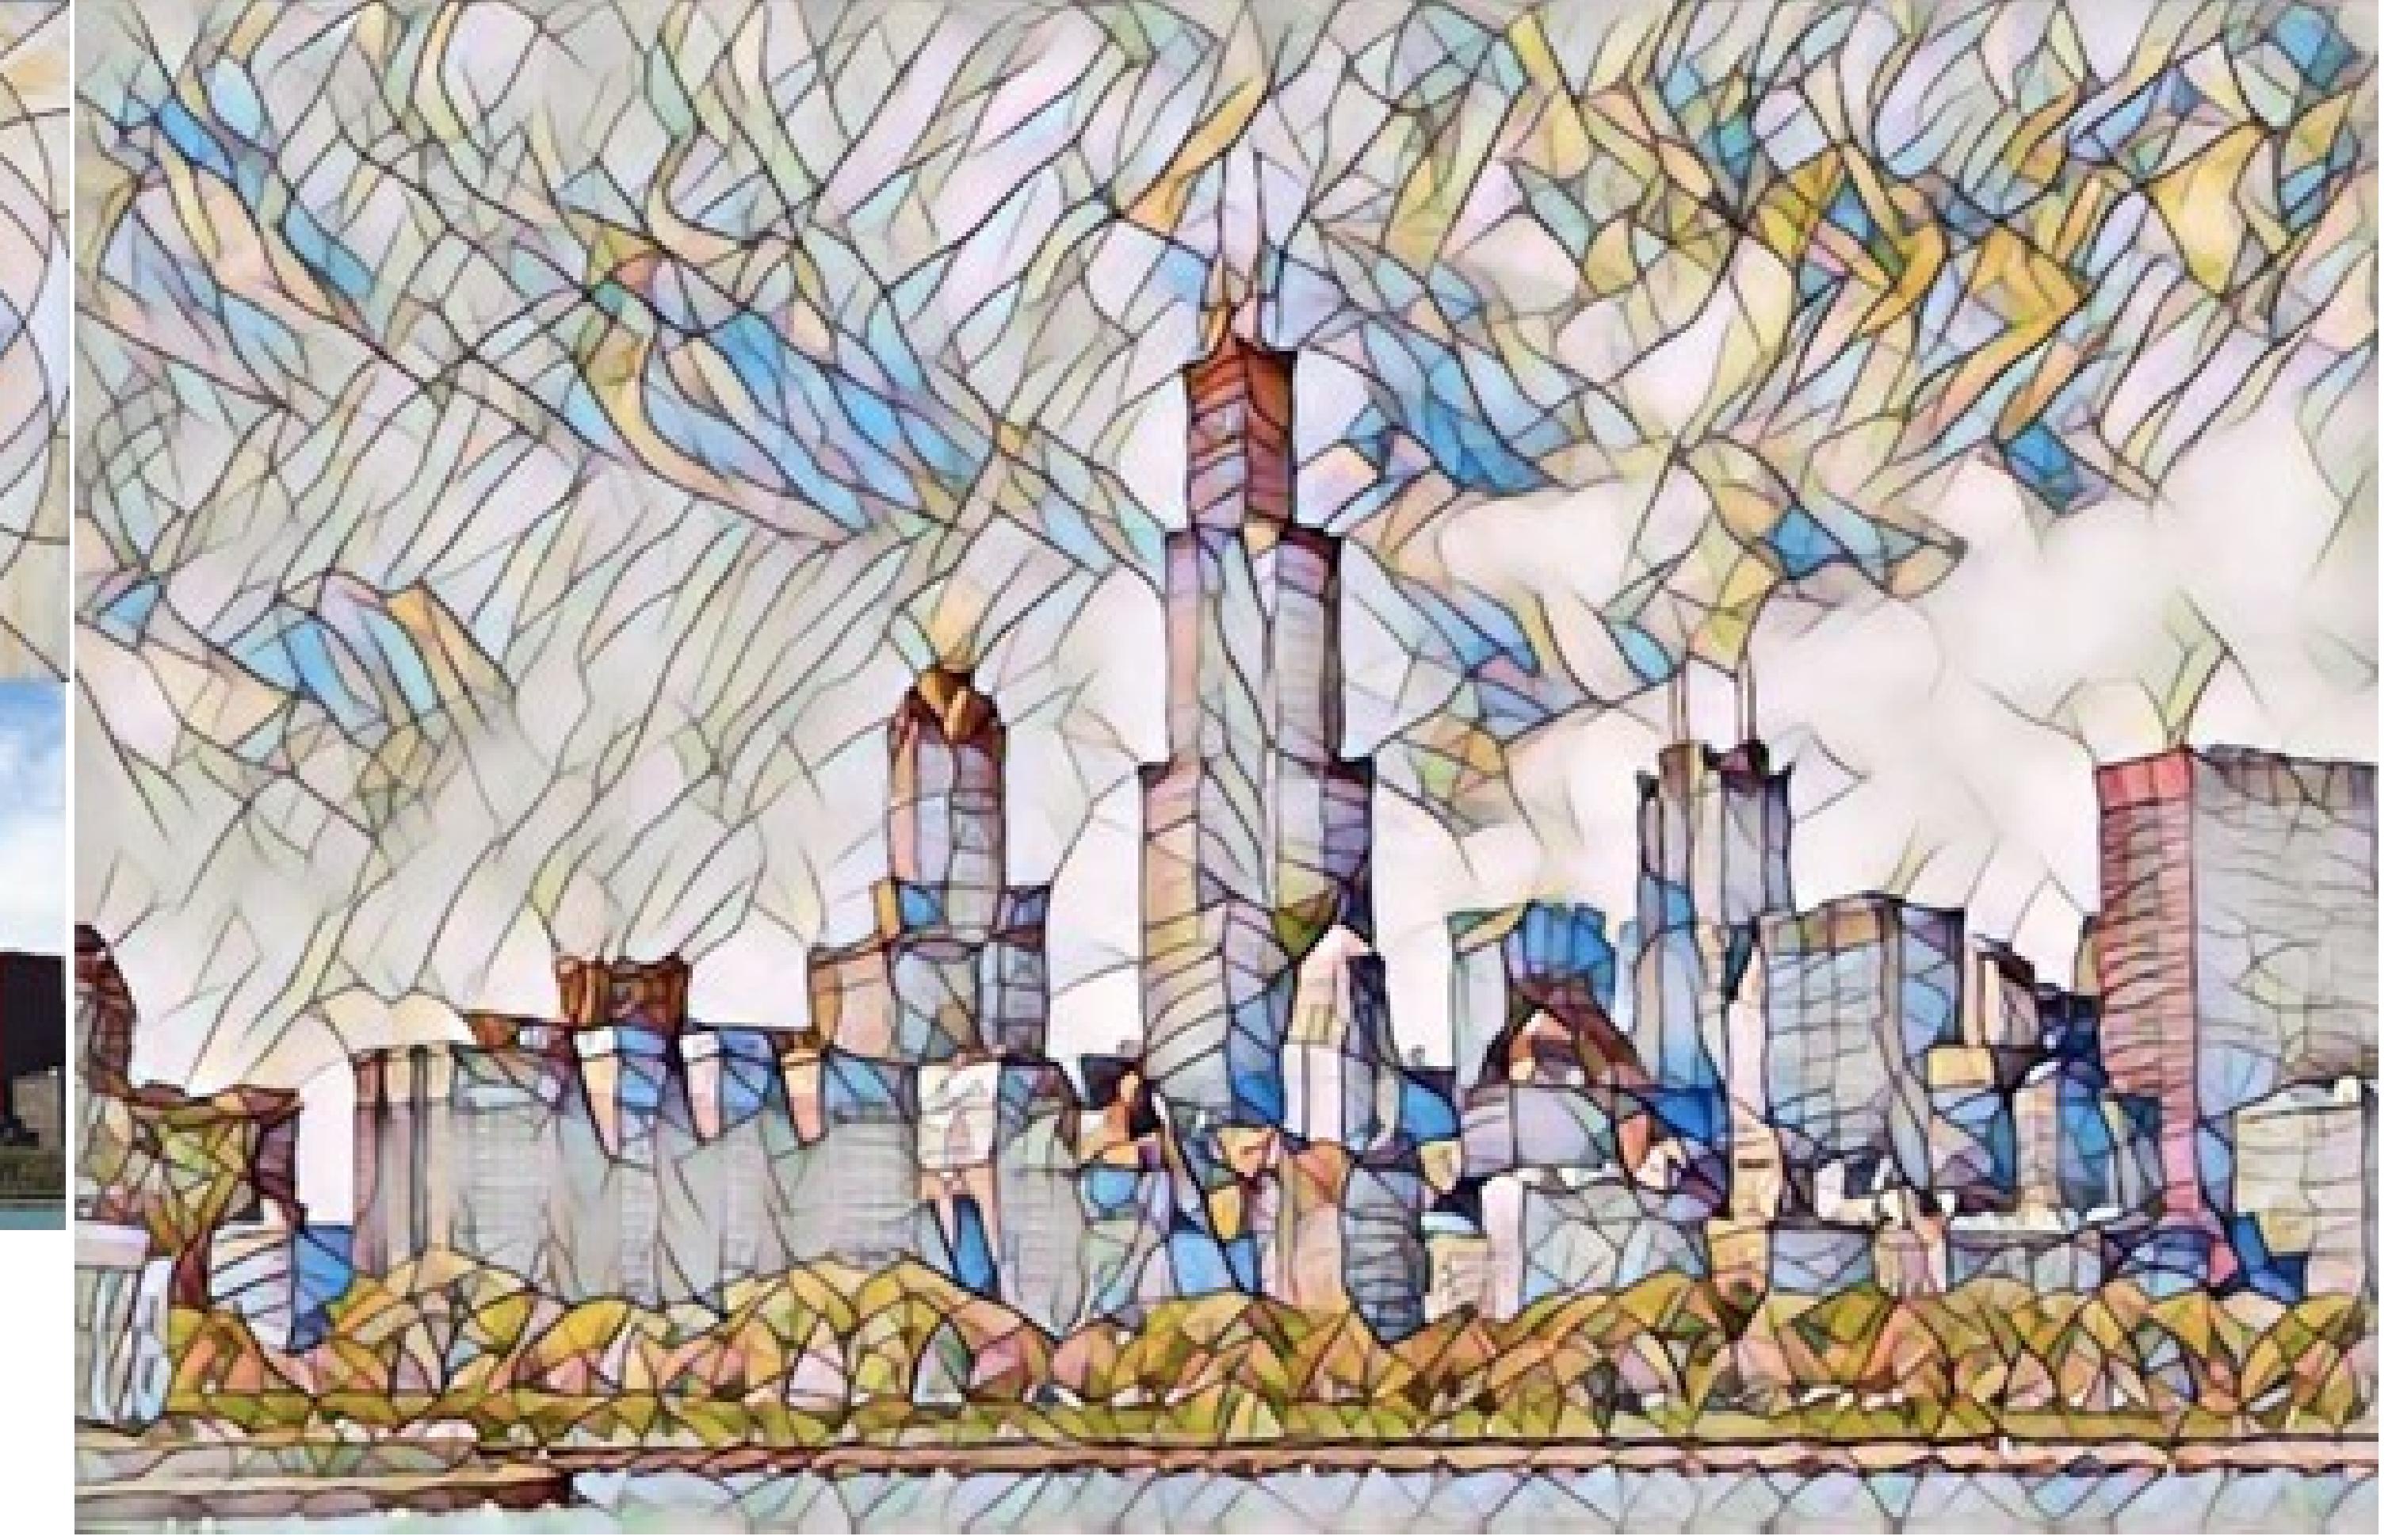
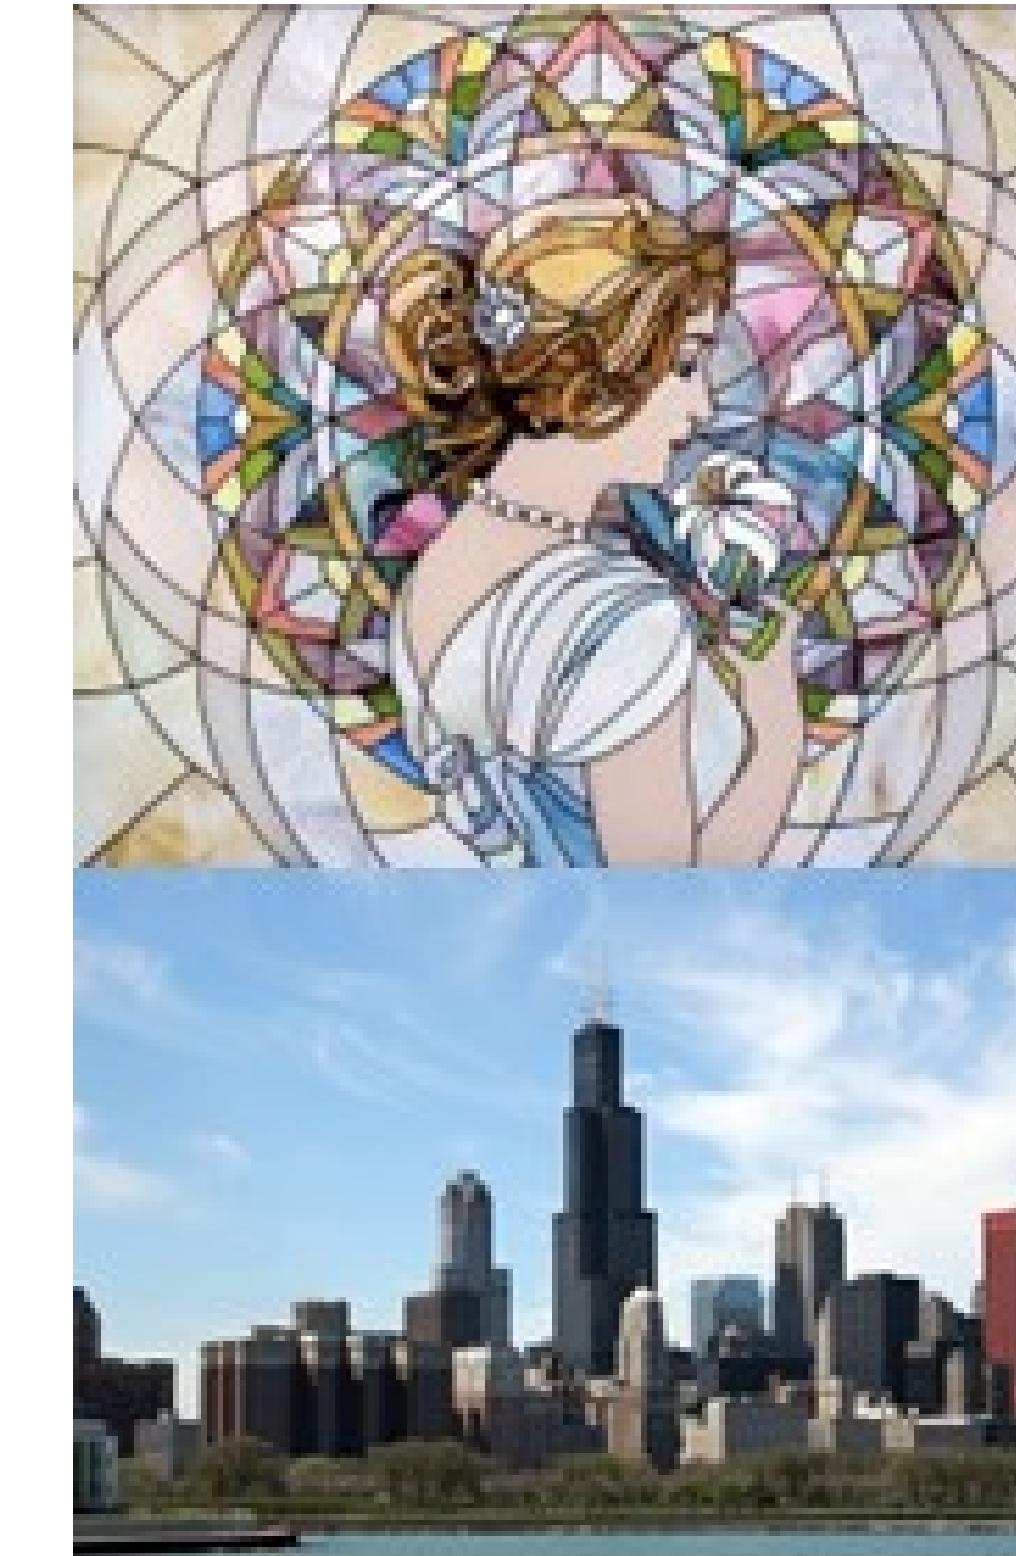
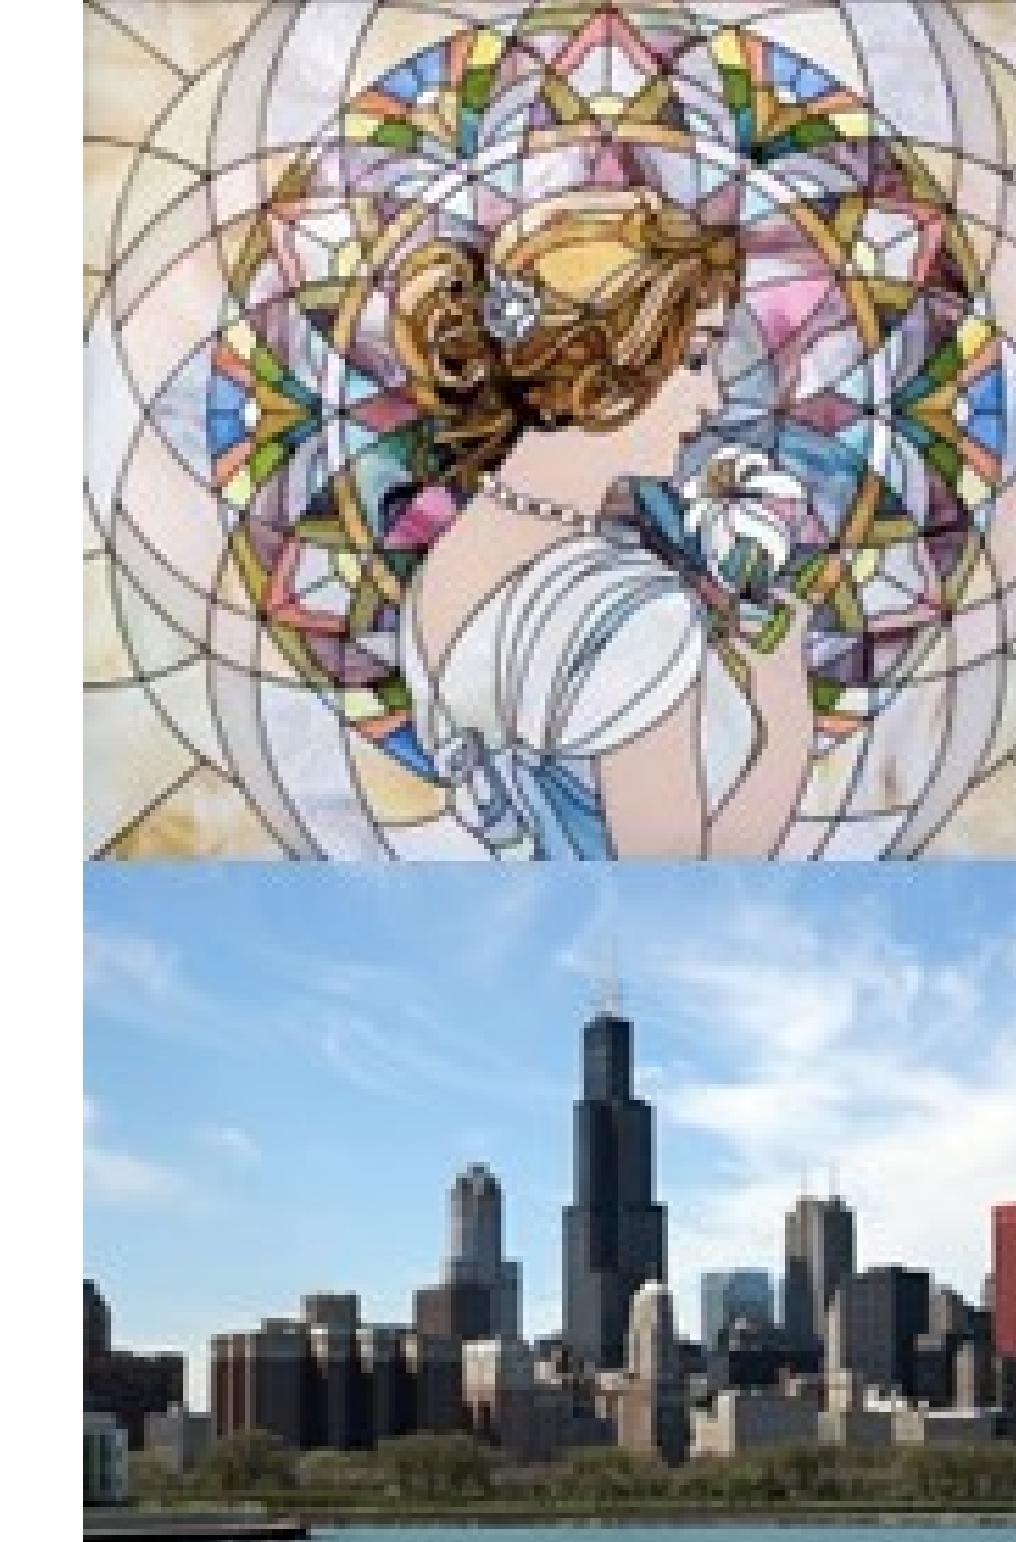
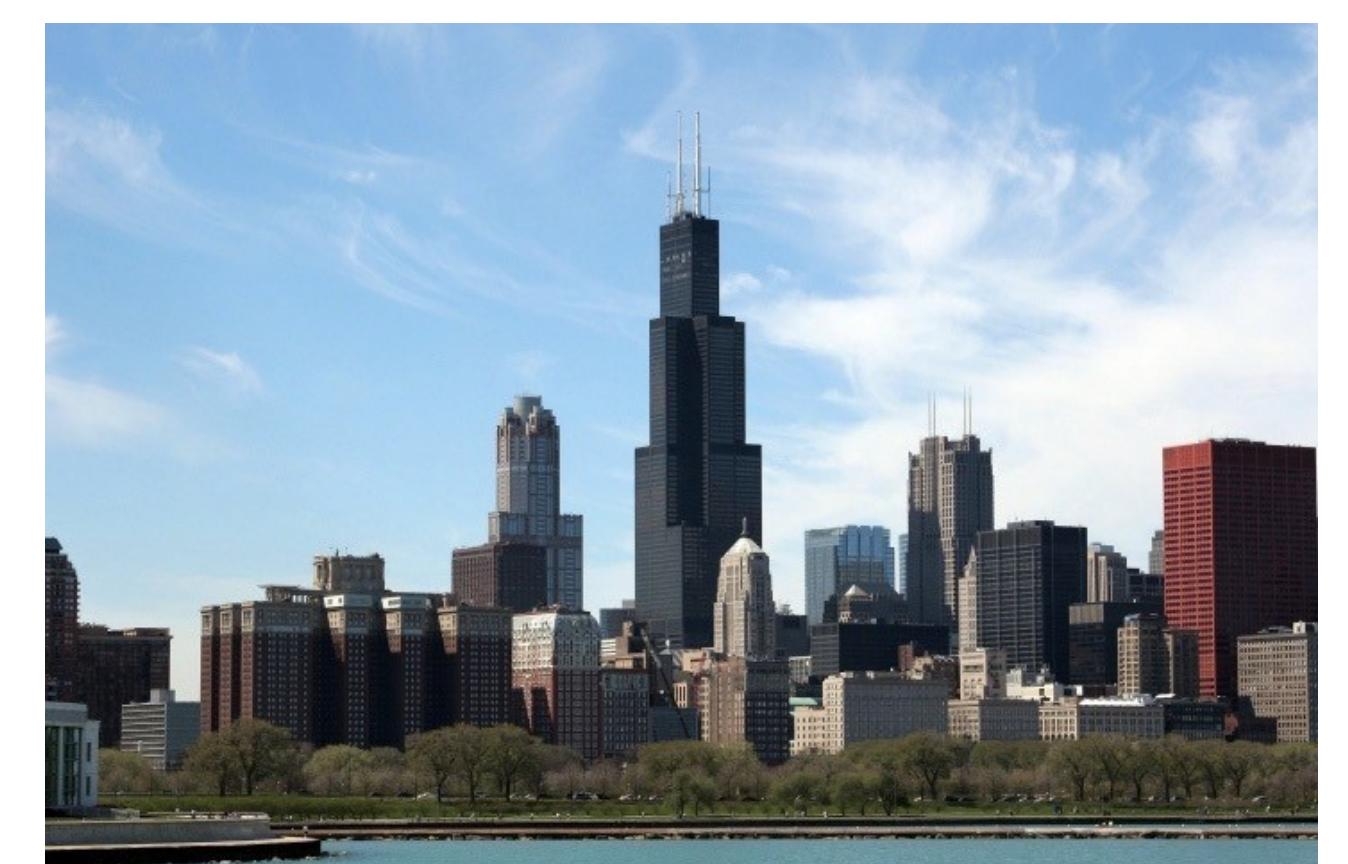
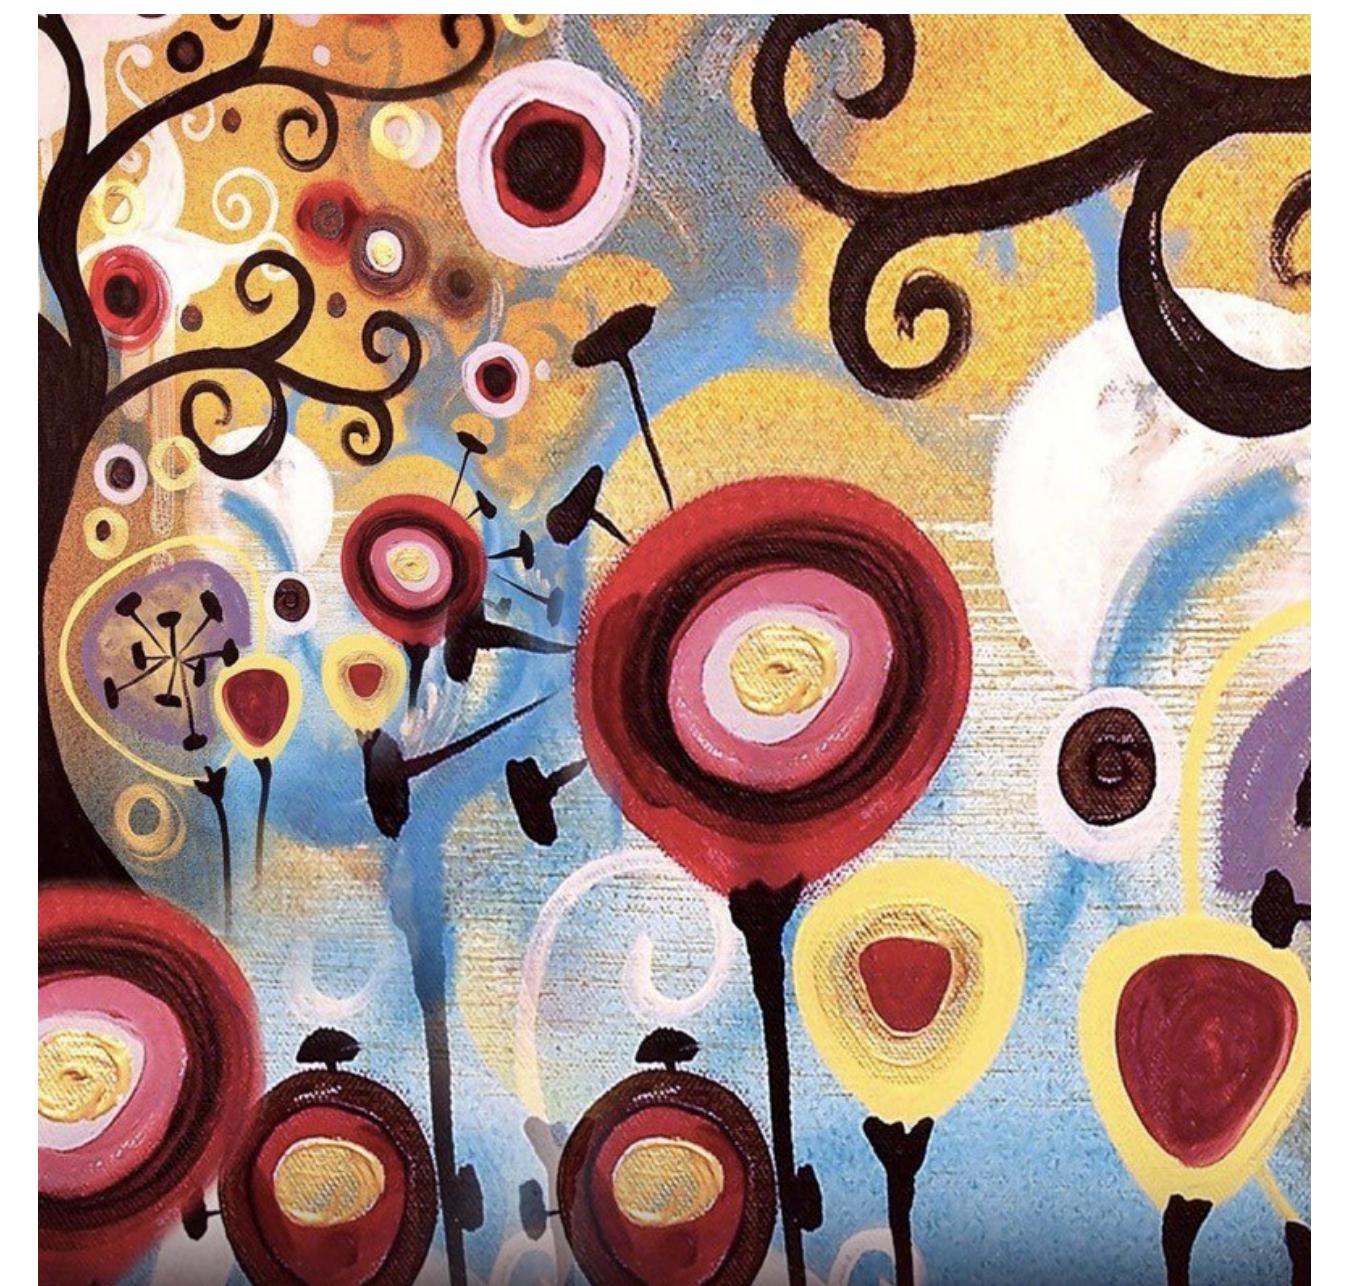
Restored Image



Credit: M. Bertalmio, G. Sapiro, V. Caselles, C. Ballester: *Image Inpainting*, SIGGRAPH 2000

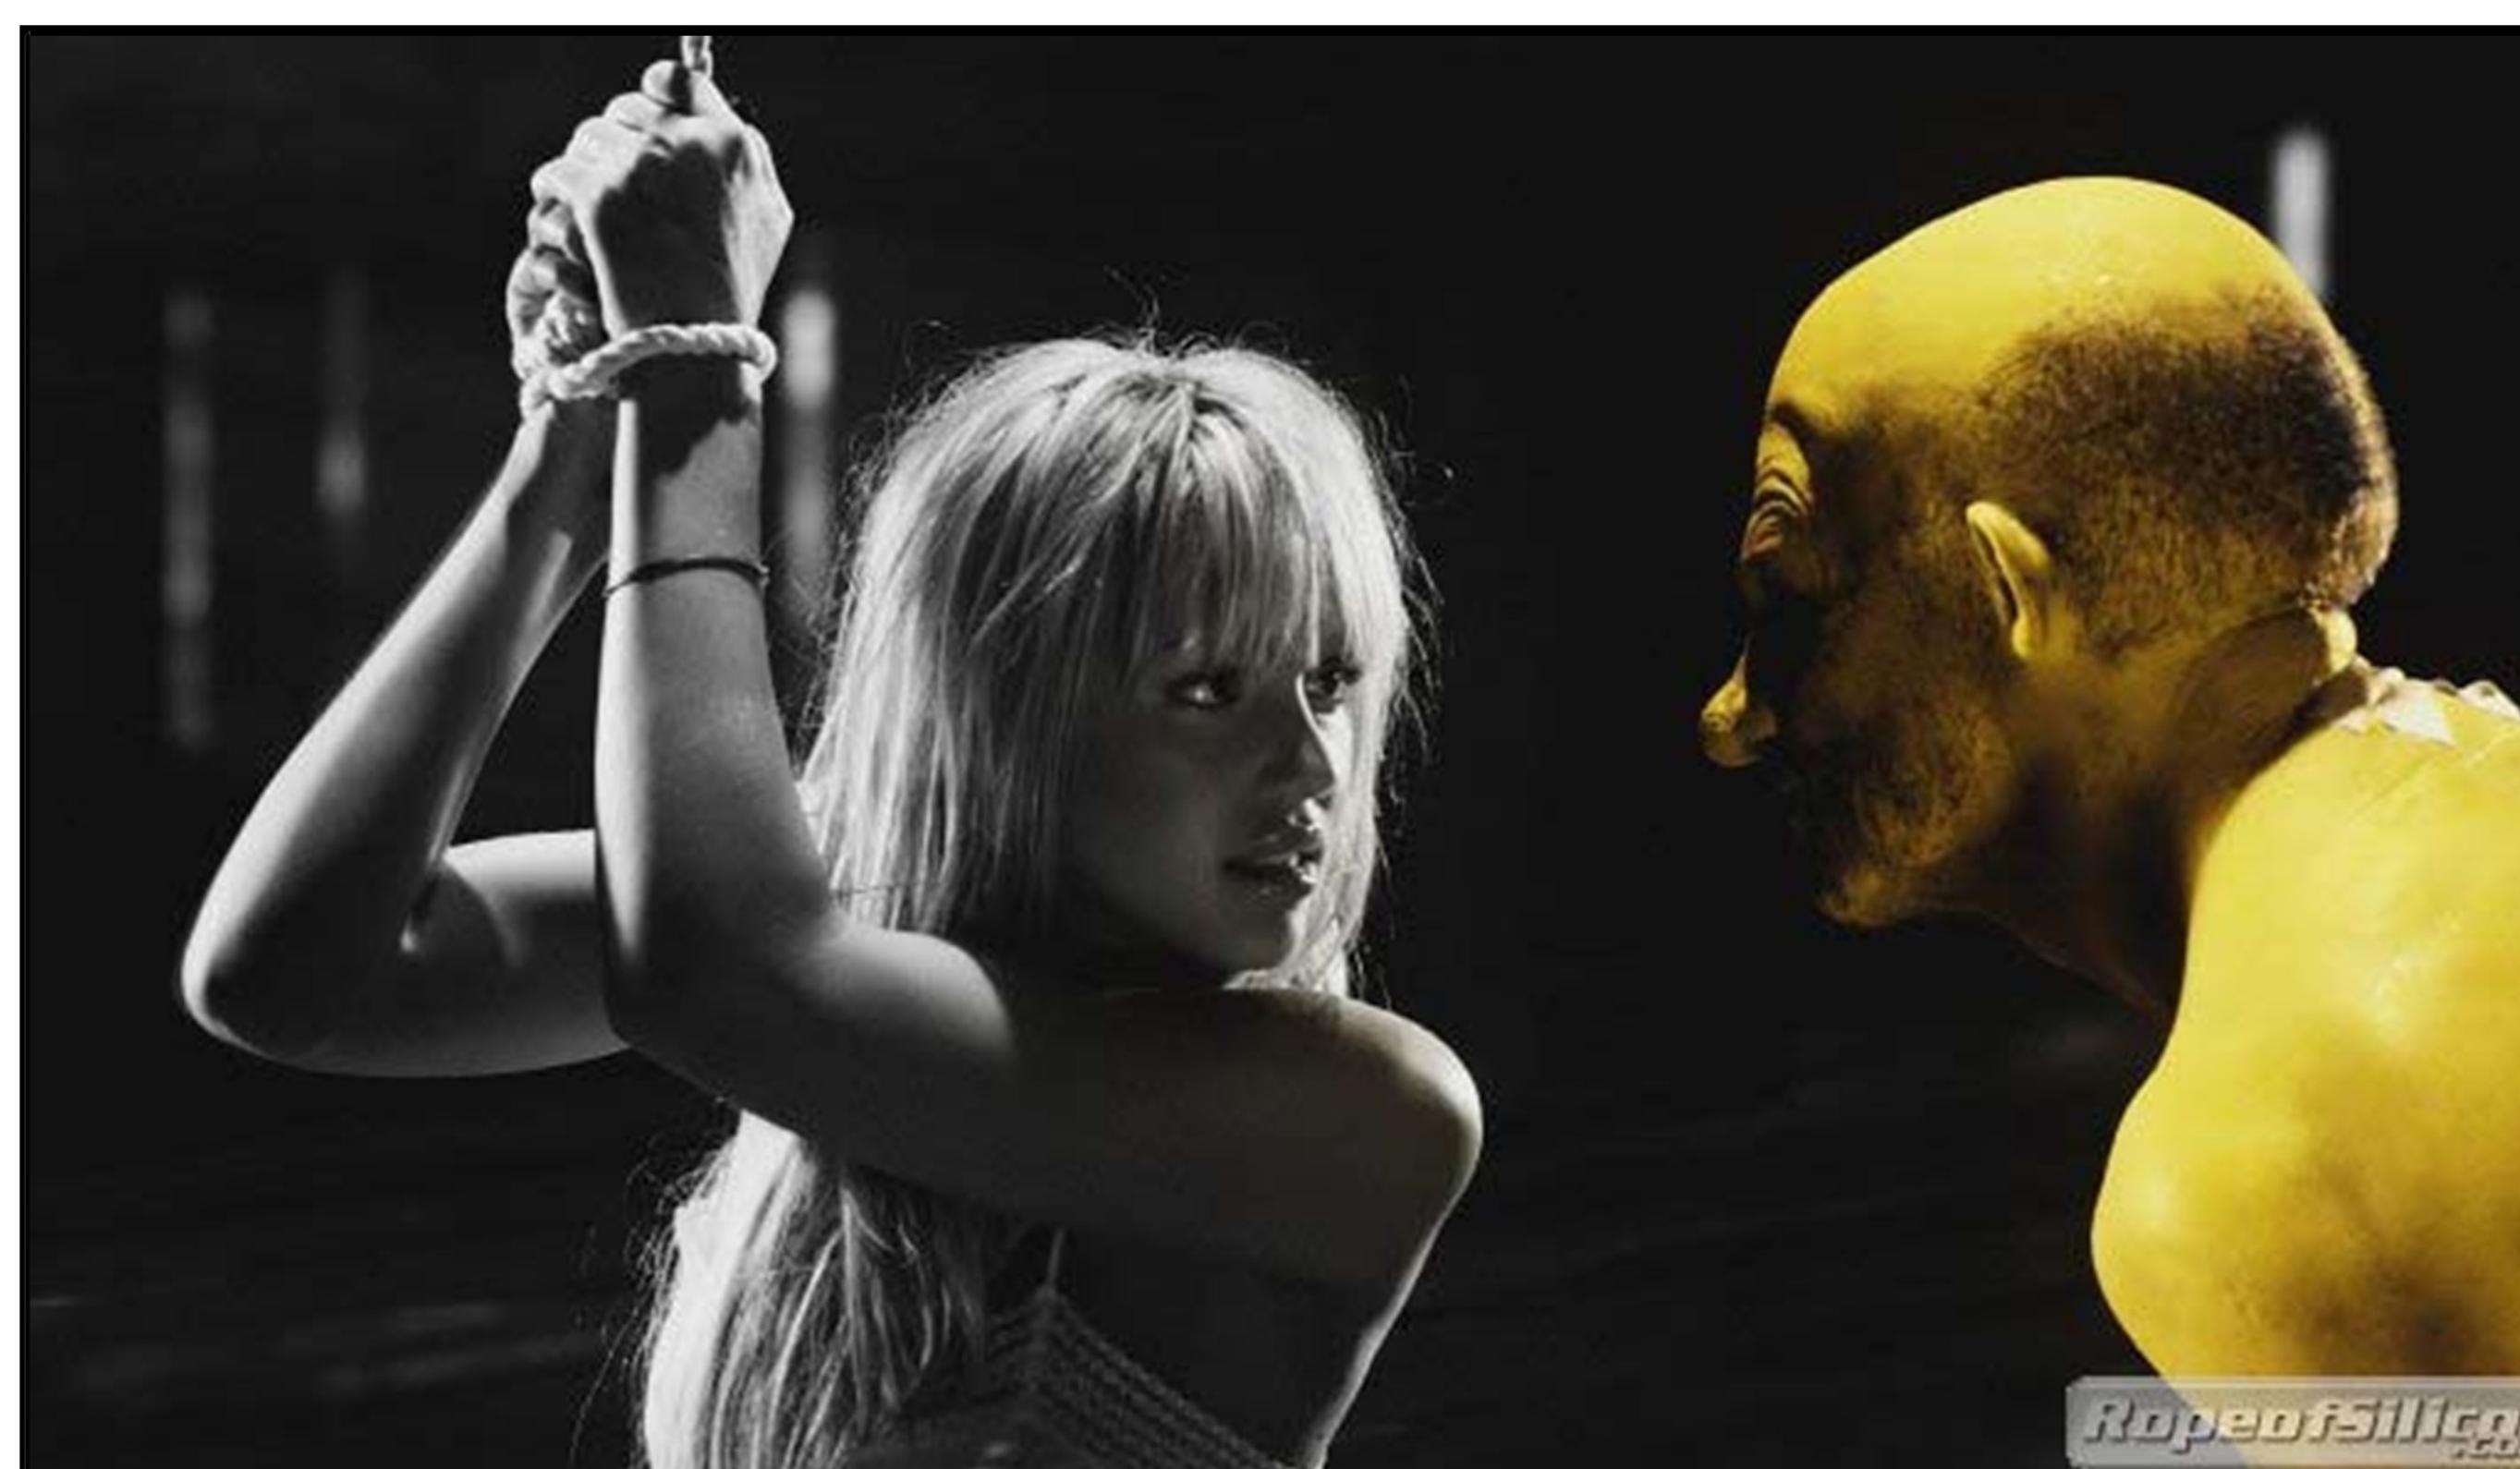
Inpainting? Reconstruct corrupted/destroyed parts of an image

Example: image style transfer



Application of image processing

Examples: Artistic (Movie Special)Effects



Application of image processing

Law Enforcement

- Number plate recognition for speed cameras or automated toll systems
- Fingerprint recognition

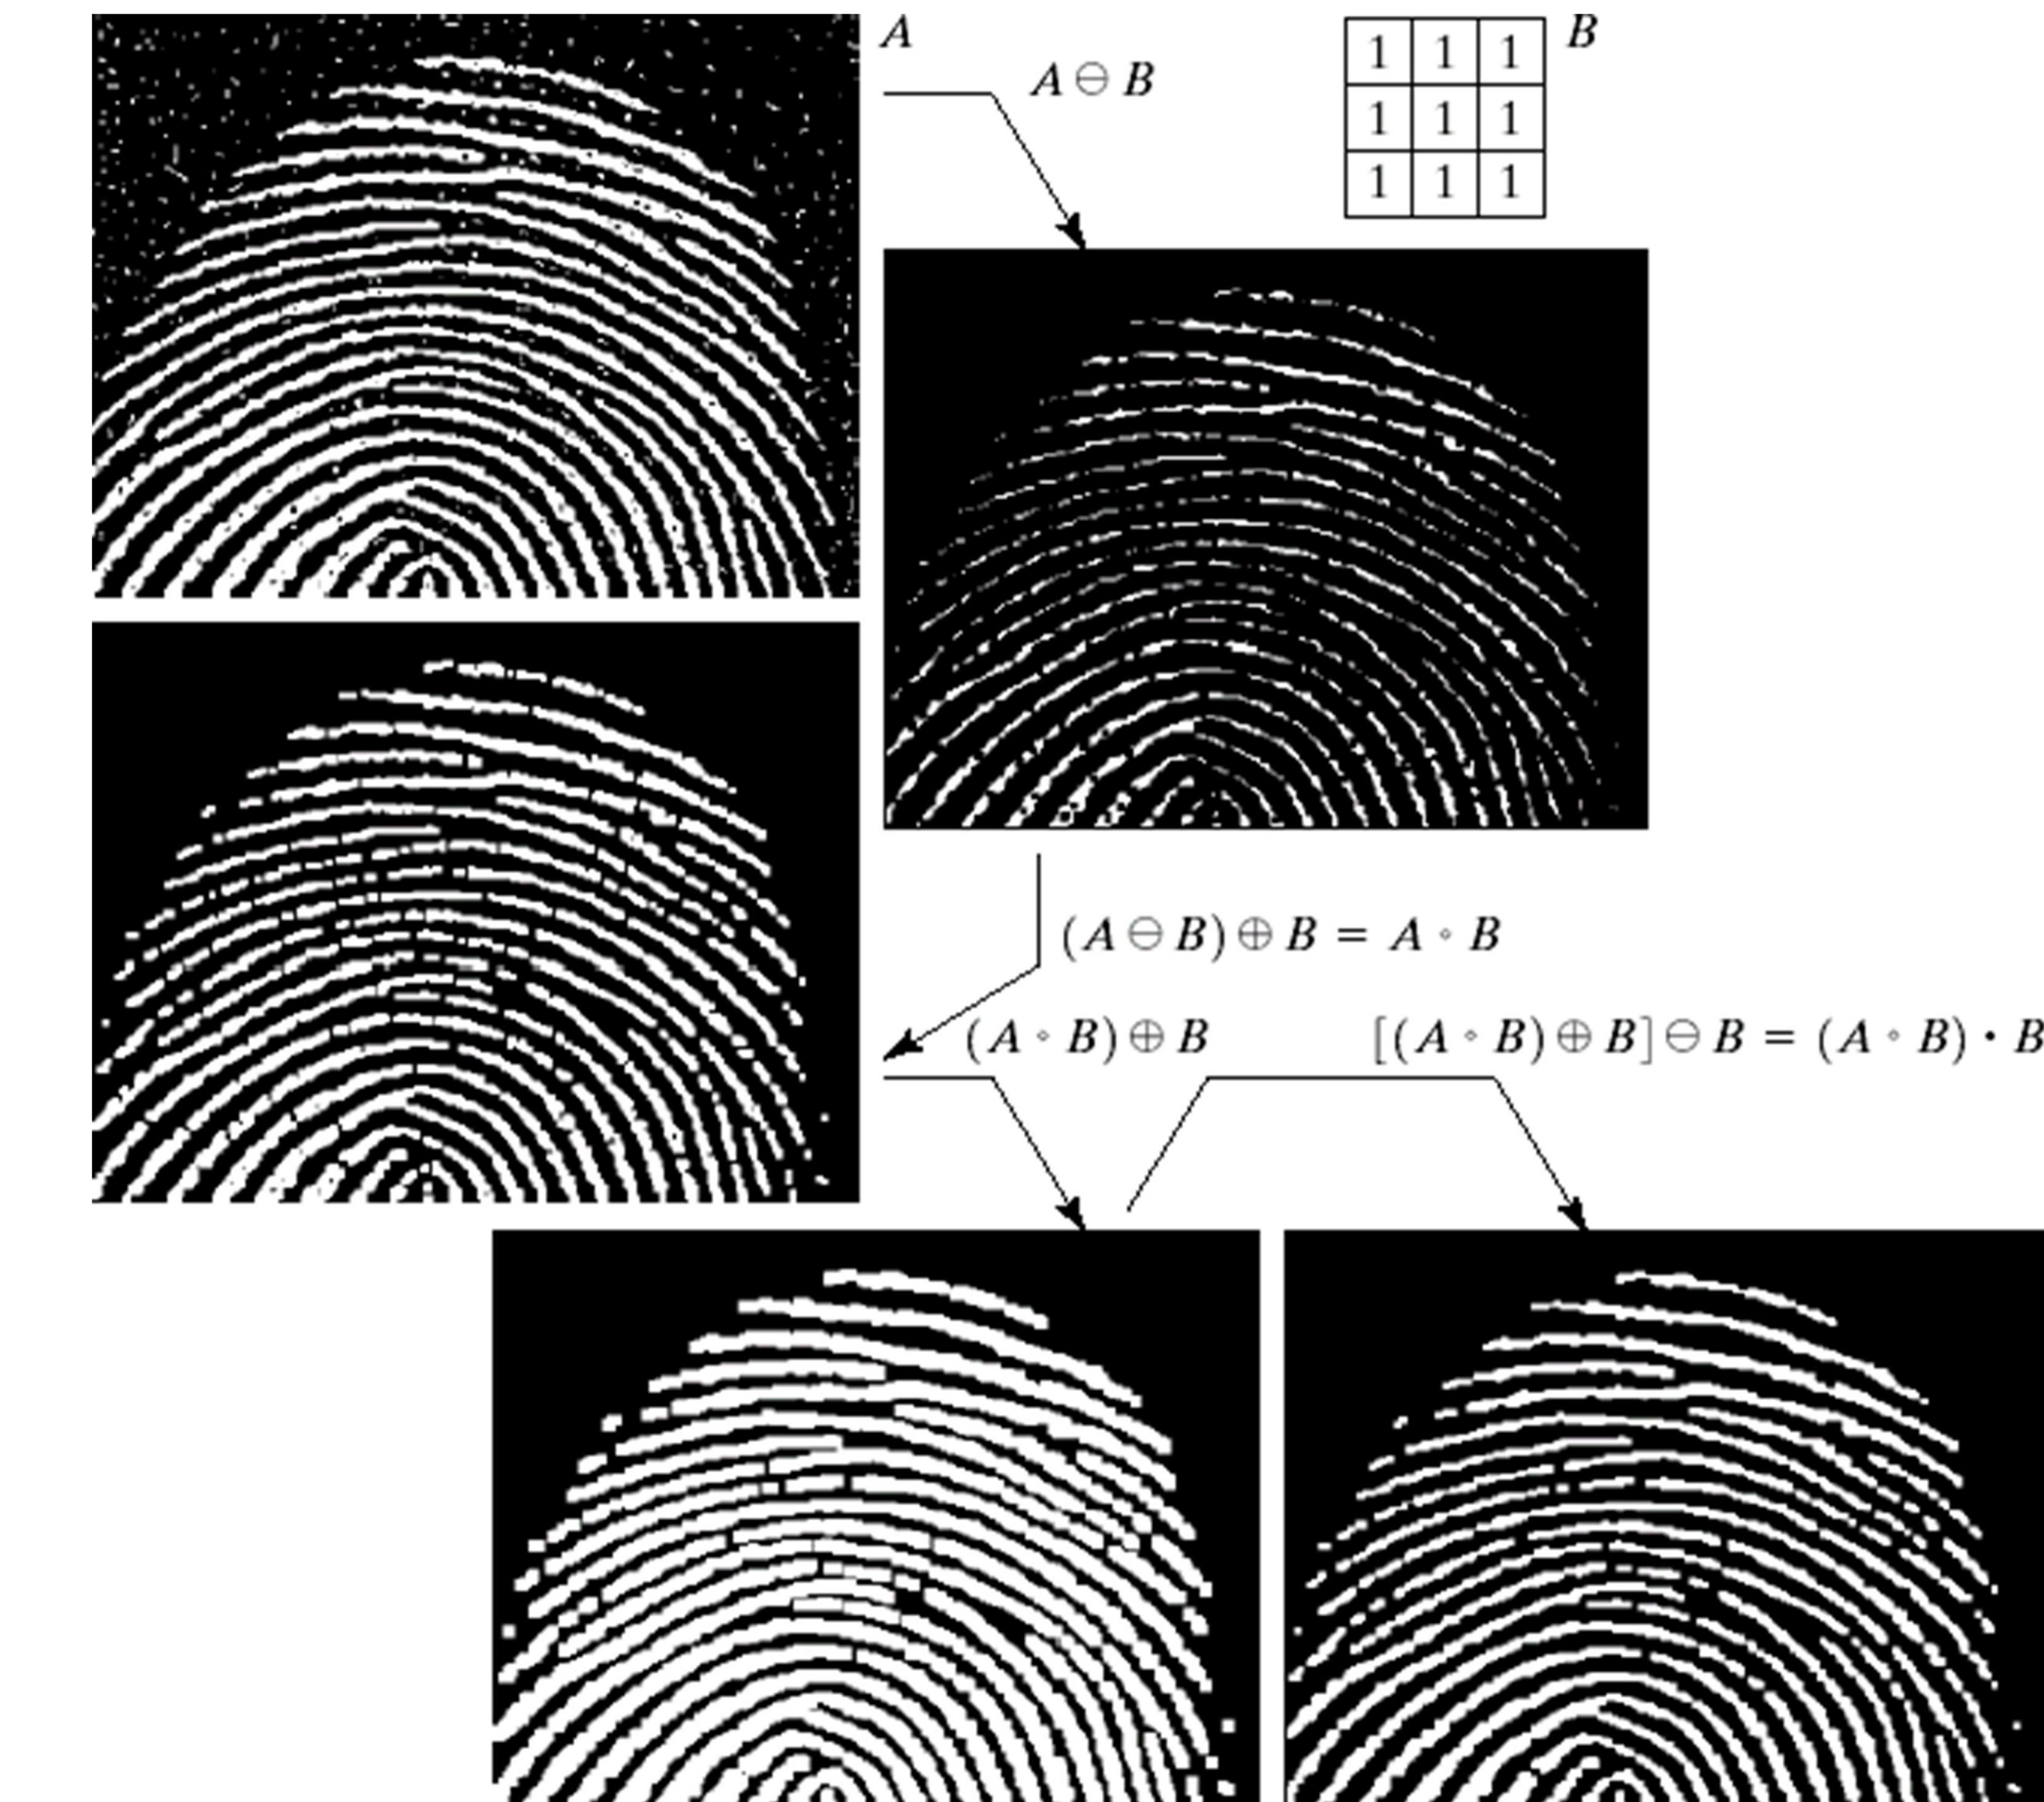
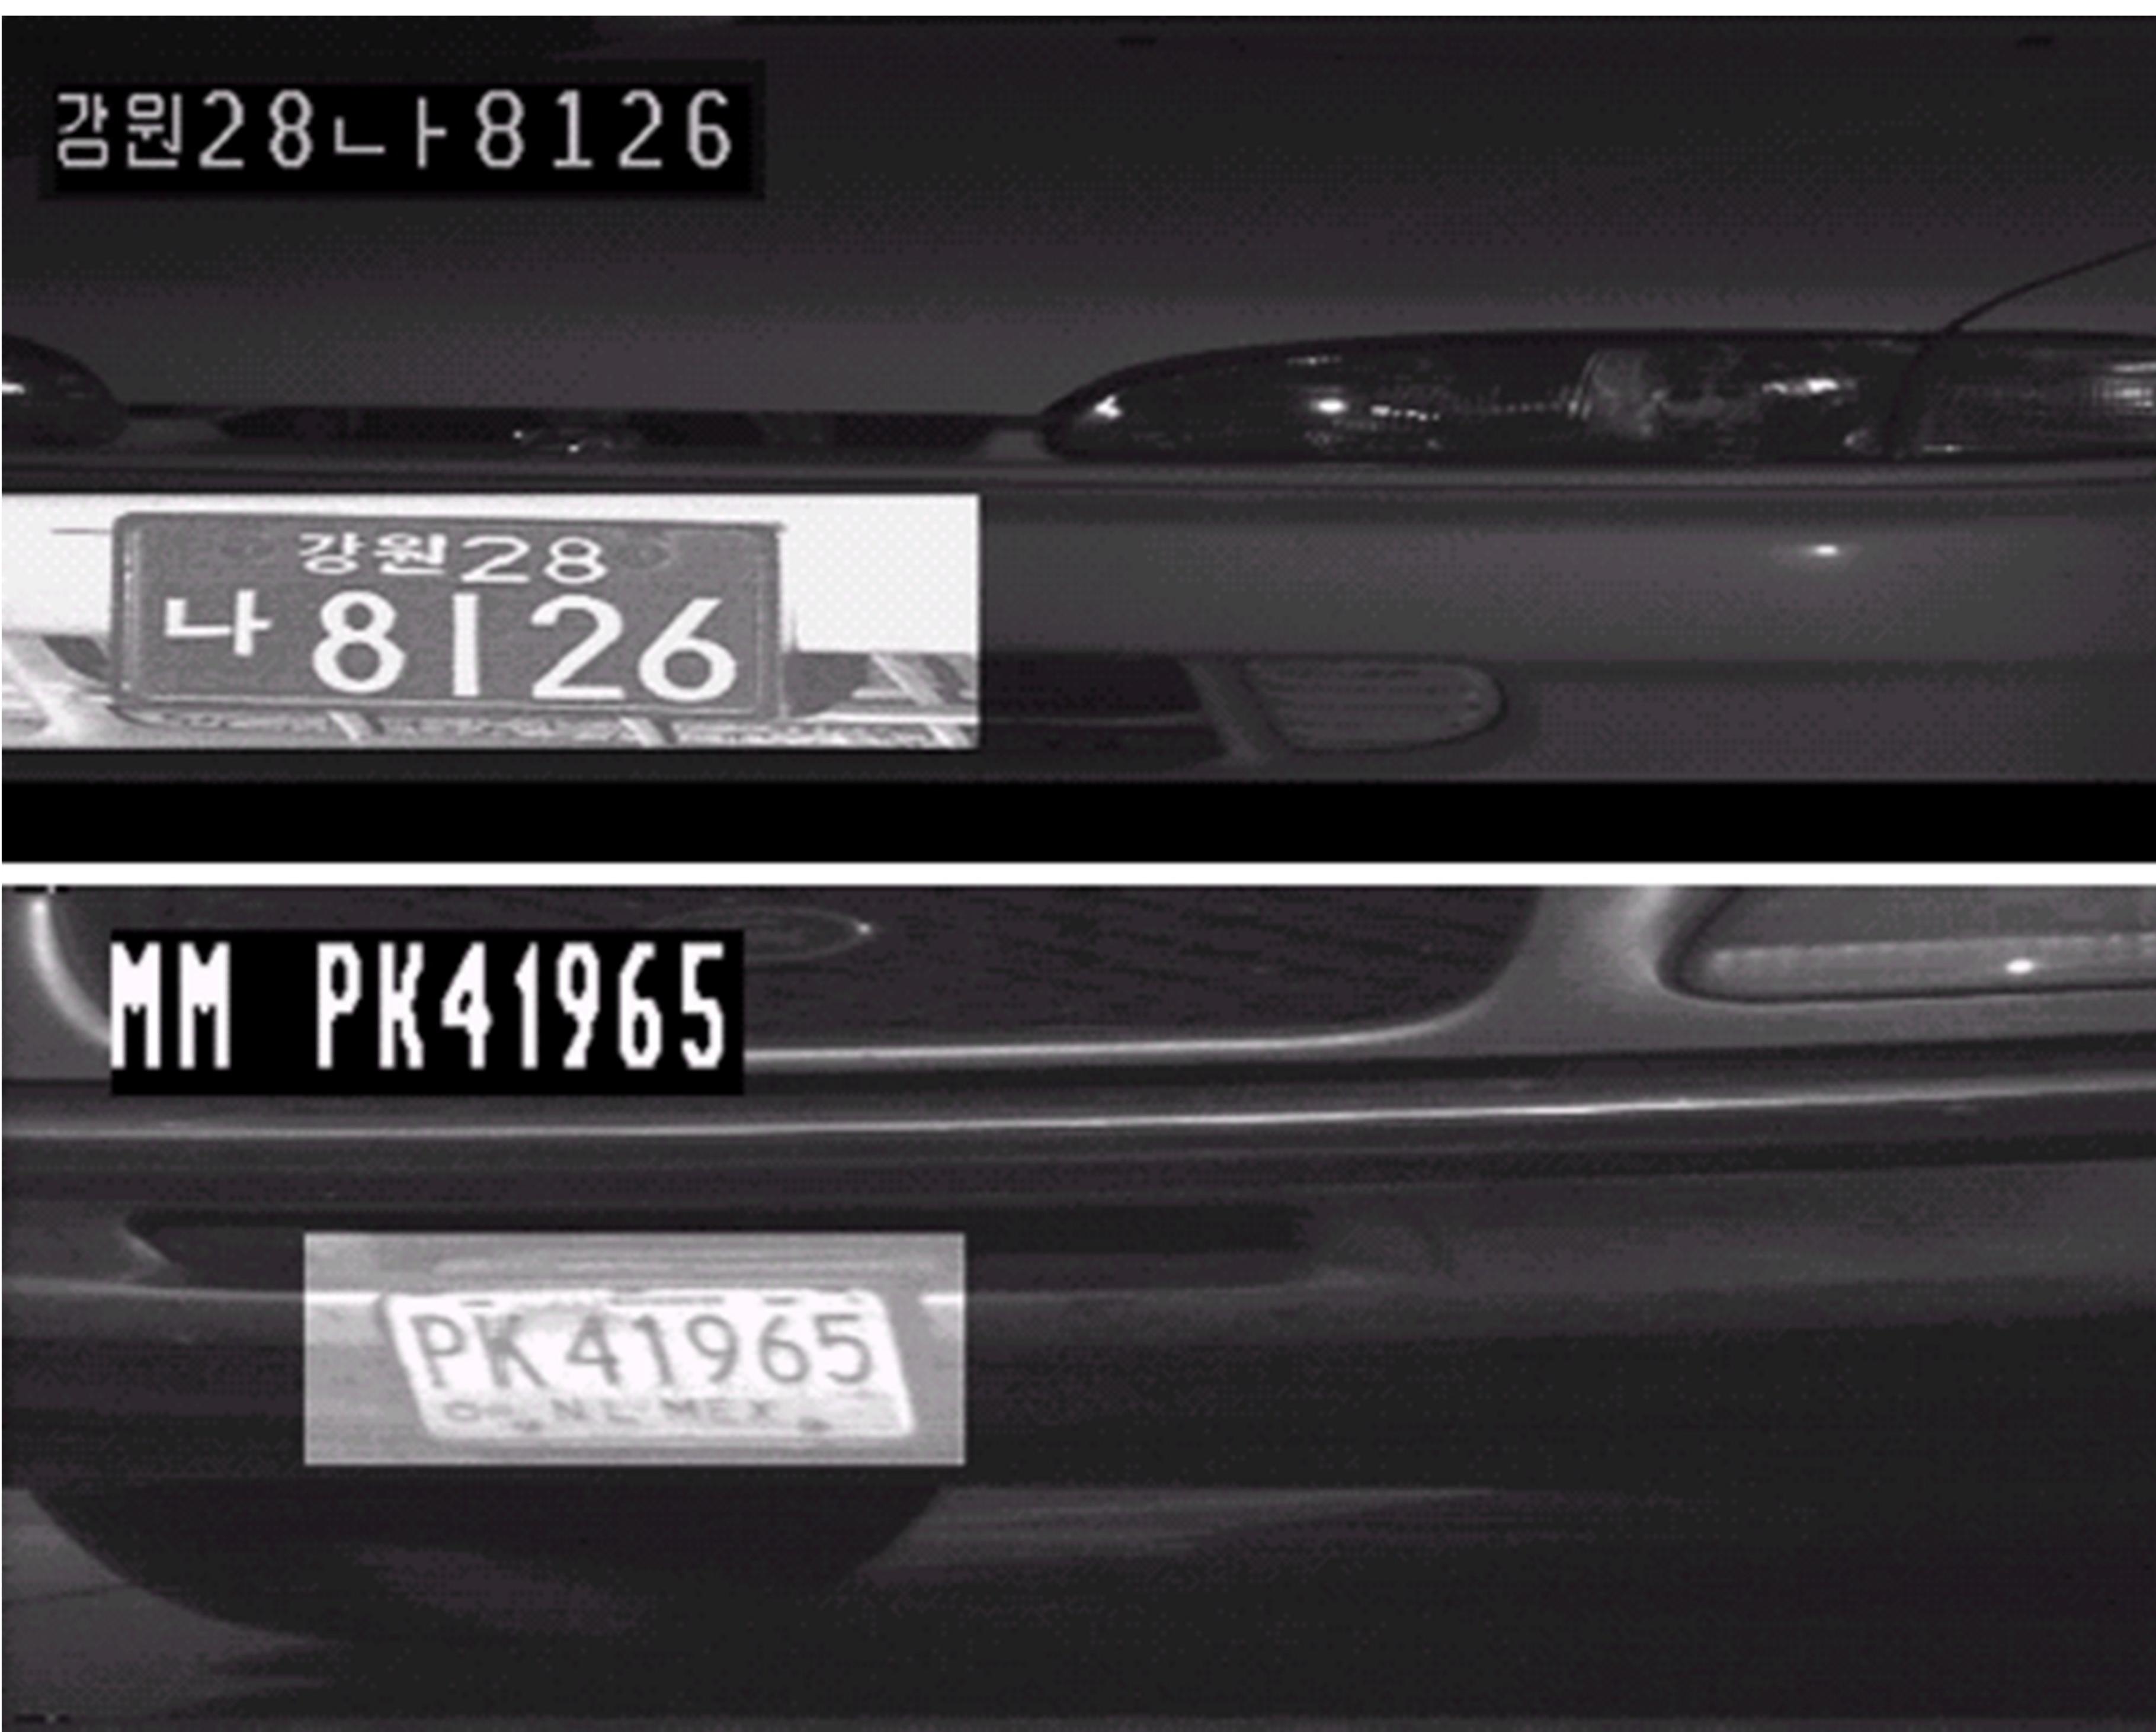
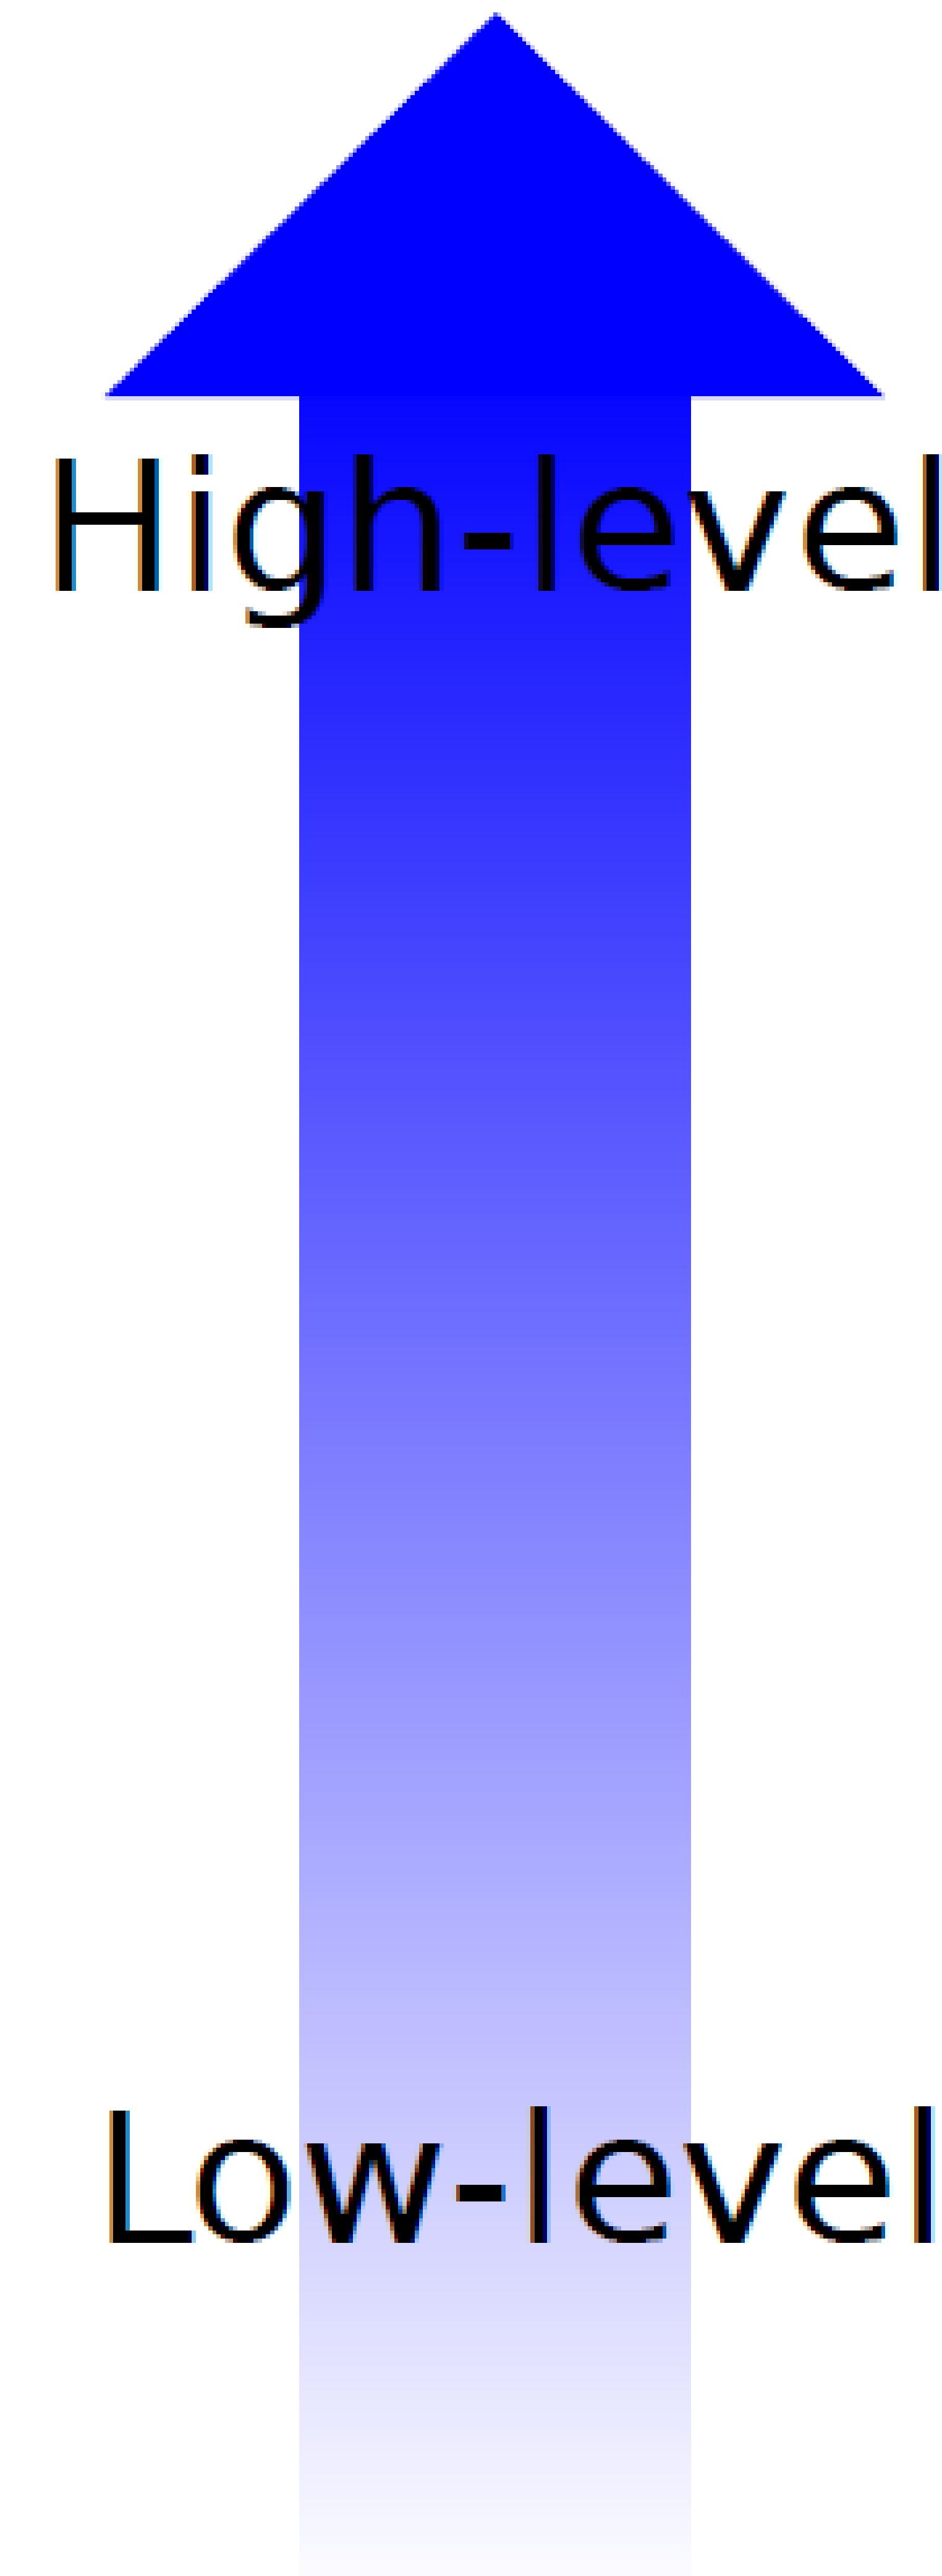


Image processing, image analysis, computer vision



Computer Vision

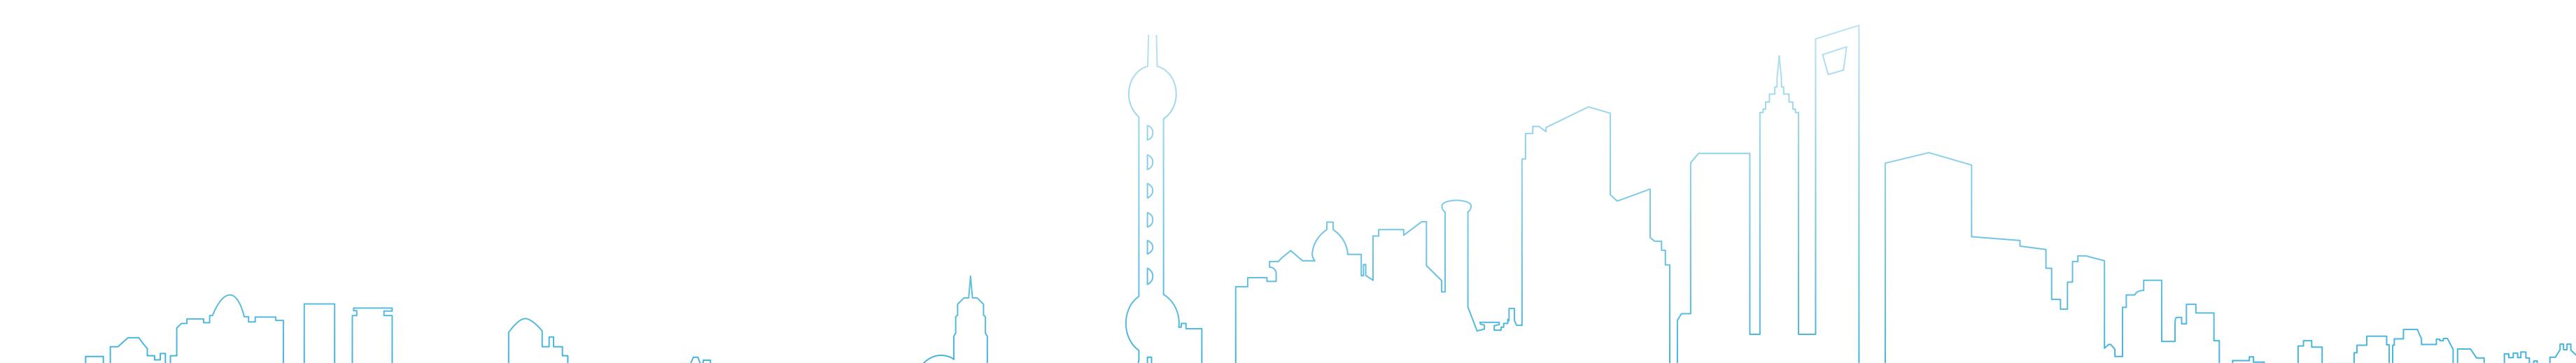
Object detection, recognition, shape analysis, tracking
Use of Artificial Intelligence and Machine Learning

Image Analysis

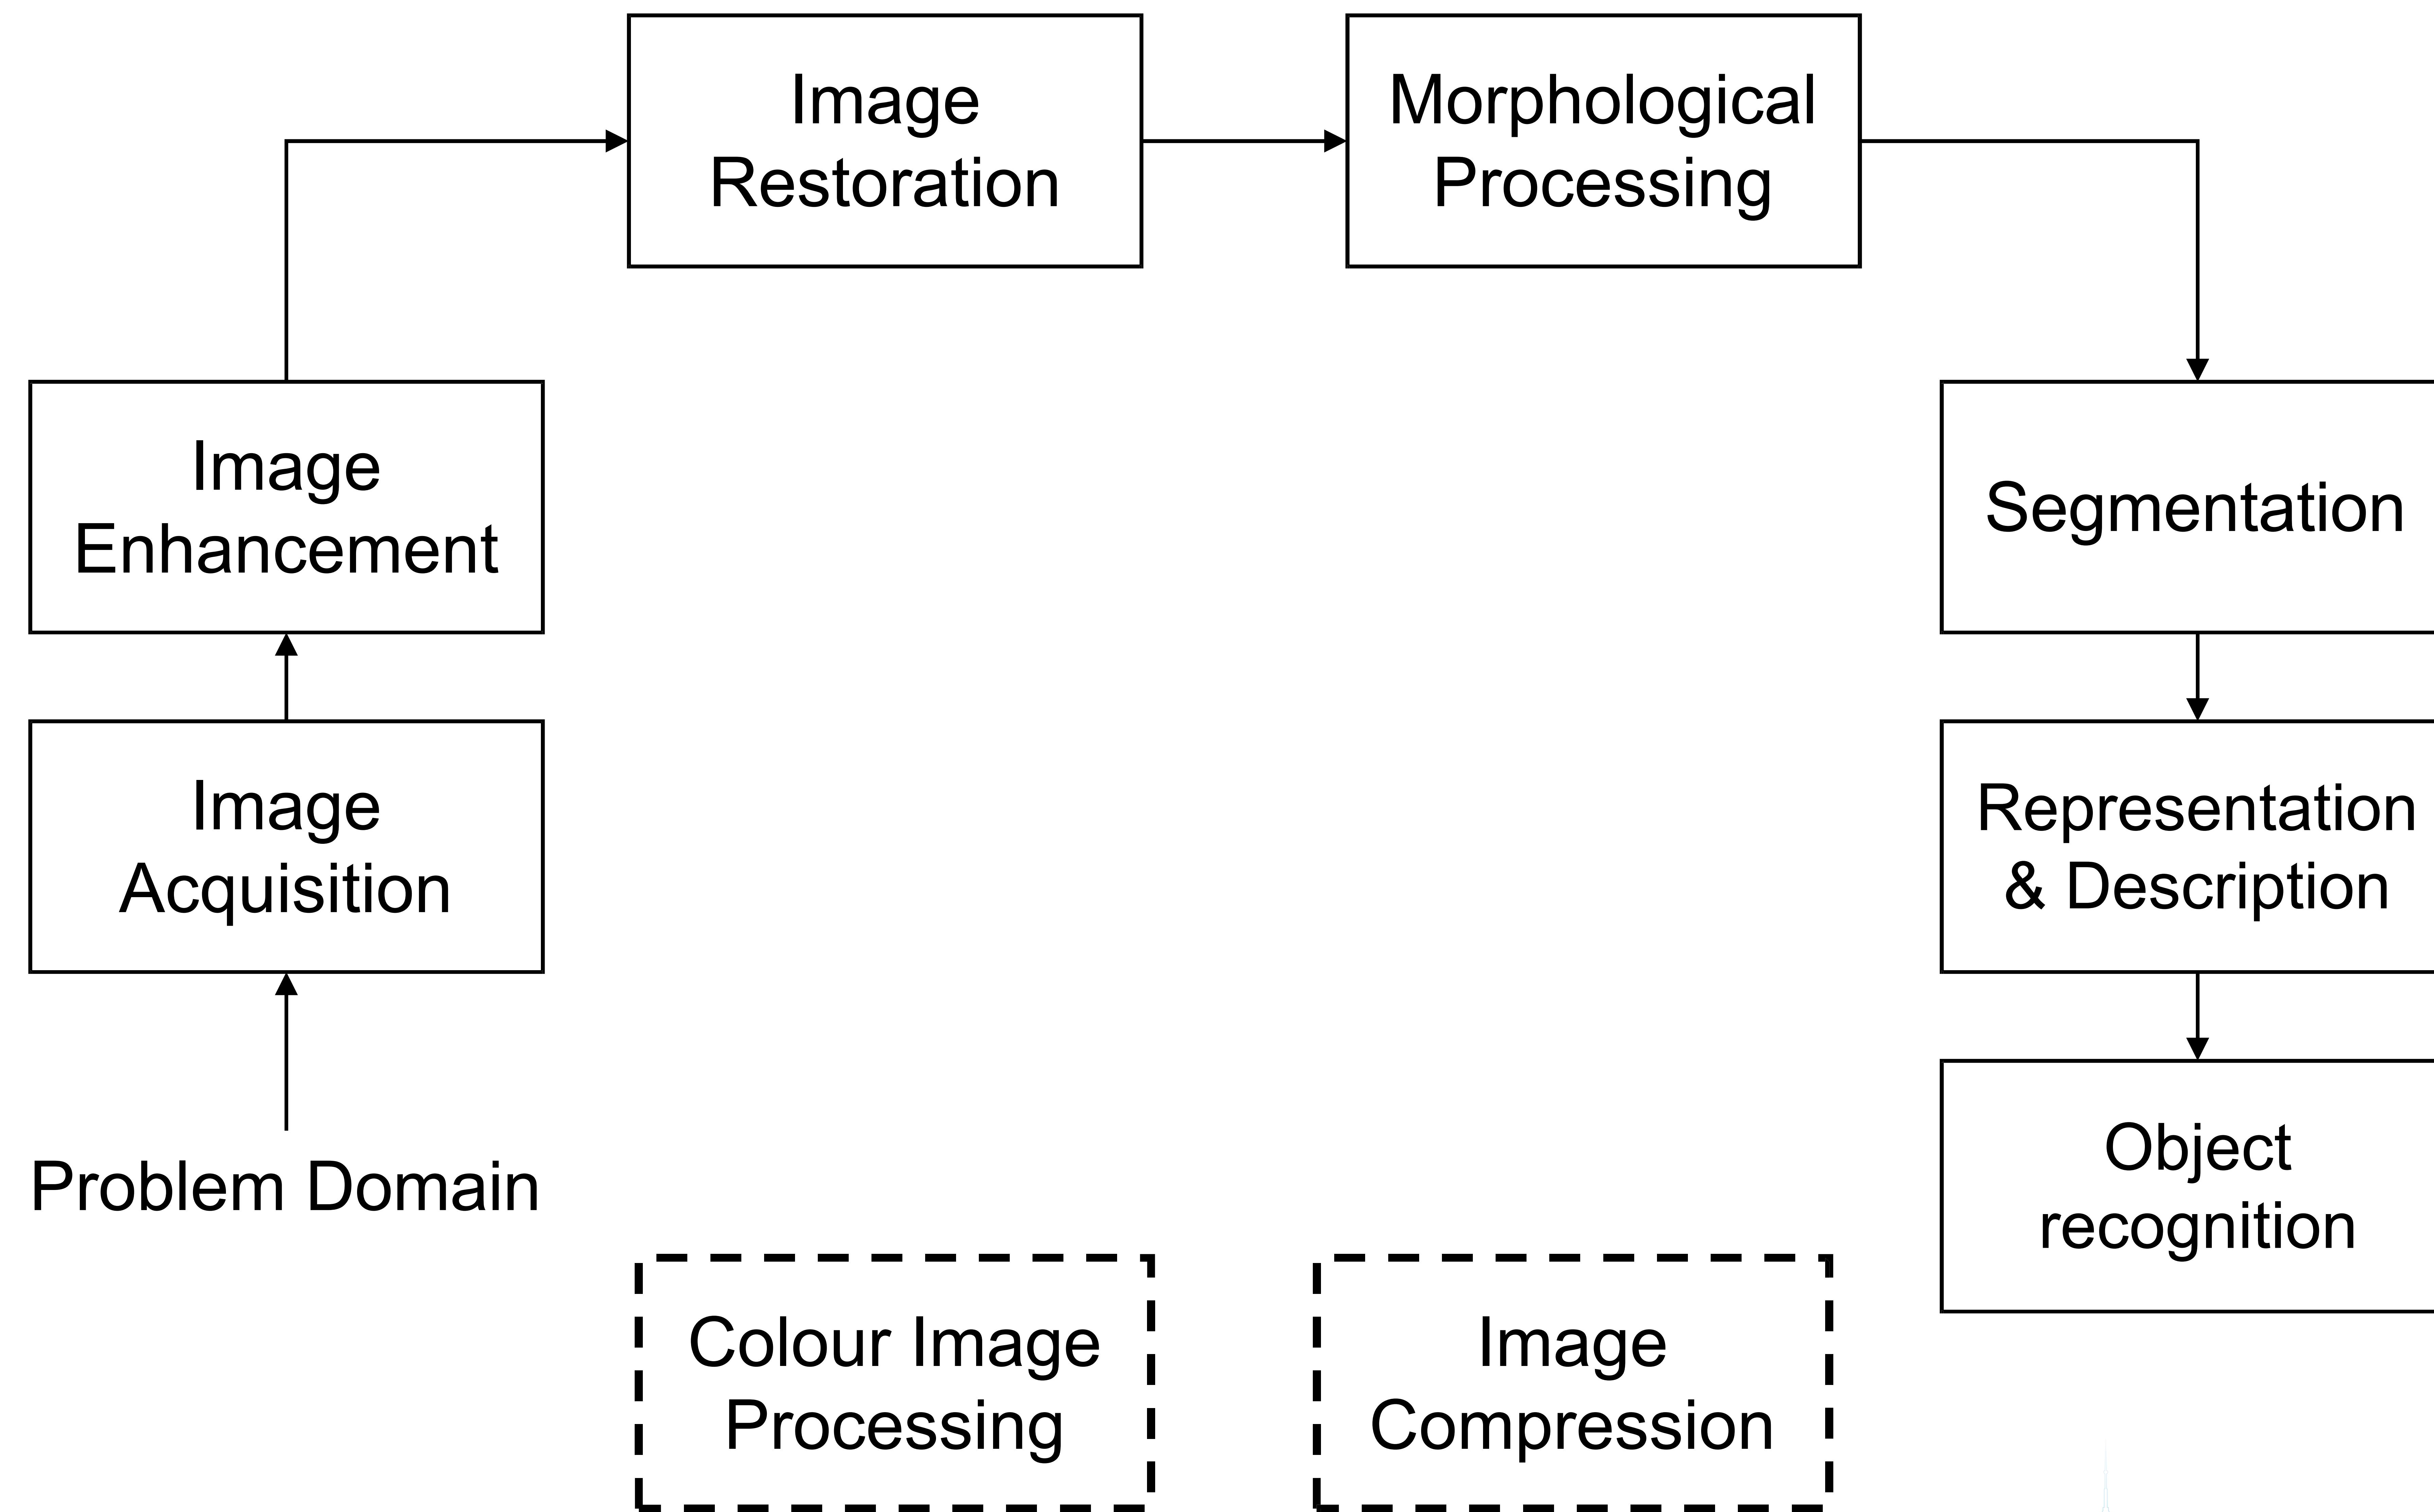
Segmentation, image registration, matching

Image Processing

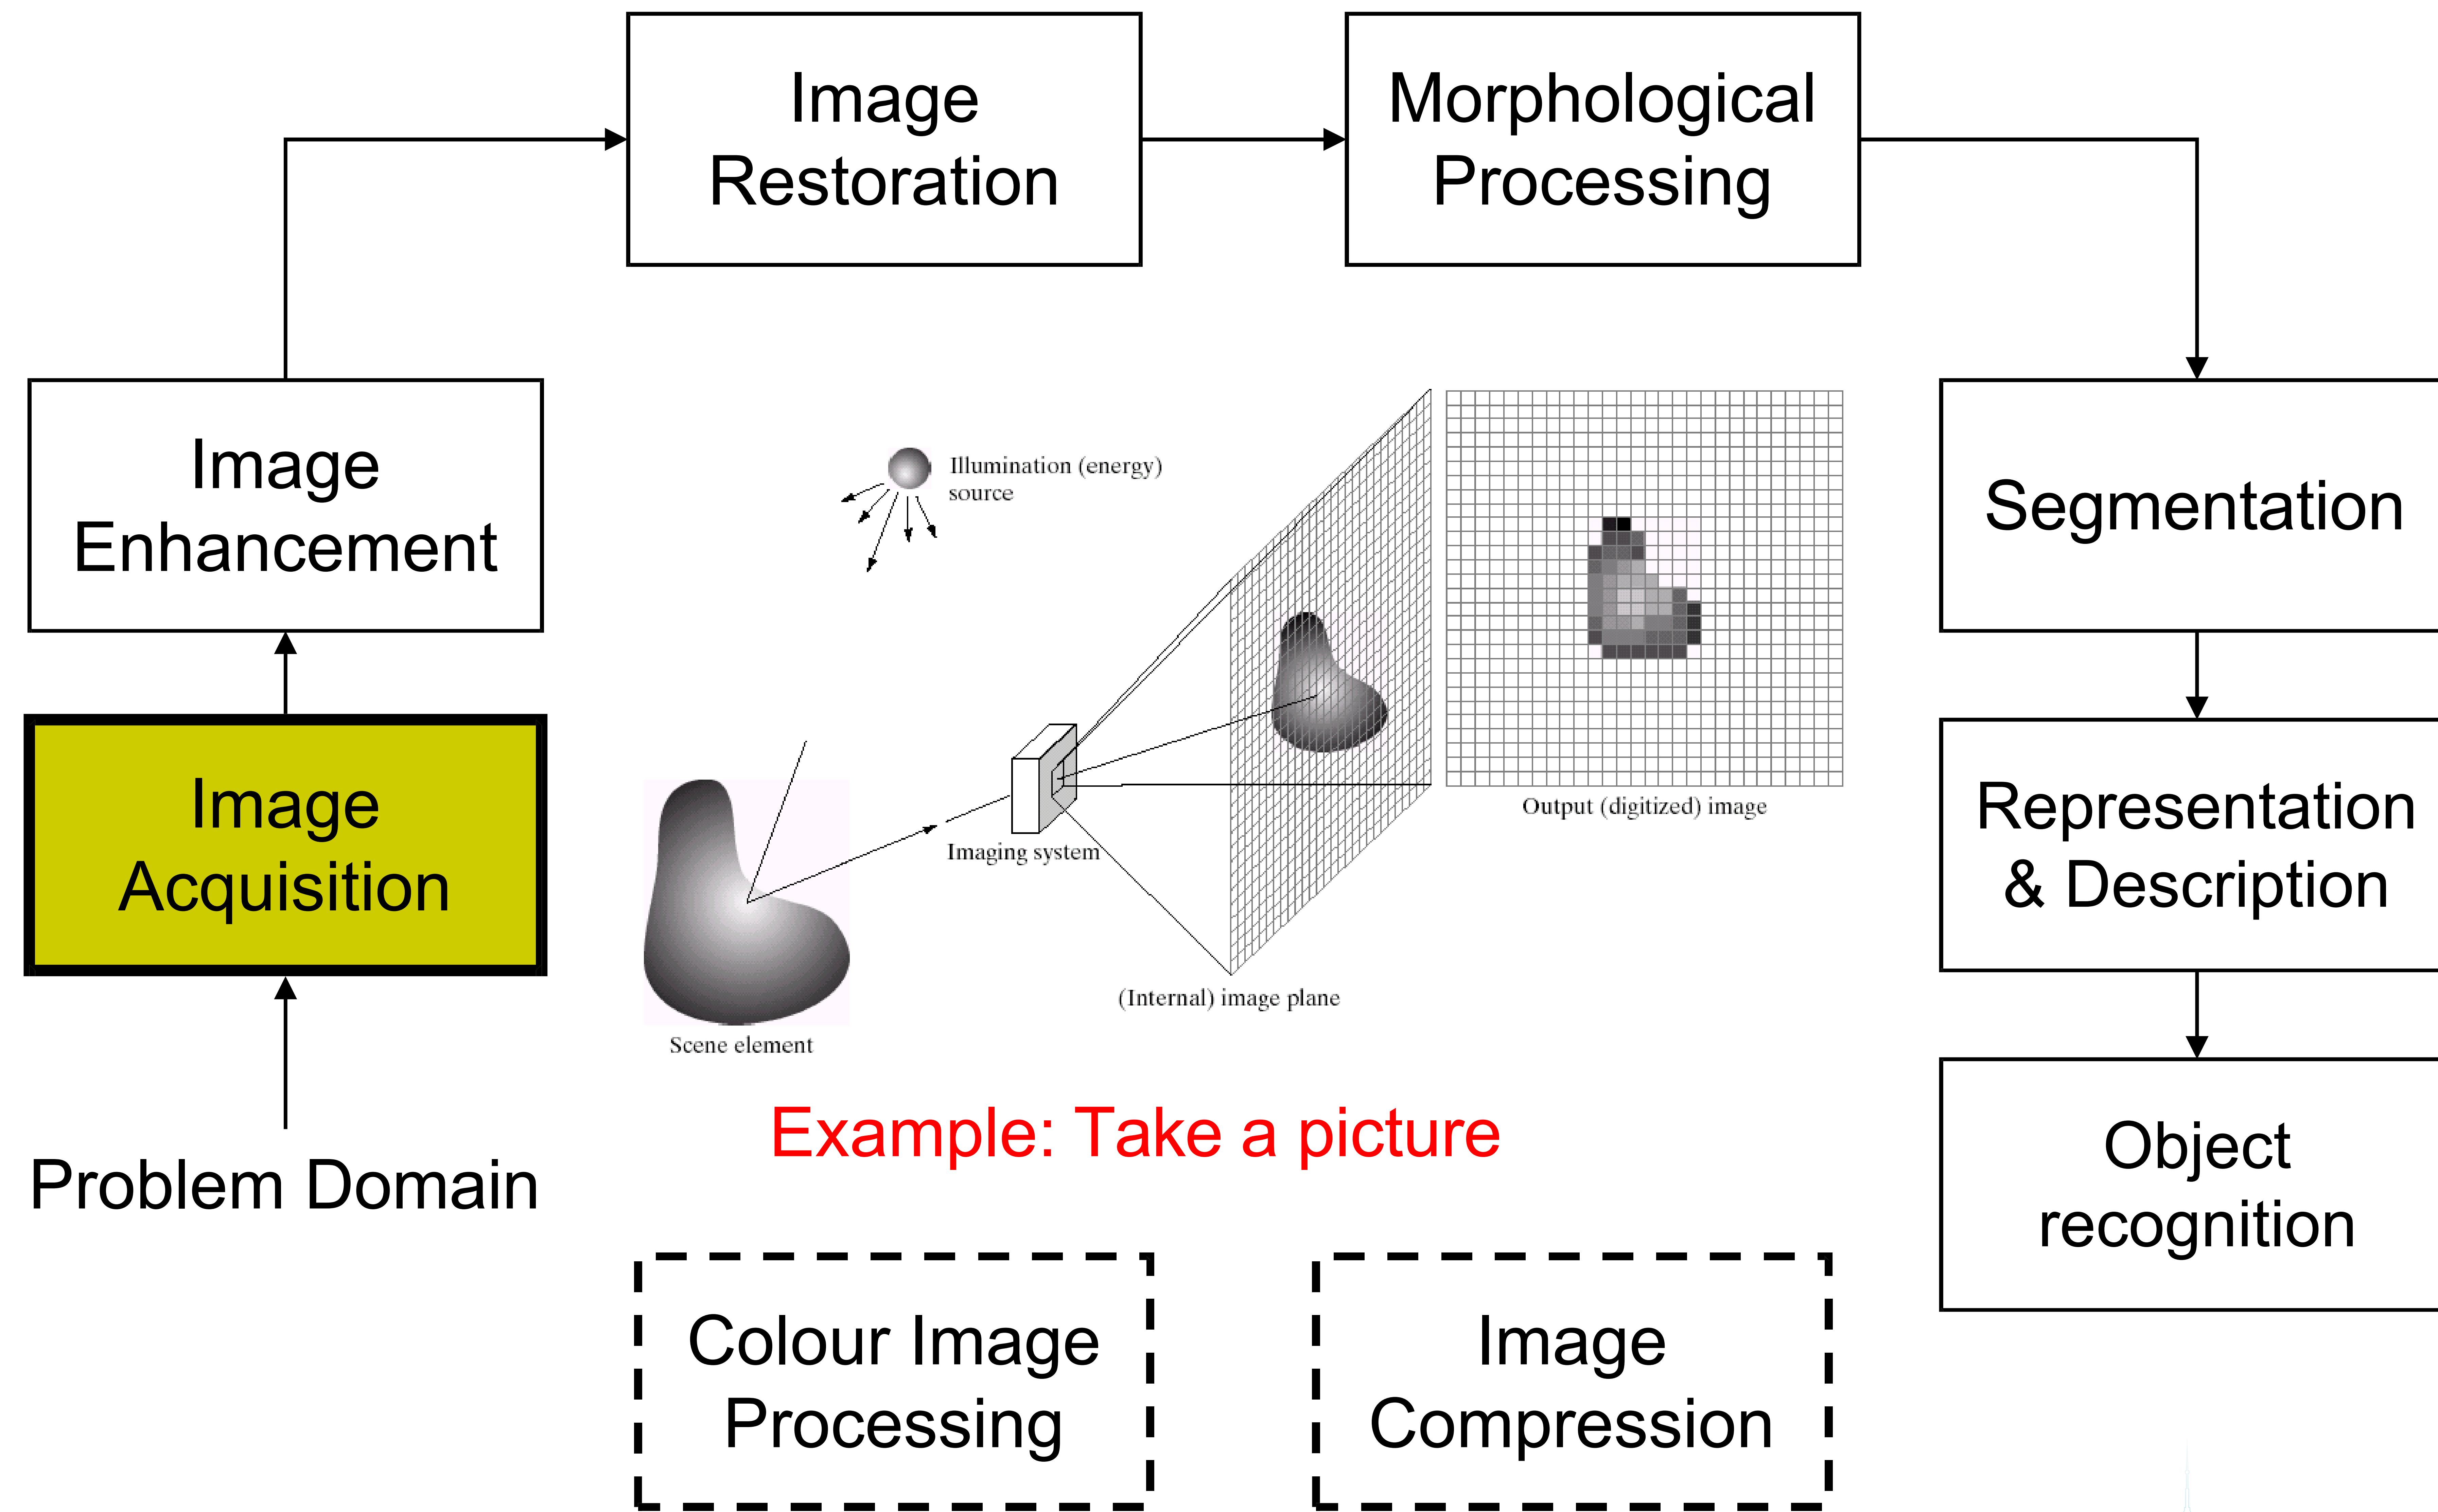
Image enhancement, noise removal, restoration,
feature detection, compression



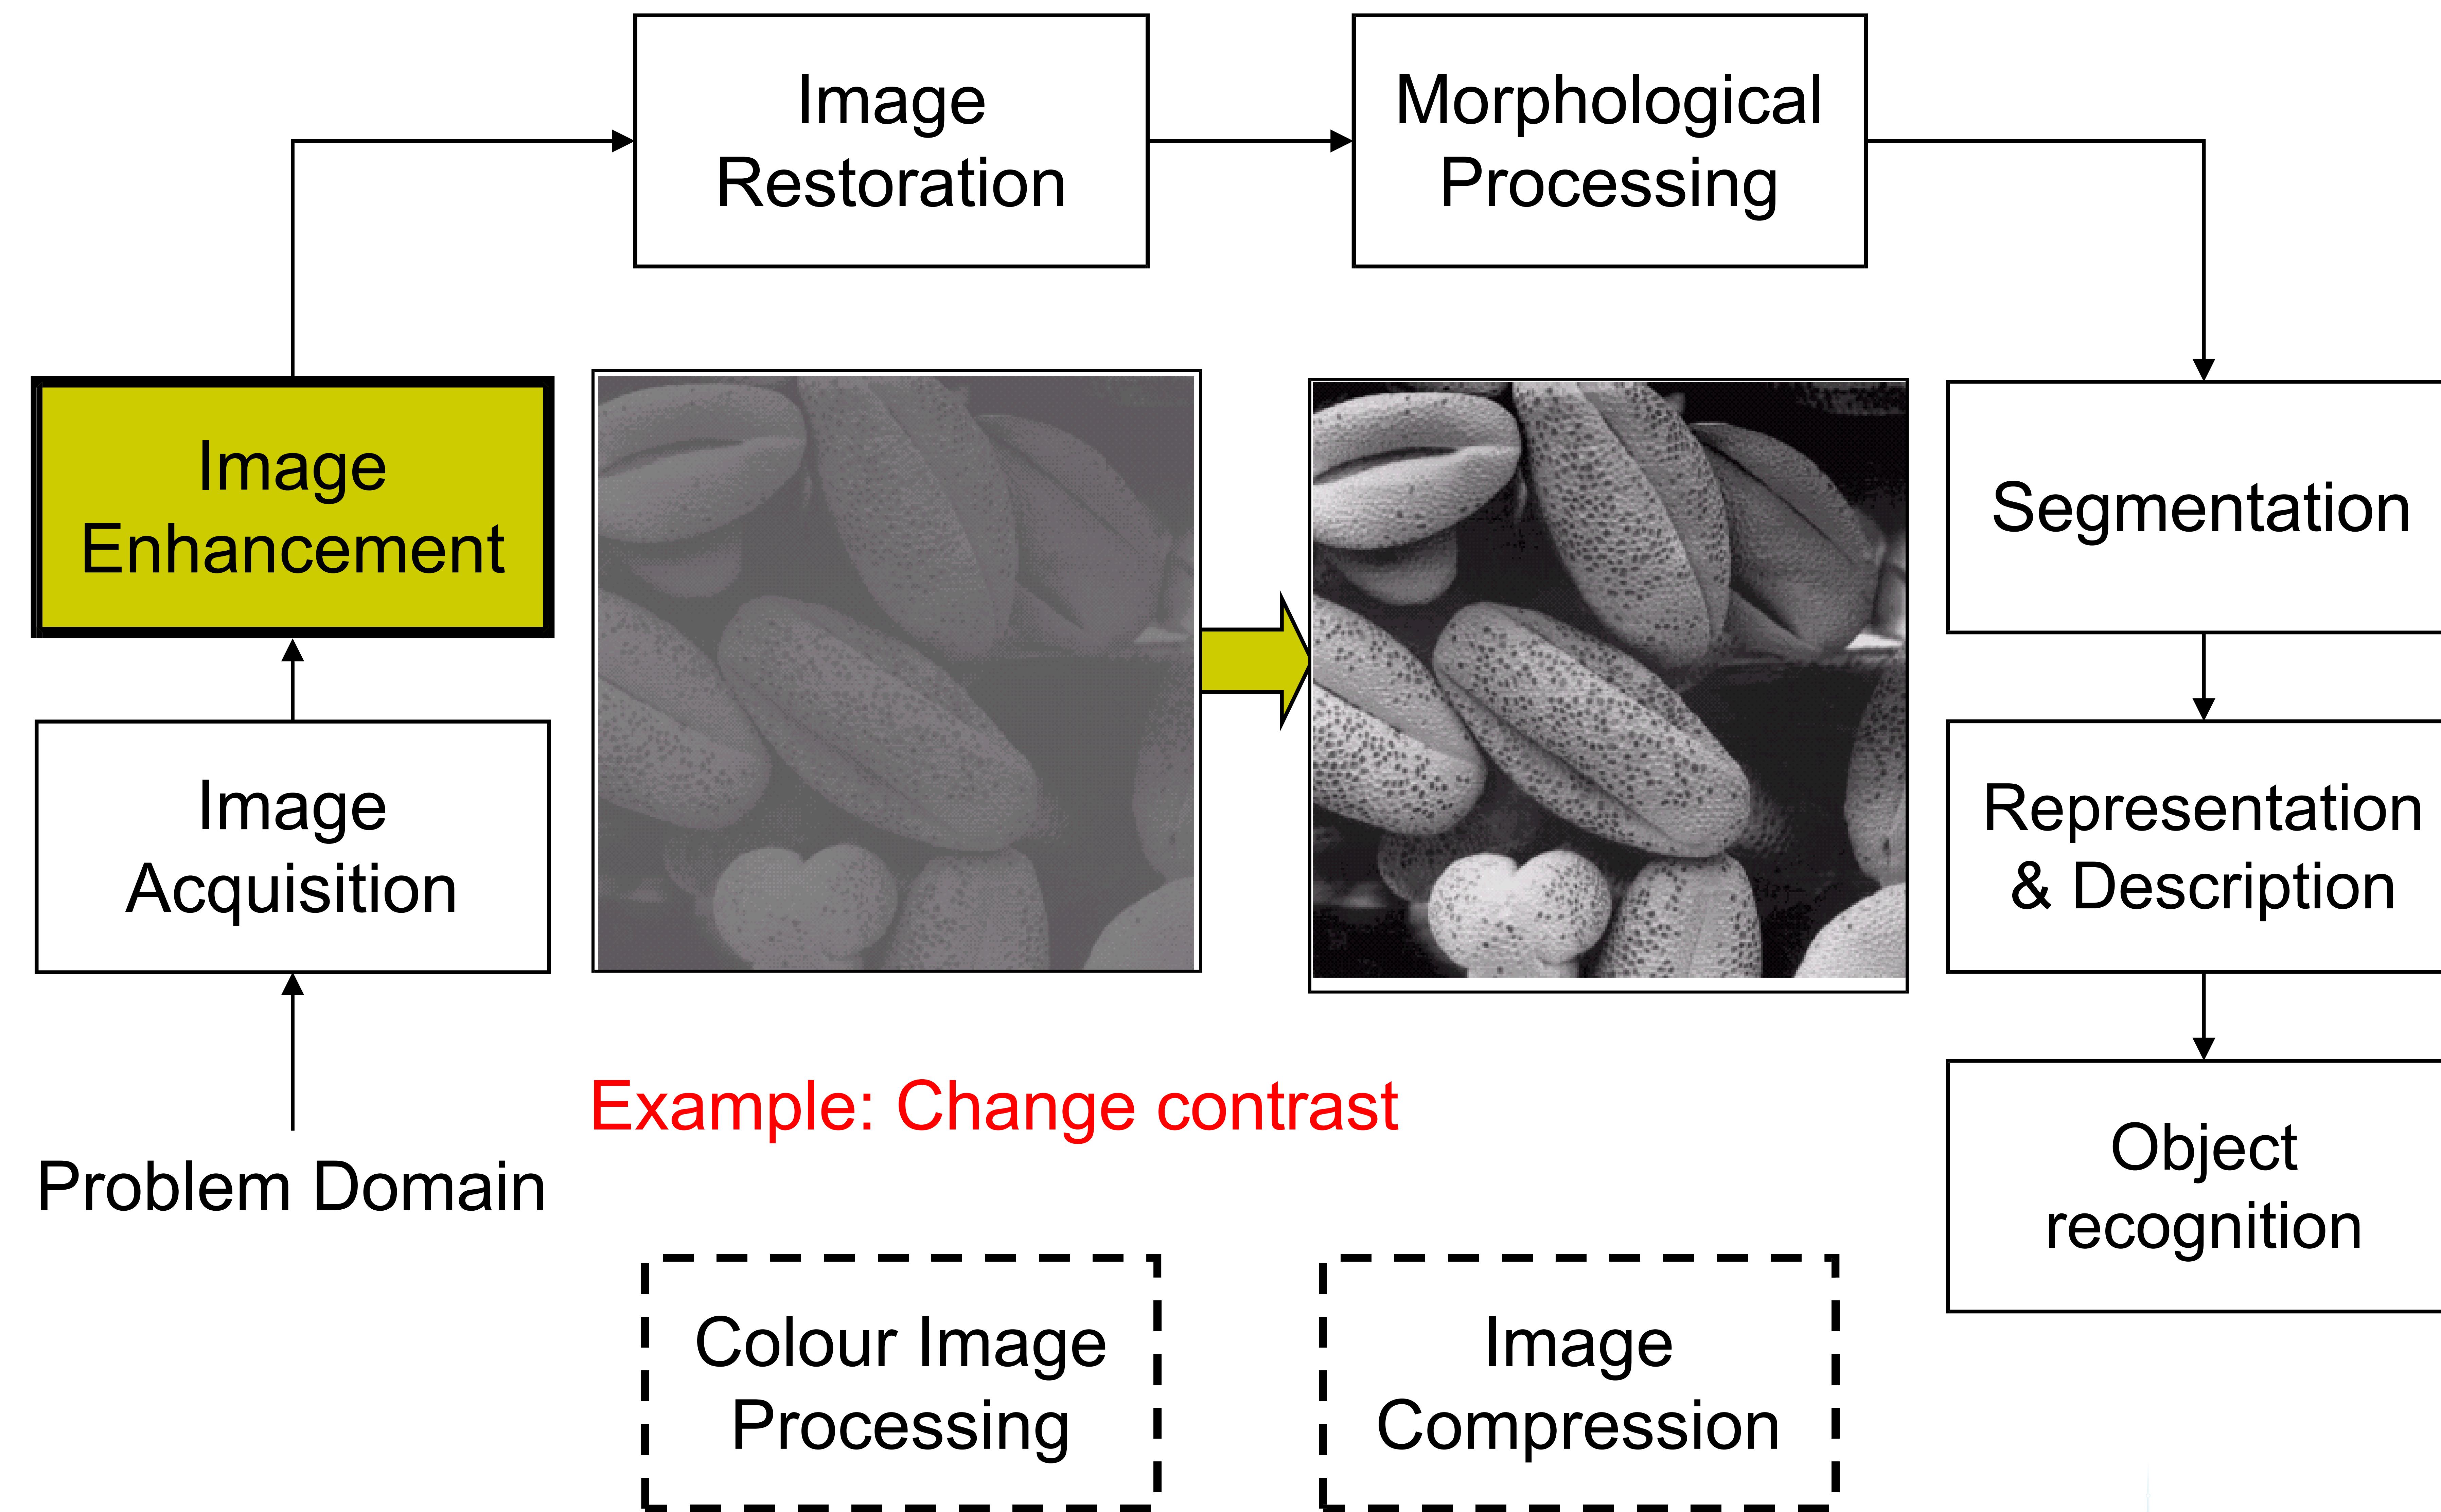
Key stages in digital image processing



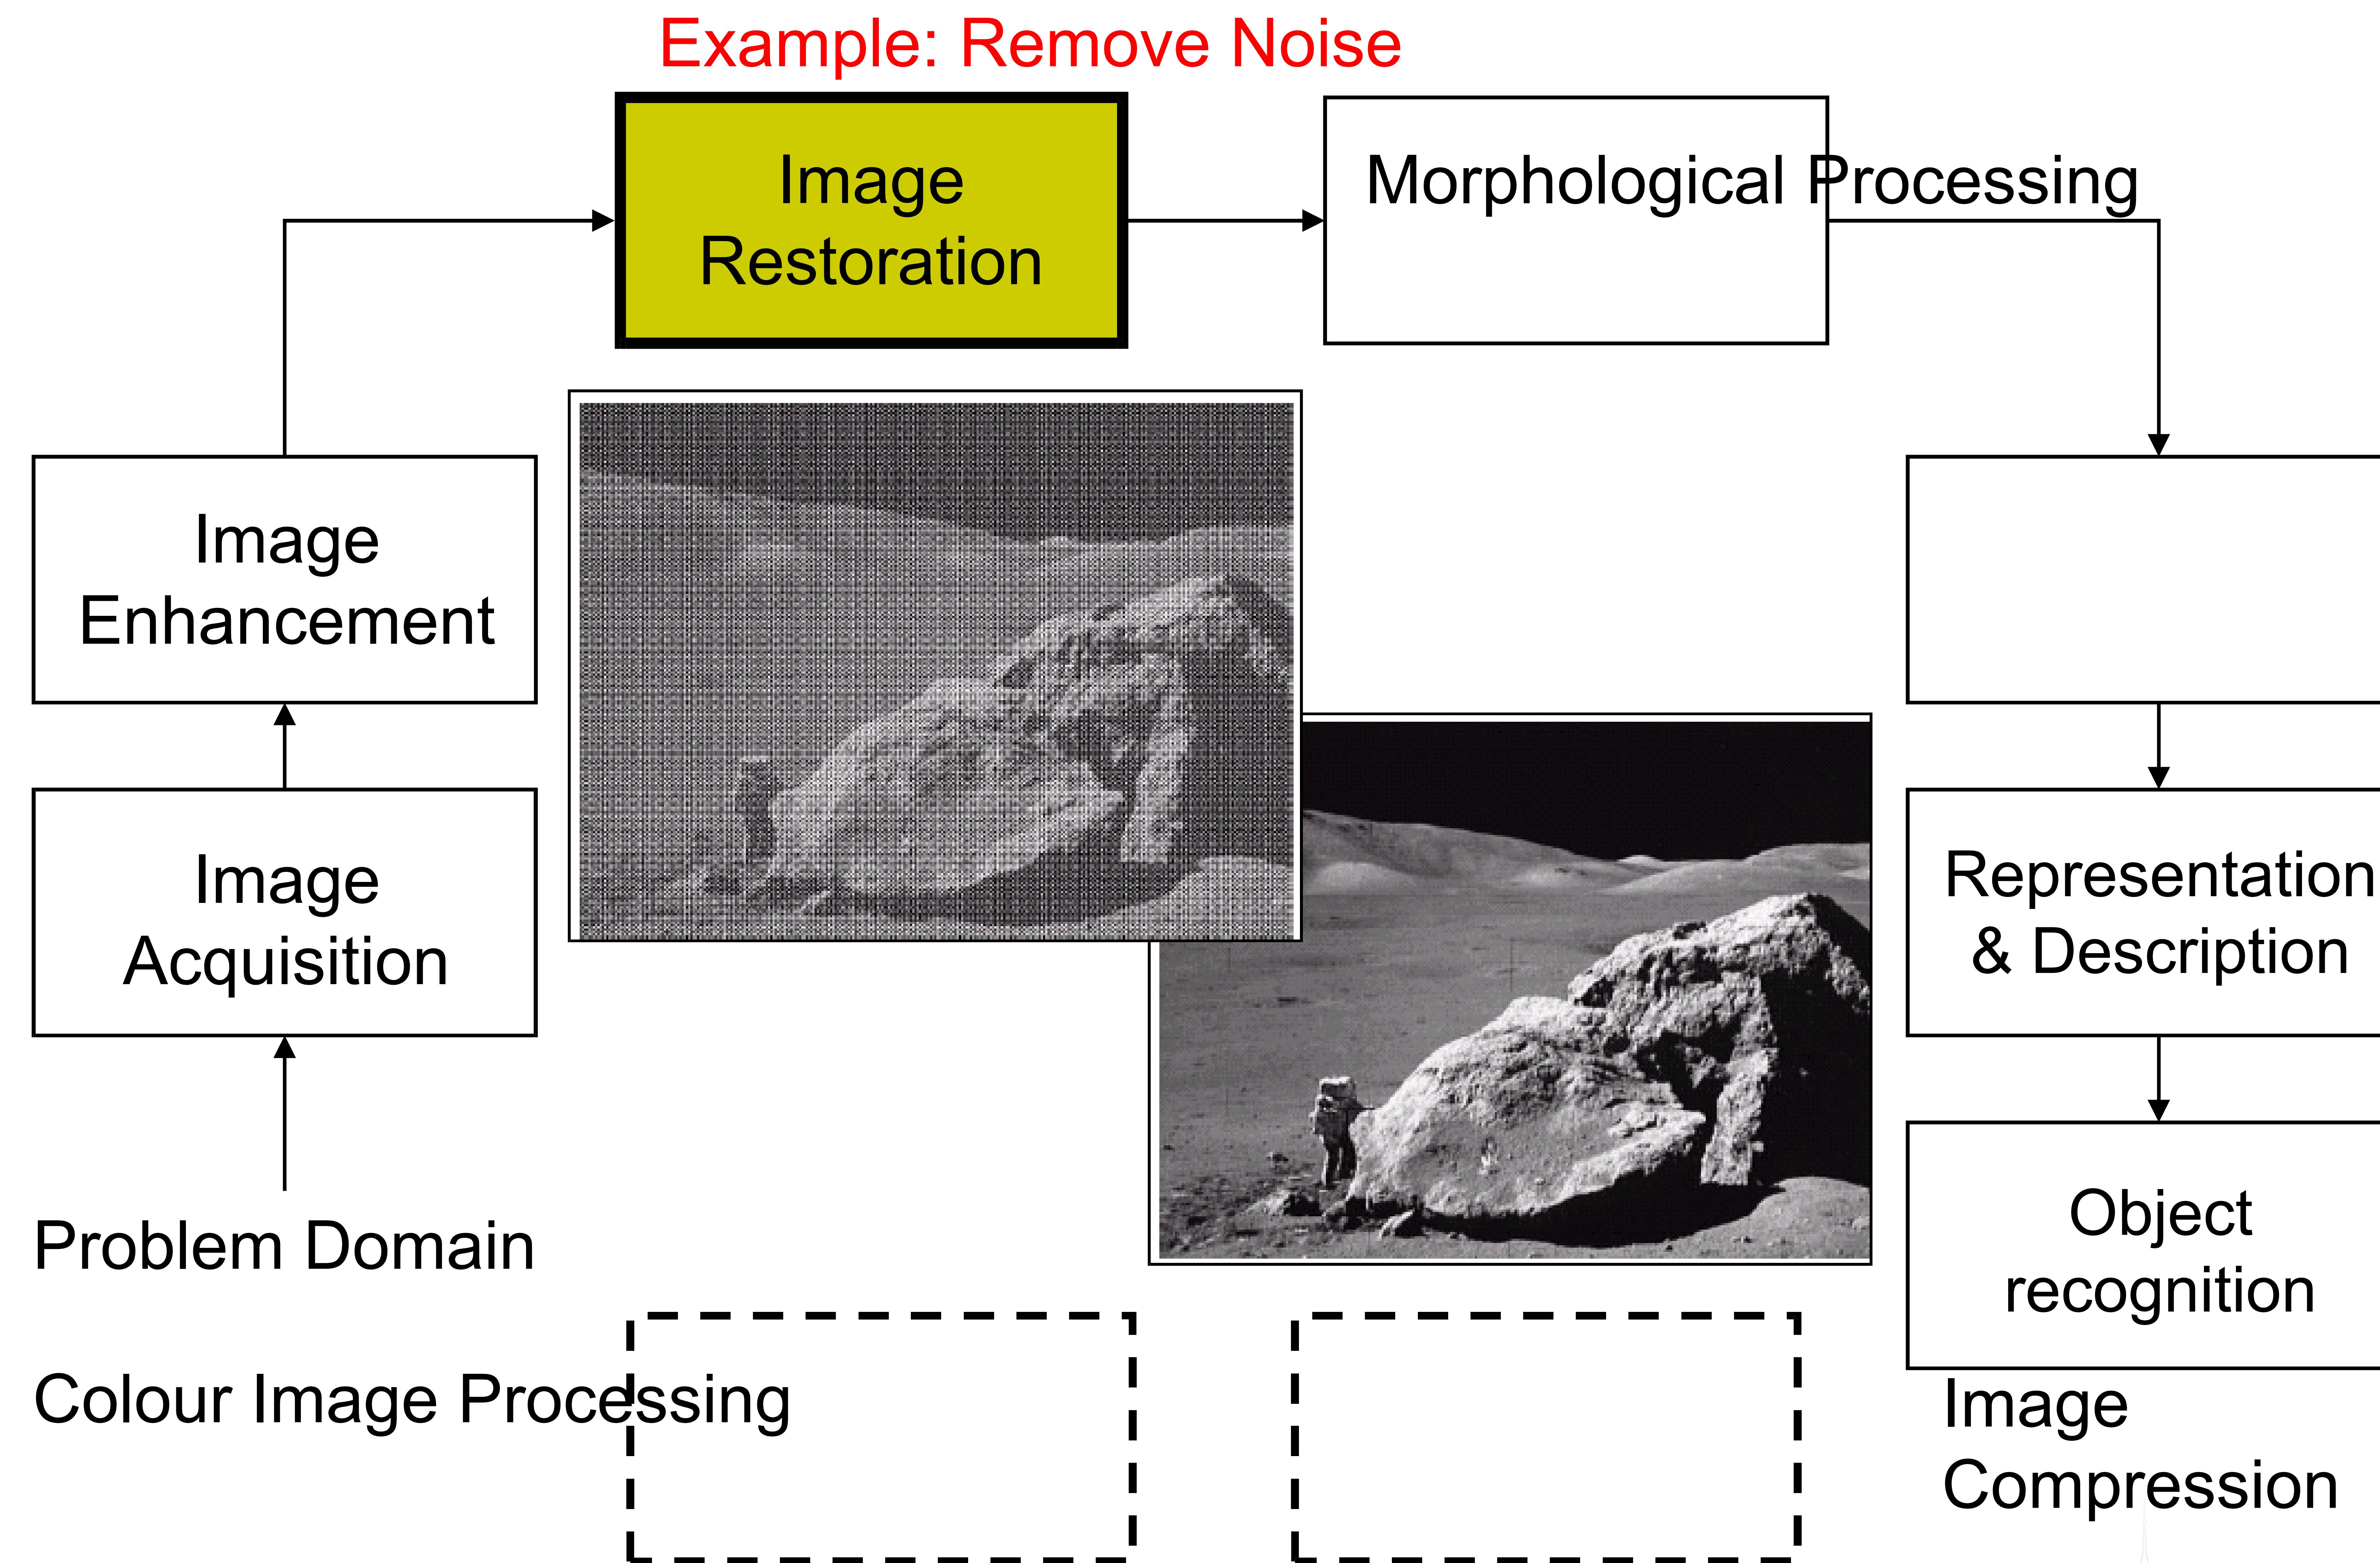
Key stages in digital image processing



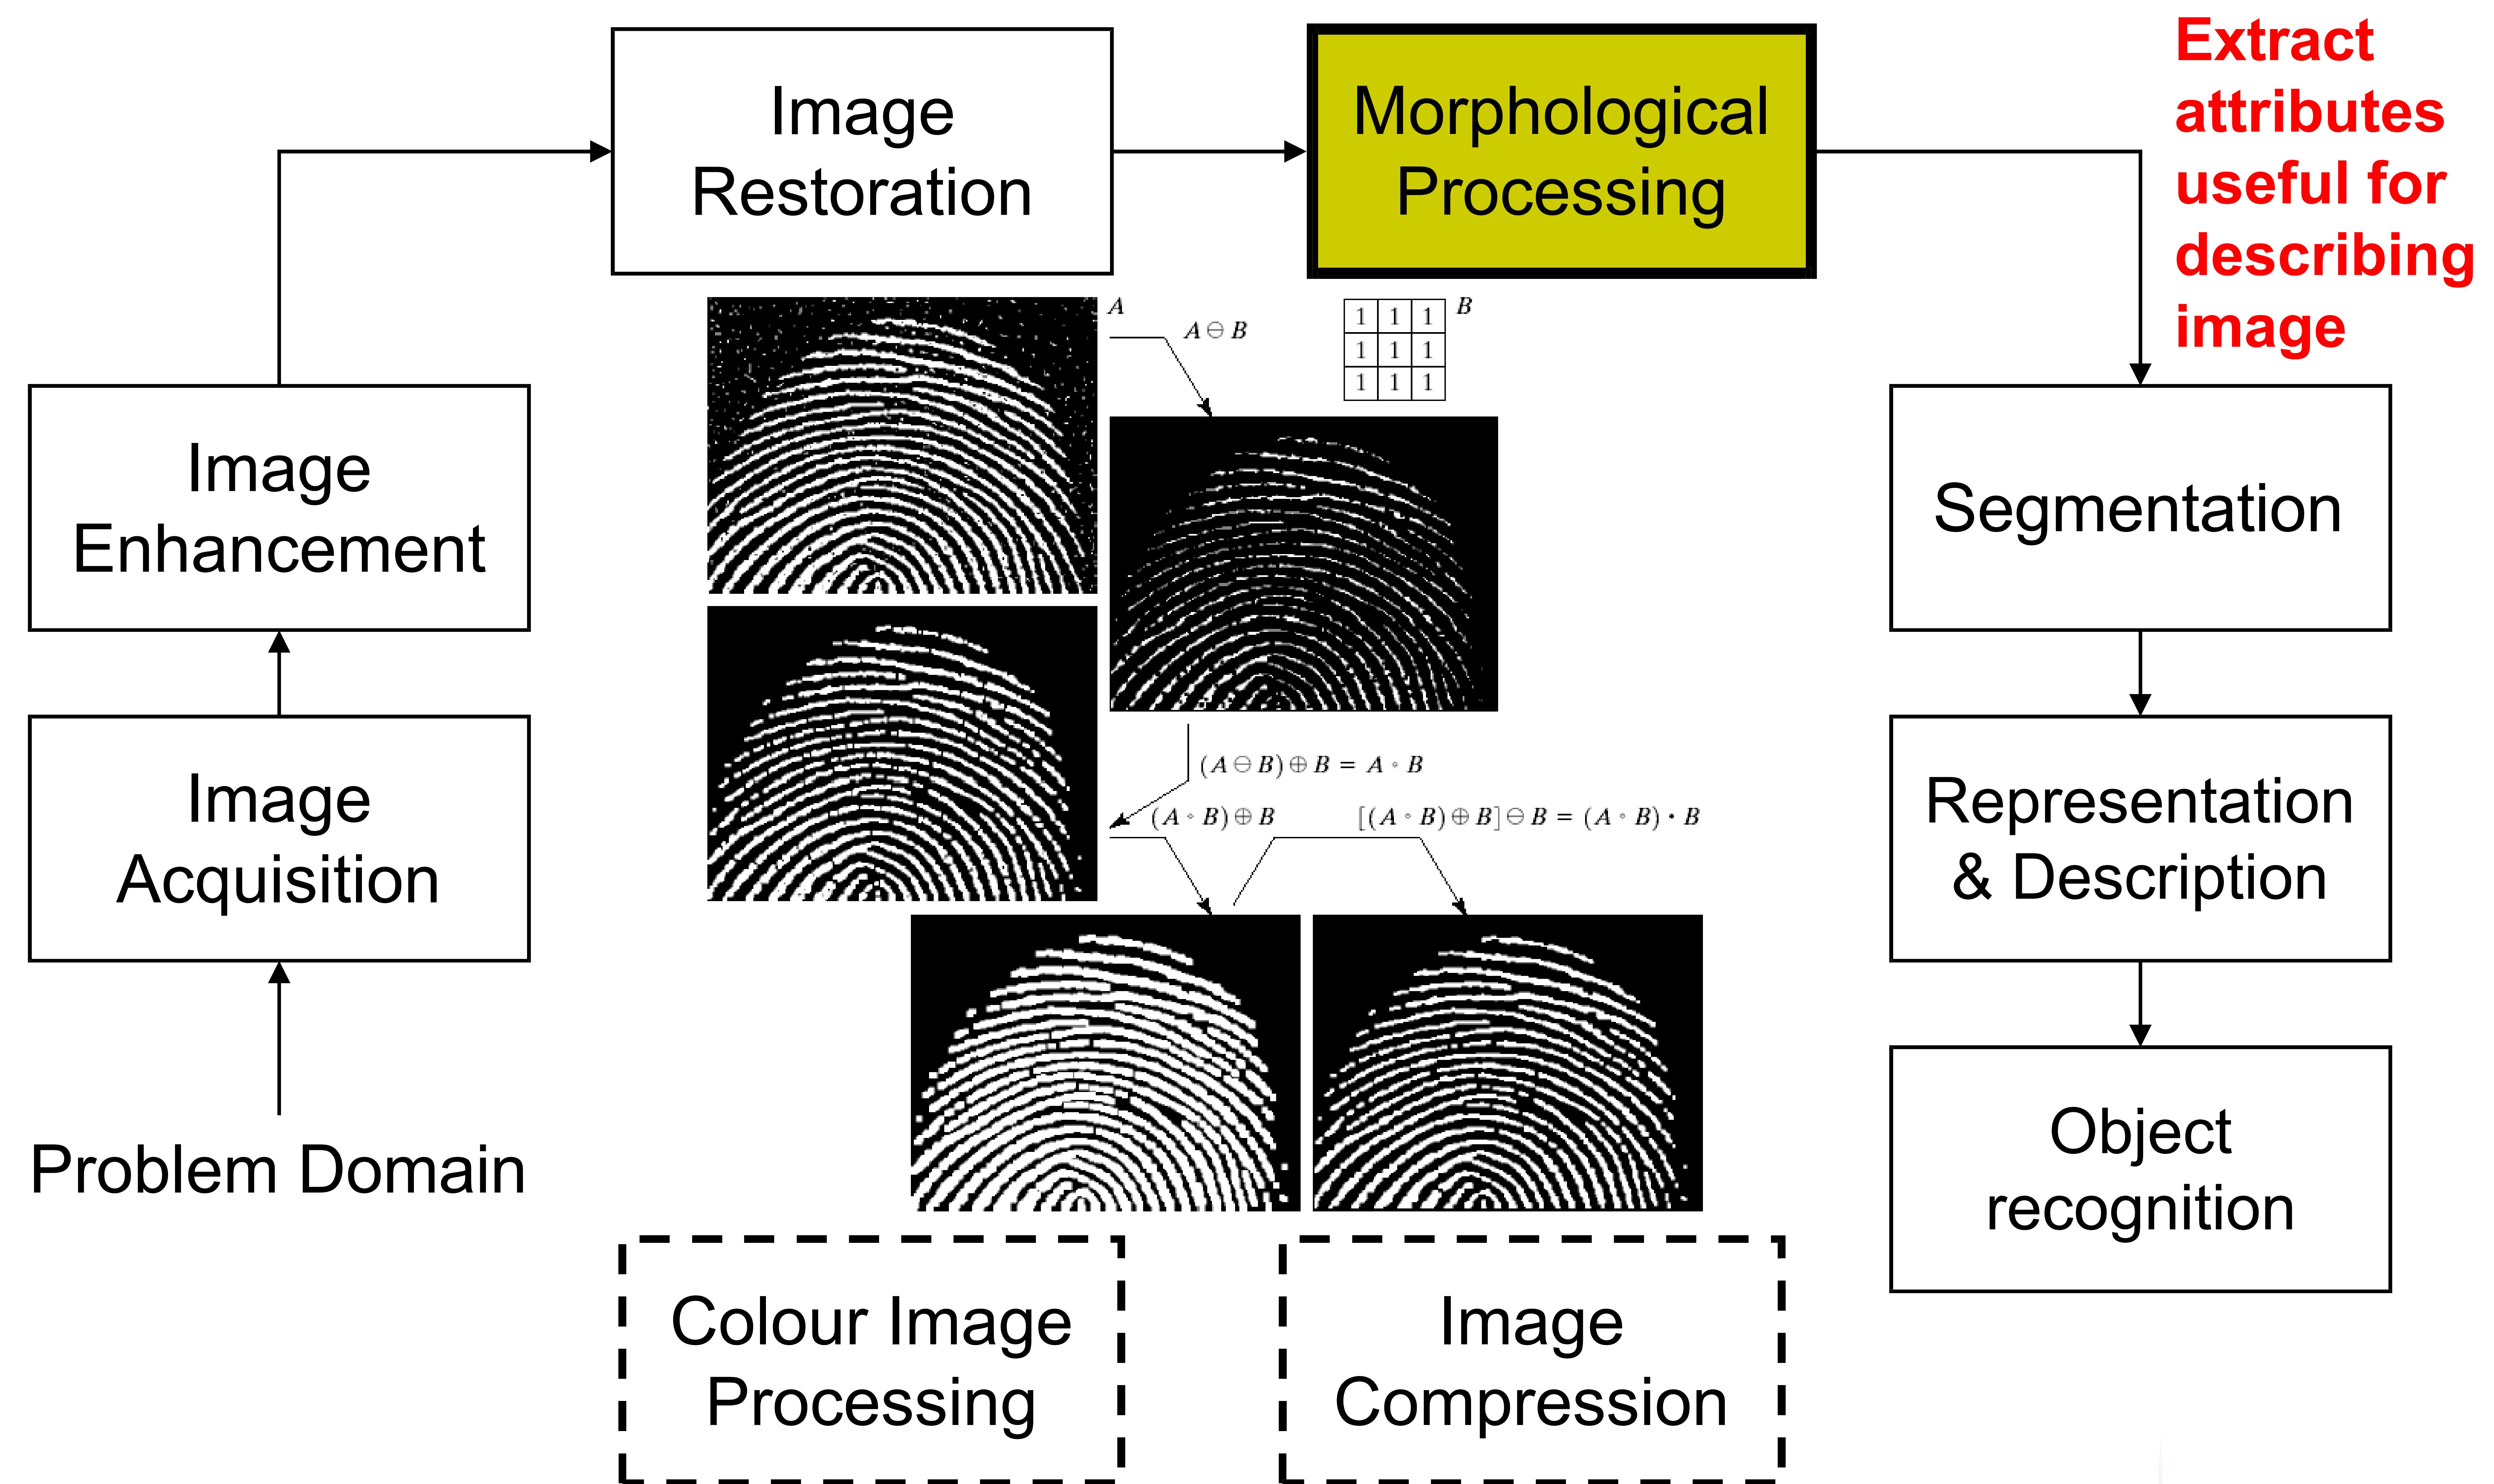
Key stages in digital image processing



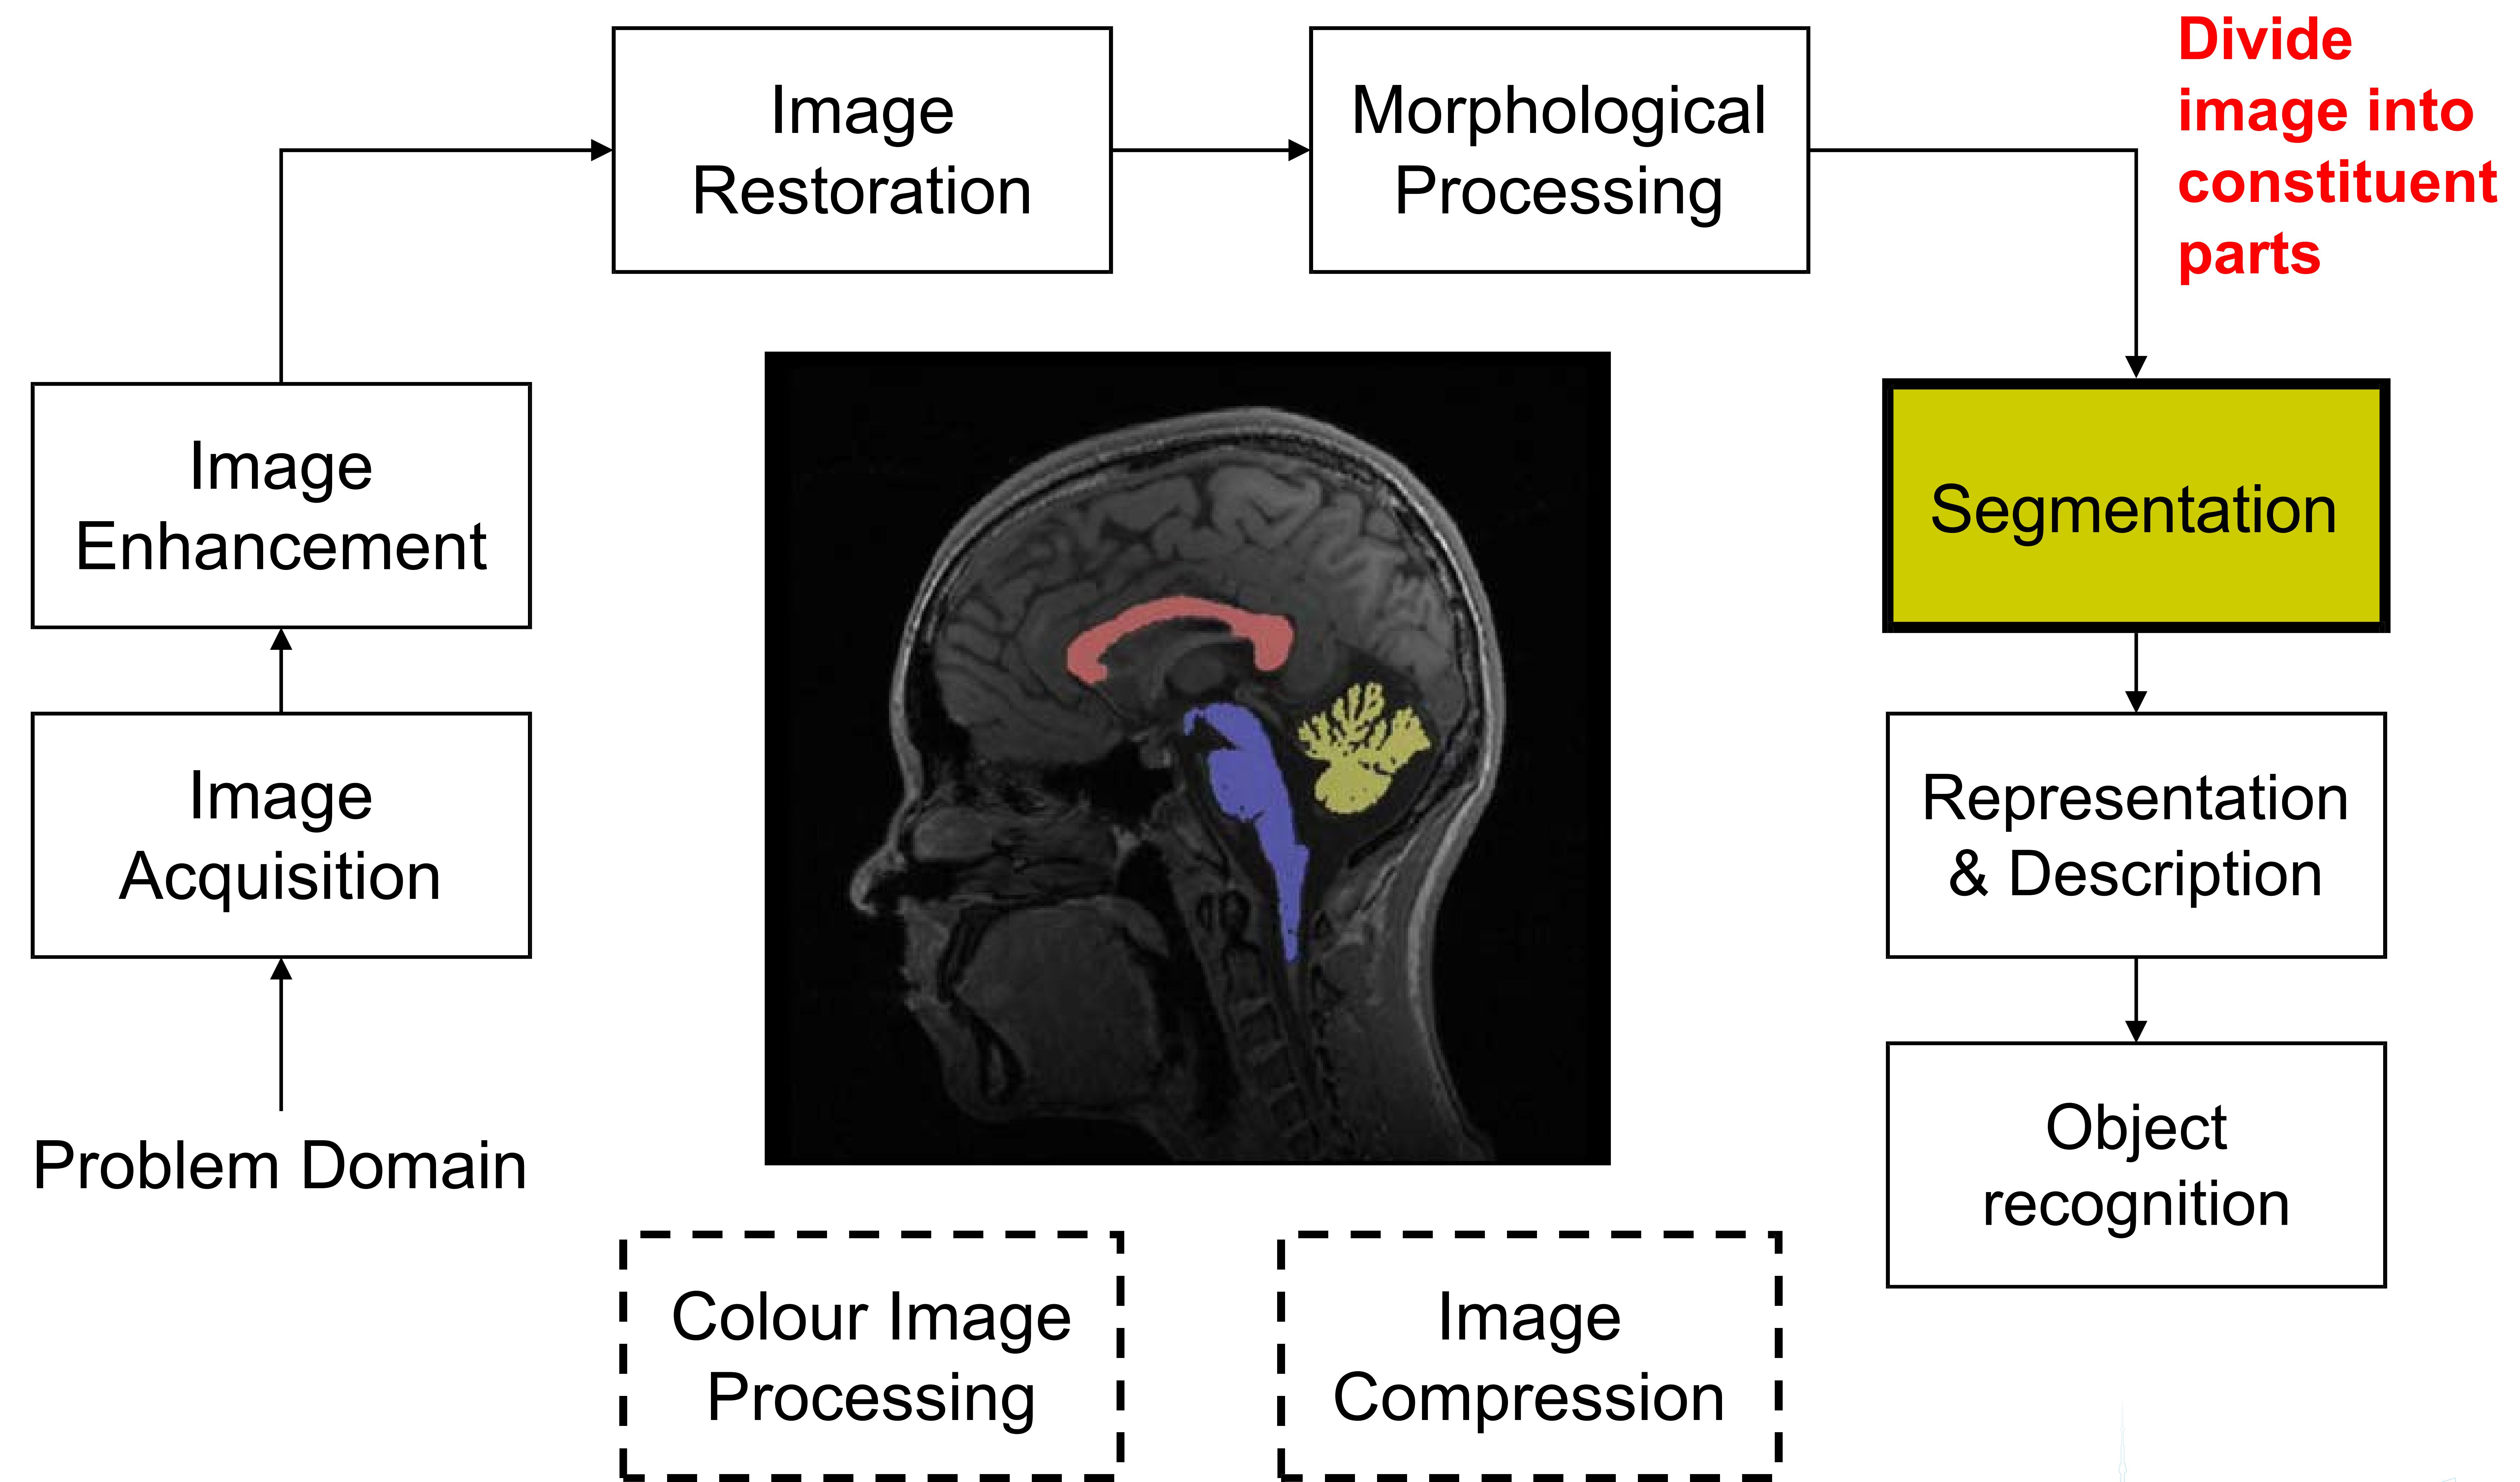
Key stages in digital image processing



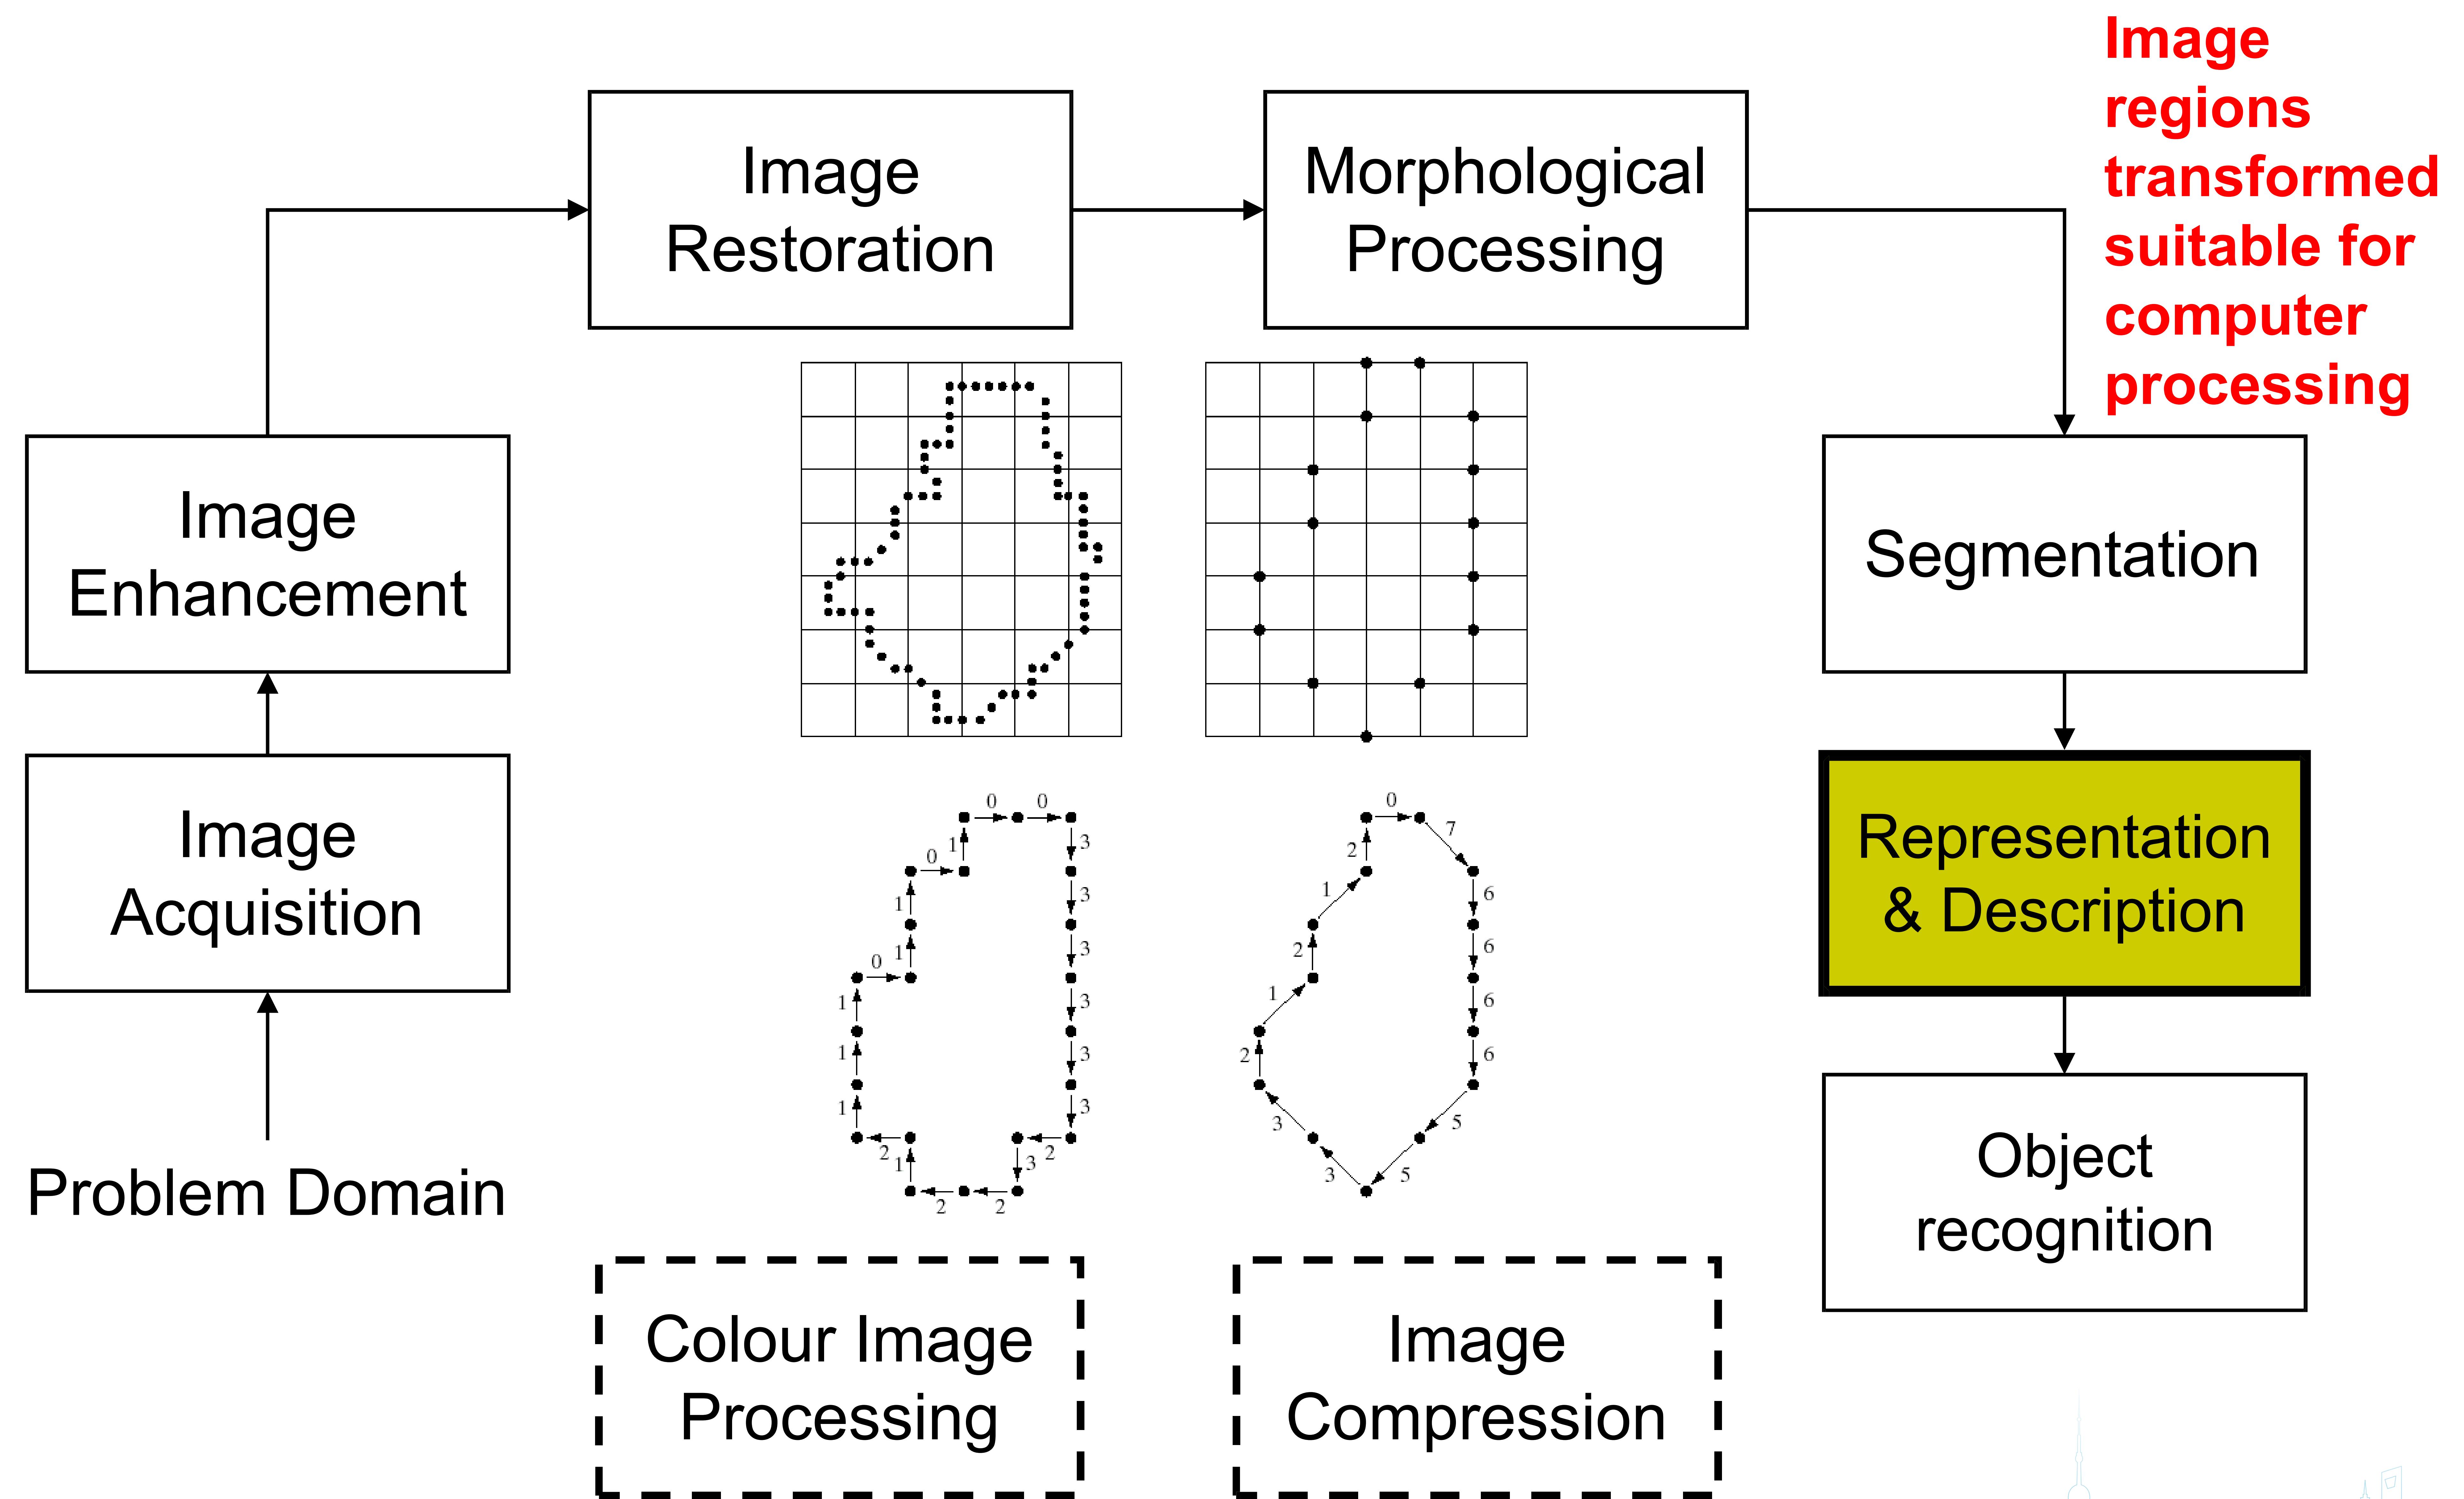
Key stages in digital image processing



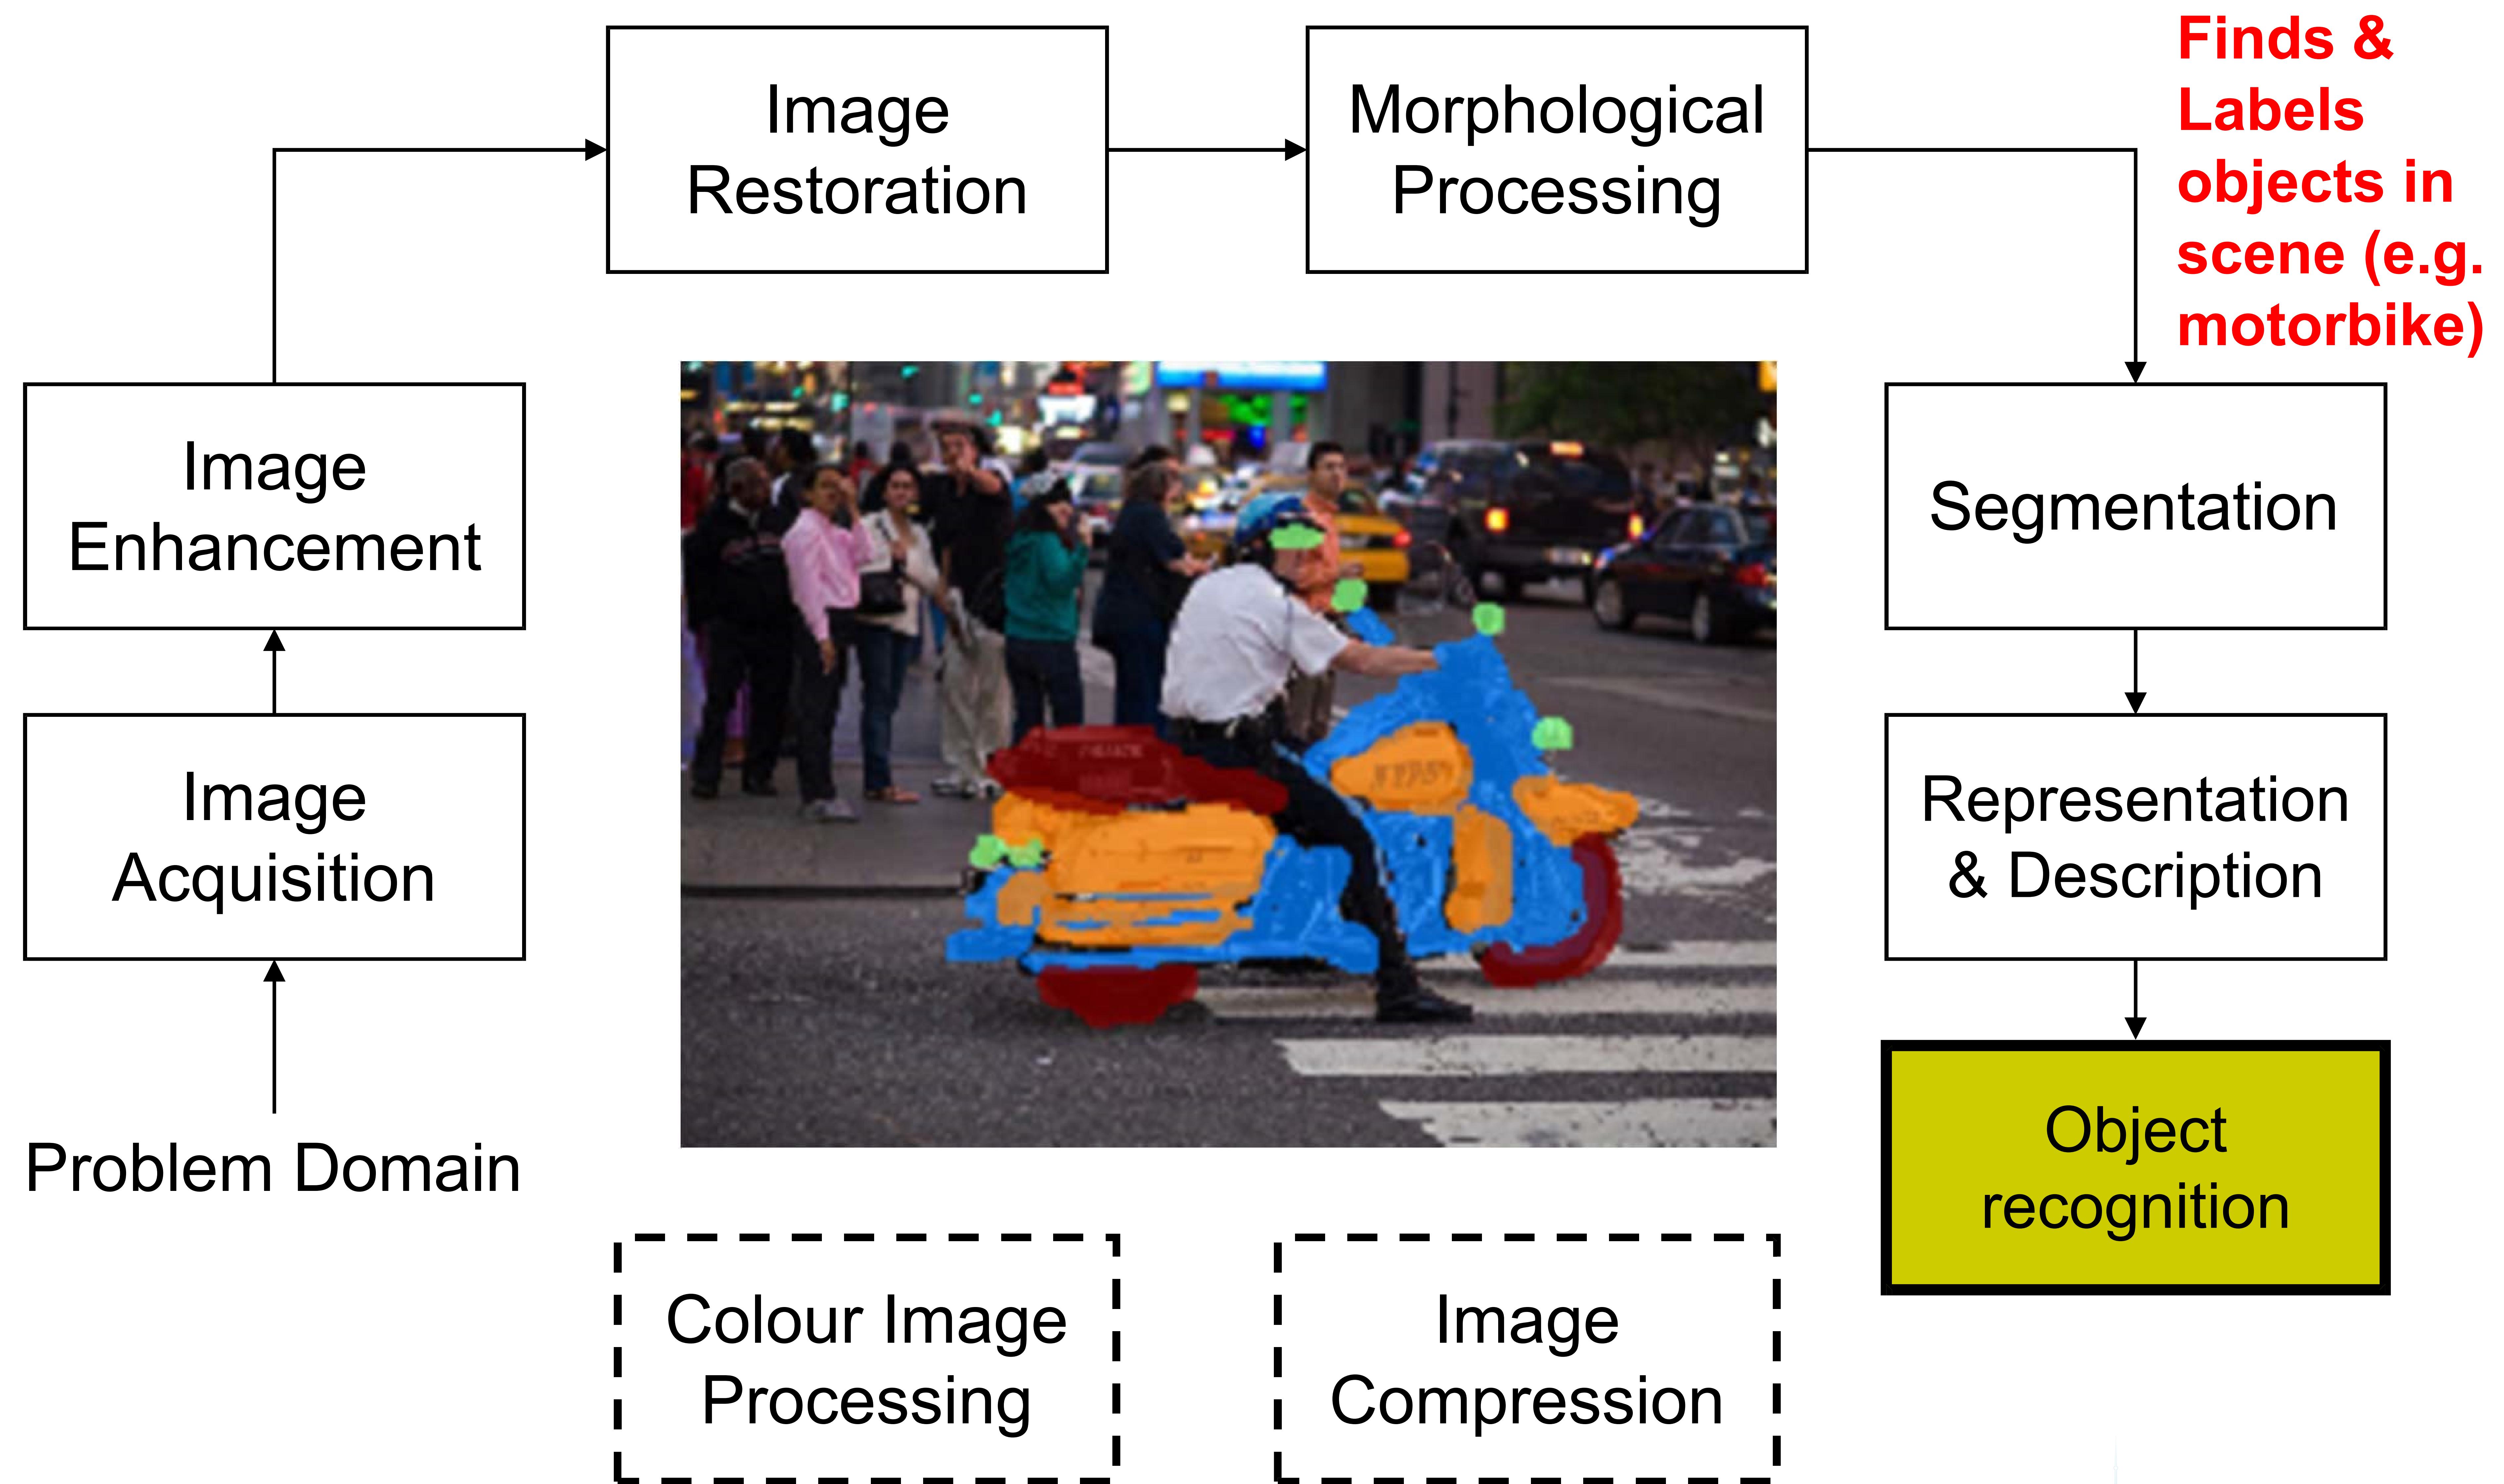
Key stages in digital image processing



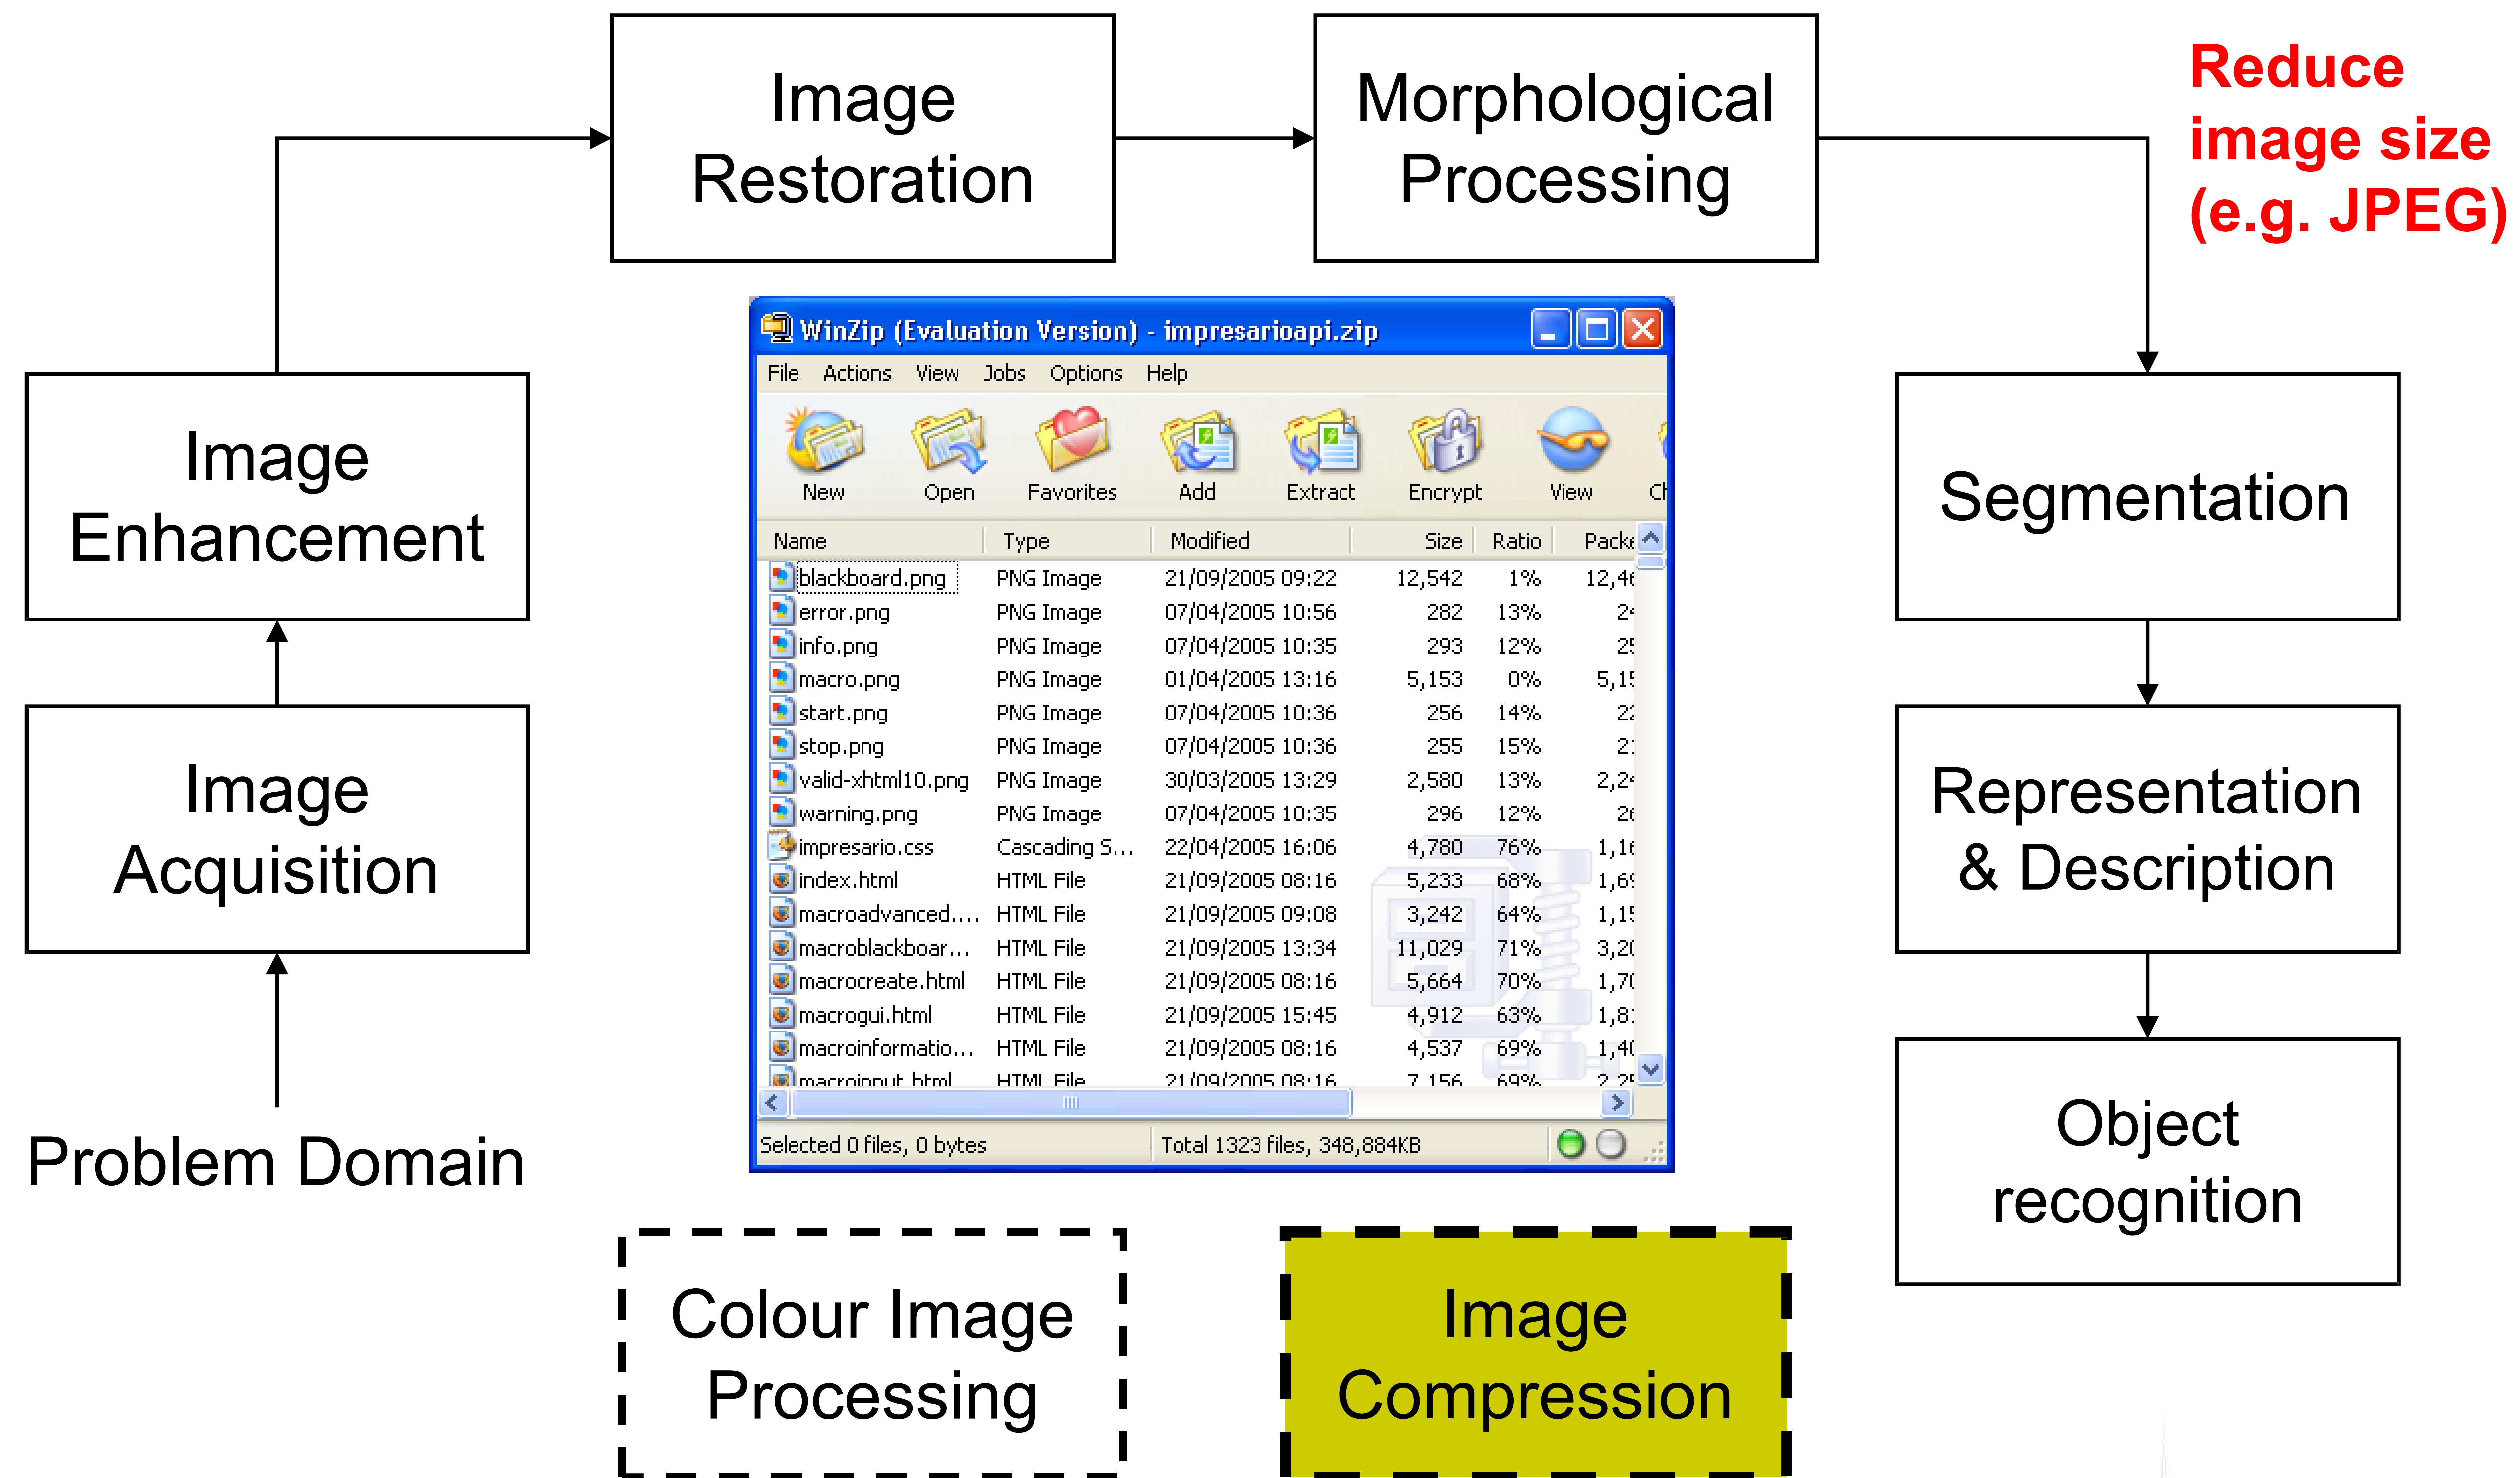
Key stages in digital image processing



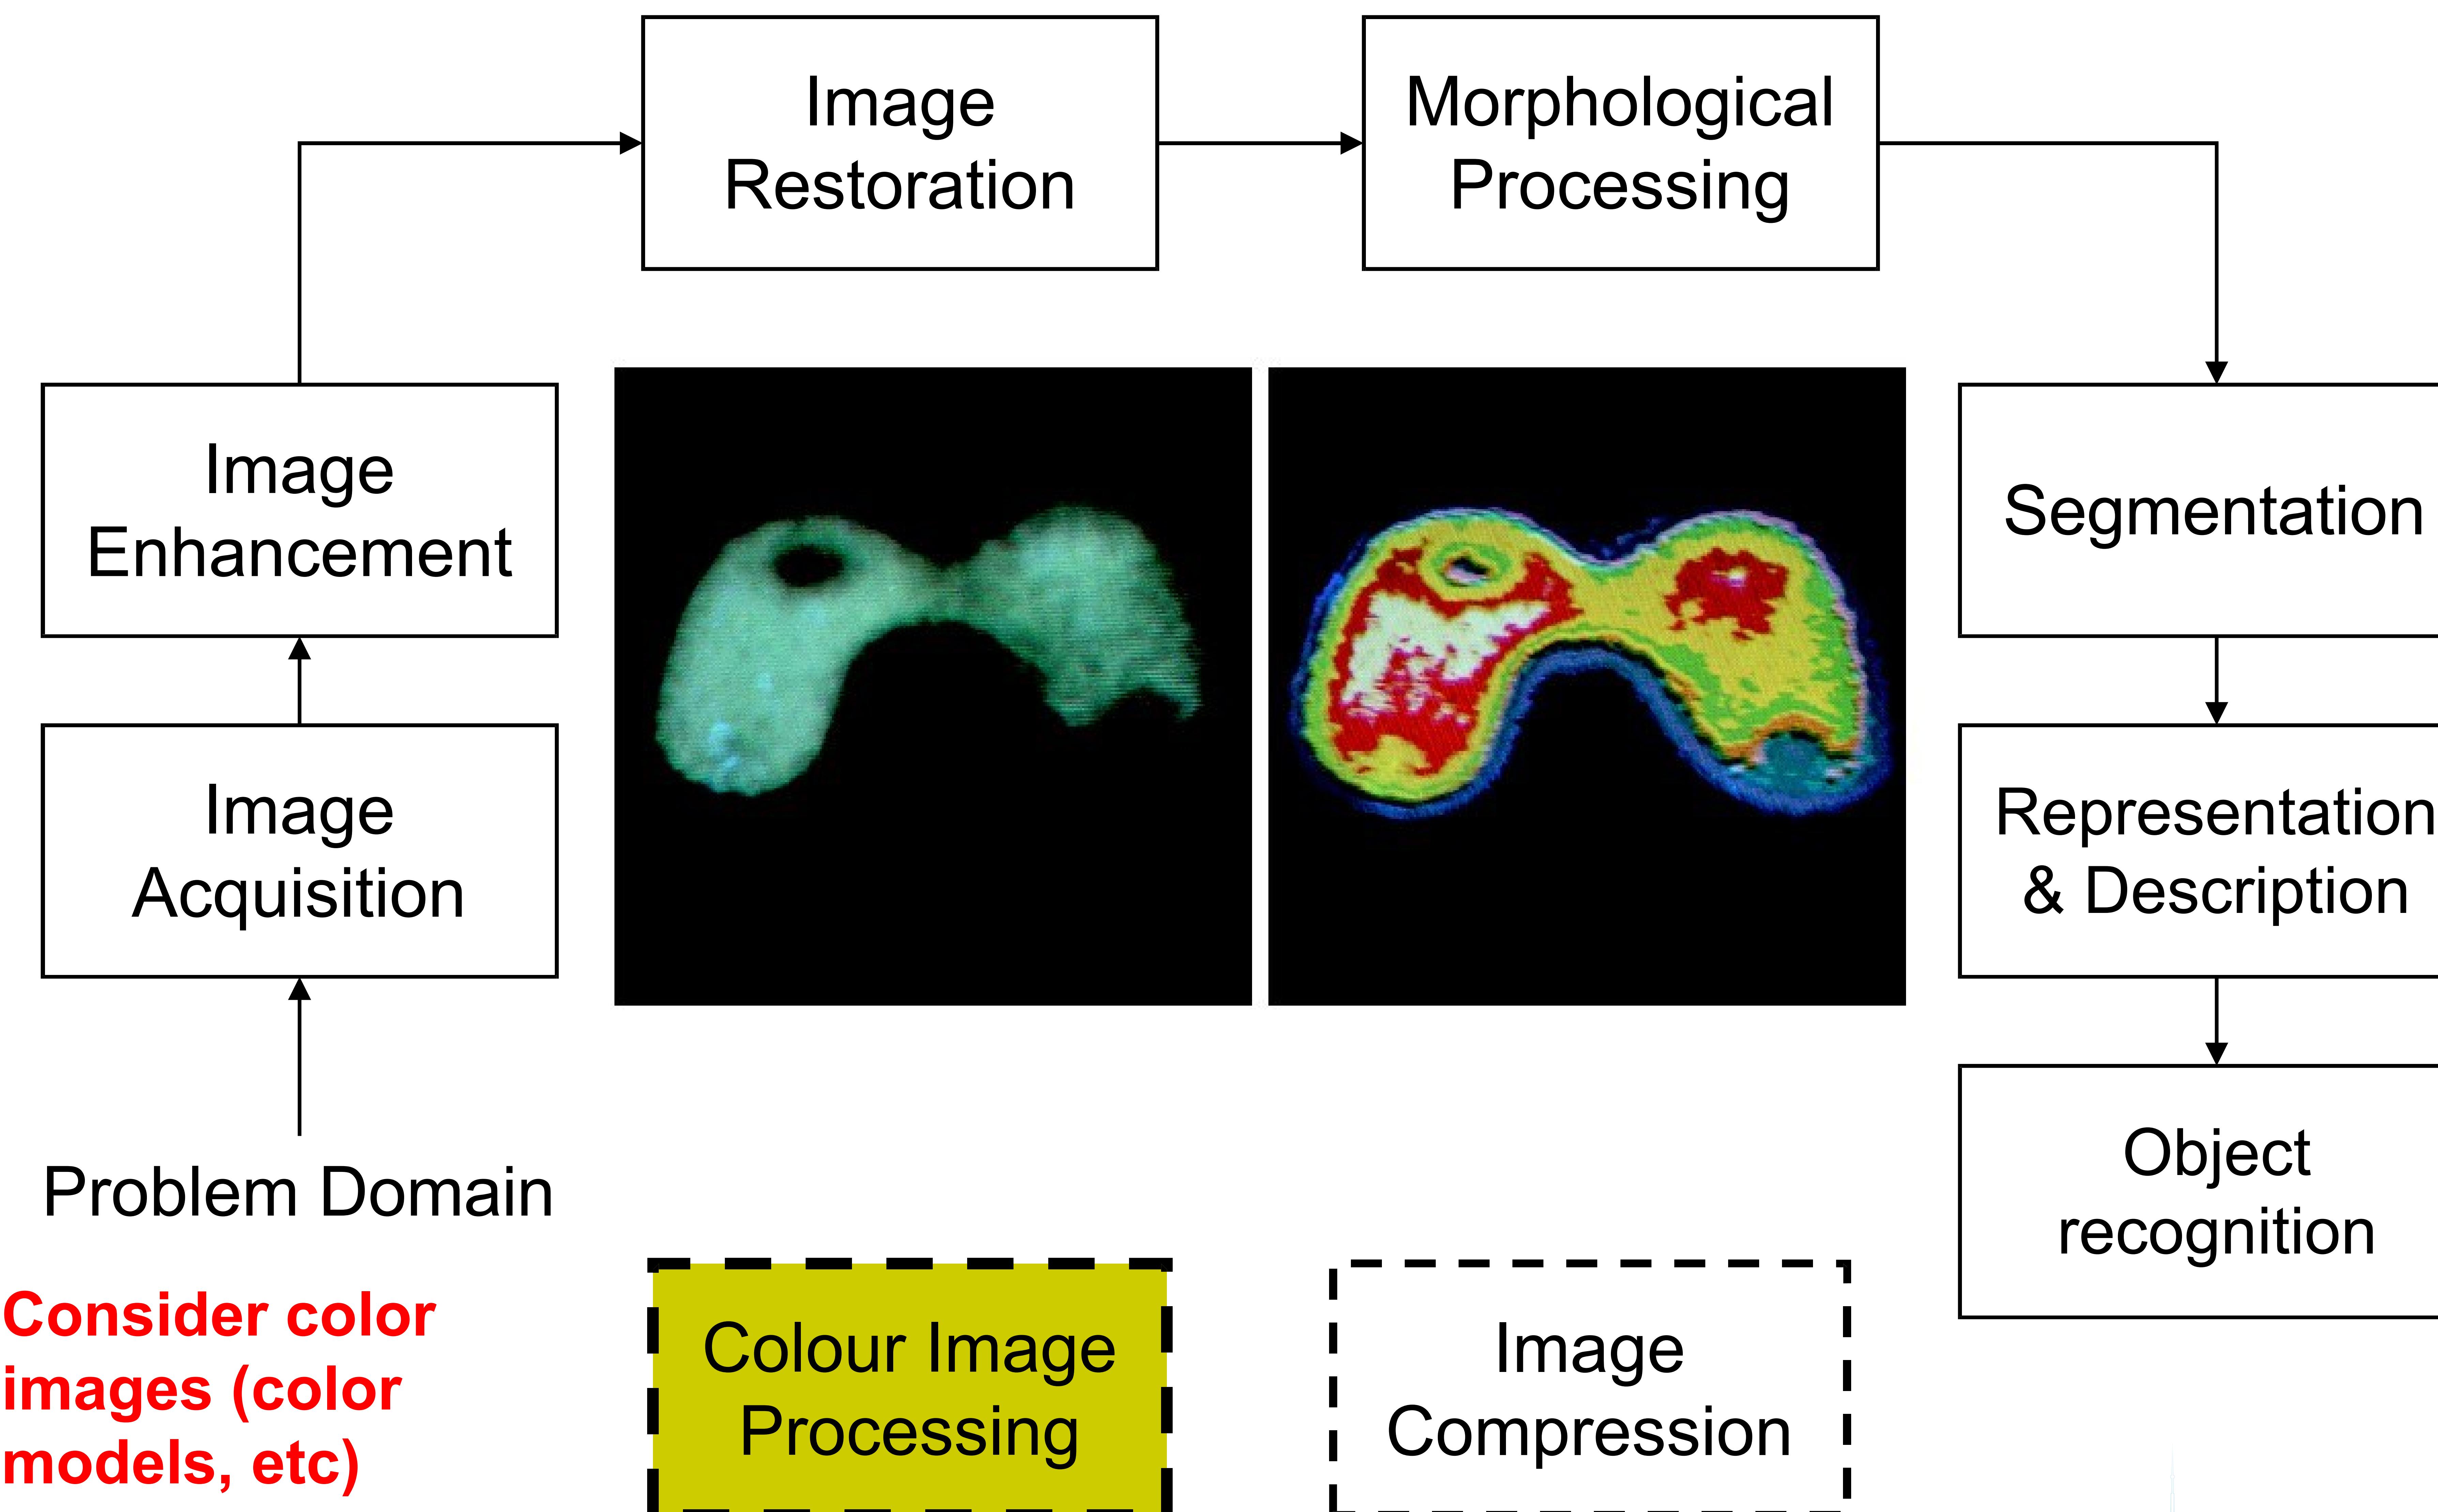
Key stages in digital image processing



Key stages in digital image processing

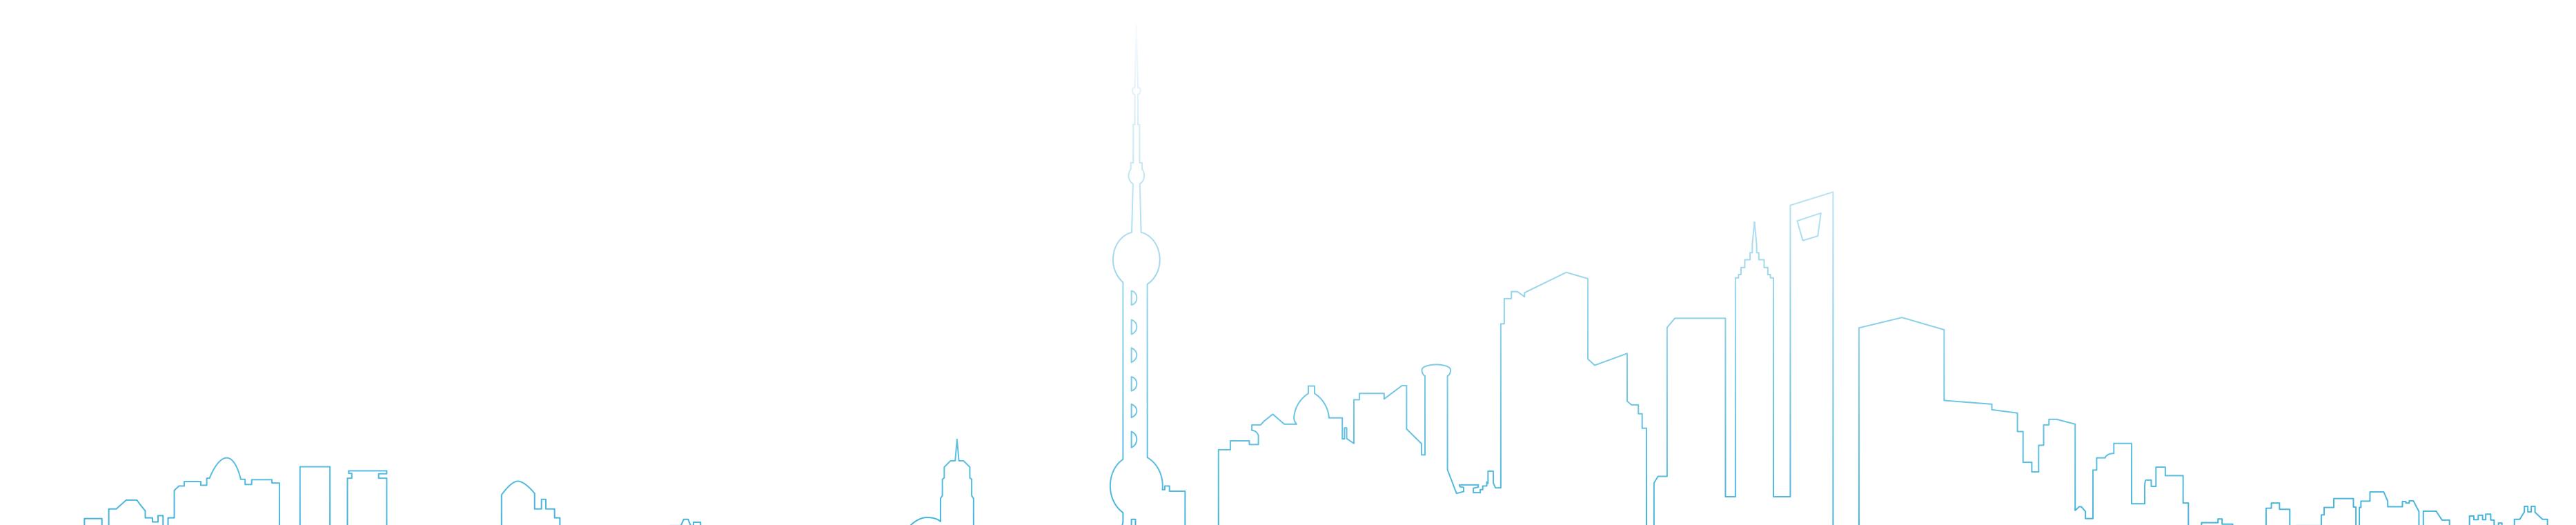


Key stages in digital image processing



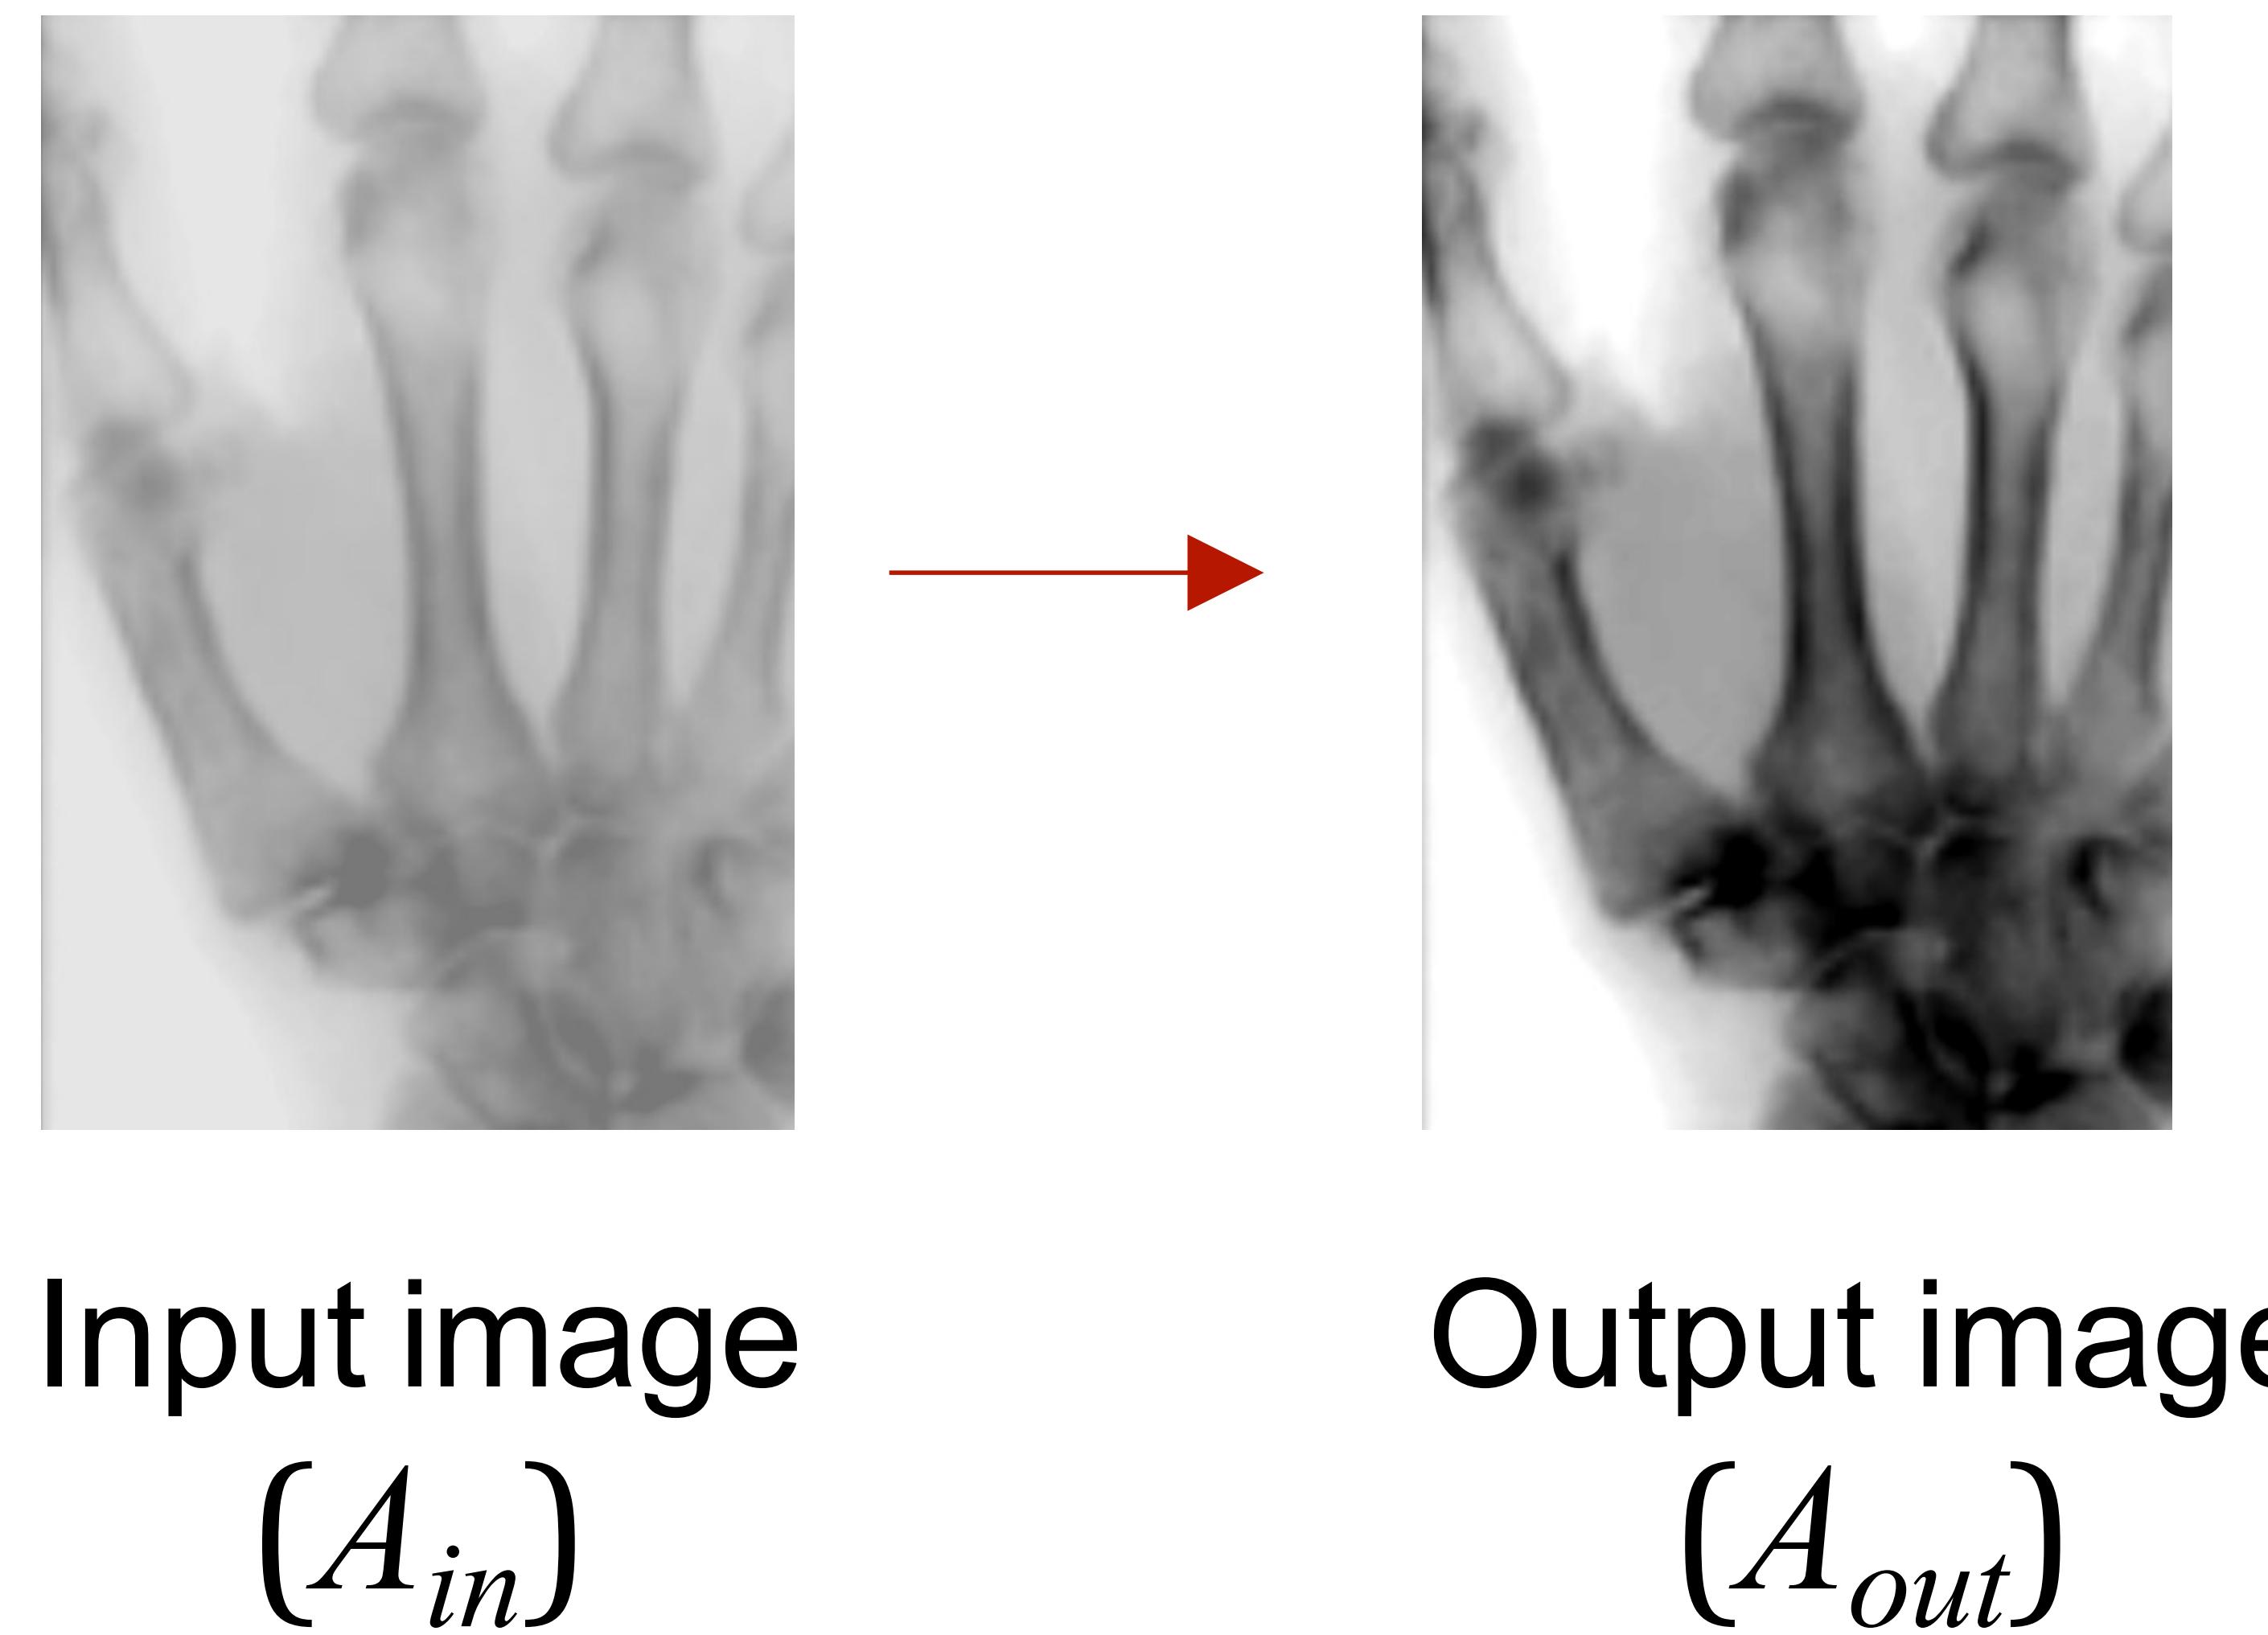
Outline: Image Data Mining

1. Introduction
2. Digital Image Processing
3. Convolutional Neural Network
4. Computer Vision Tasks



Pixel-level processing

Processing



Function $A_{out} = f(A_{in})$

Pixel Level

90	3	241	65	213
242	143	122	71	8
203	103	45	167	240
29	239	2	93	218
243	242	183	54	123

Special type of function

$$A_{out}[i,j] = f(A_{in}[i,j])$$

Pixel-level processing

- **Linear functions**
- **Nonlinear functions**

Pixel Level

90	3	241	65	213
242	143	122	71	8
203	103	45	167	240
29	239	2	93	218
243	242	183	54	123

Linear functions

Build functions in Numpy

Maths

$$= \begin{array}{|c|c|} \hline 20 & 22 \\ \hline 1 & 25 \\ \hline \end{array}$$

Nested list implementation

```
I_list = [[20,22],[1,25]]
```

$$= + 2$$

```
for row in range(2):
    for col in range(2):
        I_list[row][col] += 2
```

Submatrix:
Row[0-3], Col[5-13]

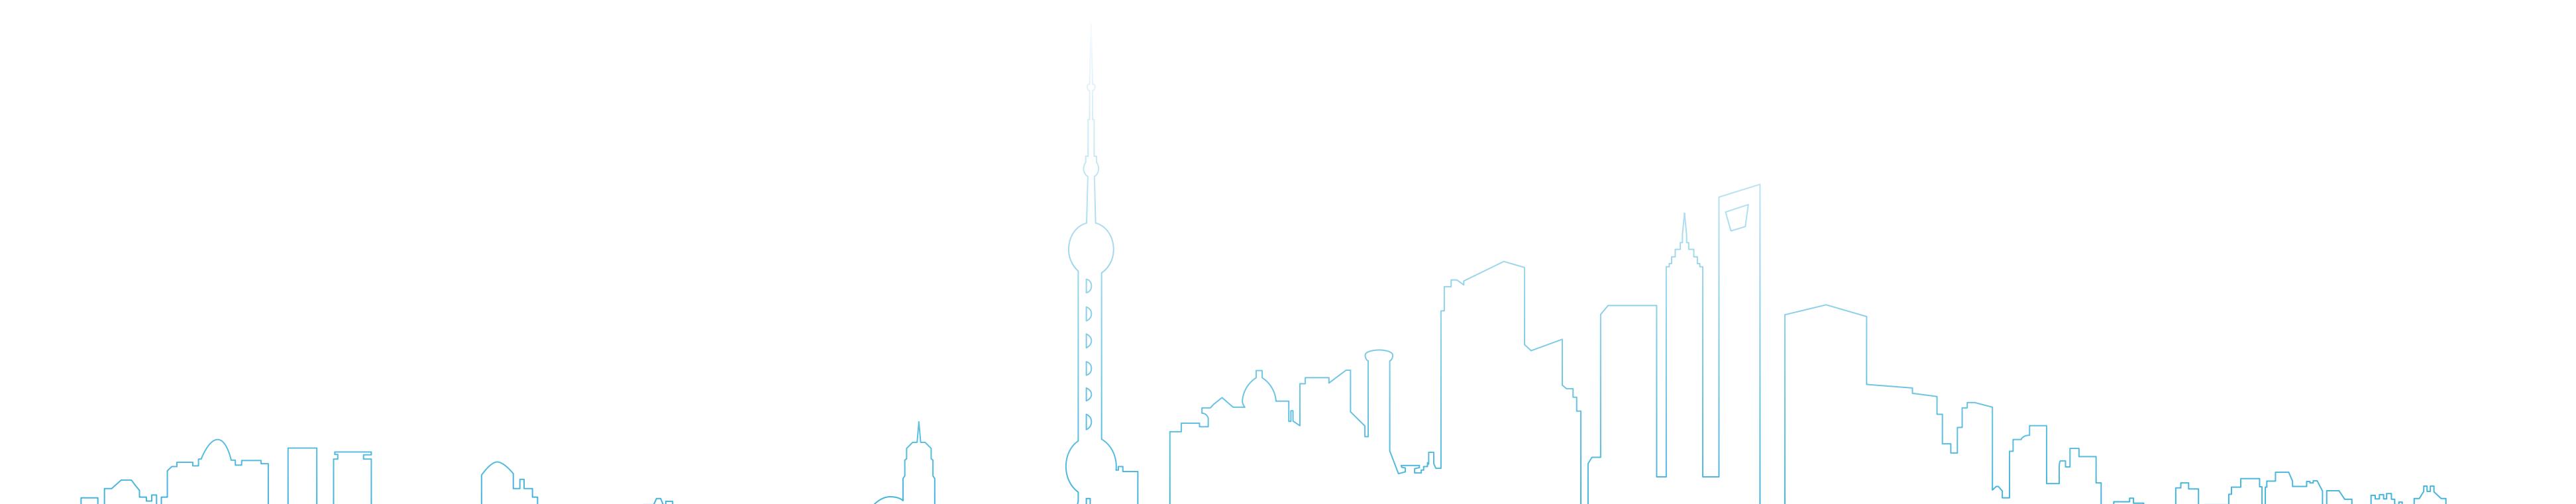
...

Numpy Implementation

```
import numpy as np
I_np = np.array([[20,22],[1,25]])
```

```
I_np += 2
```

```
I_np[0:3, 5:13]
```



Brightness adjustment: $f=x+b$

243	222	171
203	145	137
209	132	193

123	102	51
83	25	17
89	12	73

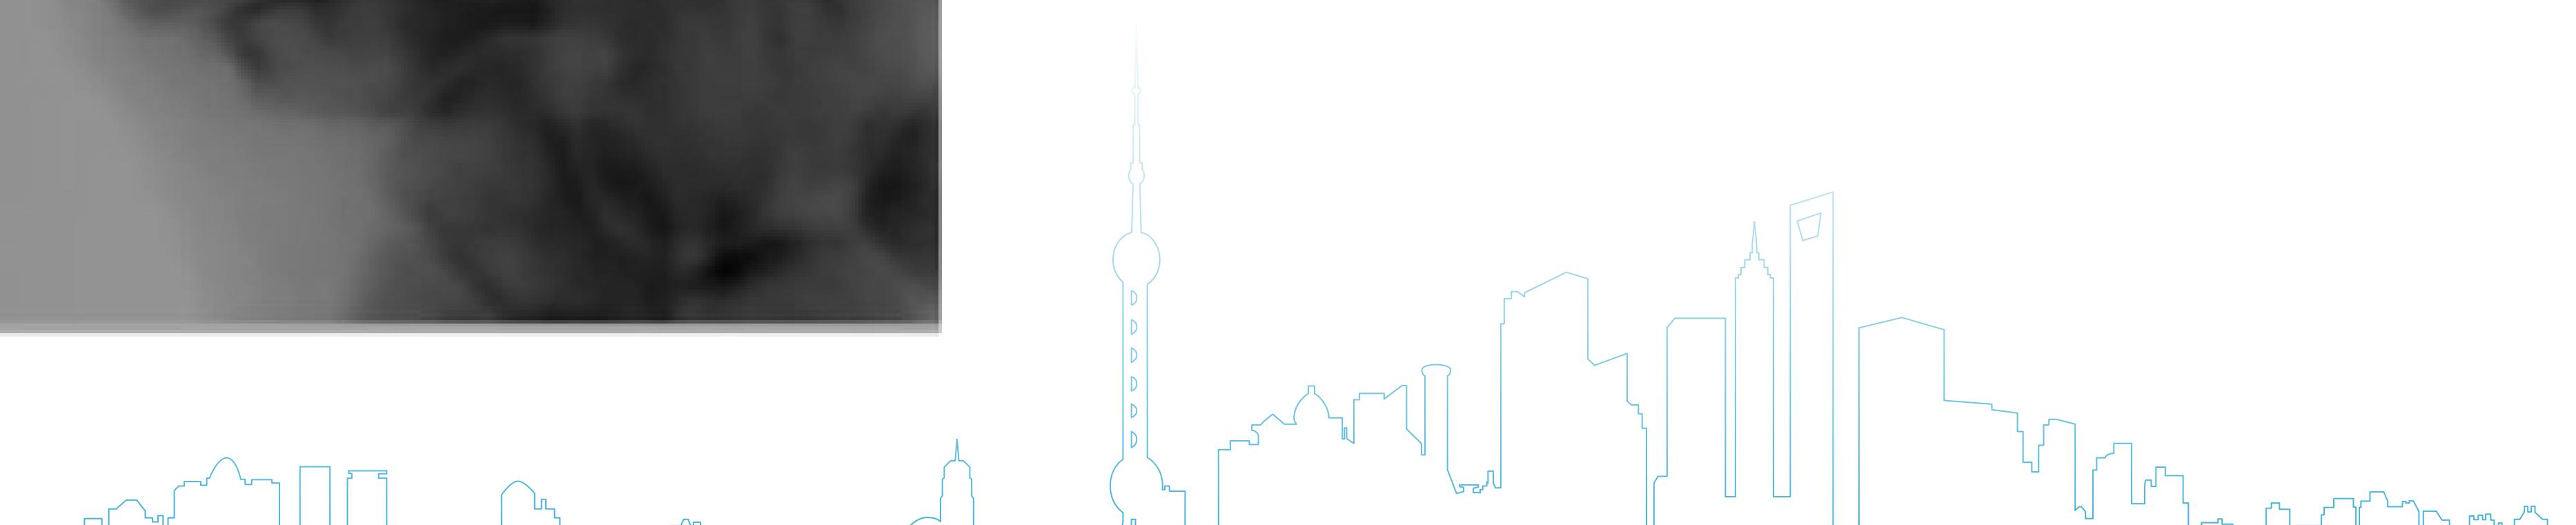
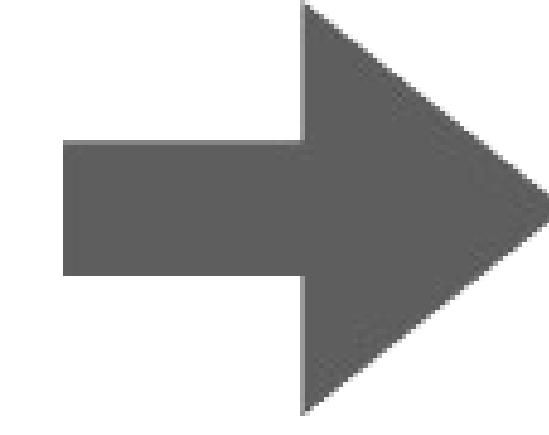
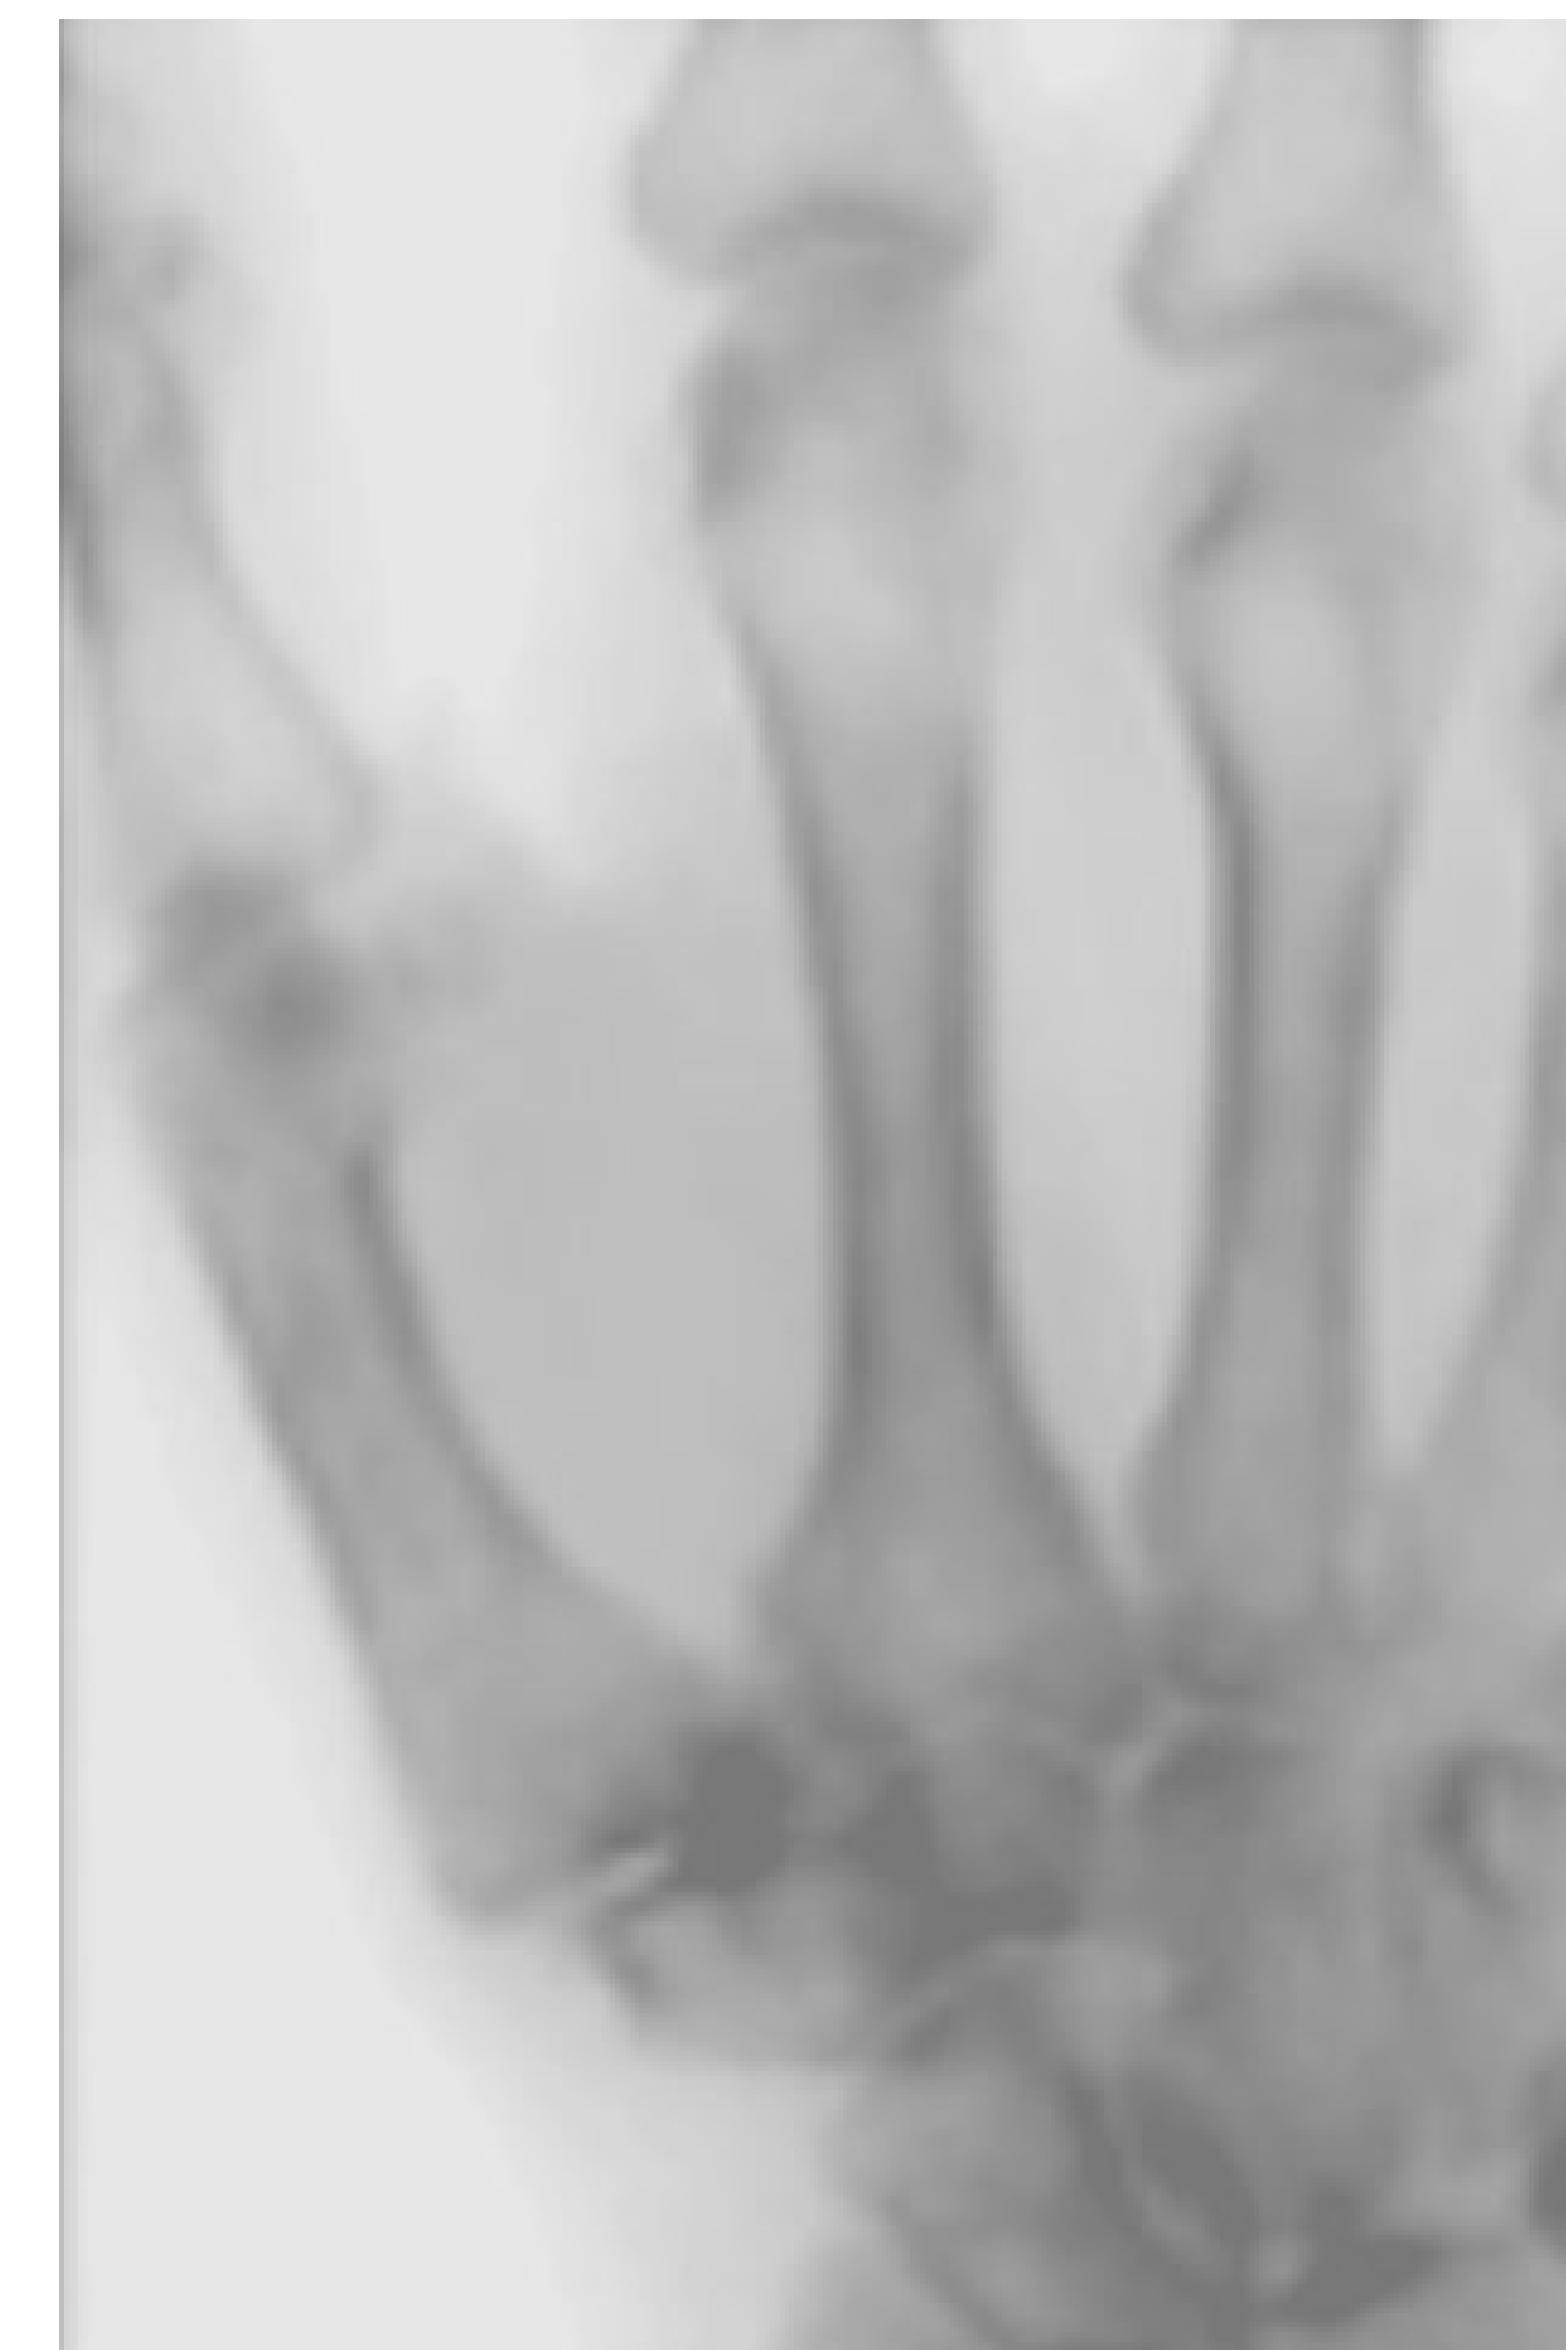
$$p \quad \rightarrow \quad p - 120$$

Input pixel value

$$p - 120$$

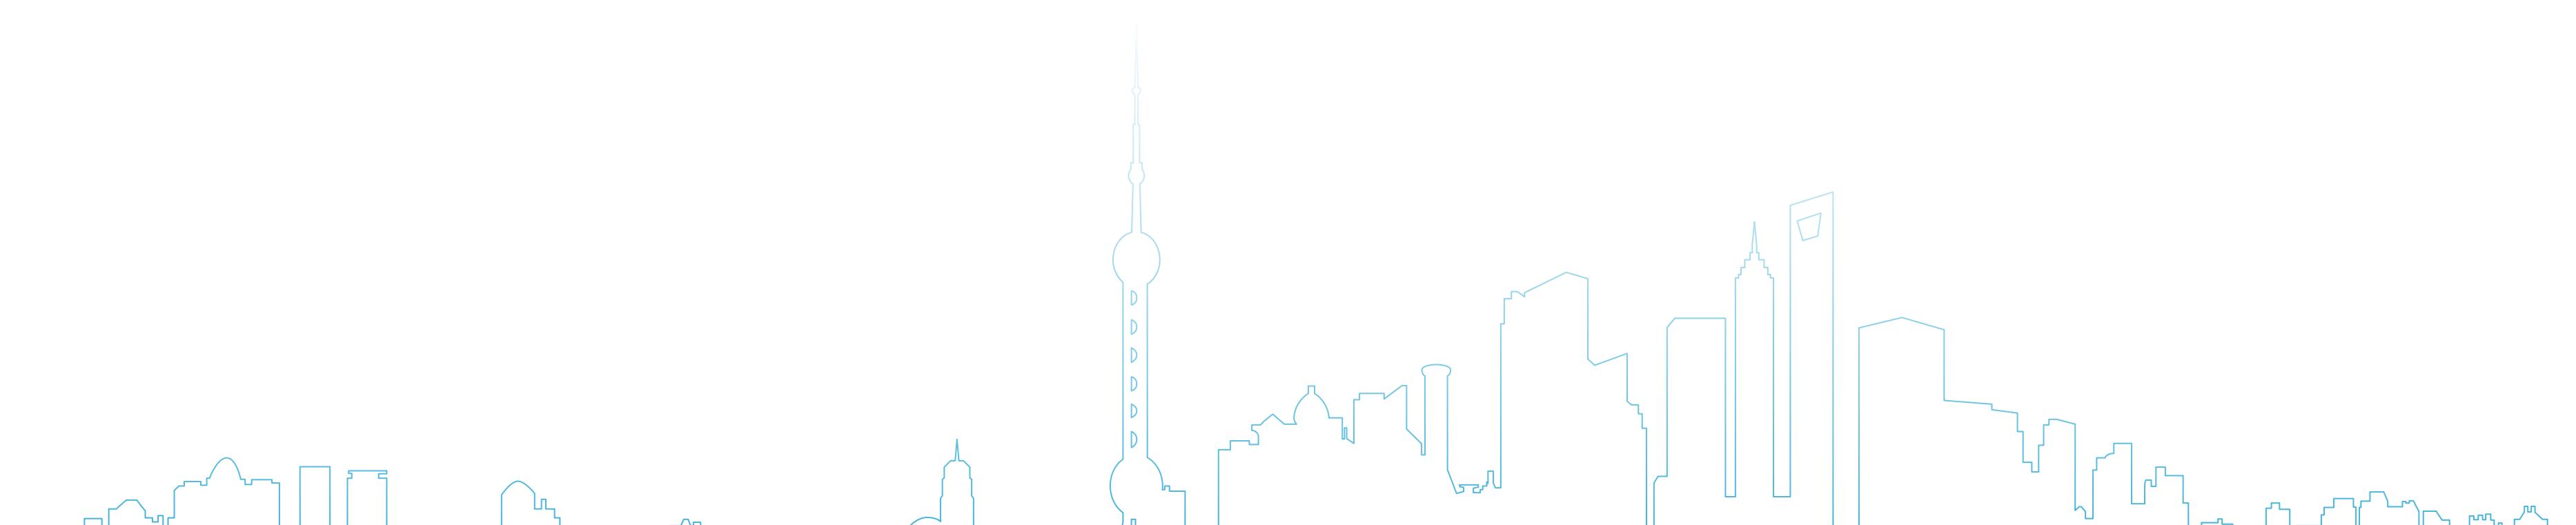
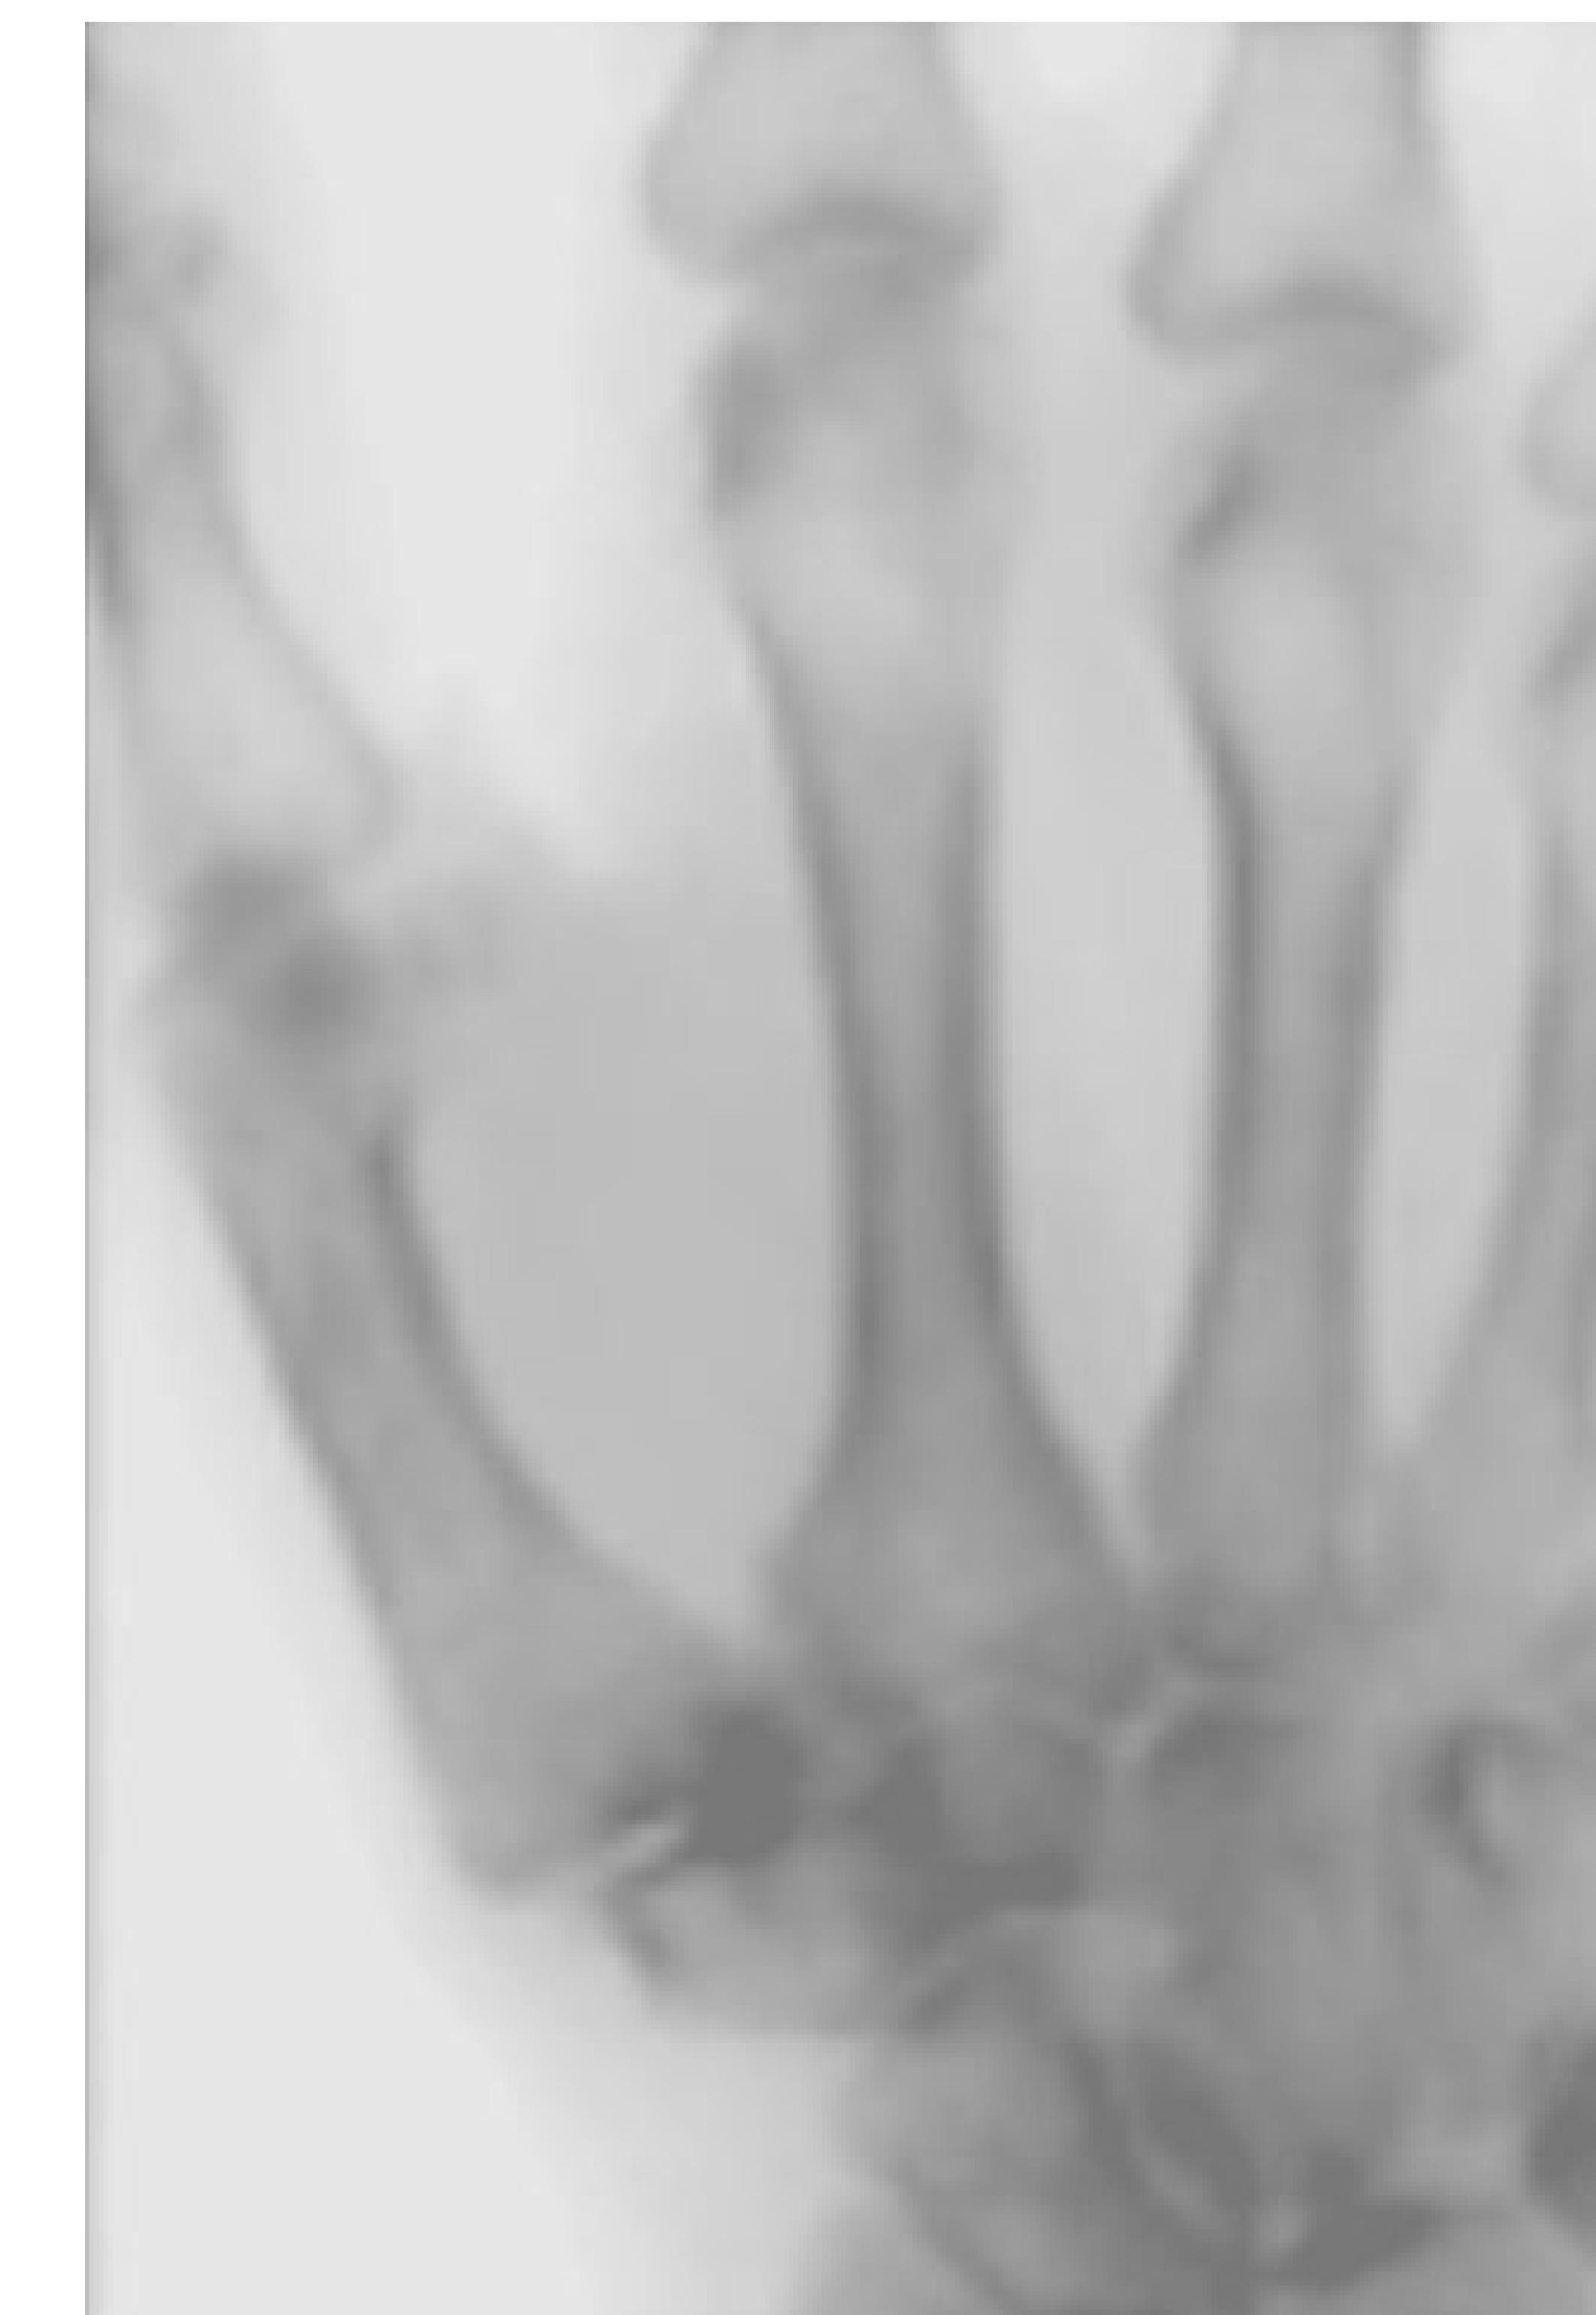
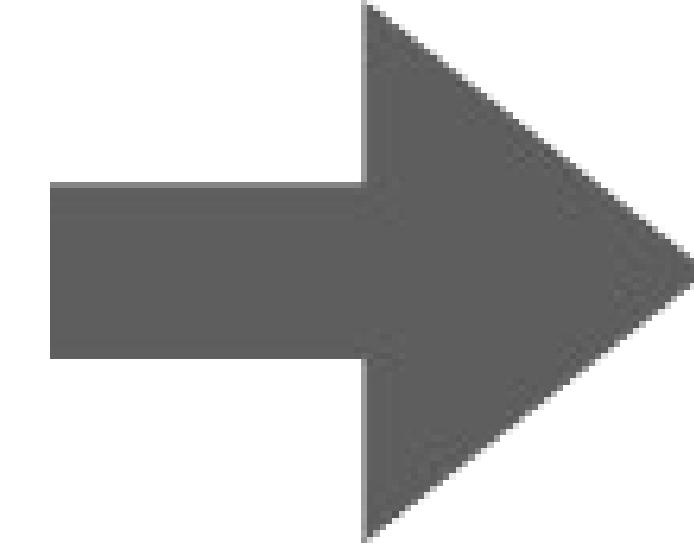
Output pixel value

```
I = imread('image.png')  
I = I - 120
```



Contrast adjustment

High contrast: $(I_{max} - I_{min})$ is bigger



Contrast adjustment: $f(x) = ax$

Scale pixel value (image range is different)

121	111	85
101	72	68
104	66	96

p



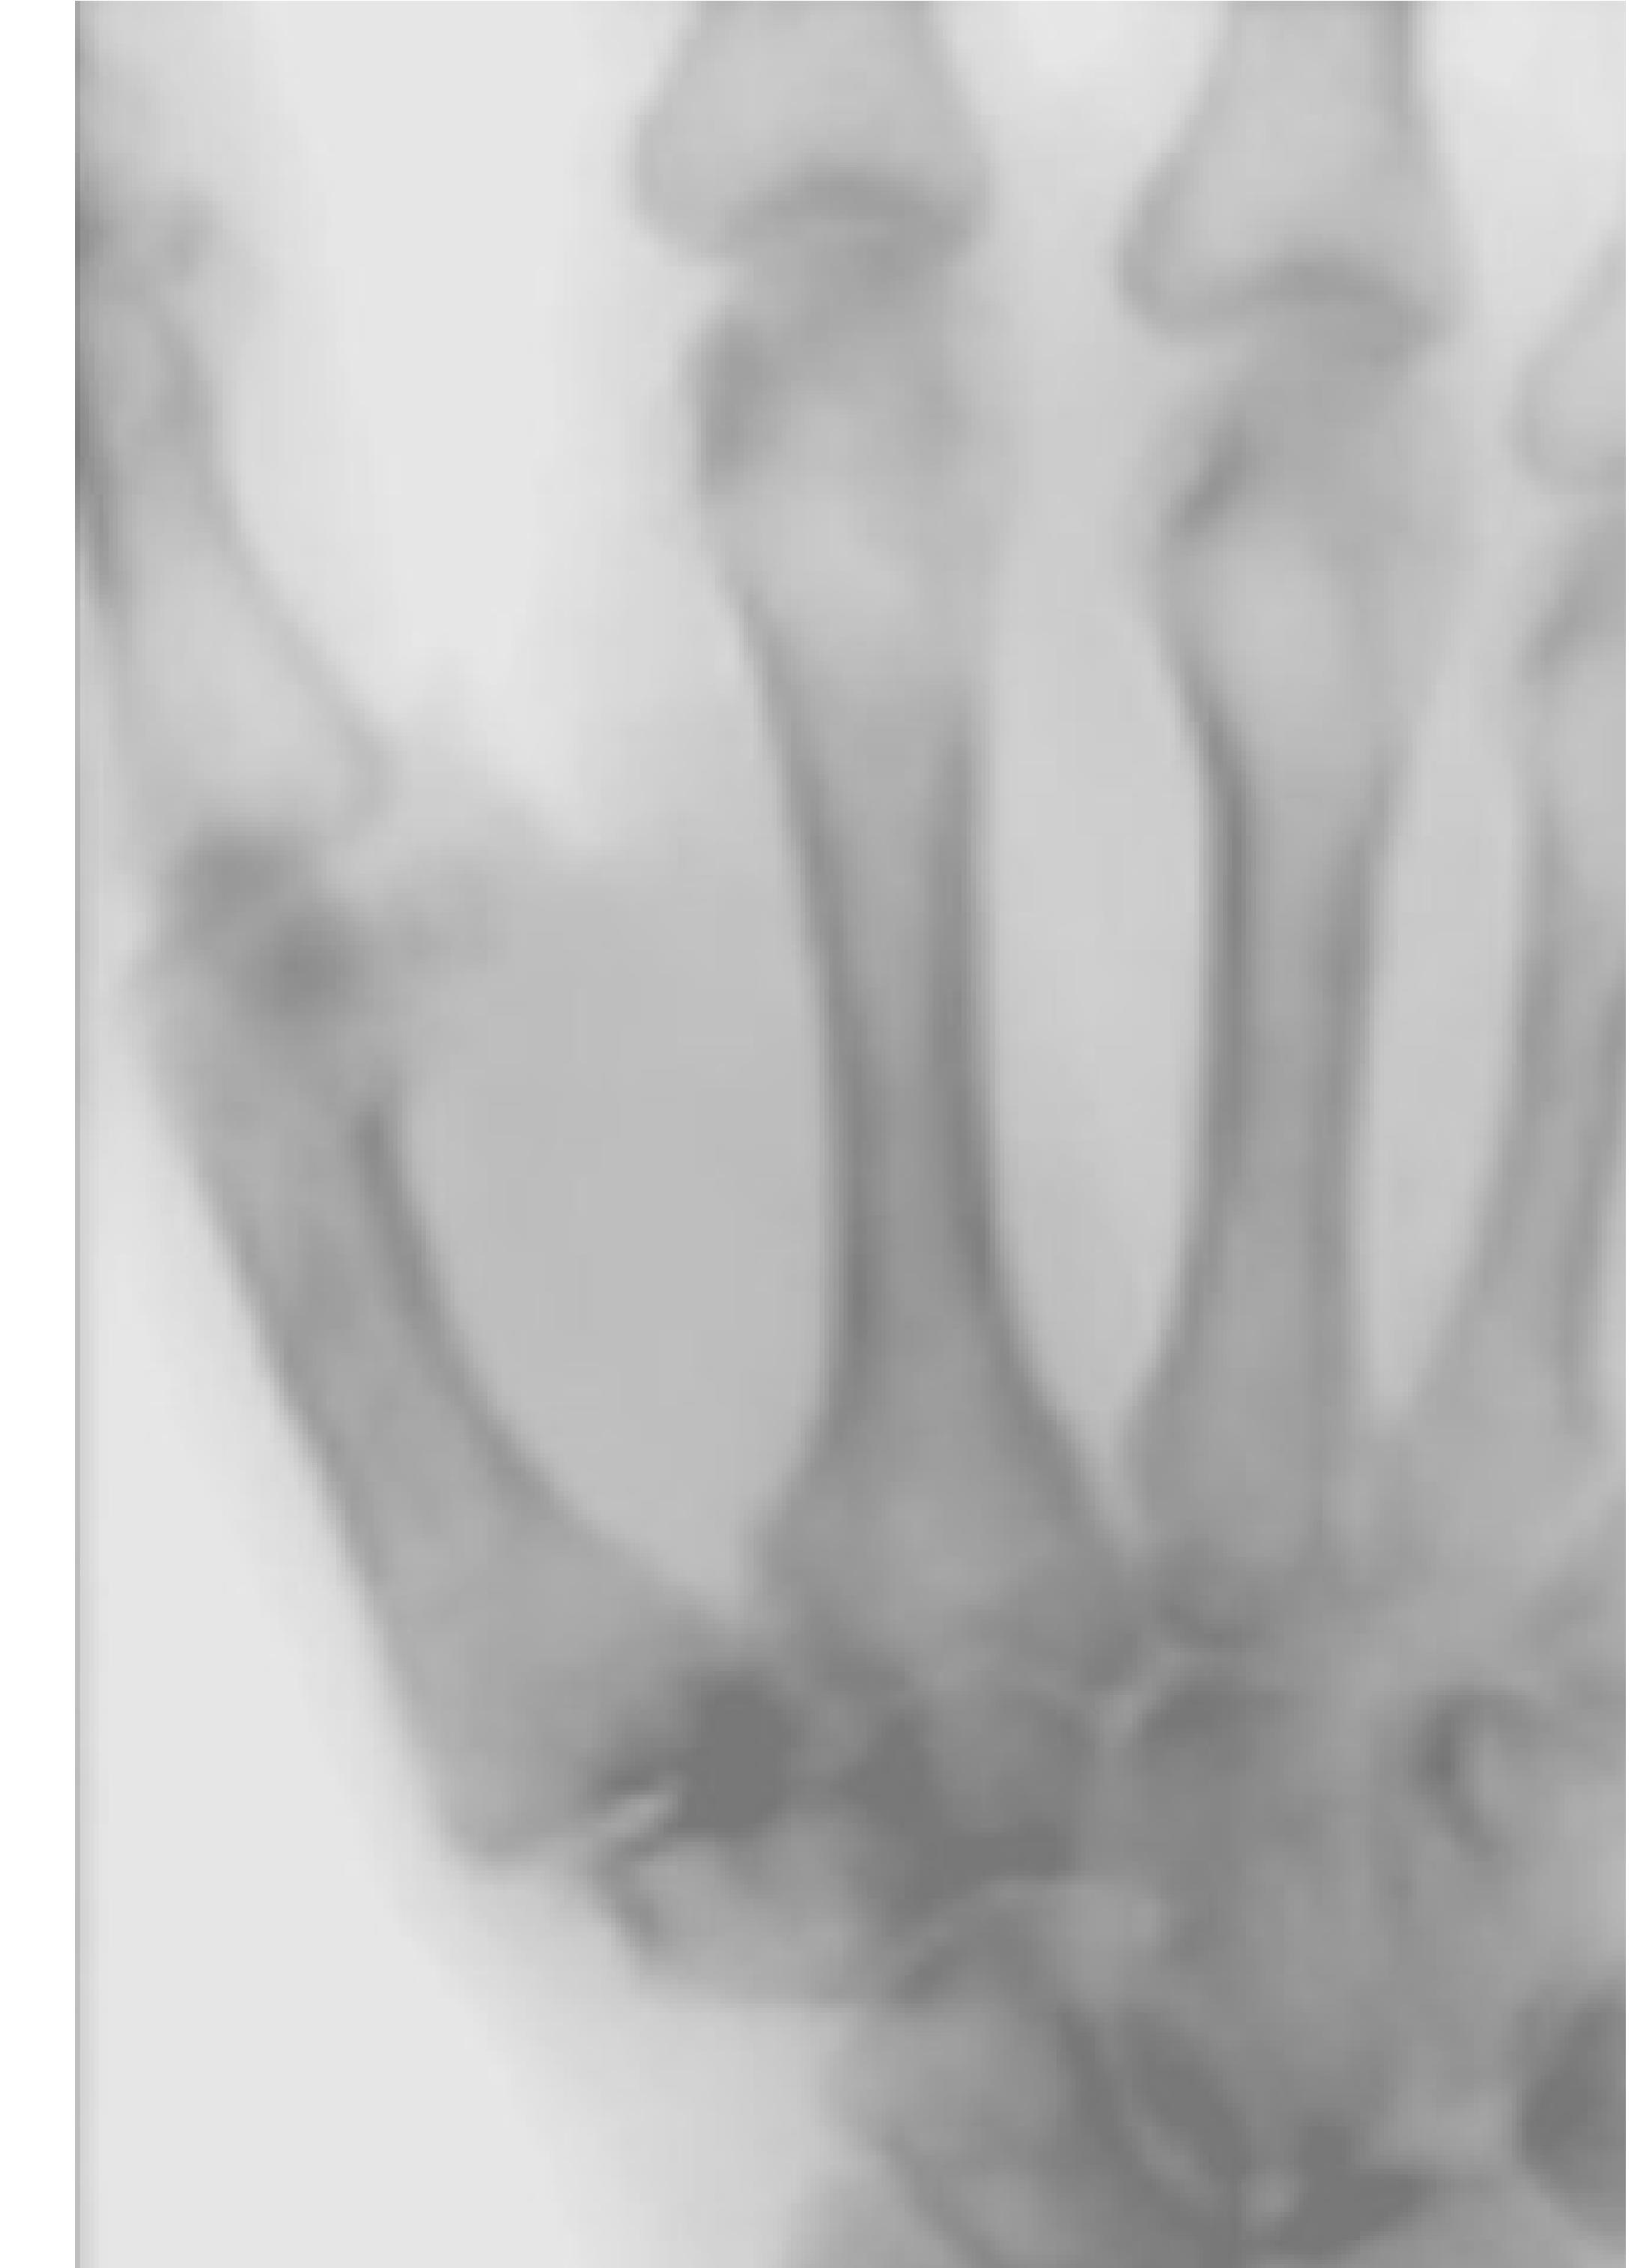
$p * 2$

Input pixel value

(max-min = 55)

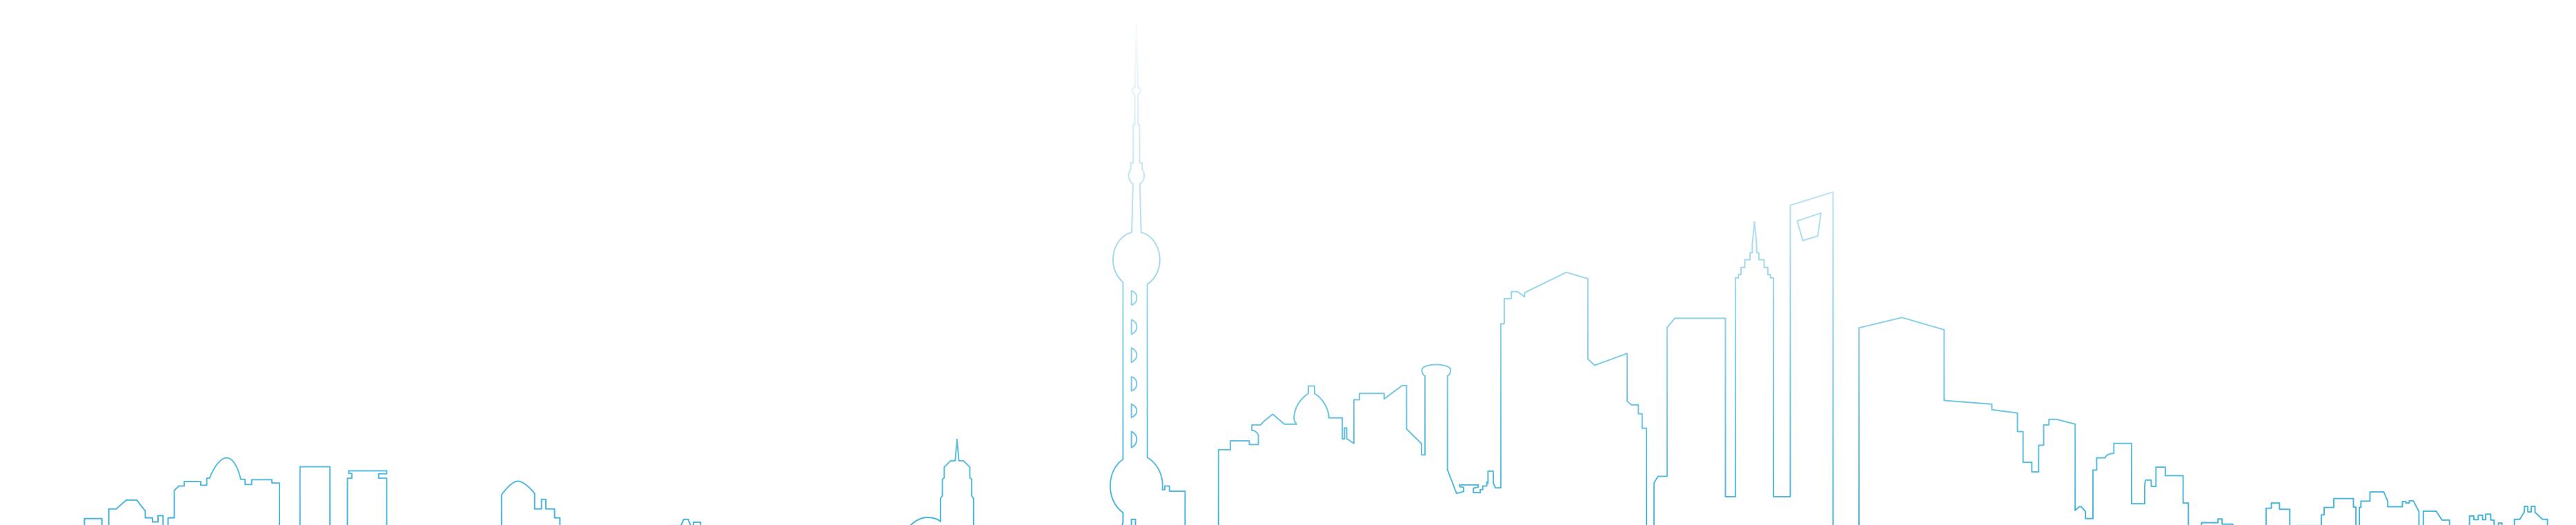
Output pixel value

(max-min = 110)

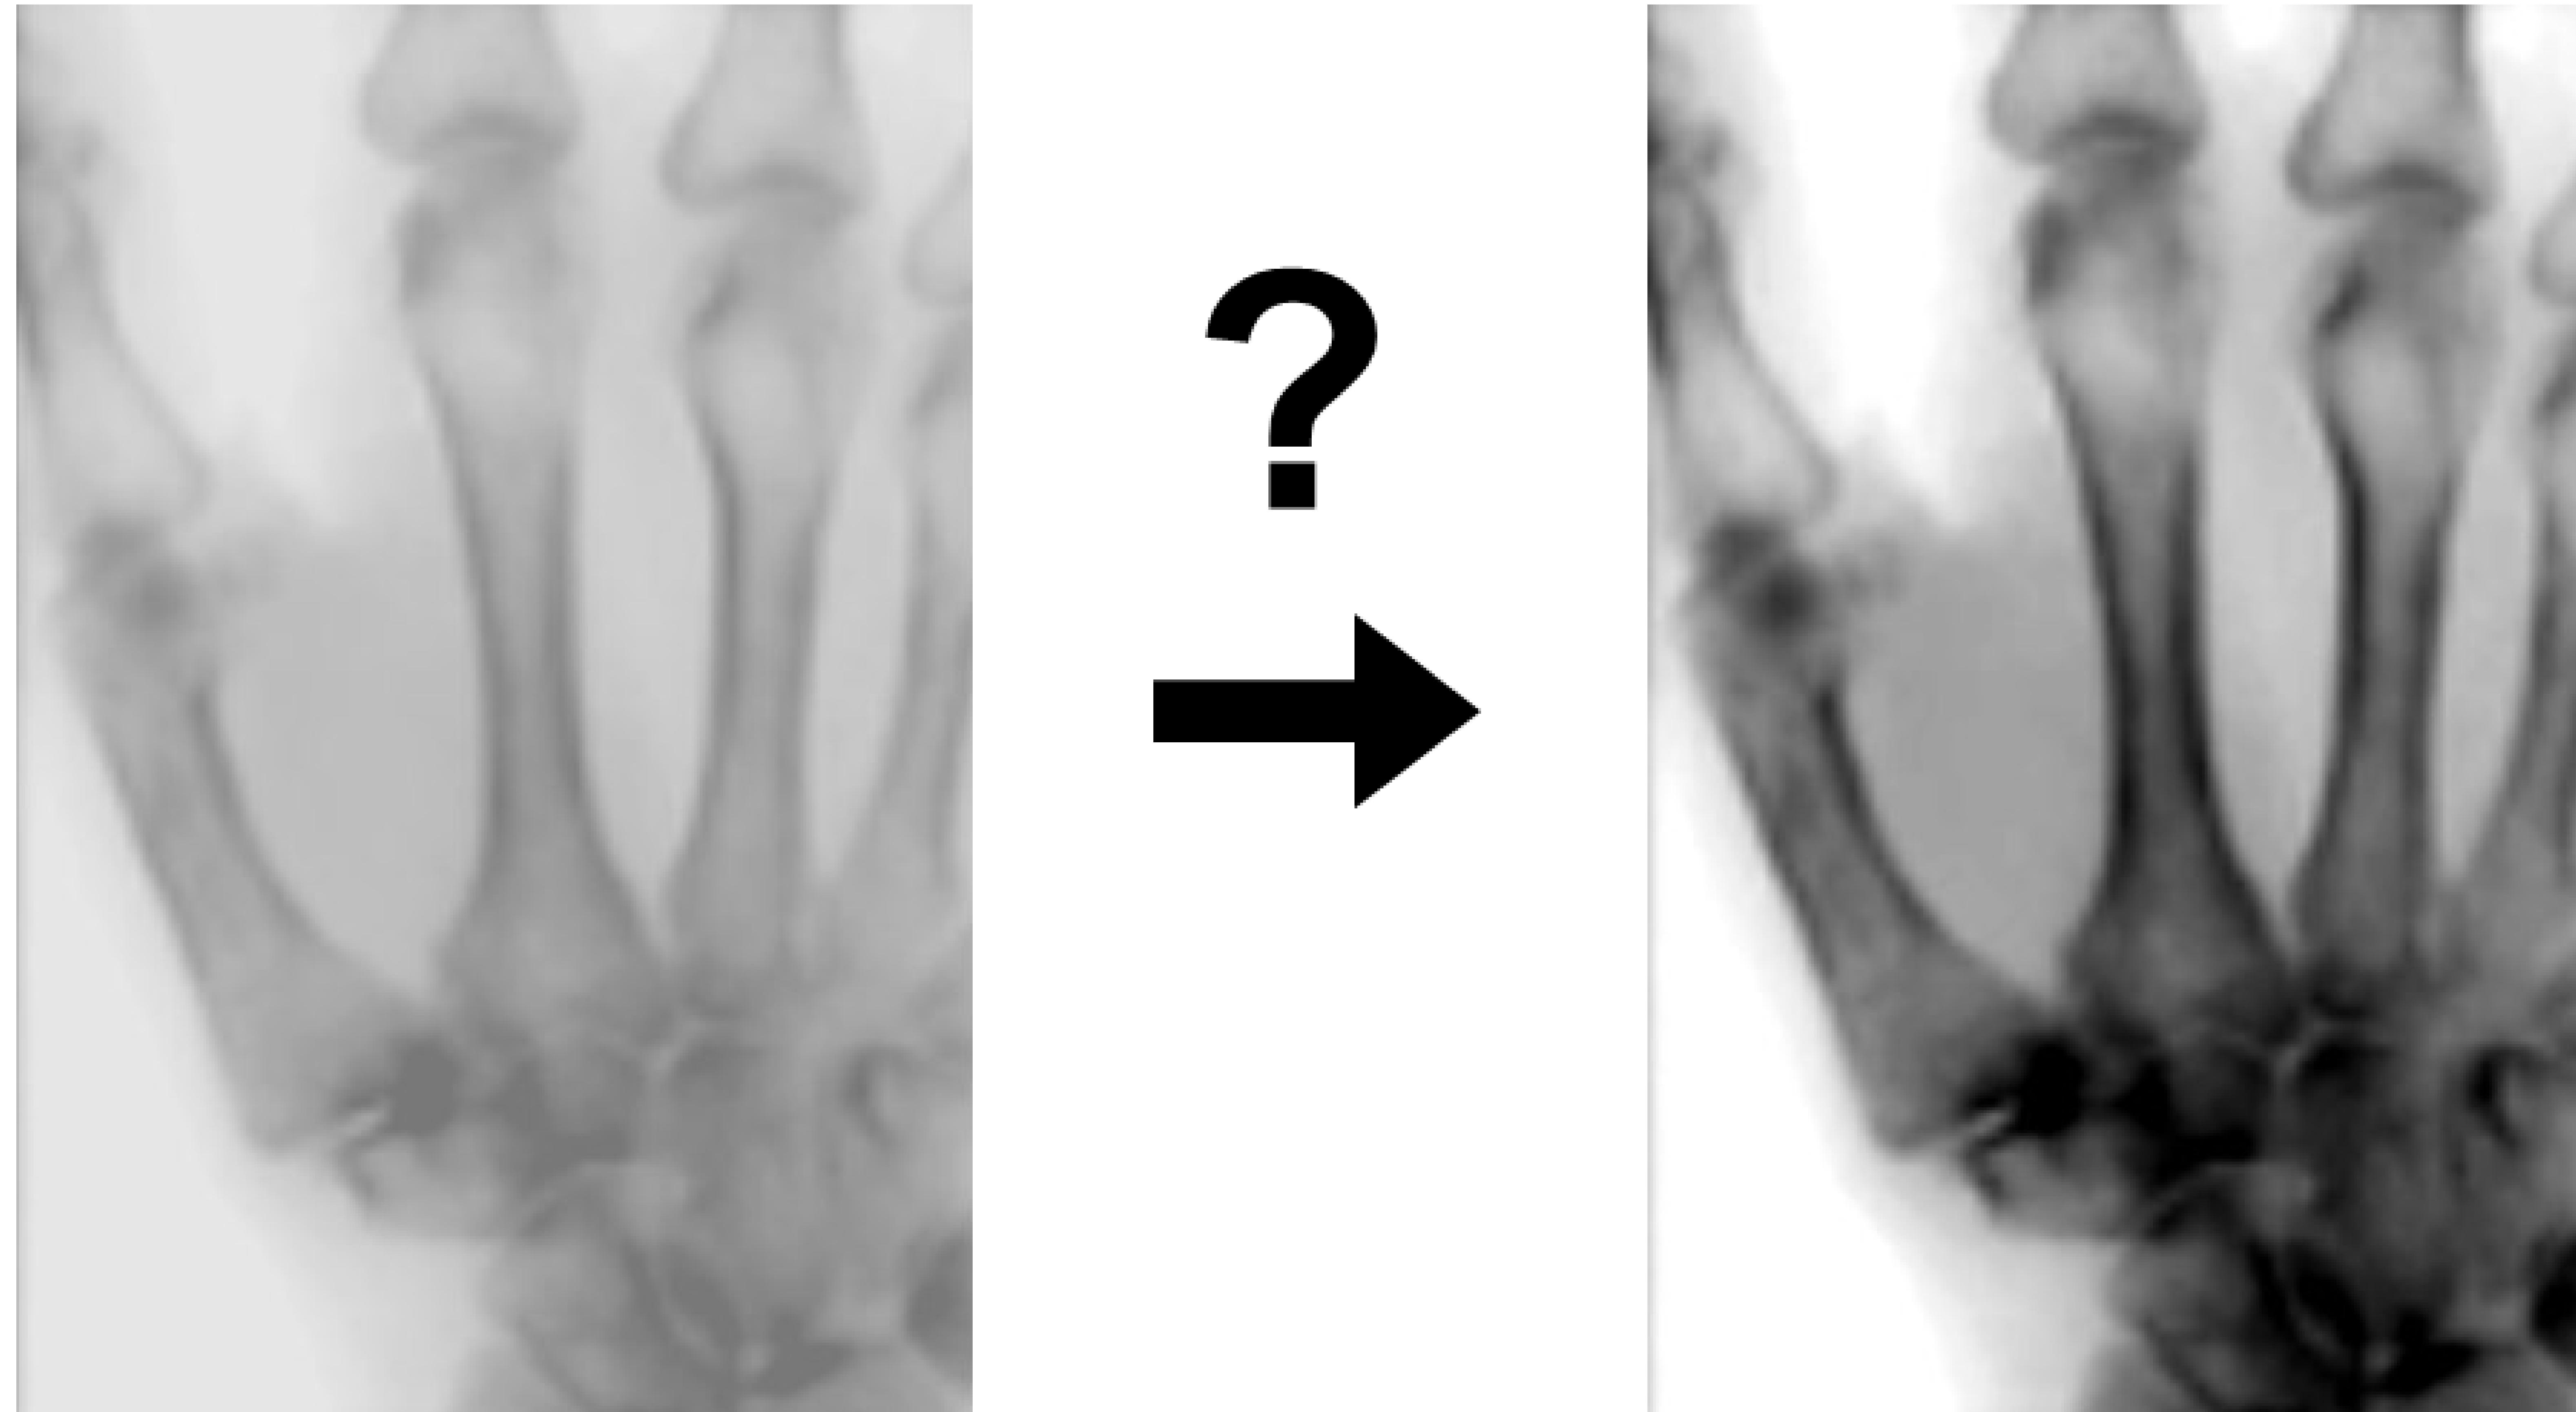


```
I = imread('image.png')
```

```
I = I * 2
```



Application: image auto-contrast



Auto-contrast is a technique used to enhance the contrast of an image automatically. It adjusts the intensity levels of the pixels to use the full range of available color values (usually 0 to 255 in an 8-bit image)

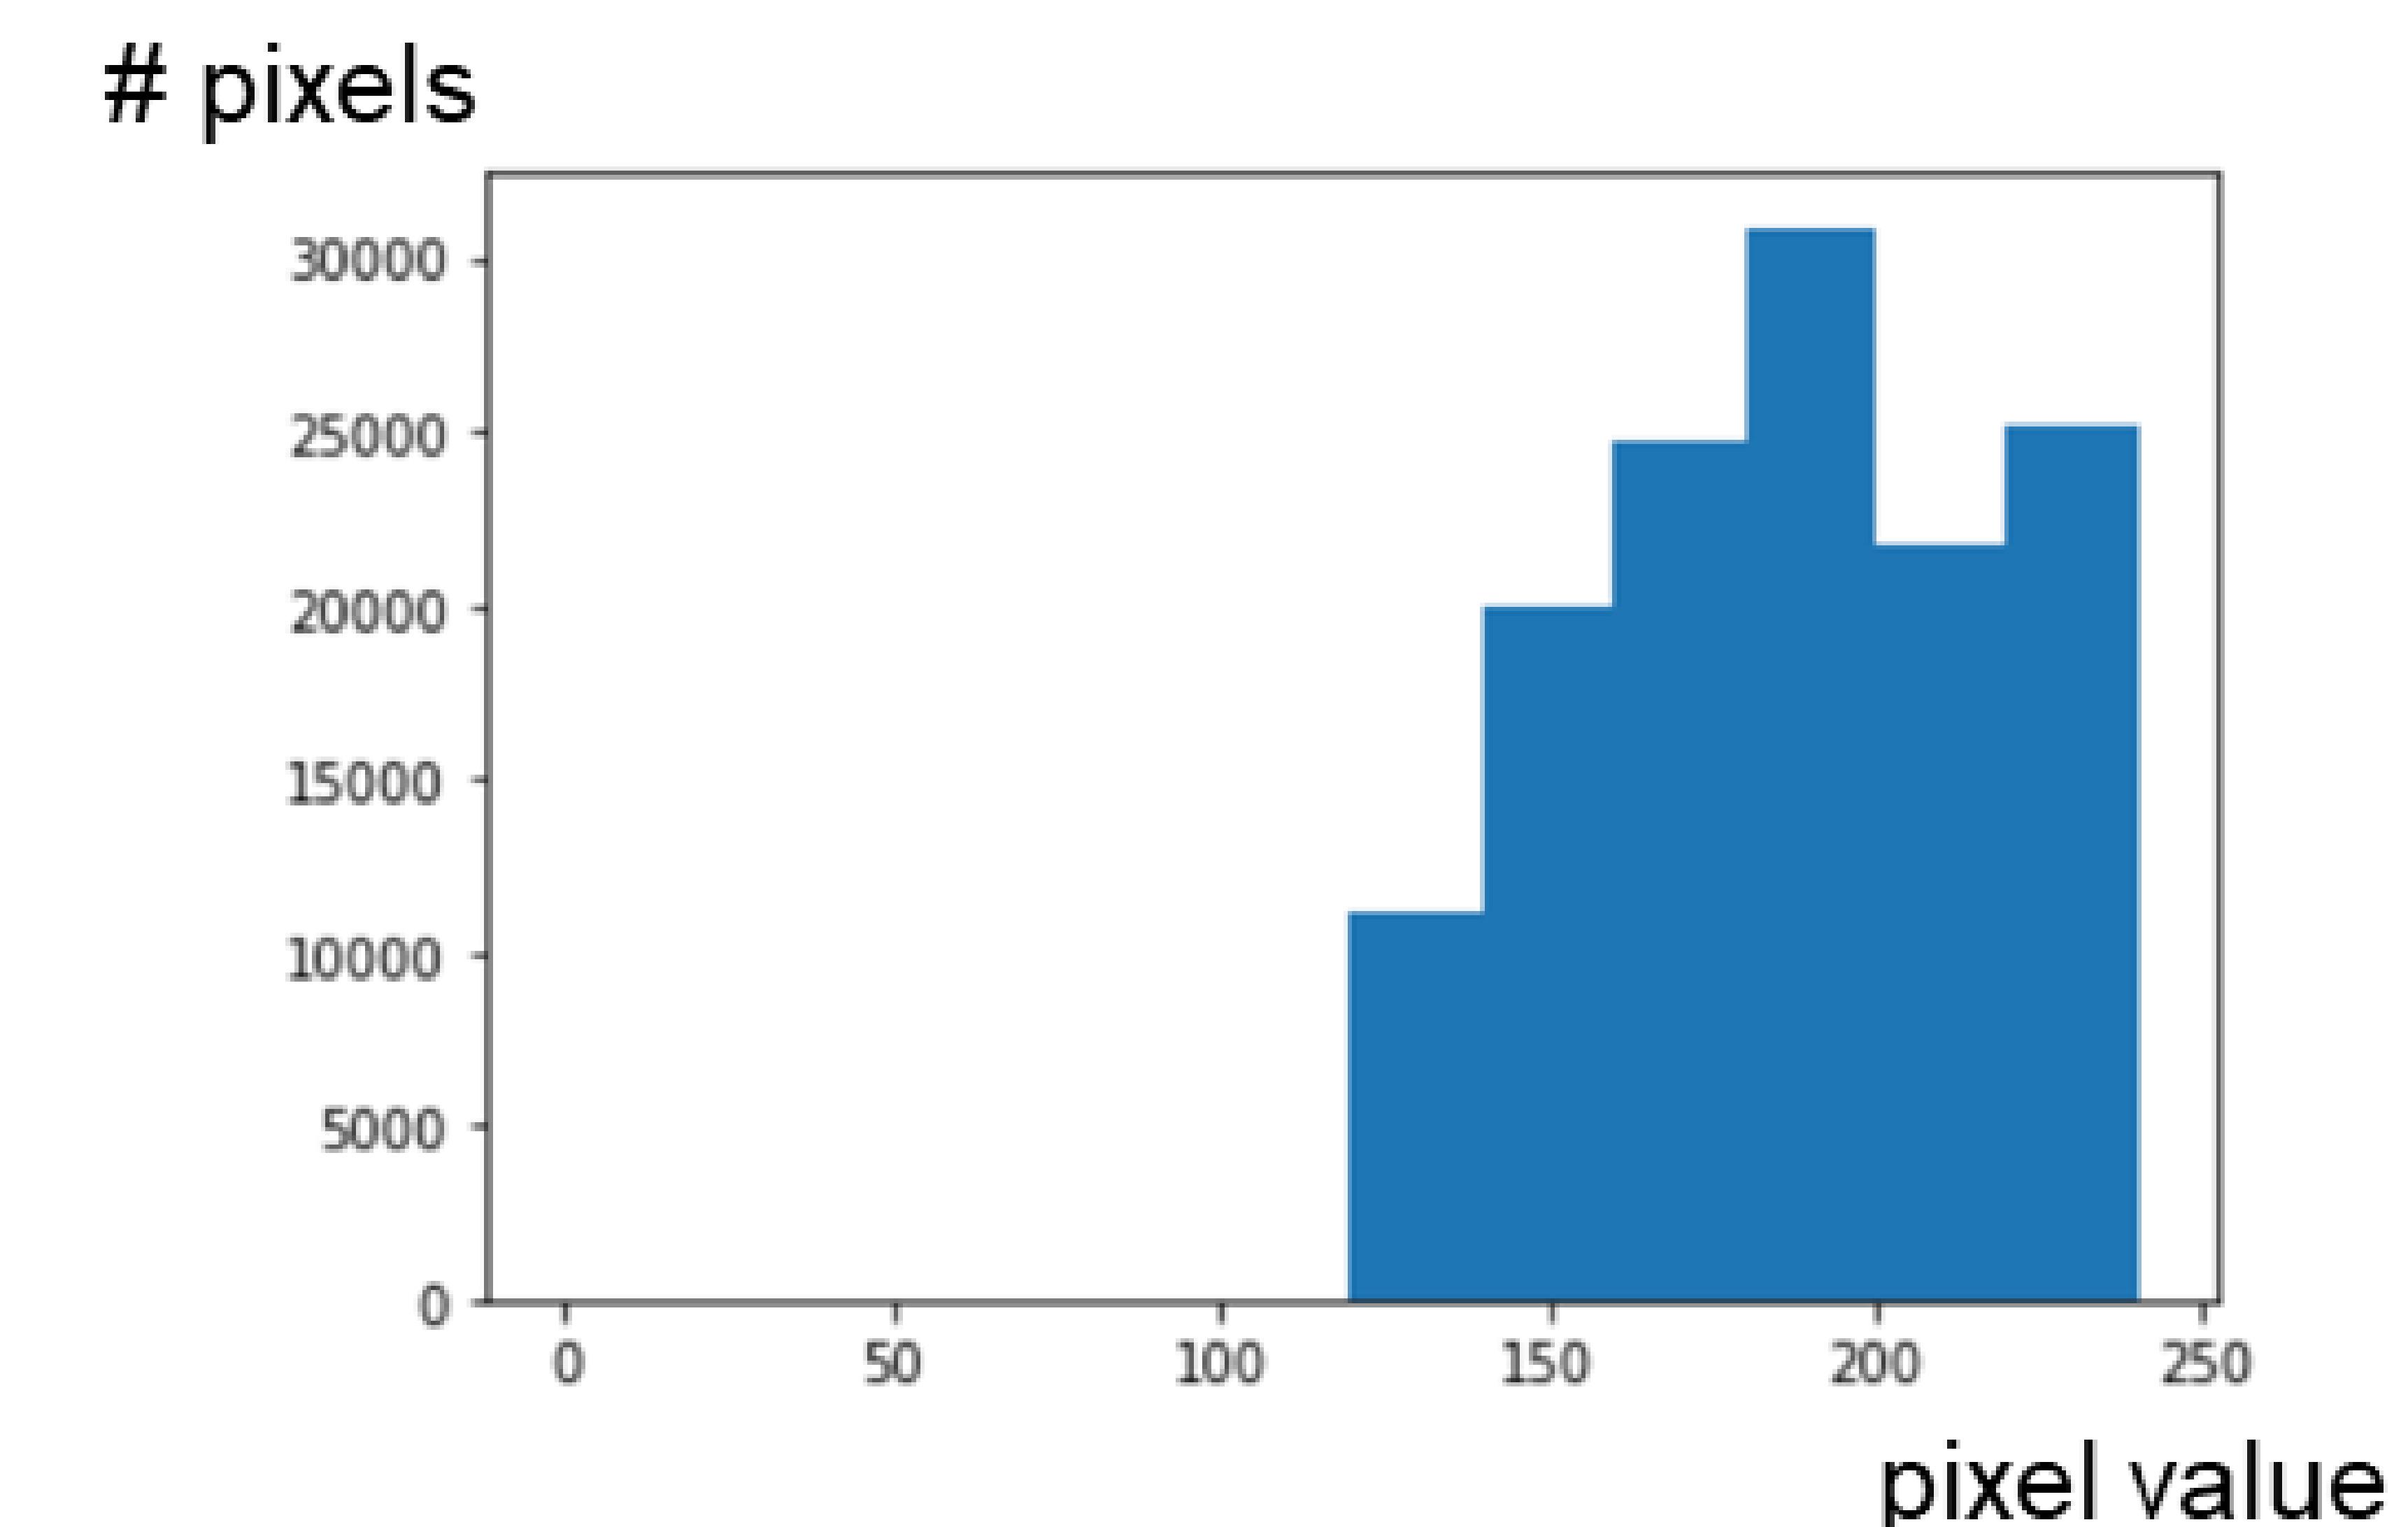
Application: image auto-contrast

Image: Statistical view



243	222	171
203	145	137
209	132	193

,



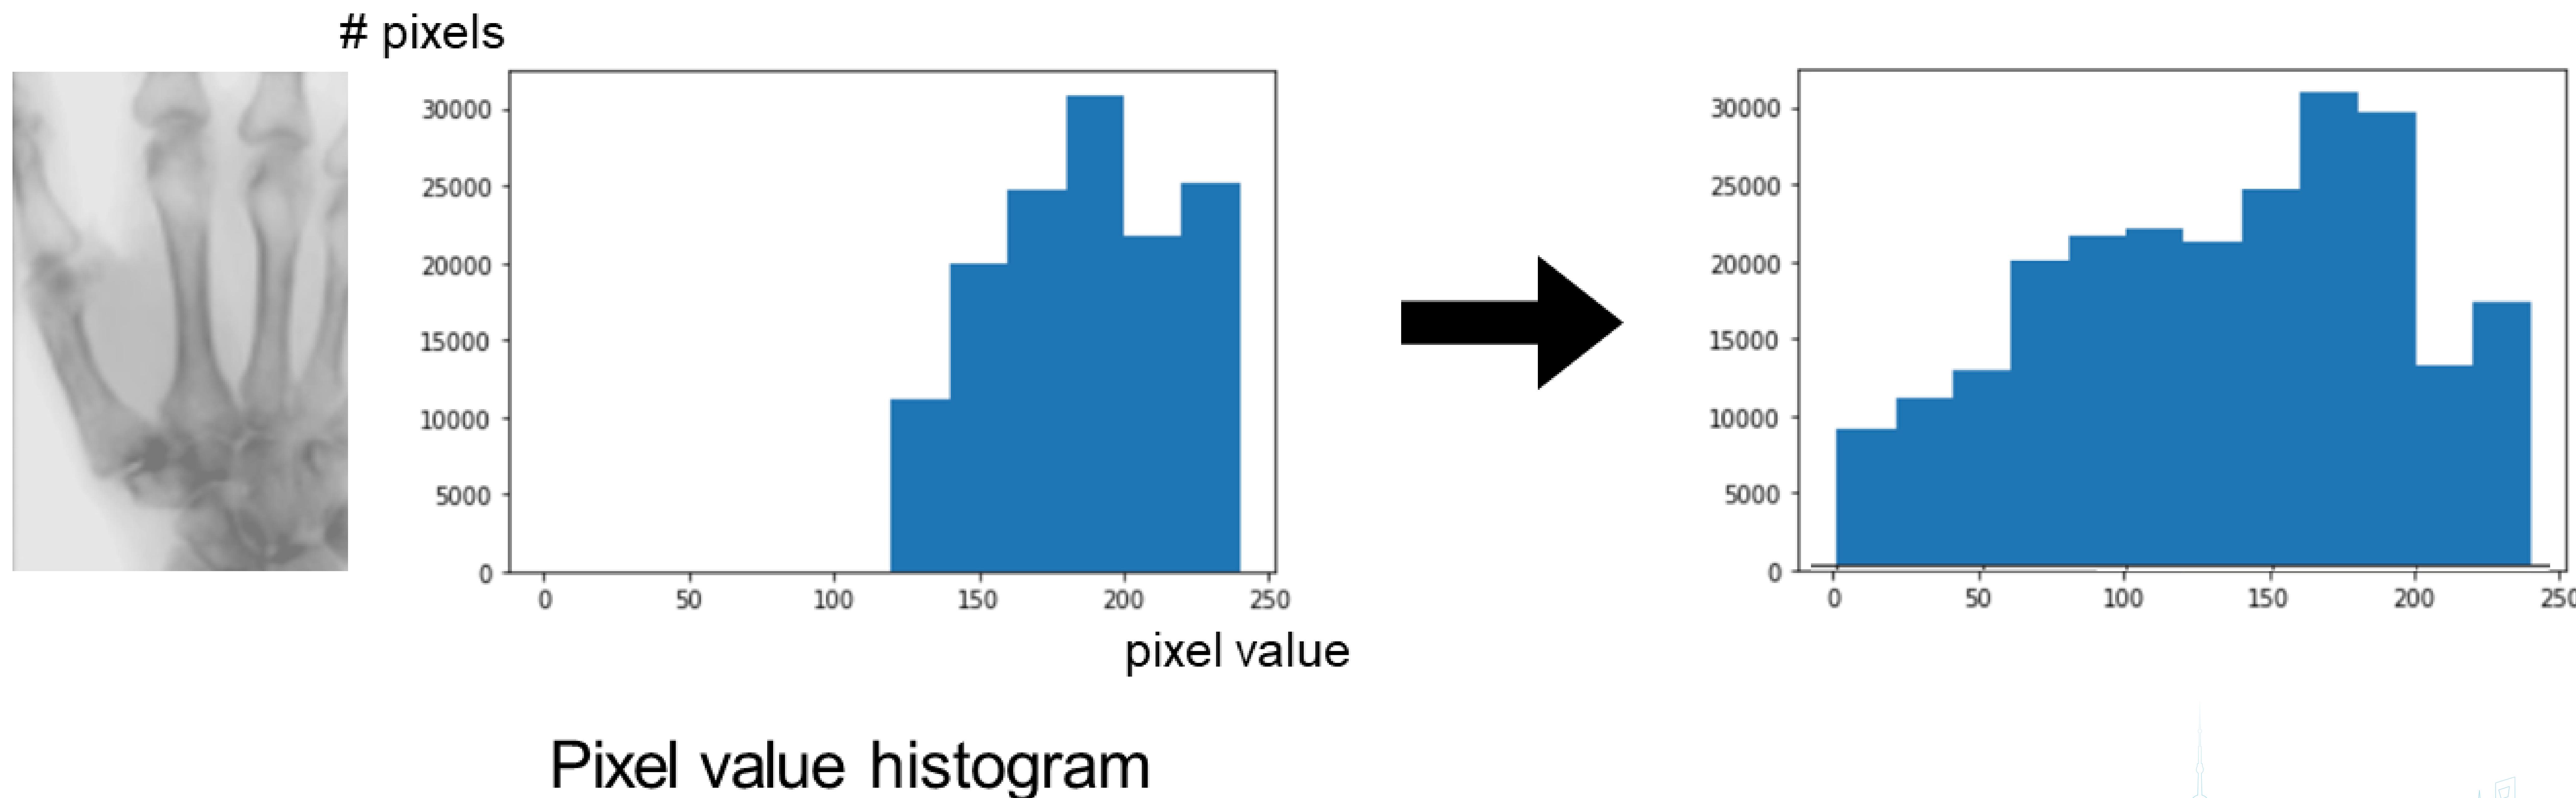
Pixel value histogram

Application: image auto-contrast

Auto-contrast: Make use of the full range

: (min, max) = 120, 230

: (min, max) = 0, 255



Application: image auto-contrast

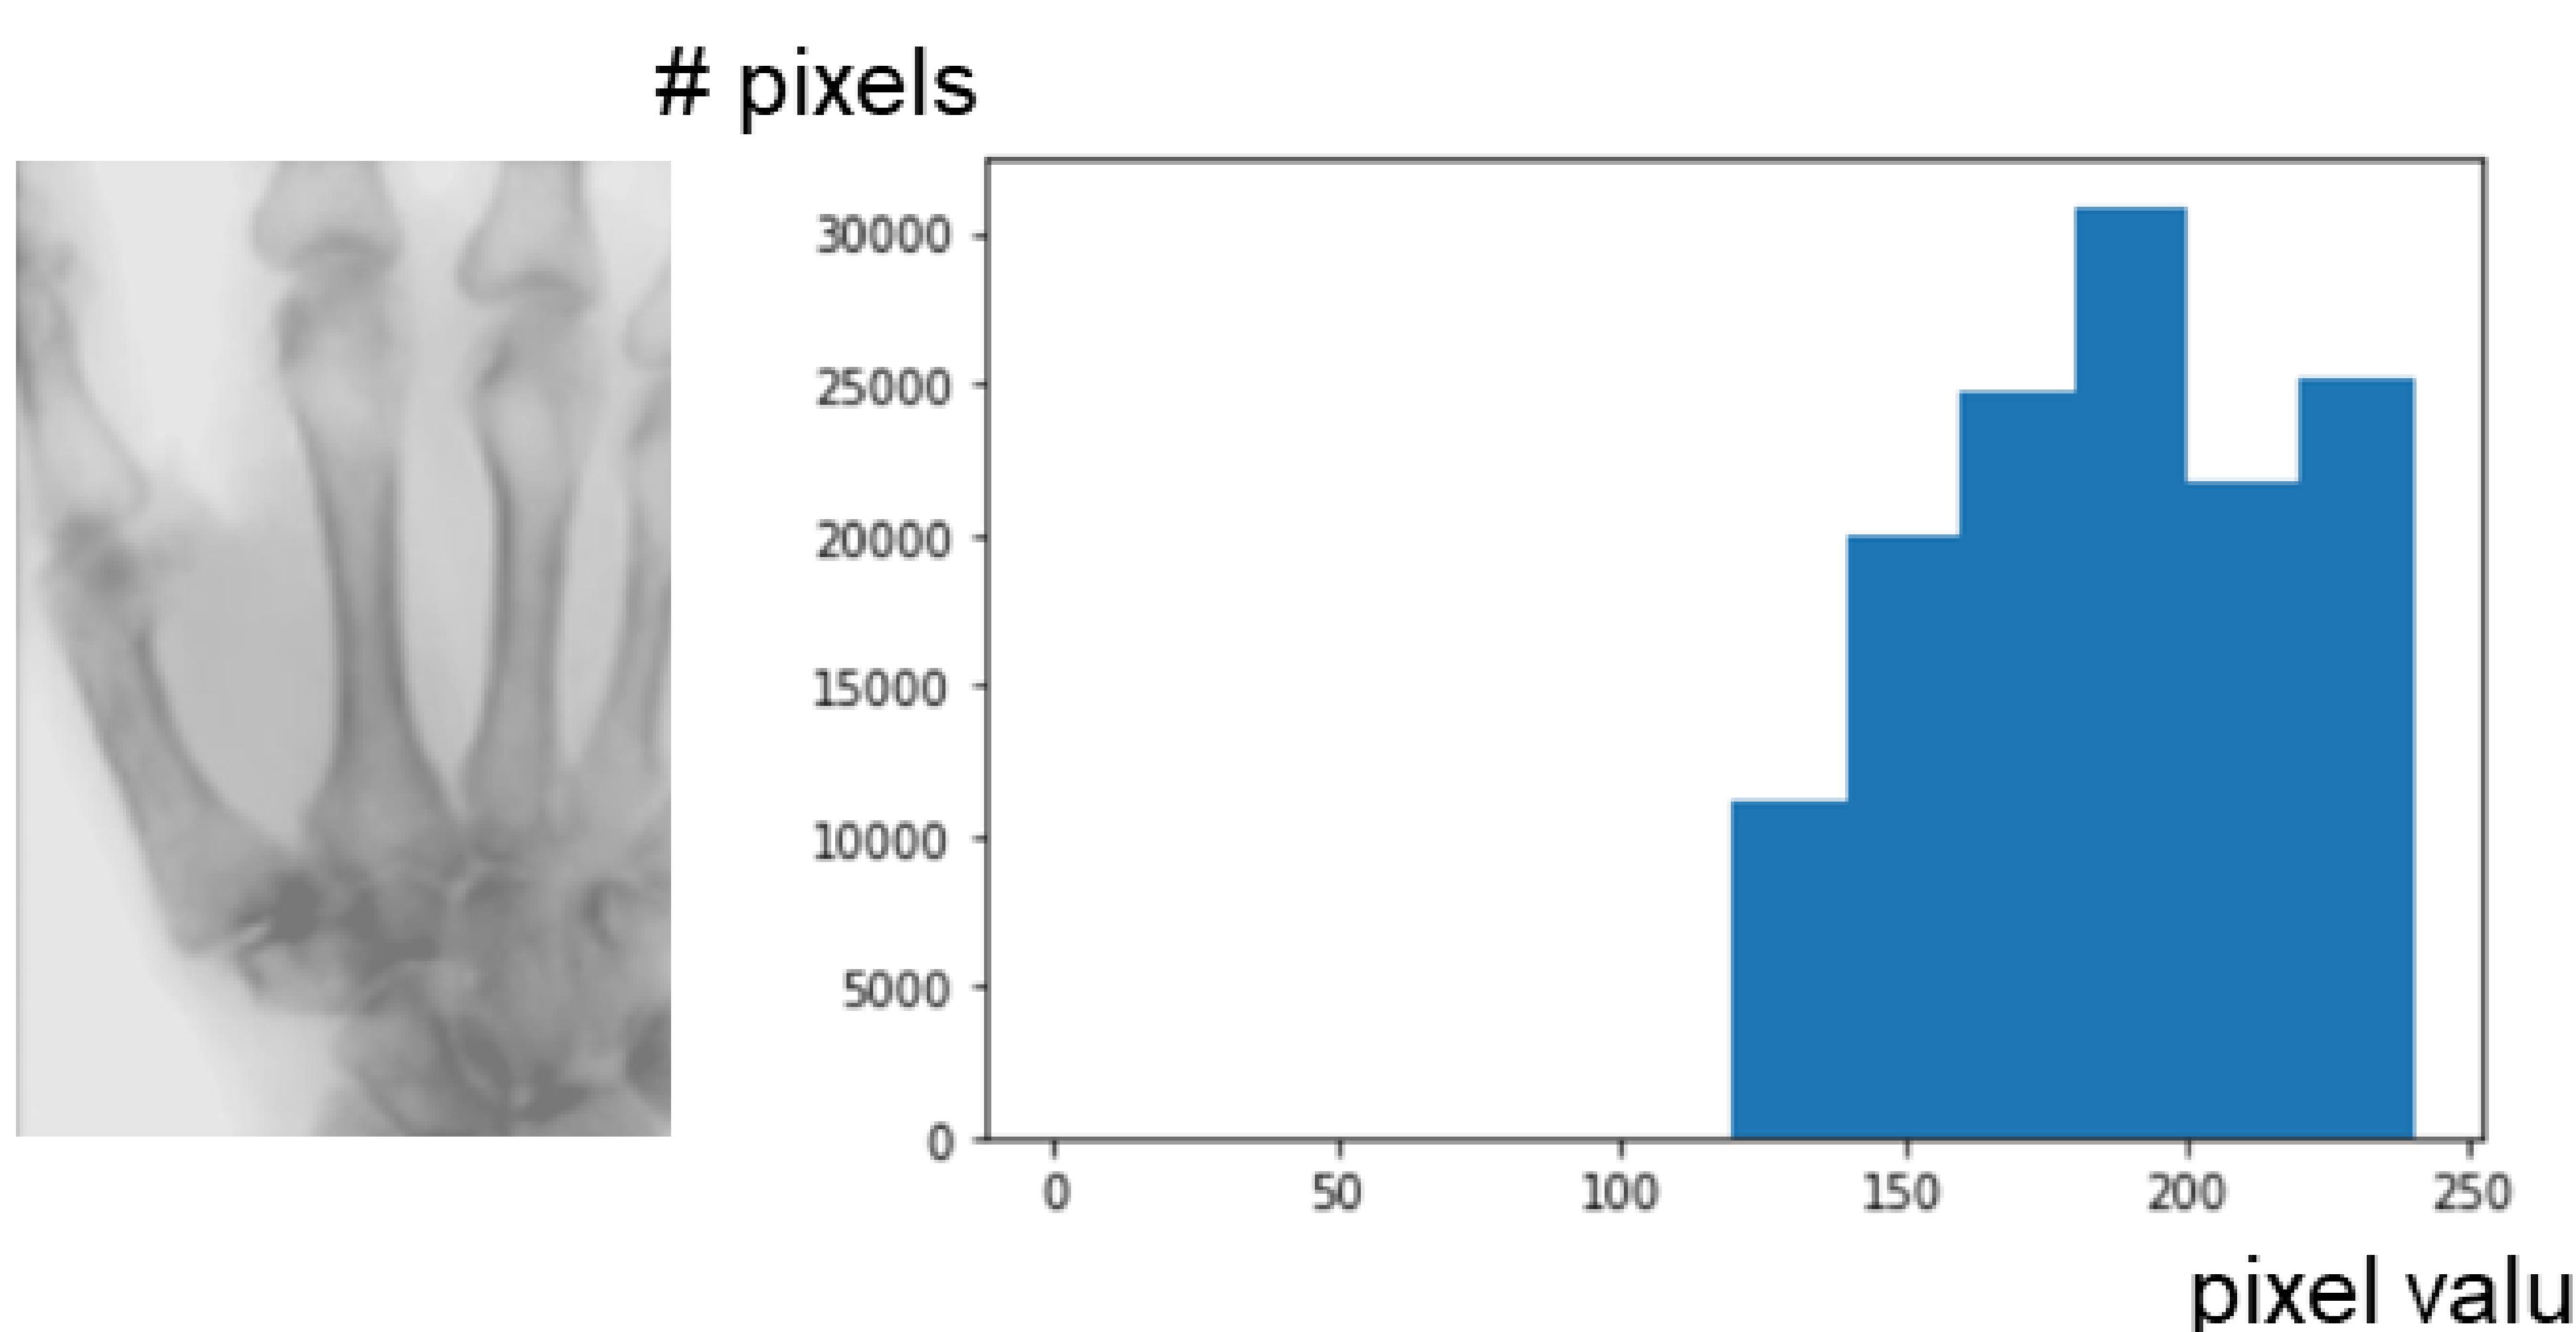
Step 1: make I_{min} to be 0

: (min, max) = 120, 230

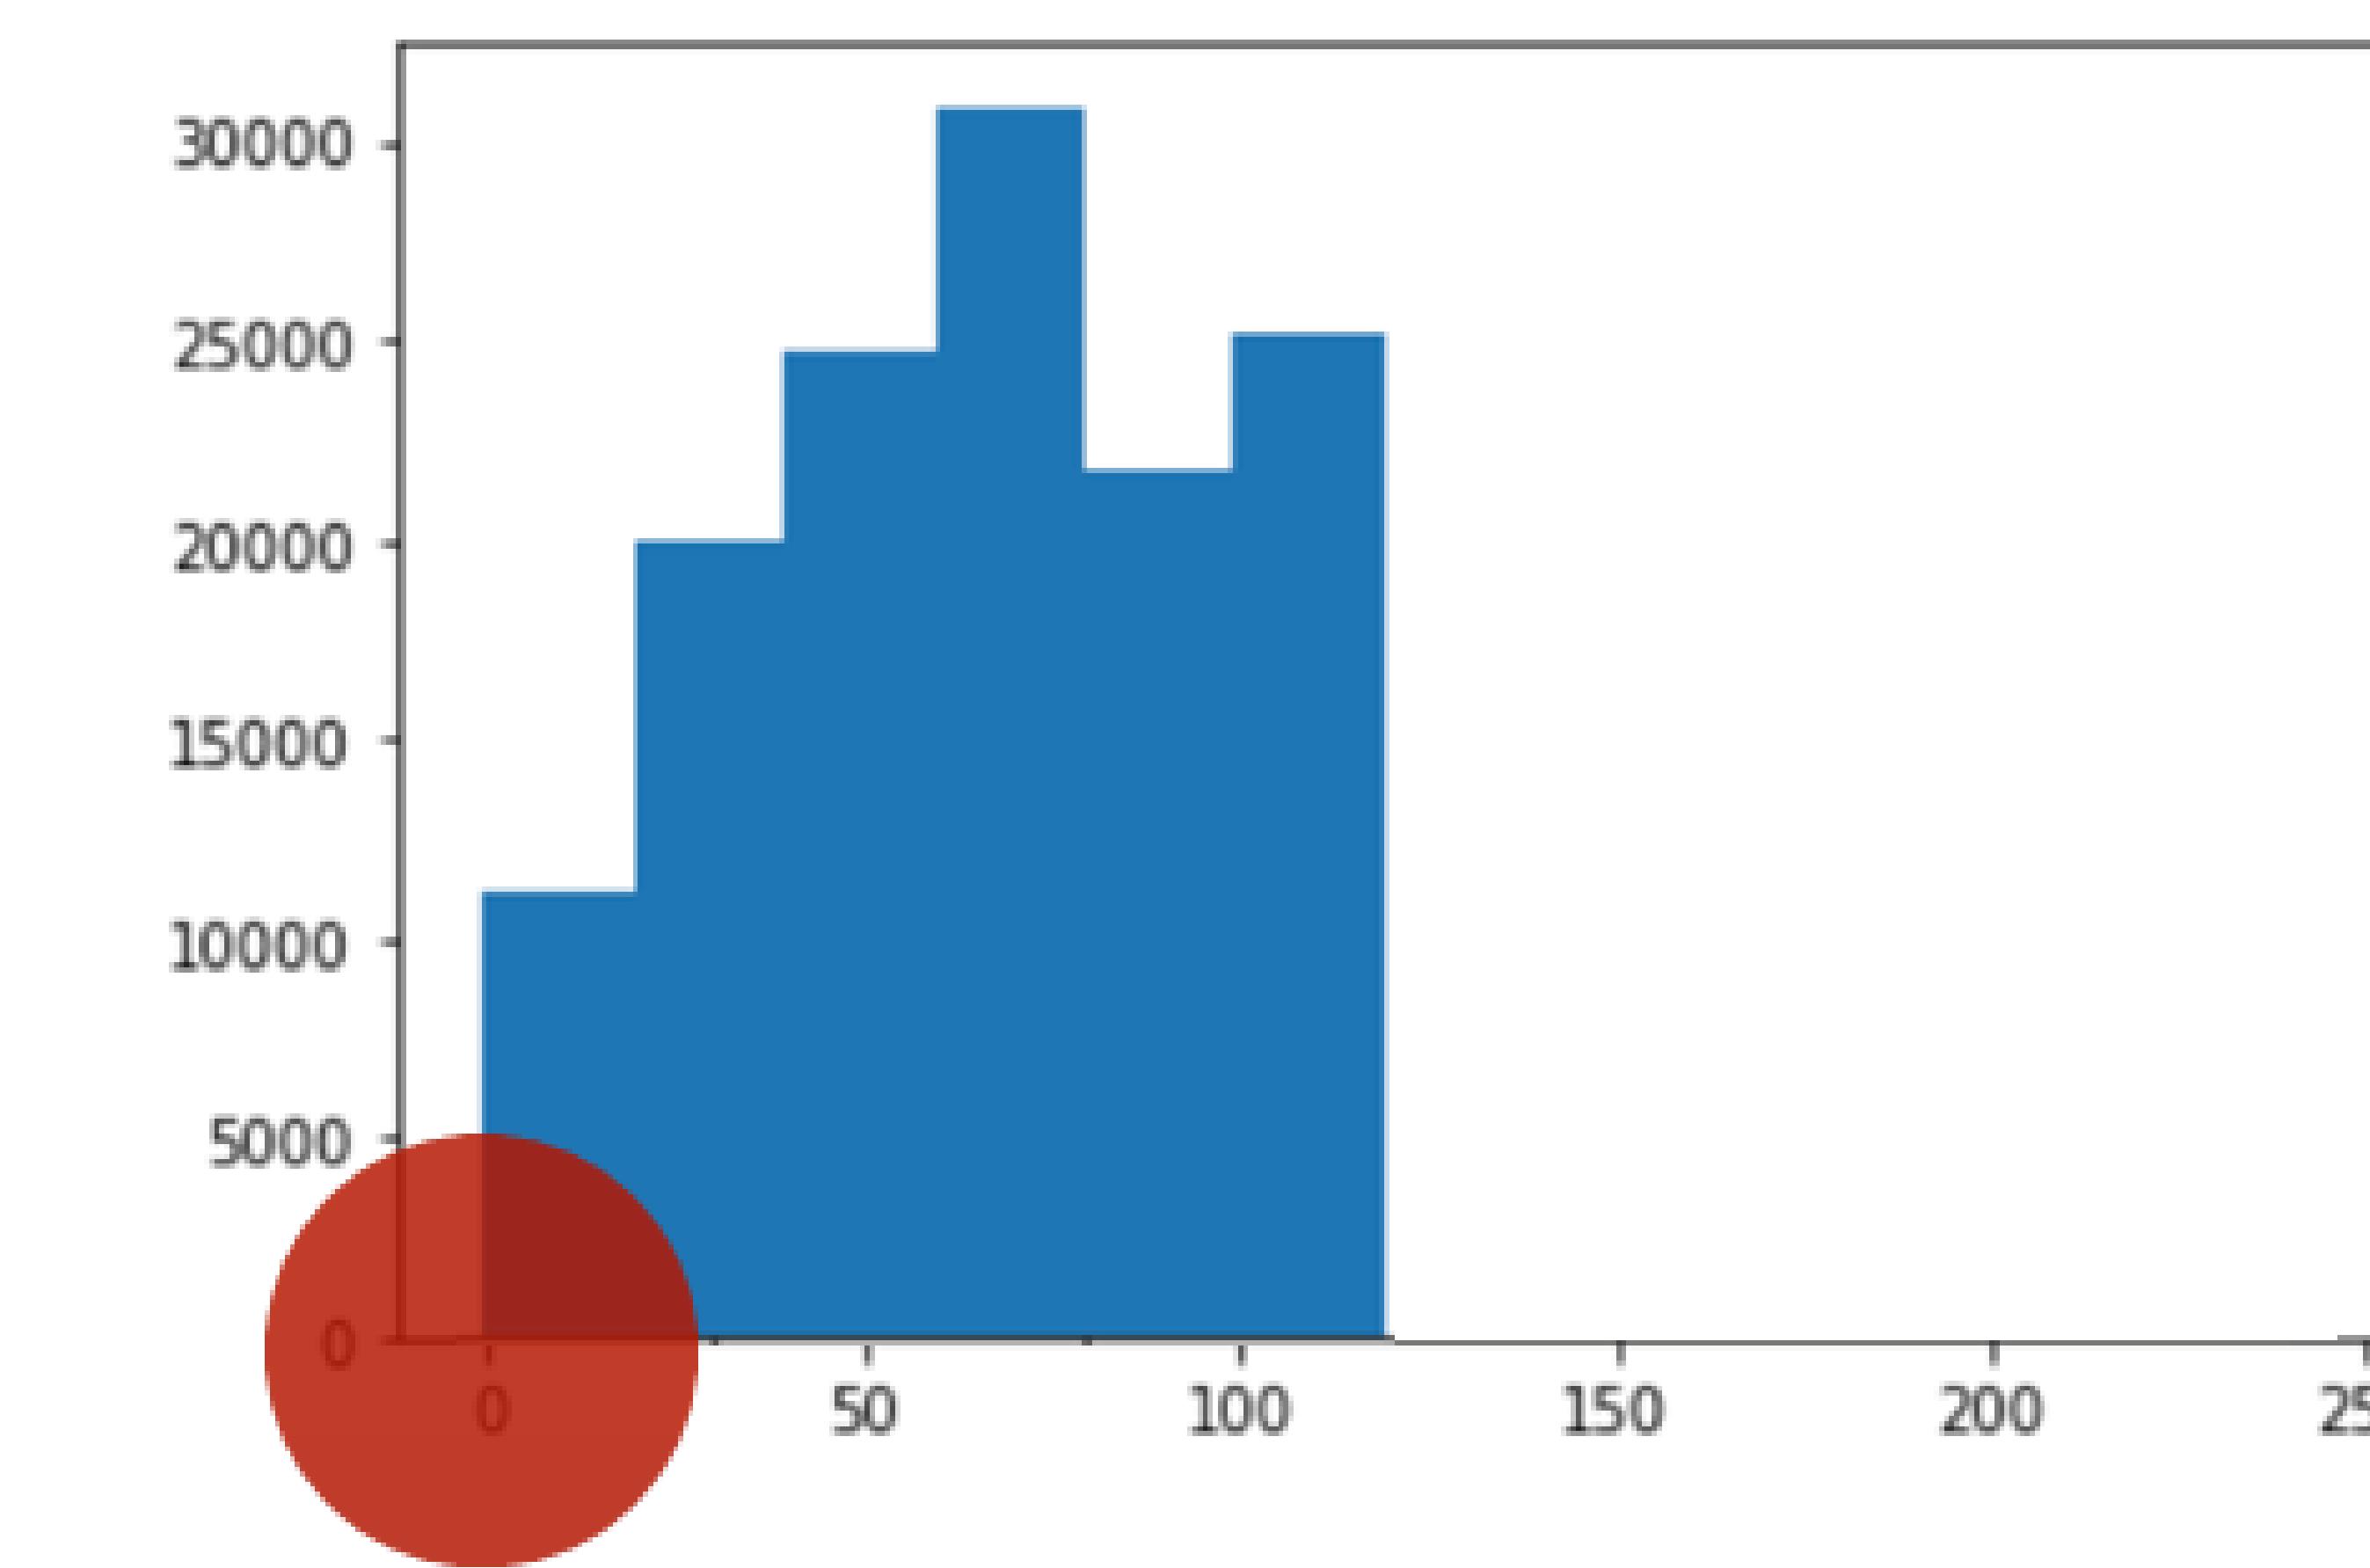
$I = I - 120$

(min, max)

0, 110



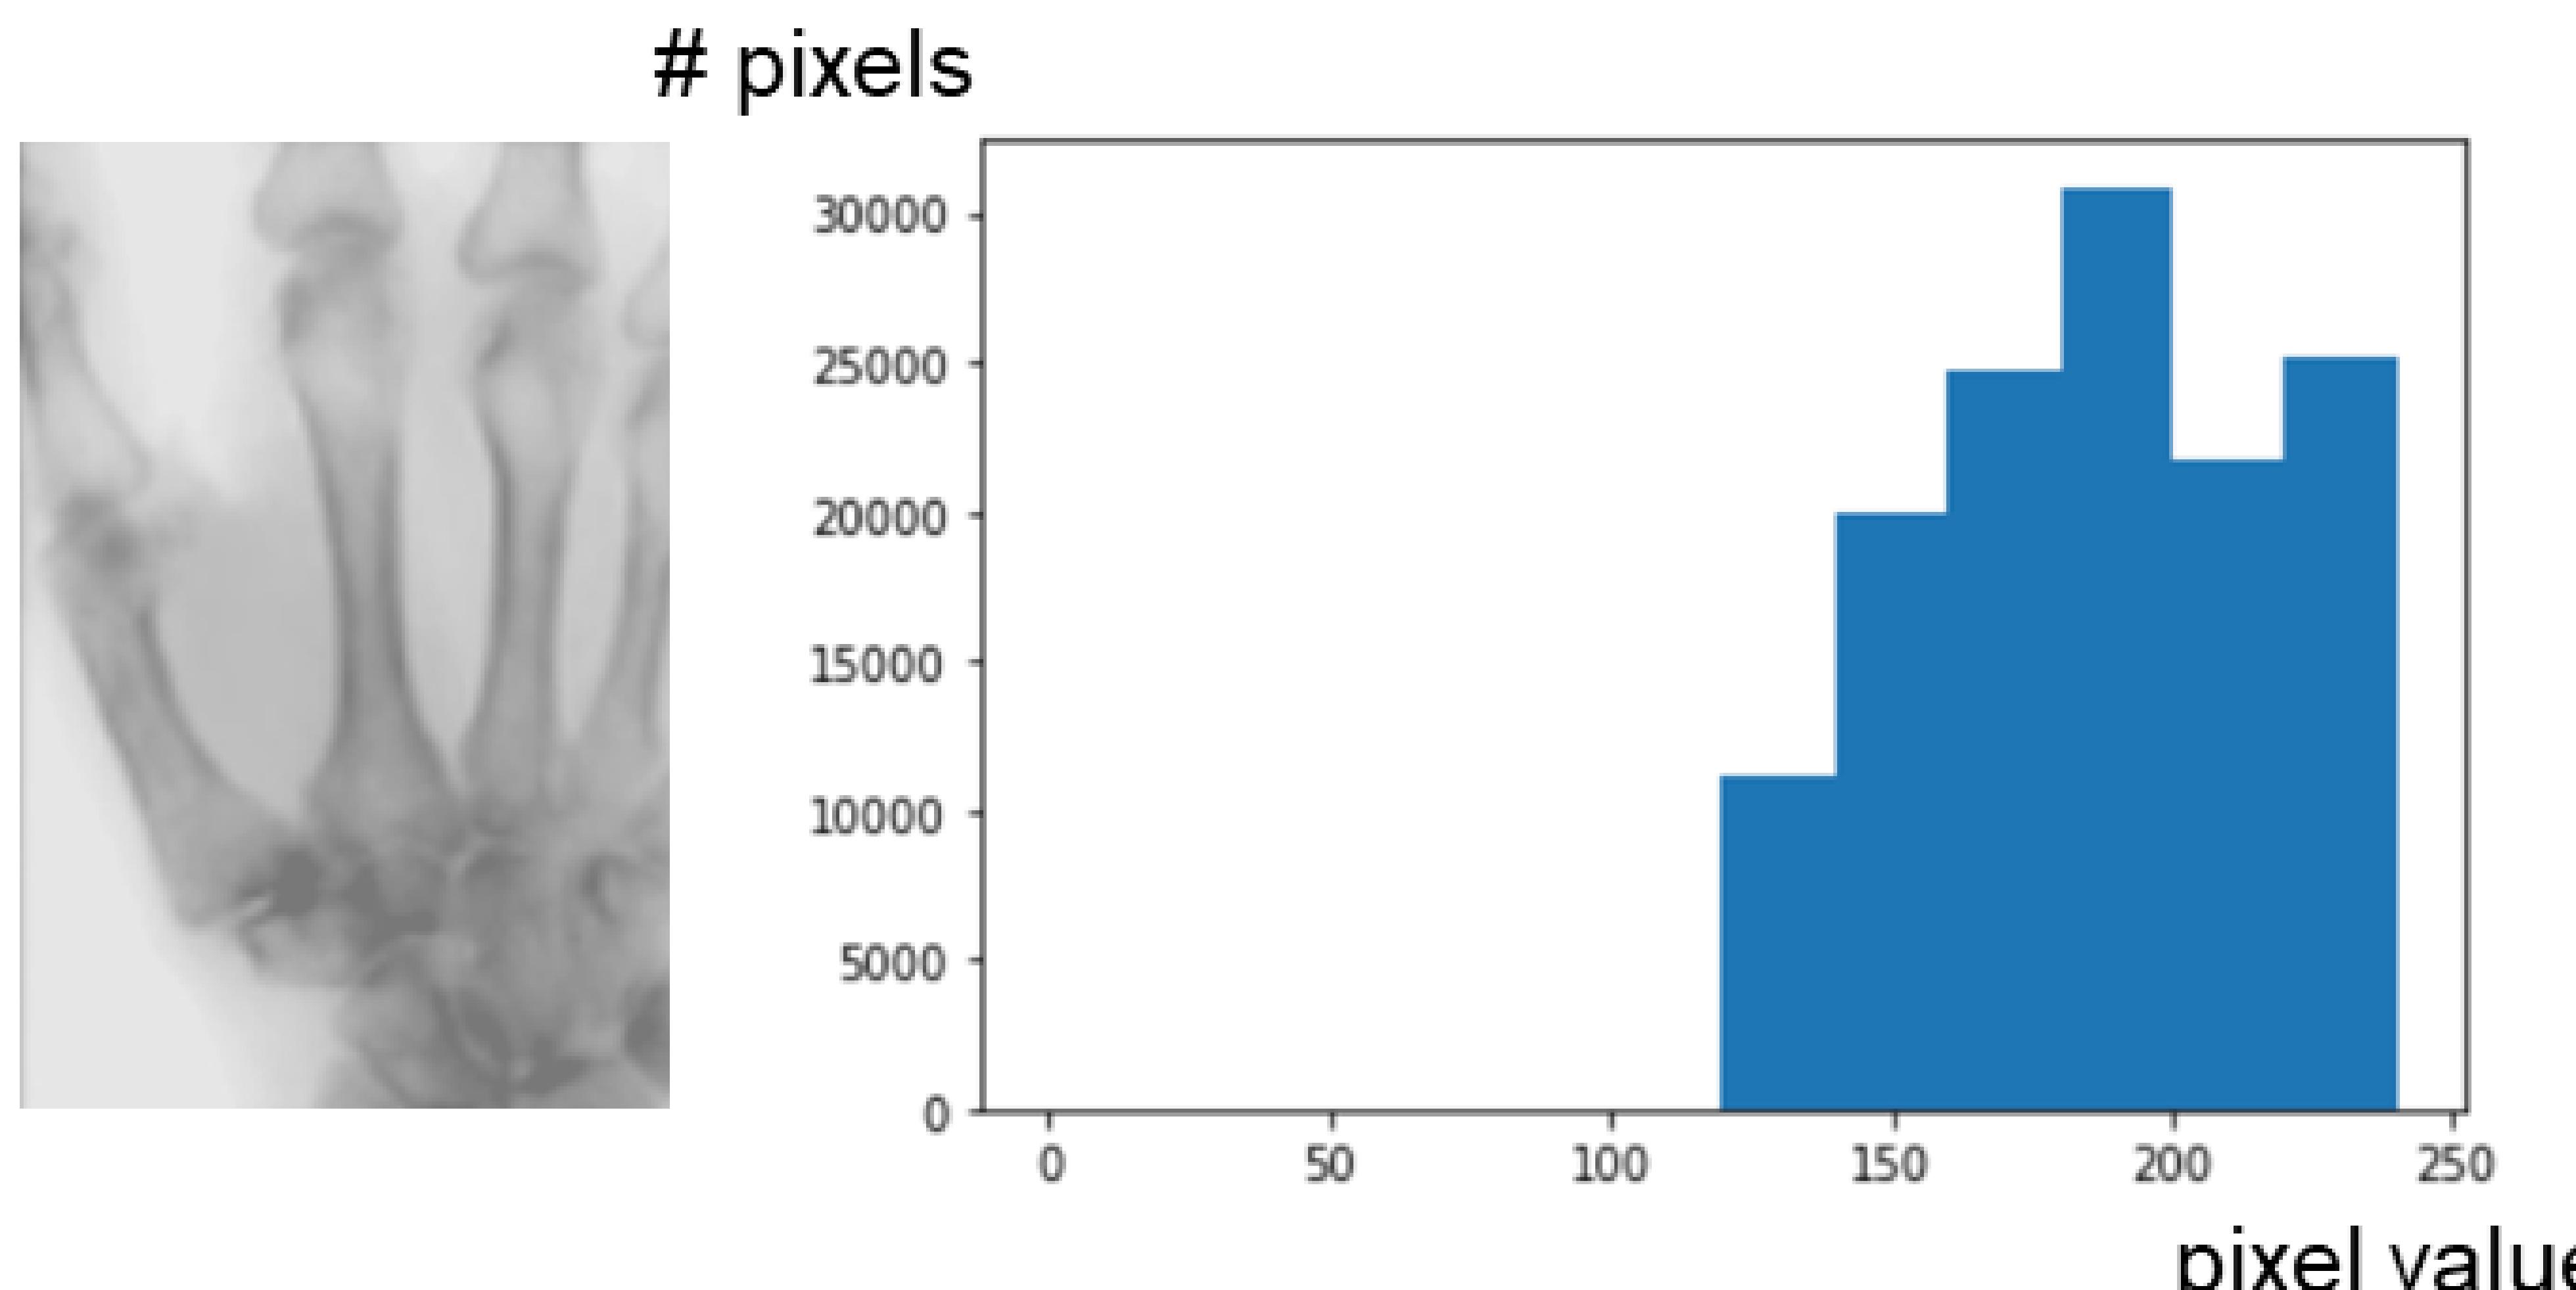
Pixel value histogram



Application: image auto-contrast

Step 2: make I_{max} to be 255

: (min, max) = 120, 230



Pixel value histogram

$I = I - 120$

$I = I / (I.max() - I.min())$

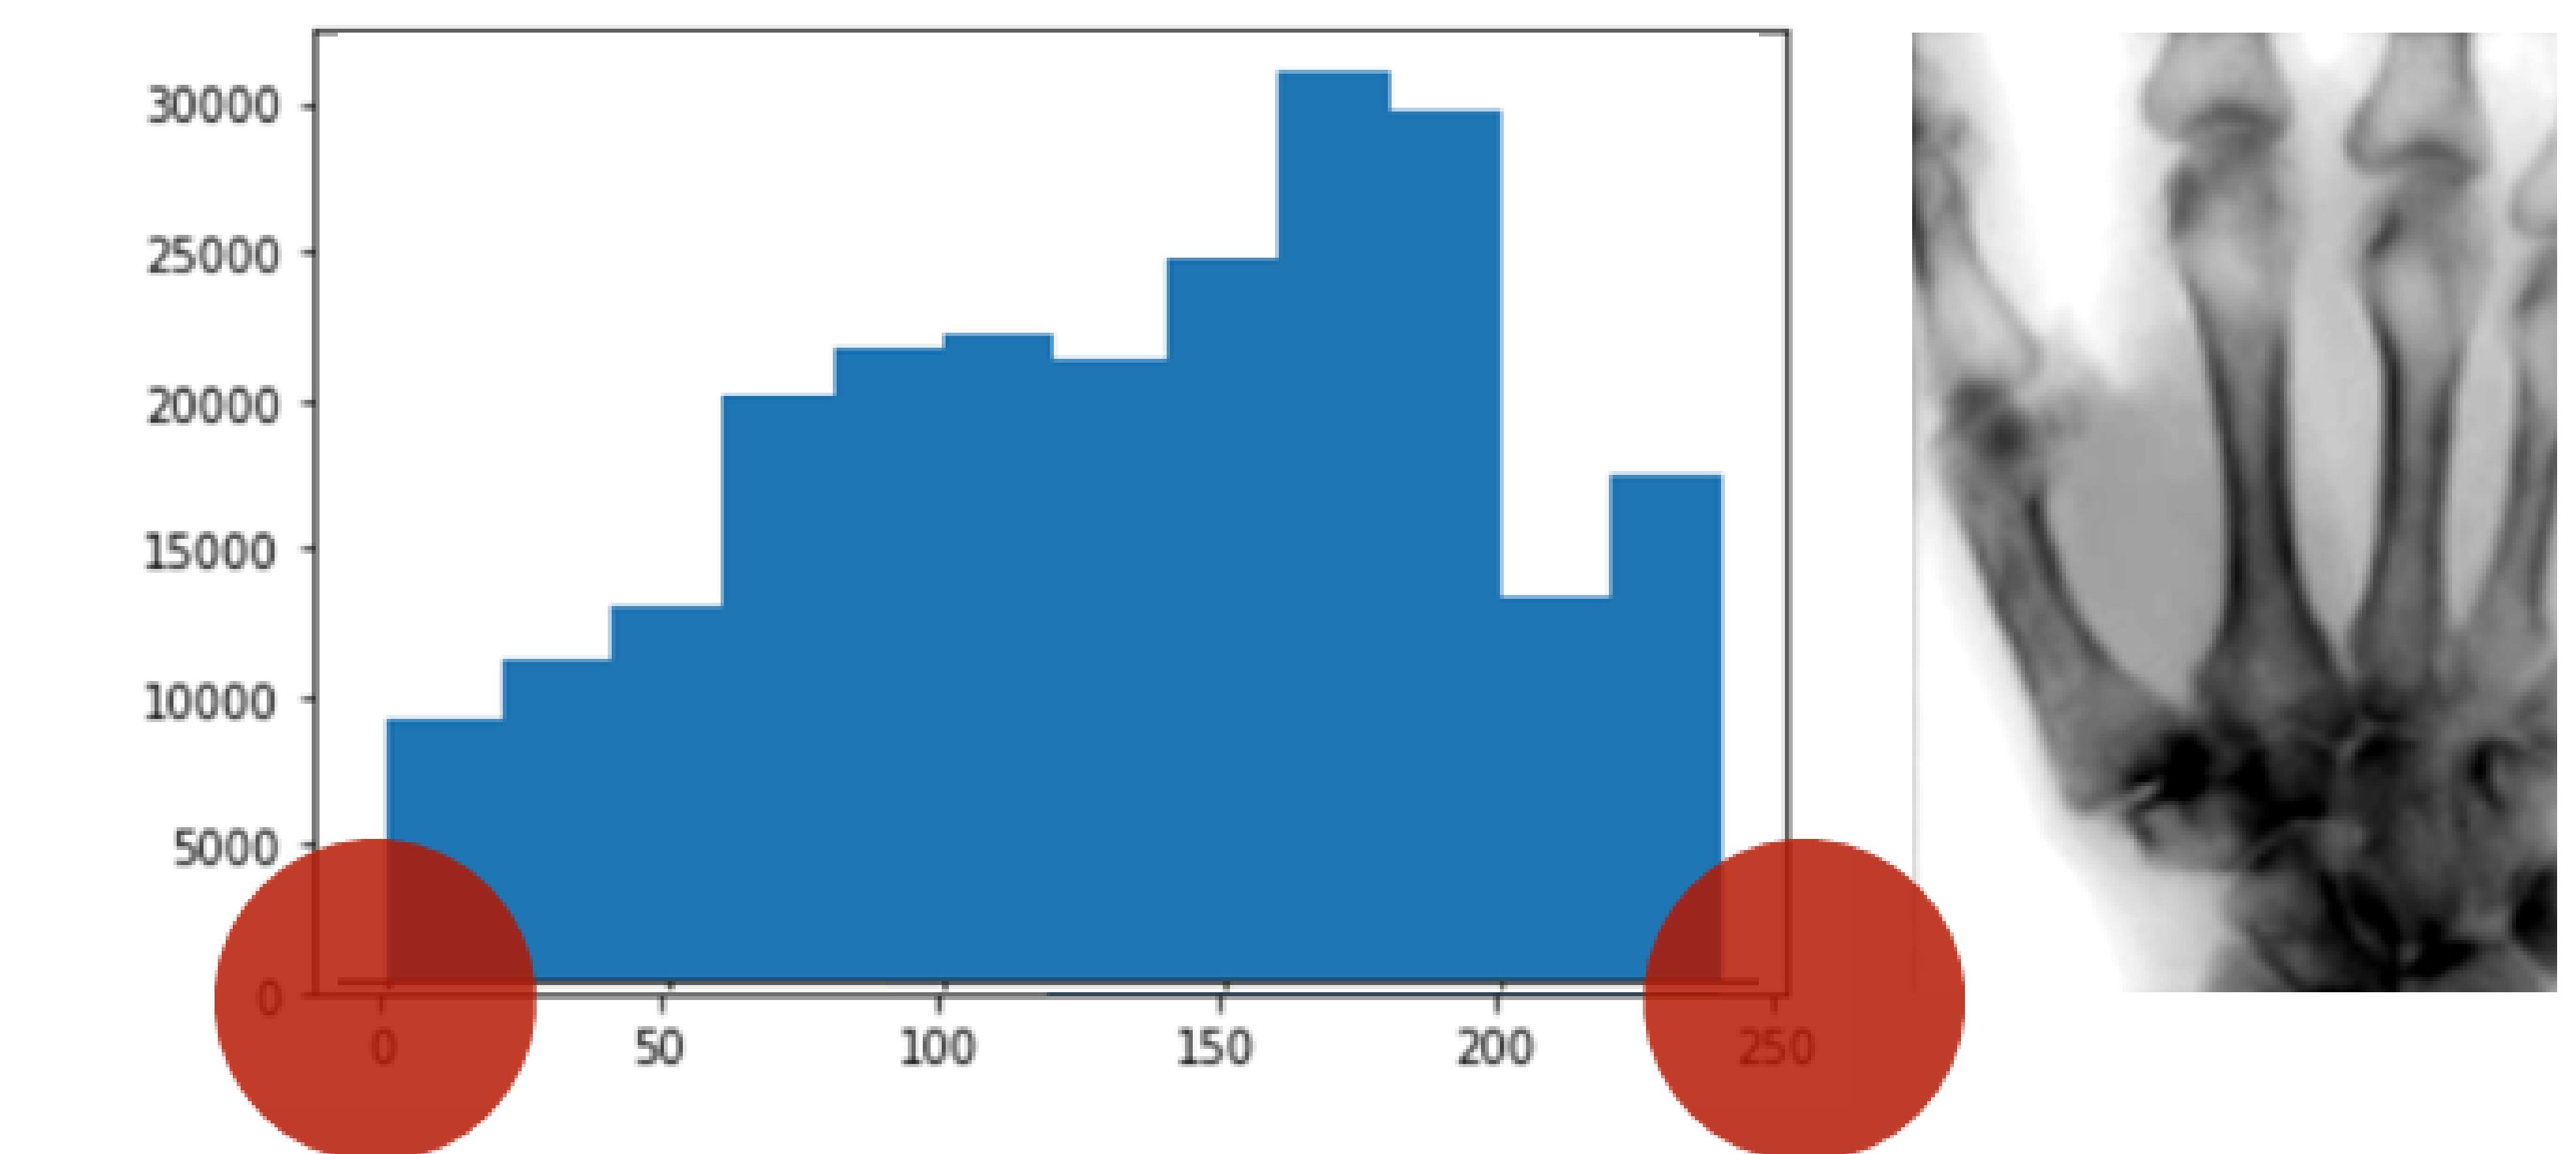
$I = I * 255$

(min, max)

0, 110

0, 1

0, 255

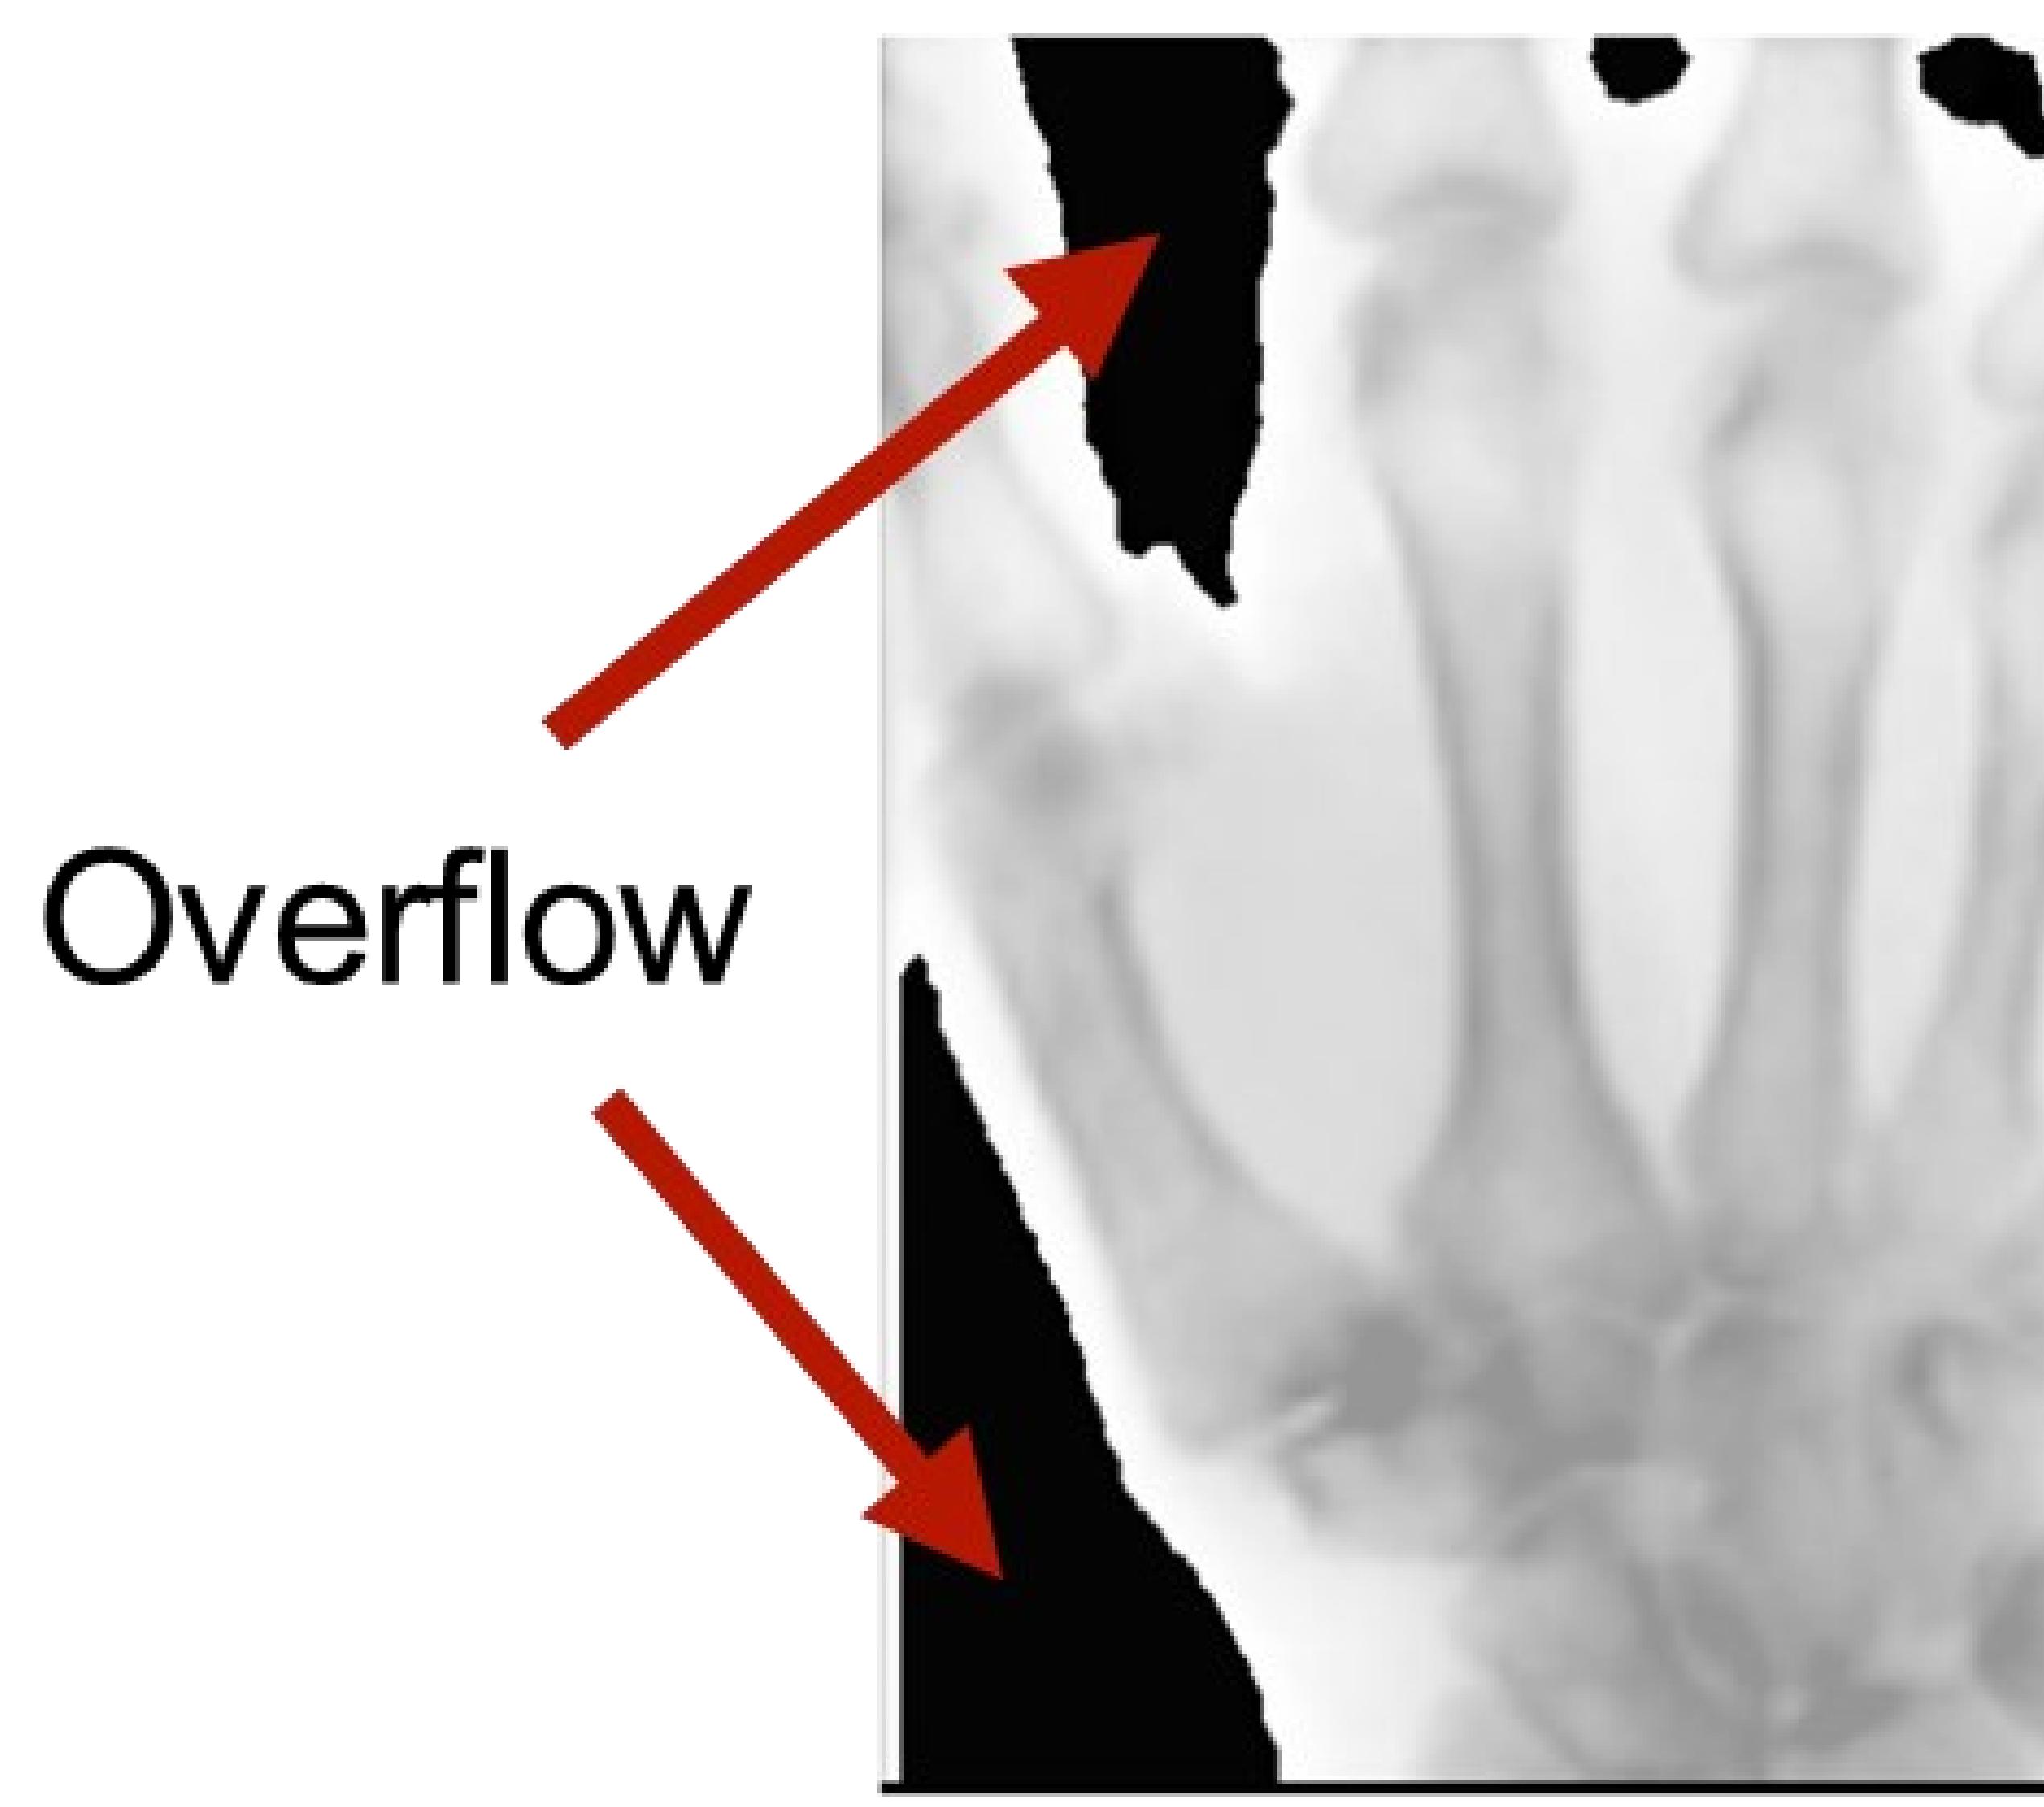


Linear functions

Implementation: uint8 arithmetic

$$100 + 200 = ?$$

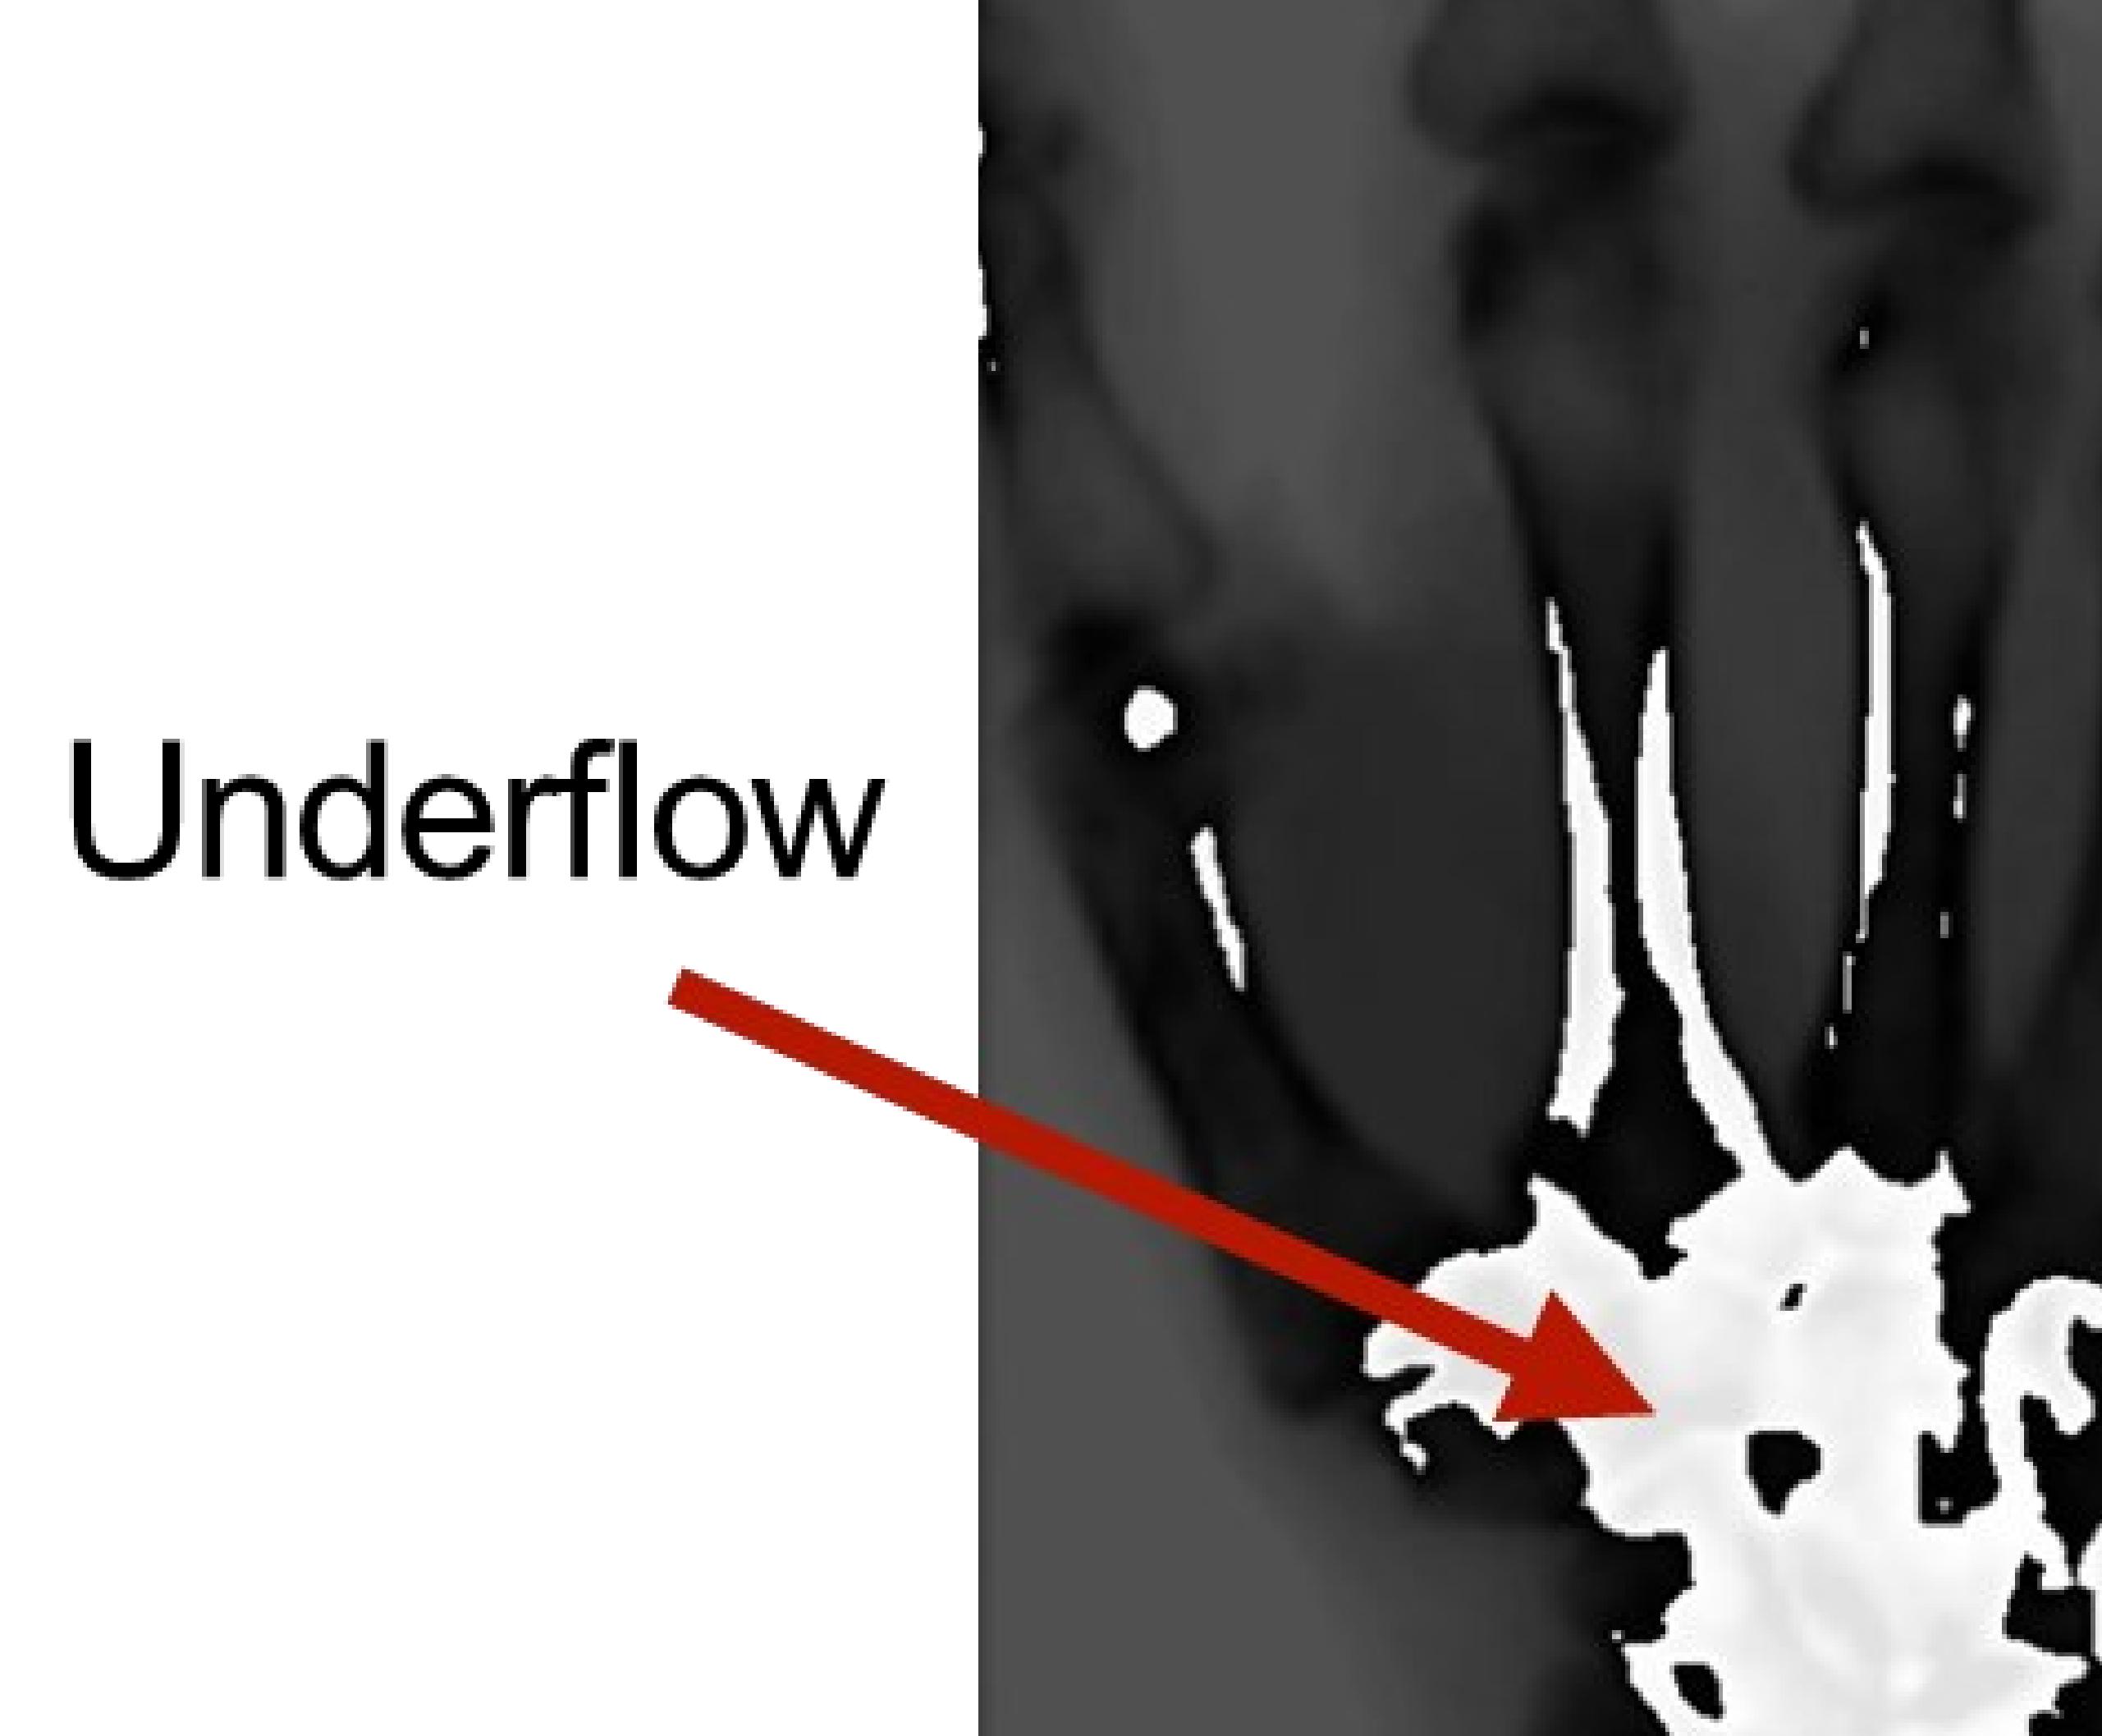
Overflow: $300 - 256 = 44$



$$I = I + 30$$

$$100 - 120 = ?$$

Underflow: $-20 + 256 = 236$



$$I = I - 150$$

Image arithmetic

- convert to “Float”
- do arithmetic
- clip it $[0,255]$

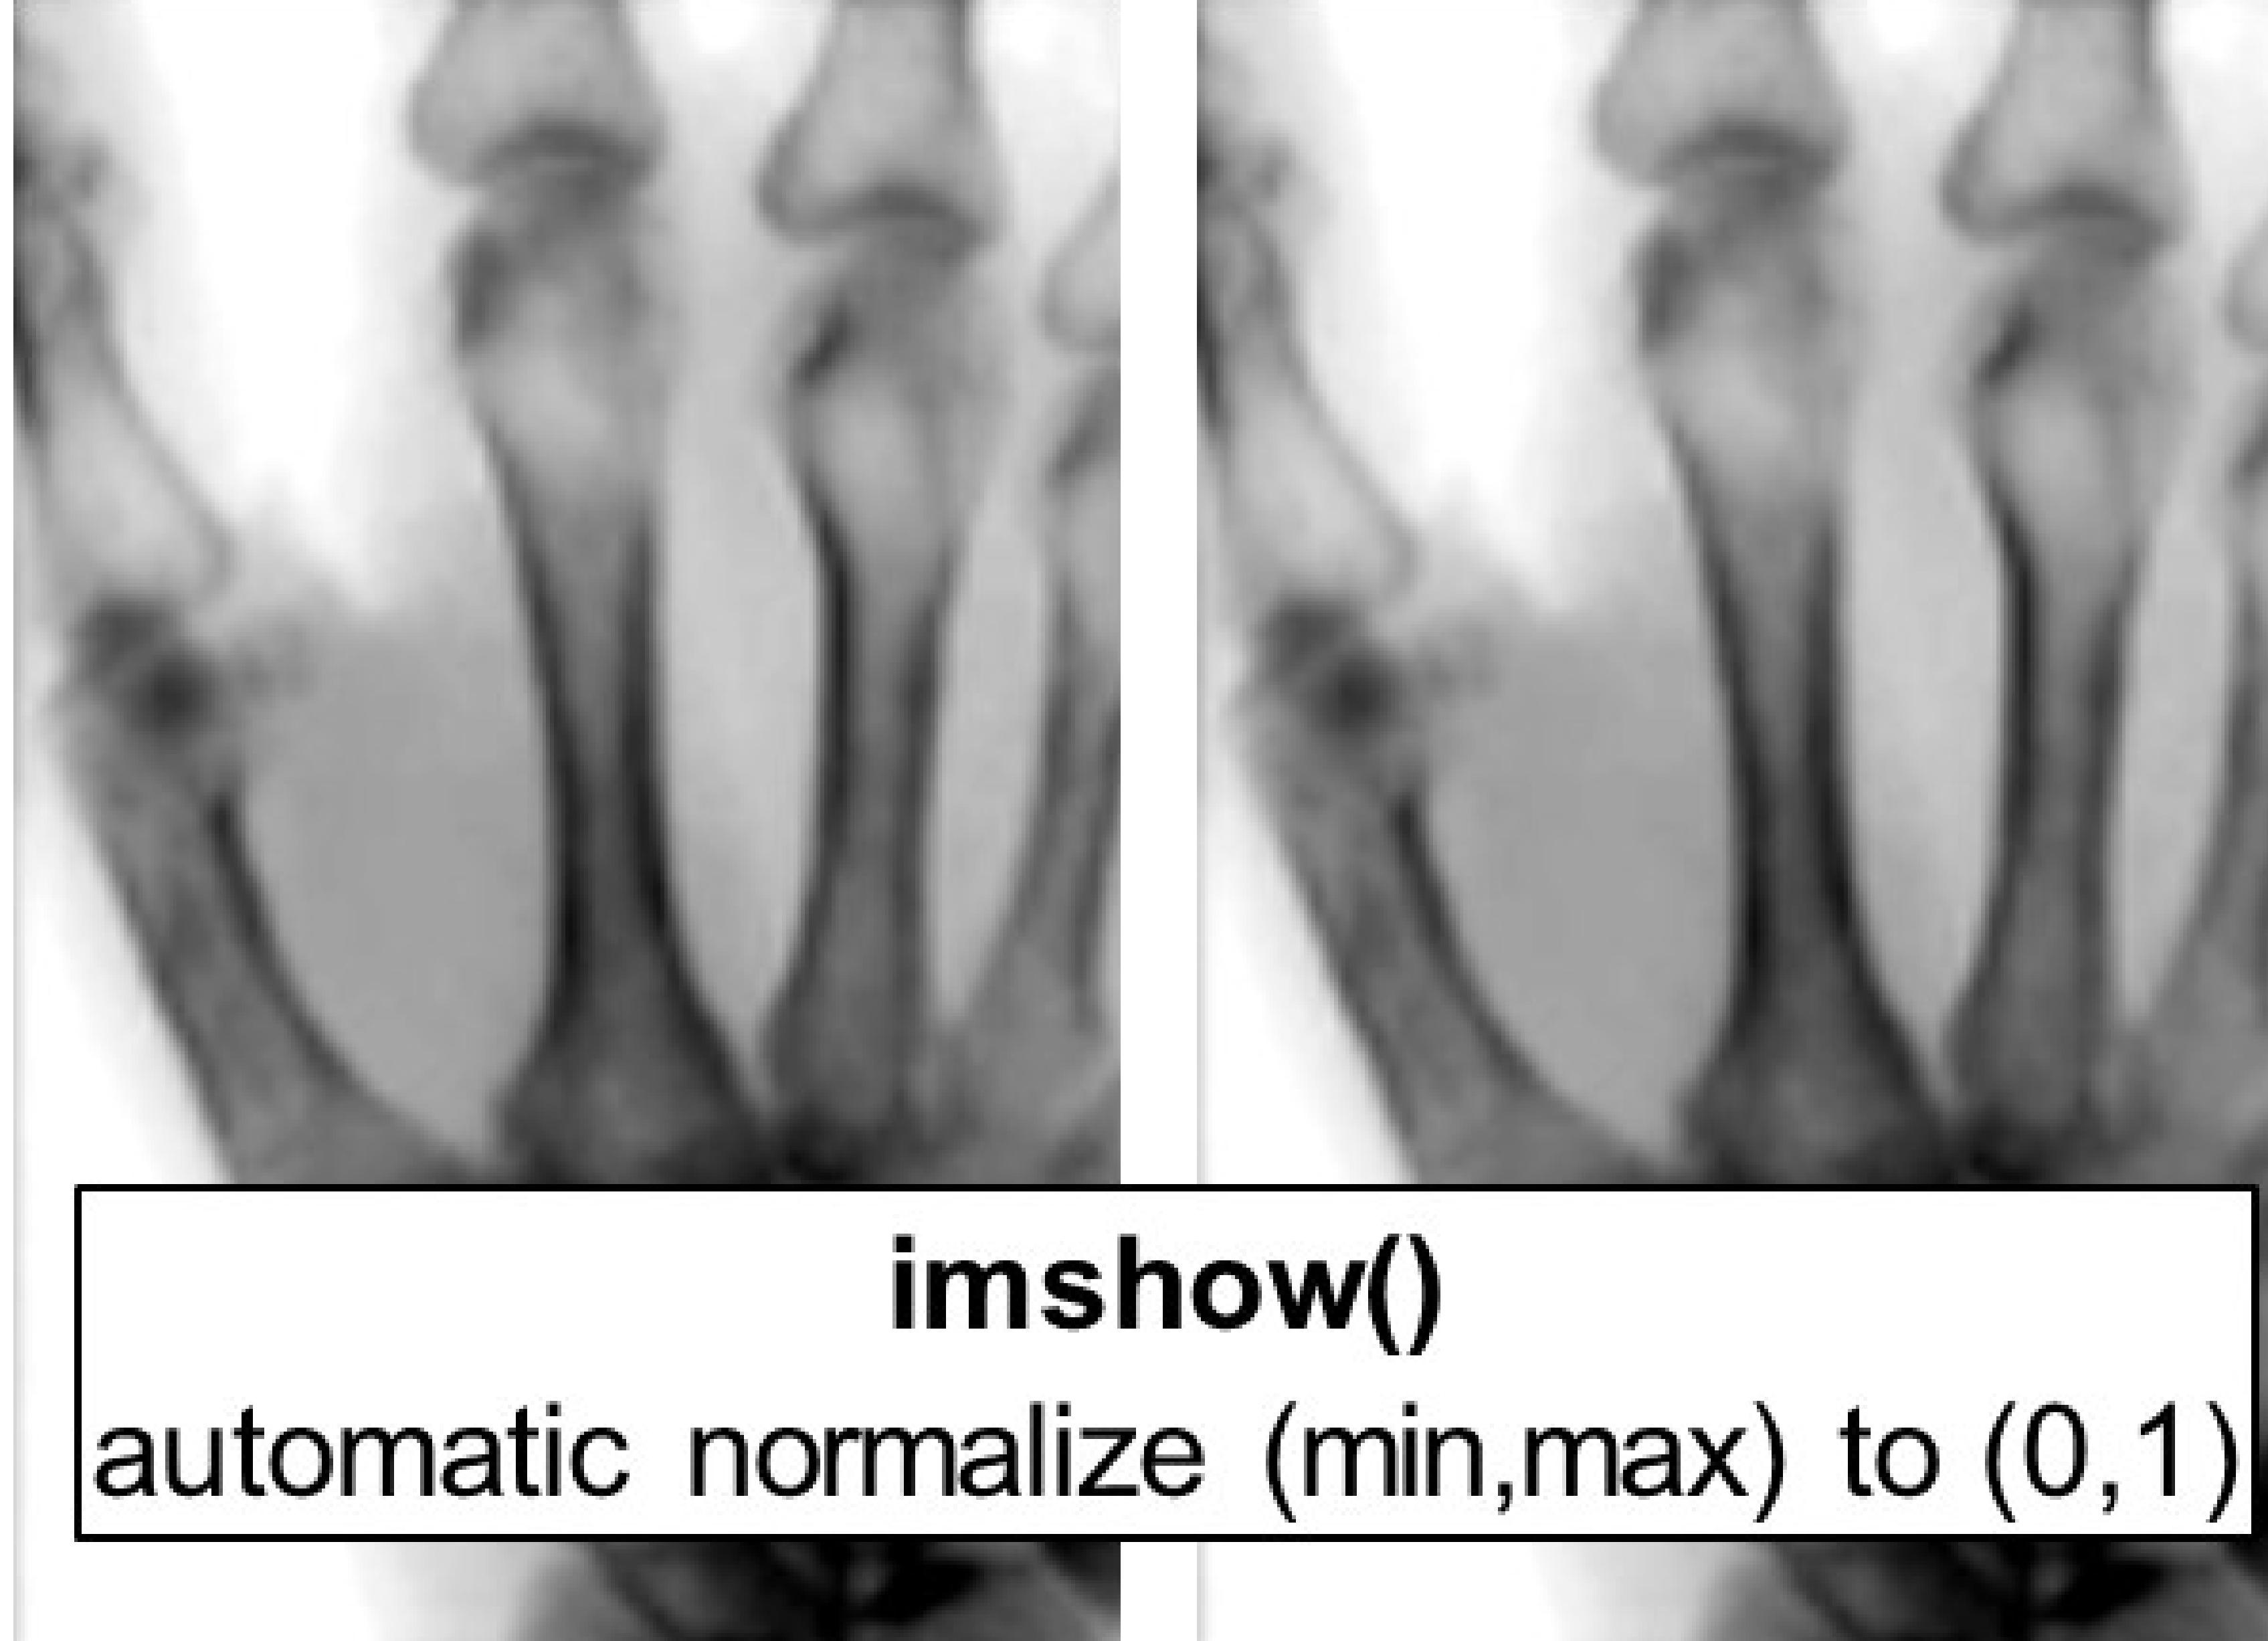
`I.astype(float)`

`np.clip(im, 0, 255).astype(np.uint8)`

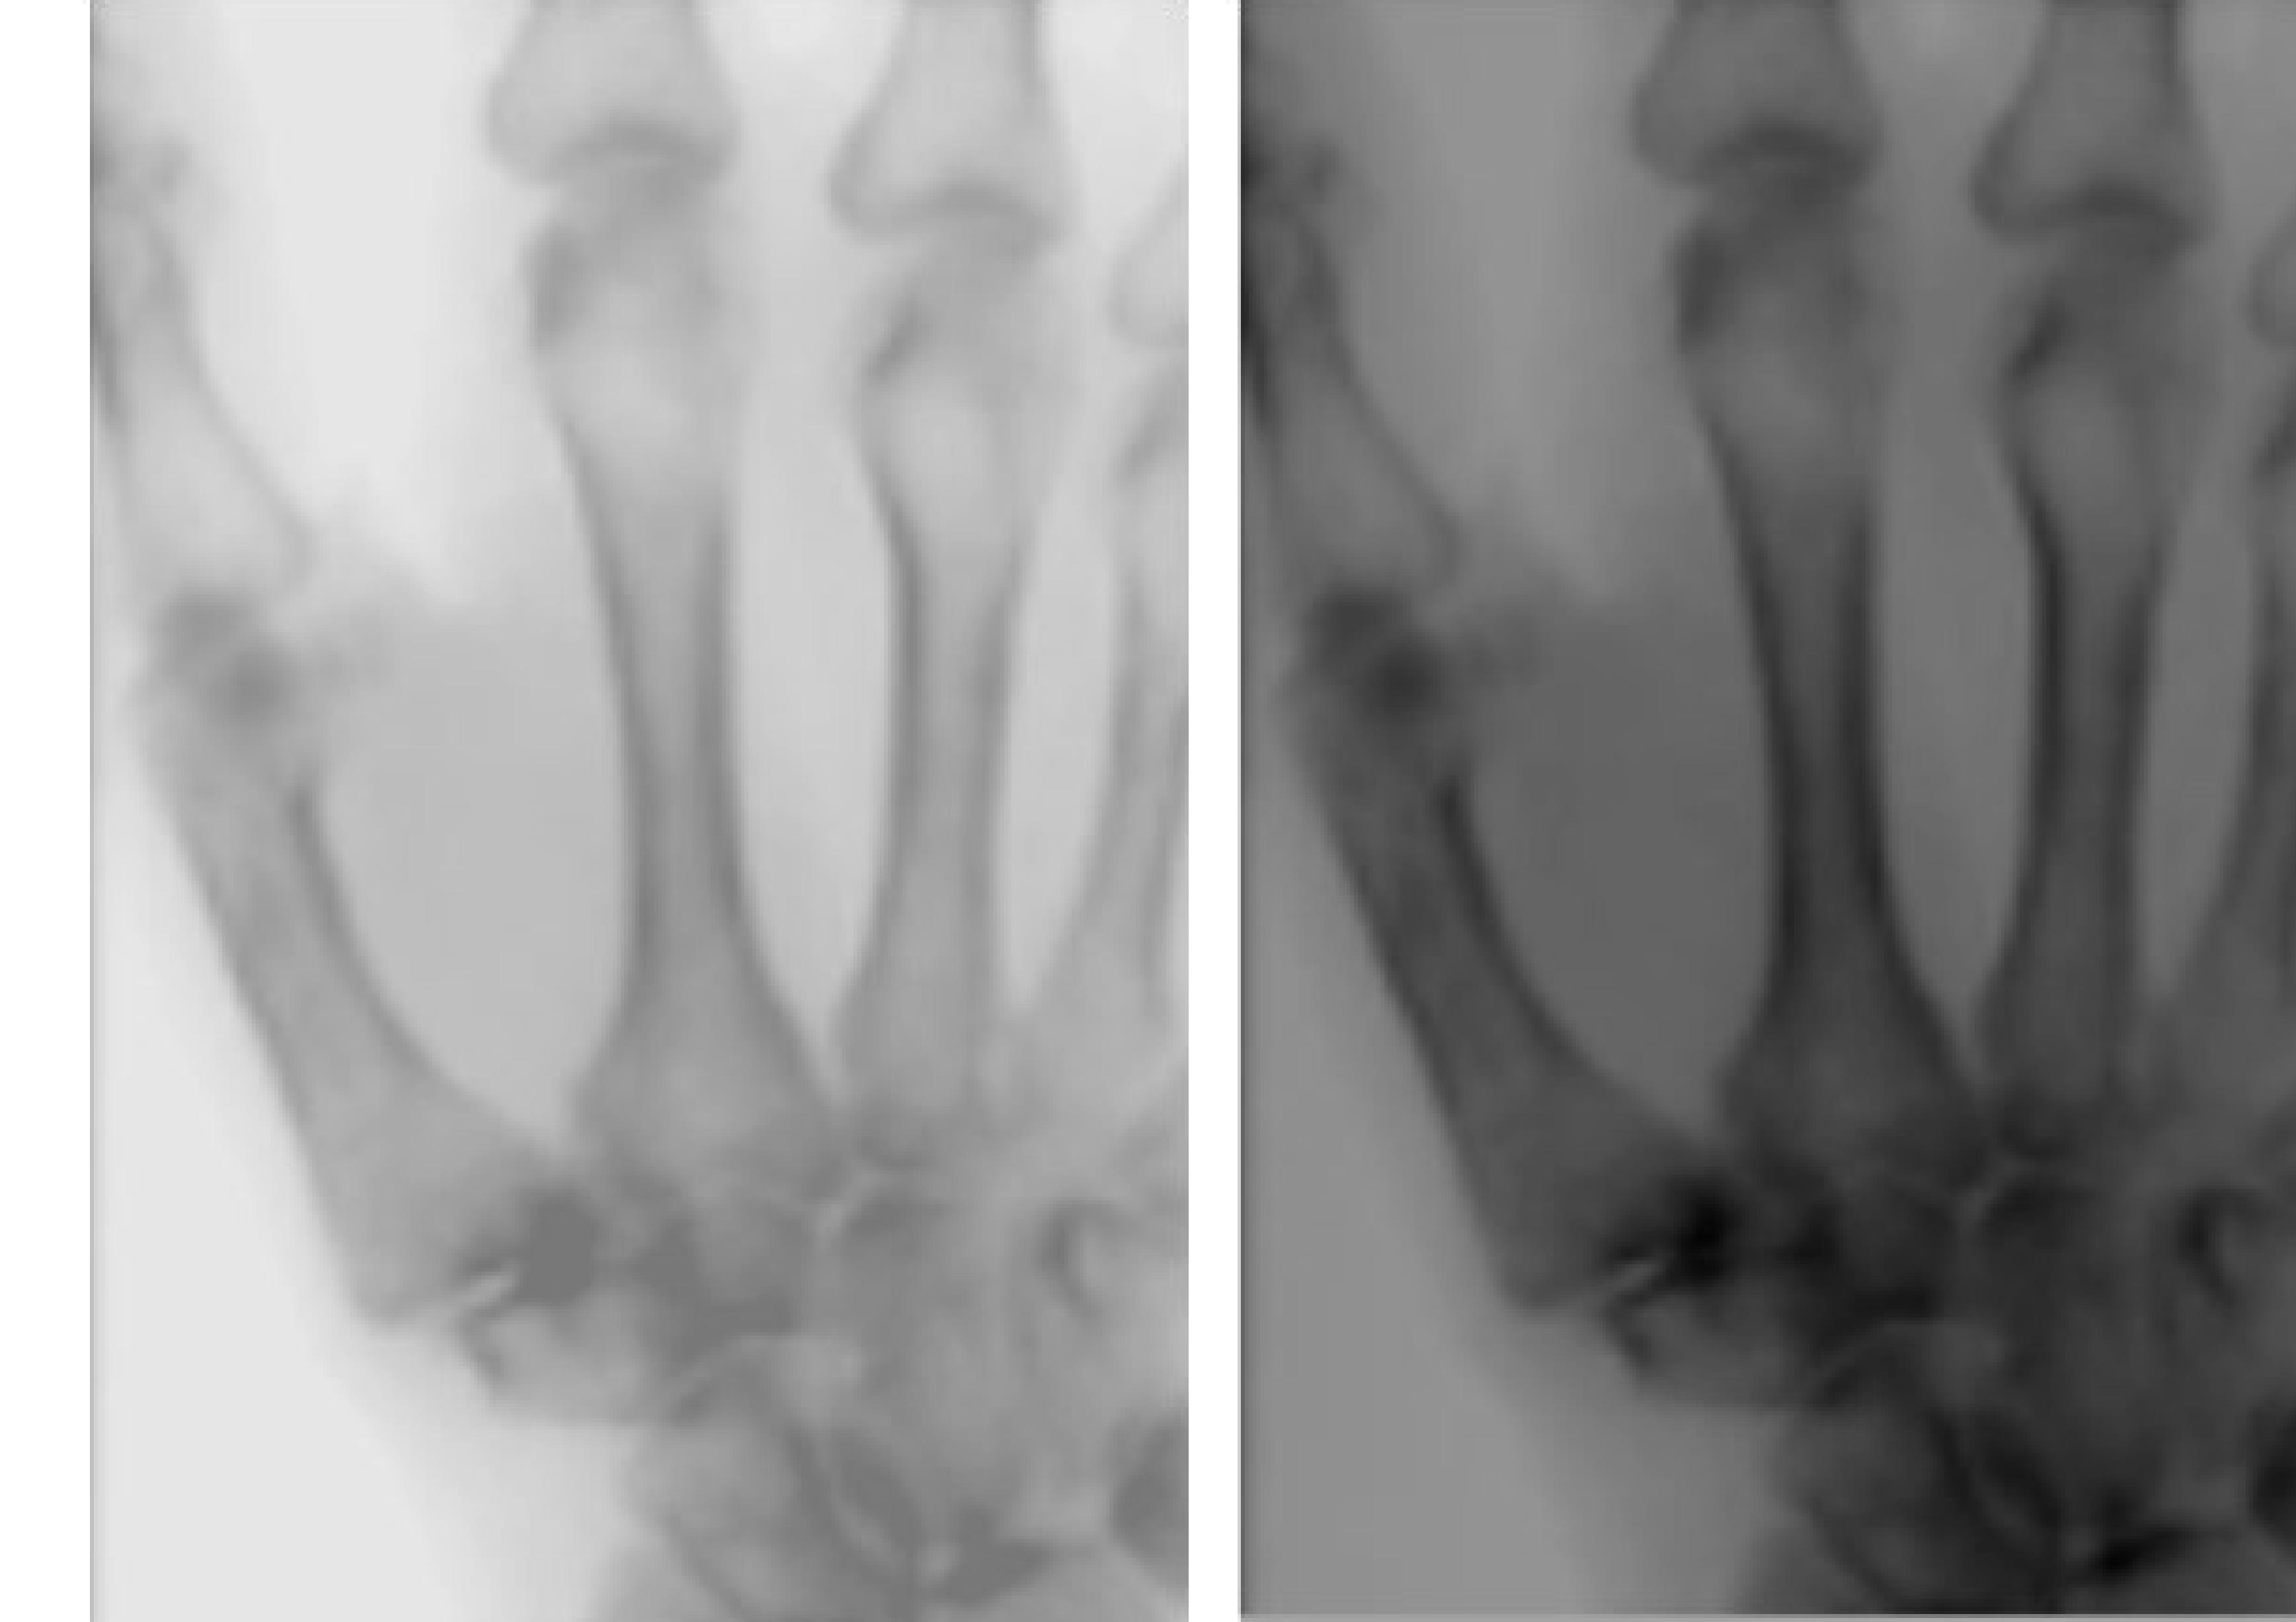
Linear functions

Implementation: plt.imshow()

```
plt.subplot(121)
plt.imshow(image, cmap='gray')
plt.axis('off')
plt.subplot(122)
plt.imshow(image - 100, cmap='gray')
plt.axis('off')
```

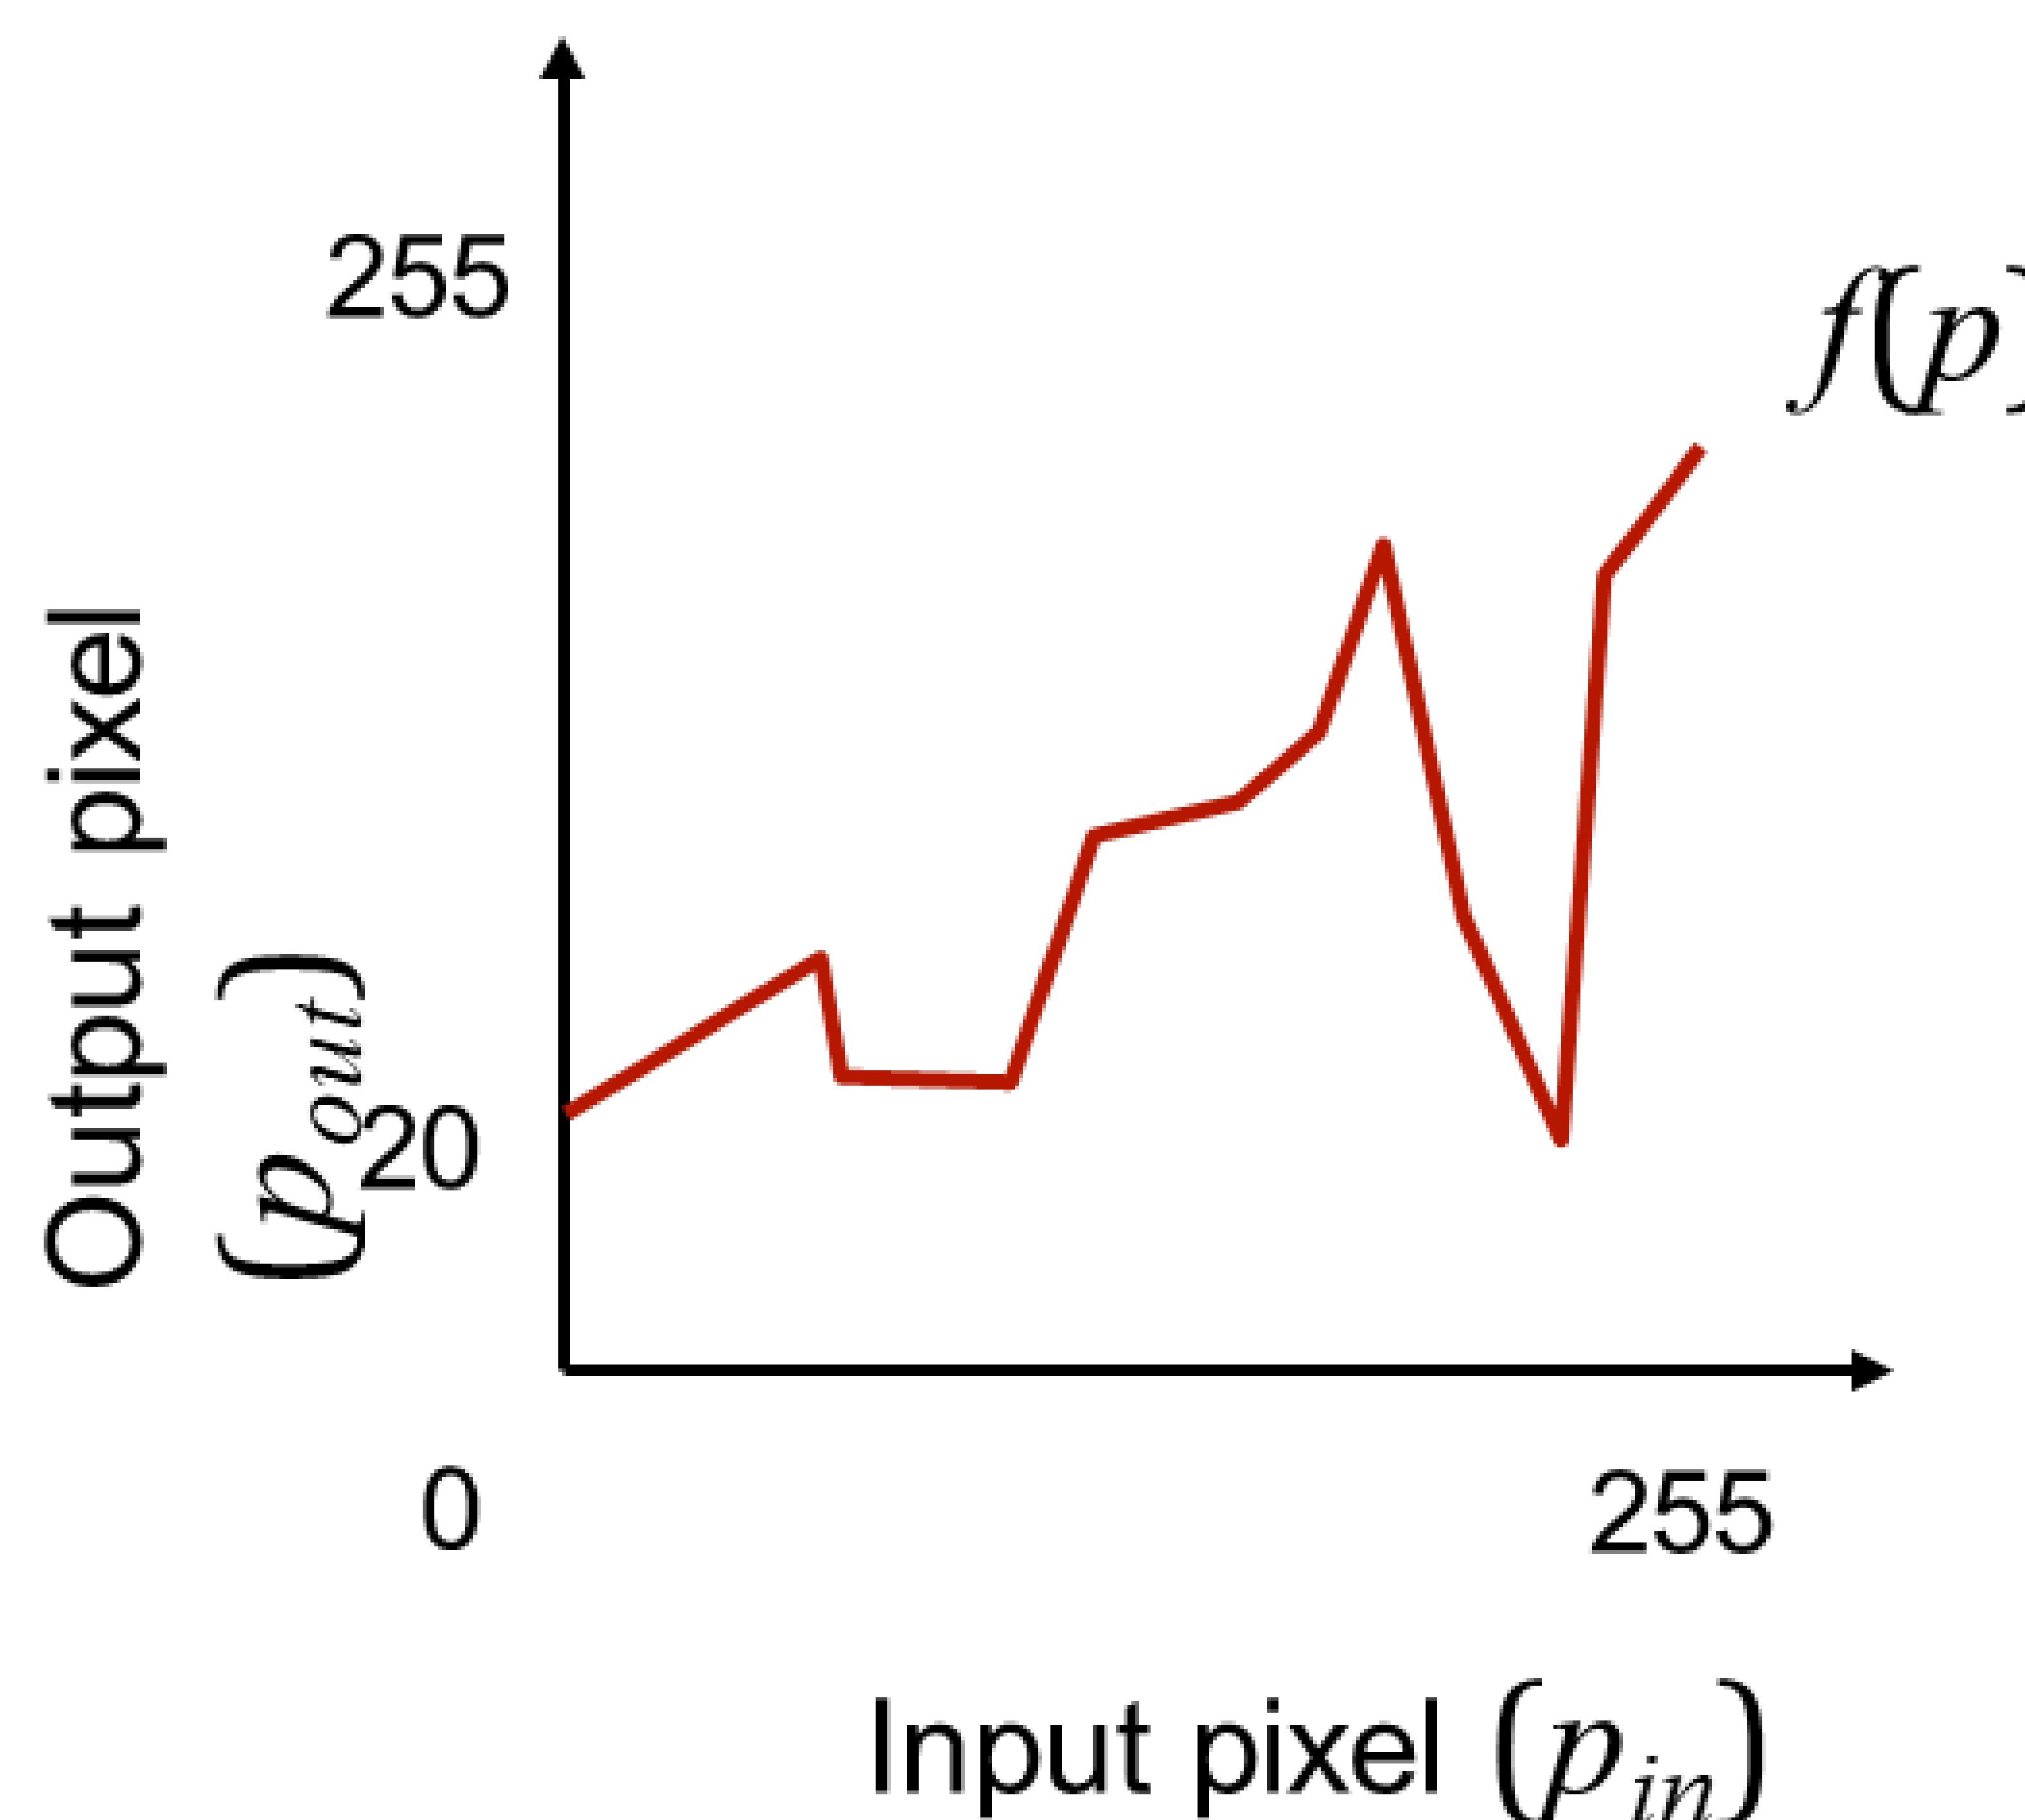


```
plt.subplot(121)
plt.imshow(image, cmap='gray', vmin=0, vmax=255)
plt.axis('off')
plt.subplot(122)
plt.imshow(image - 100, cmap='gray', vmin=0, vmax=255)
plt.axis('off')
```

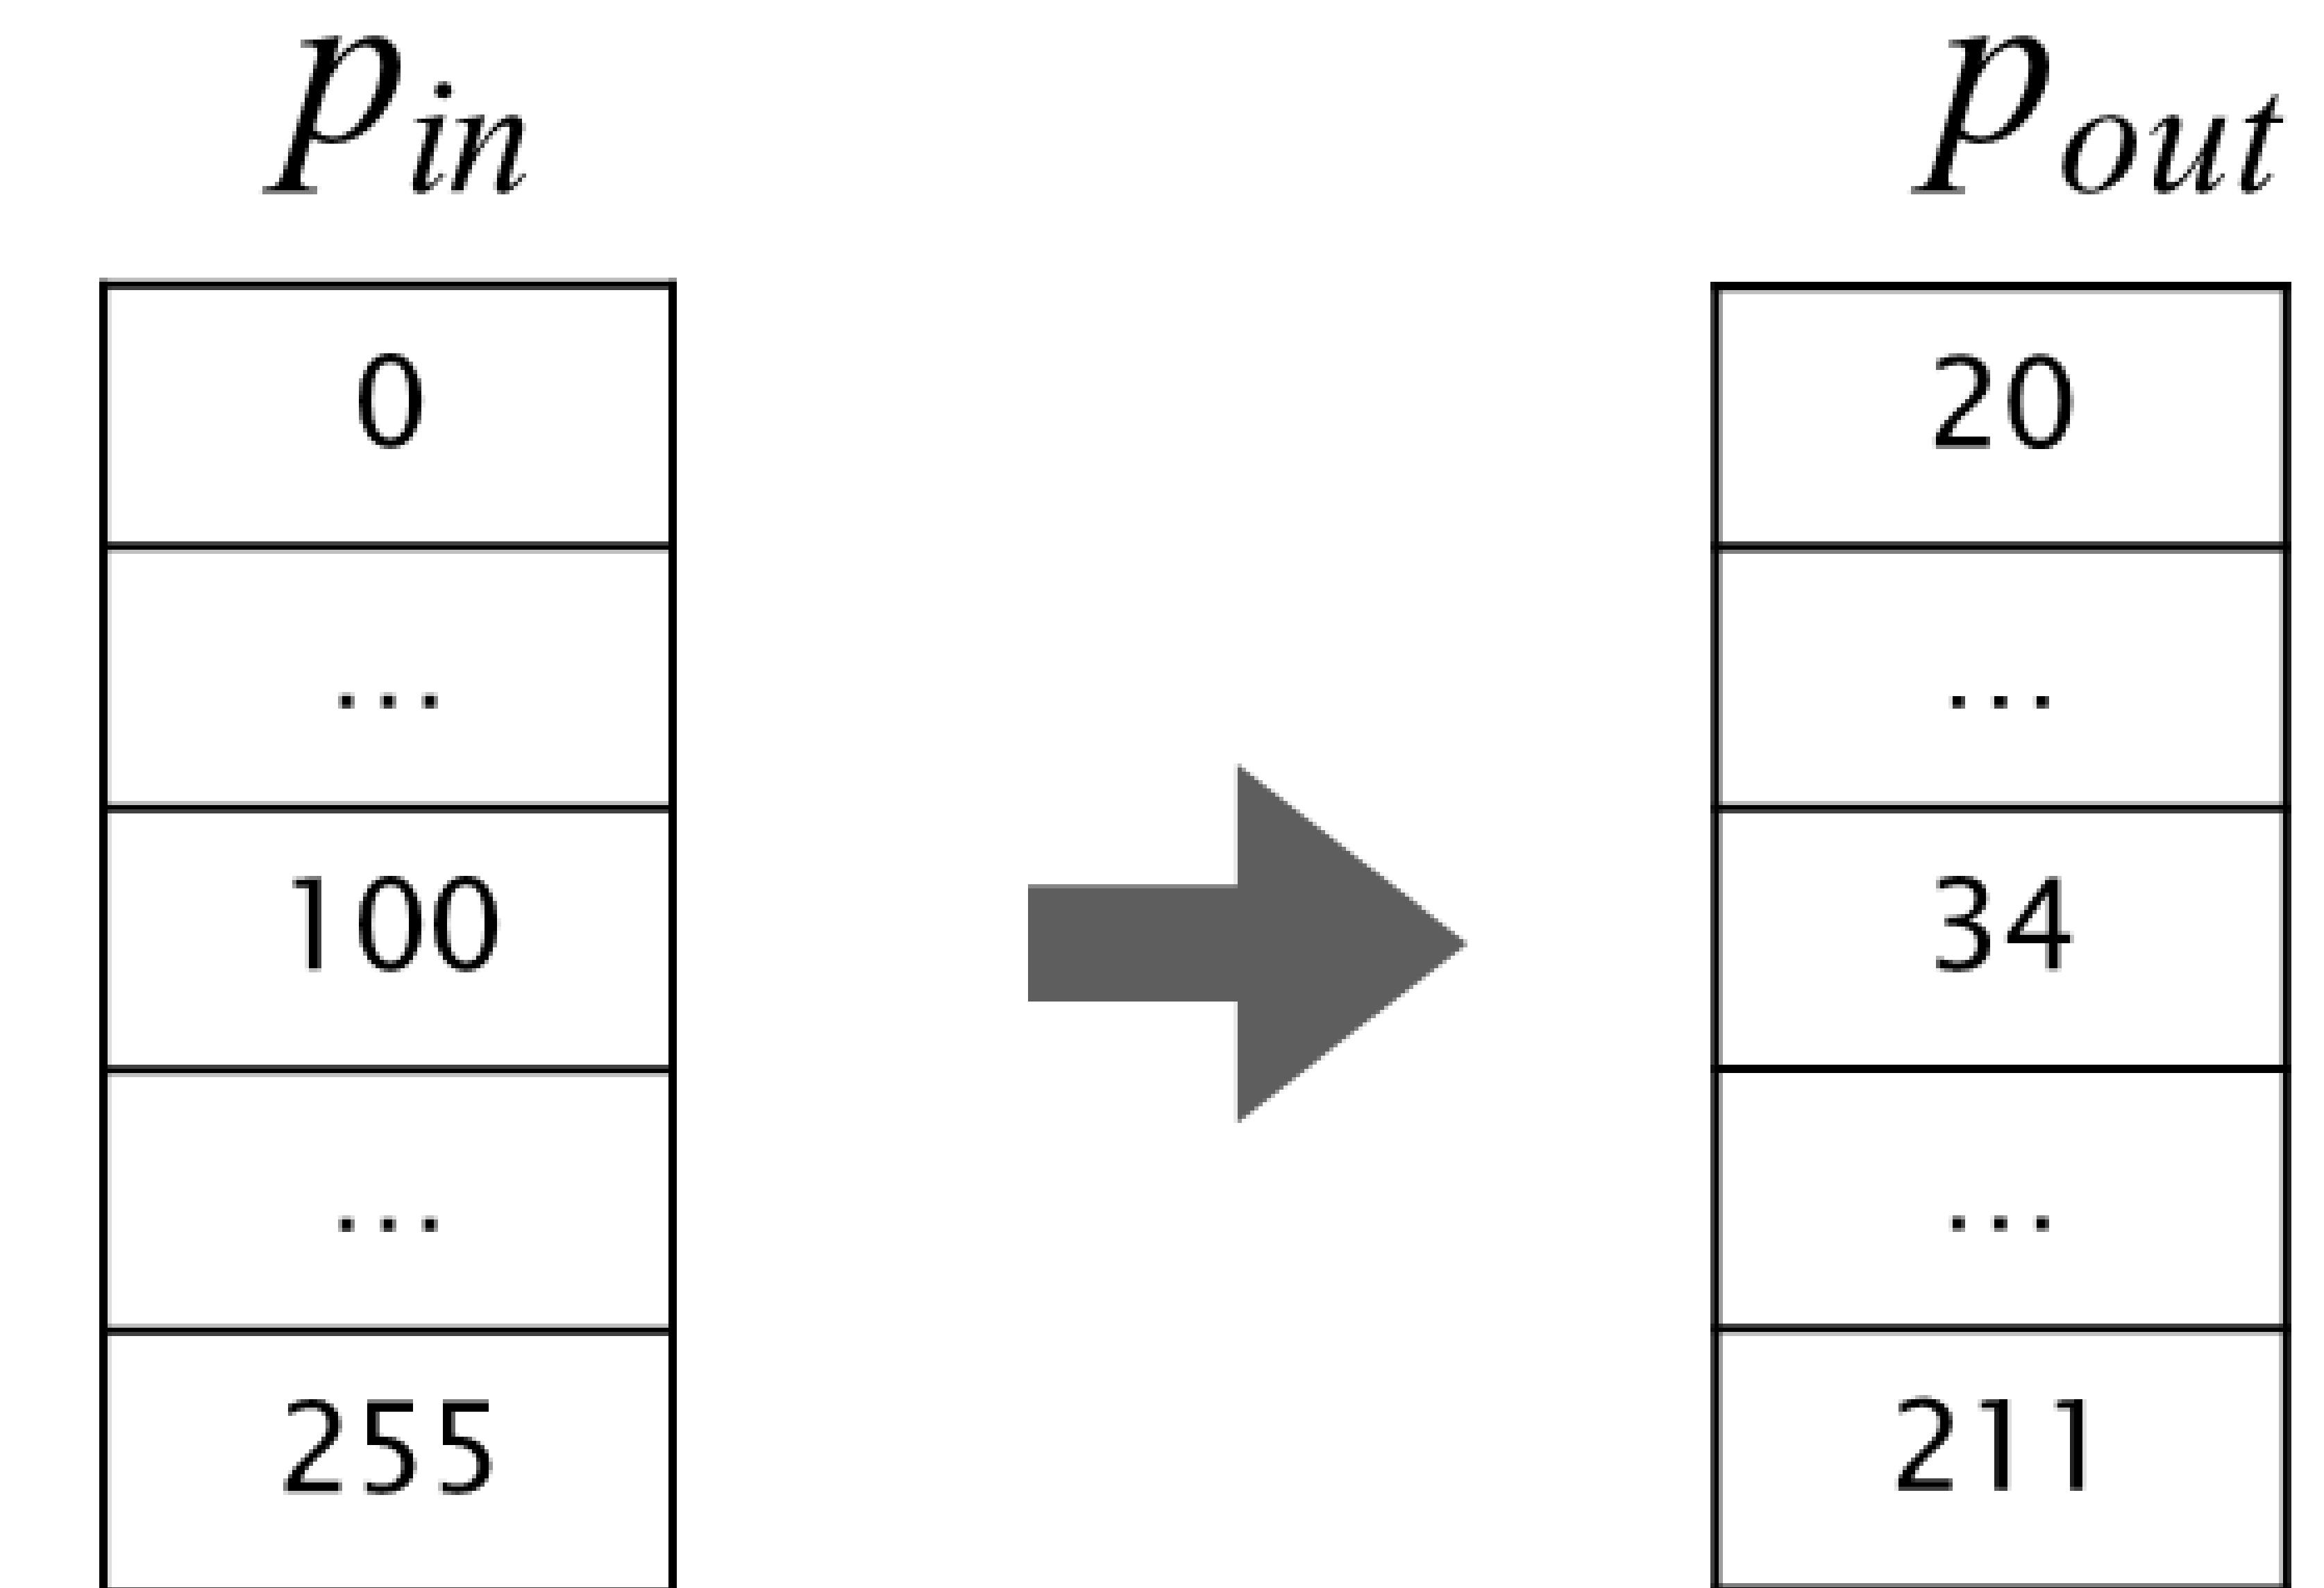


Linear transfer functions

Visualization

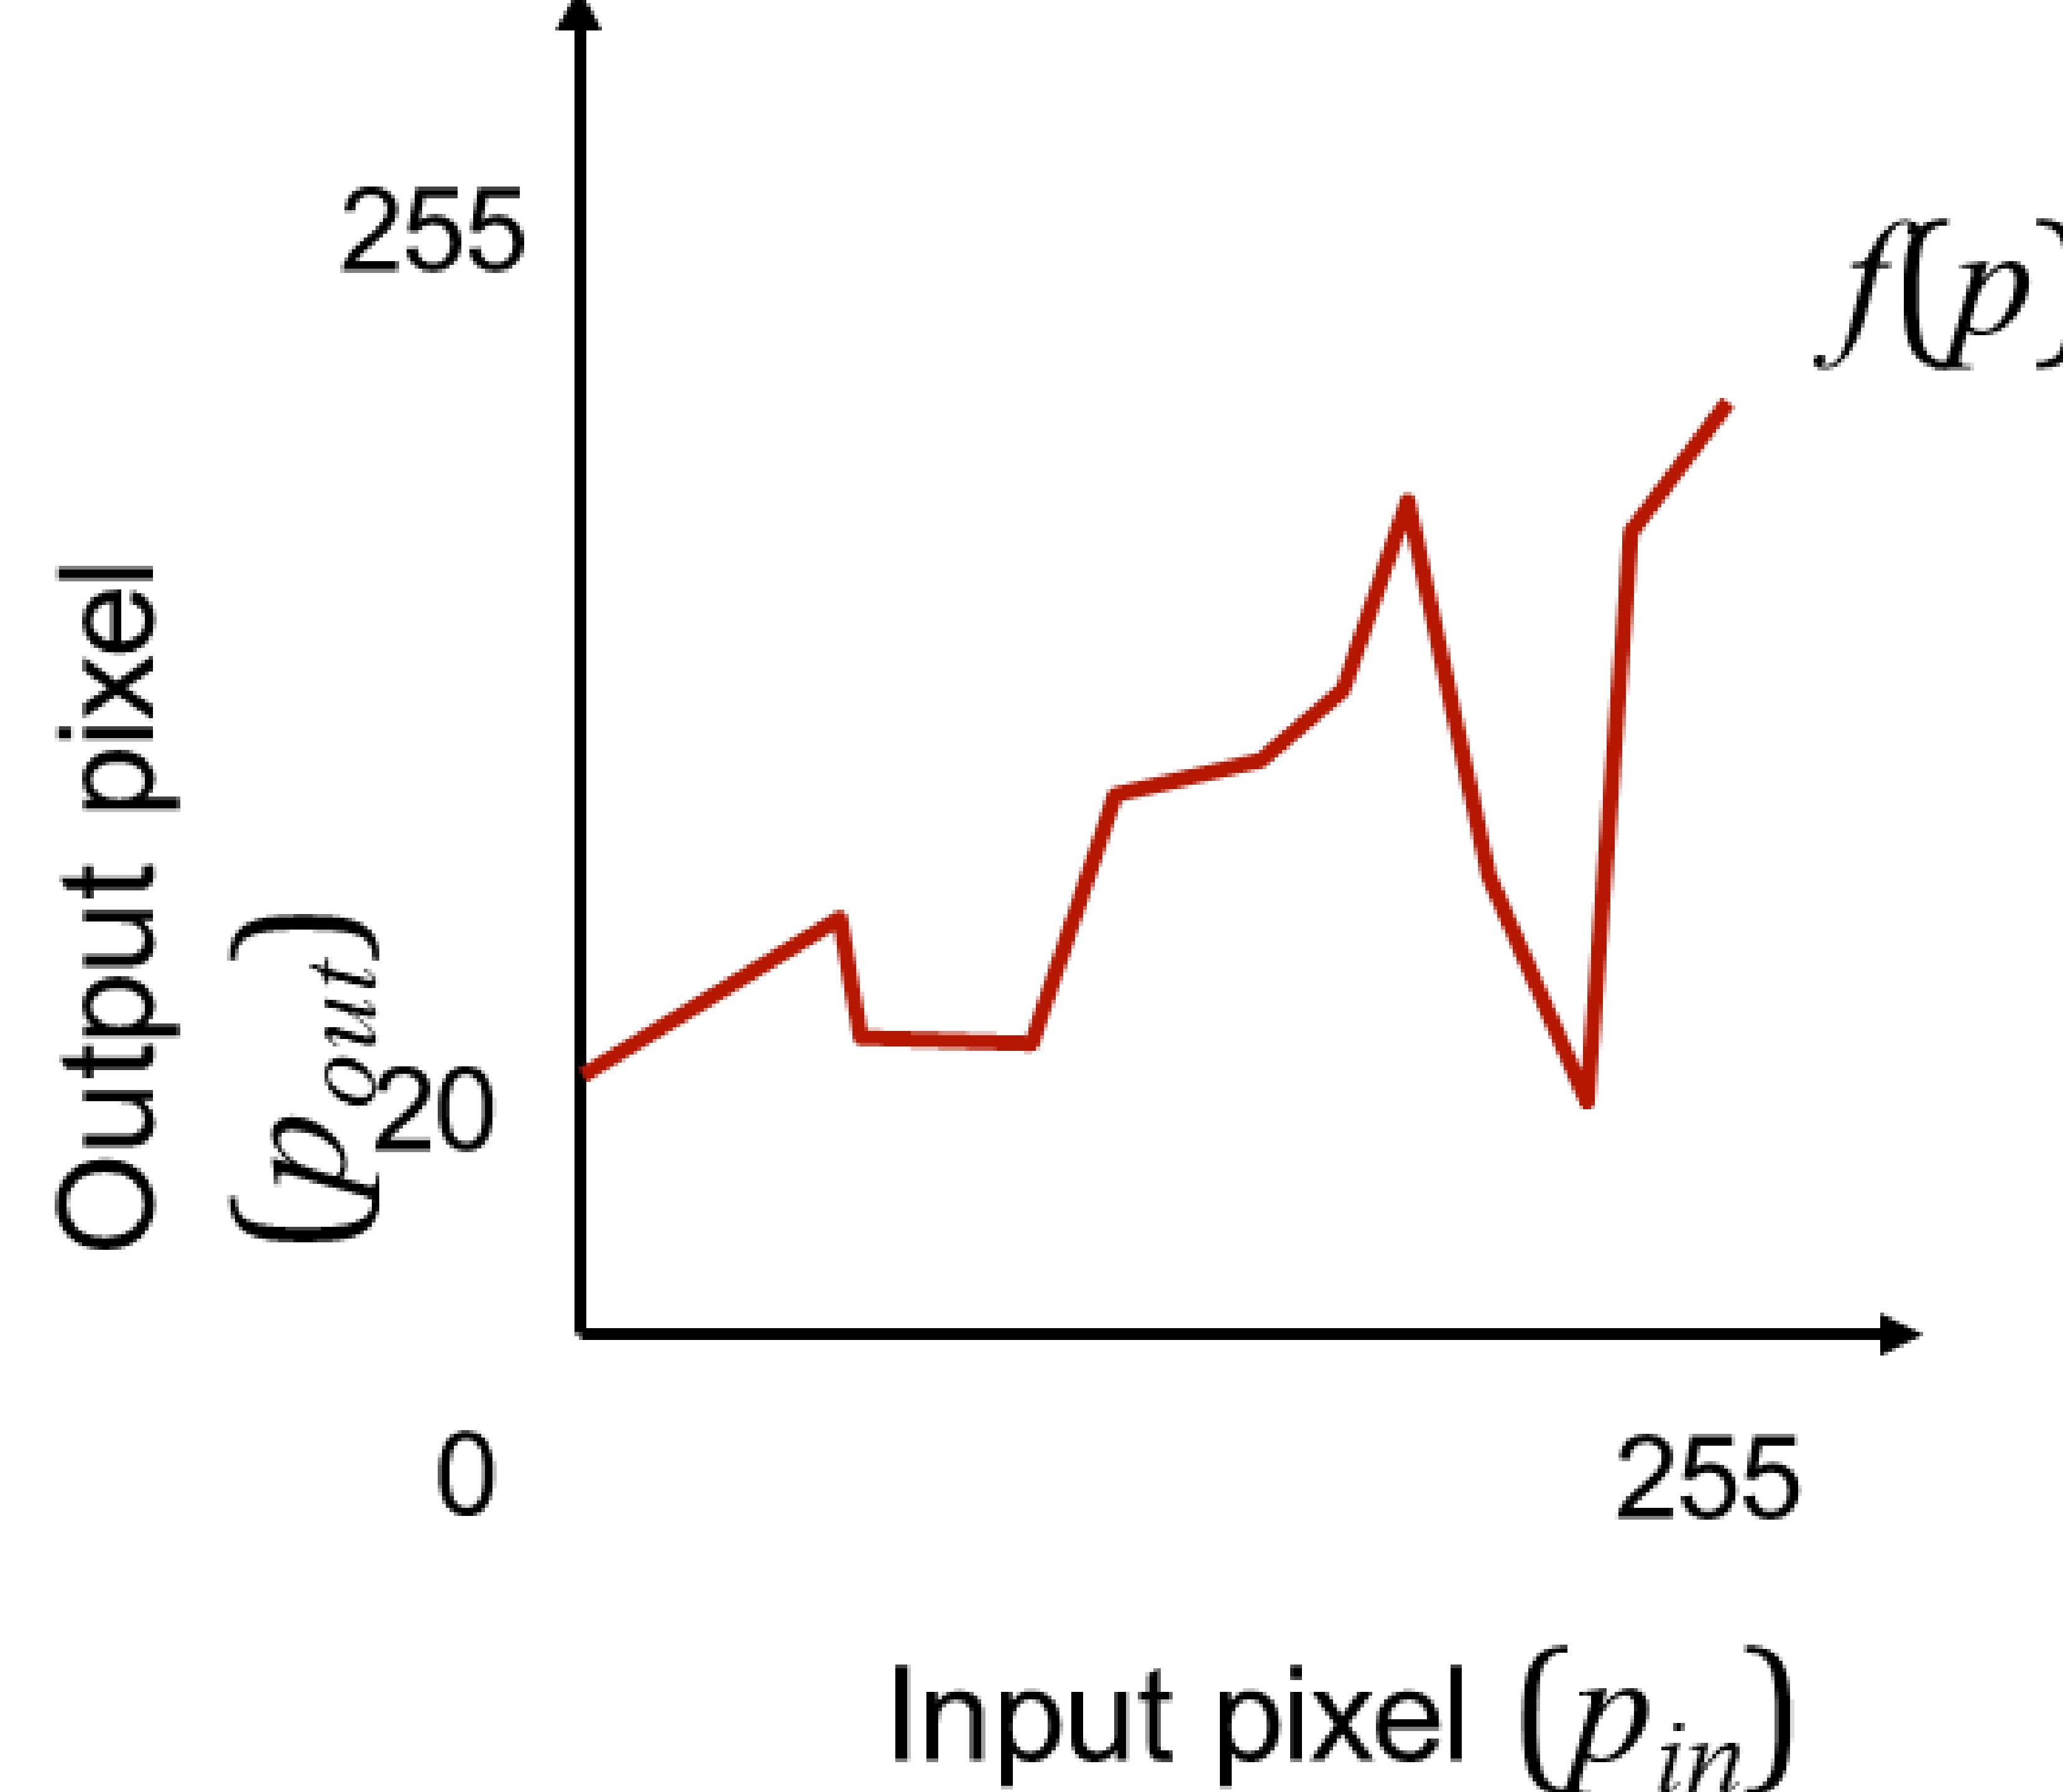


Example: an arbitrary function

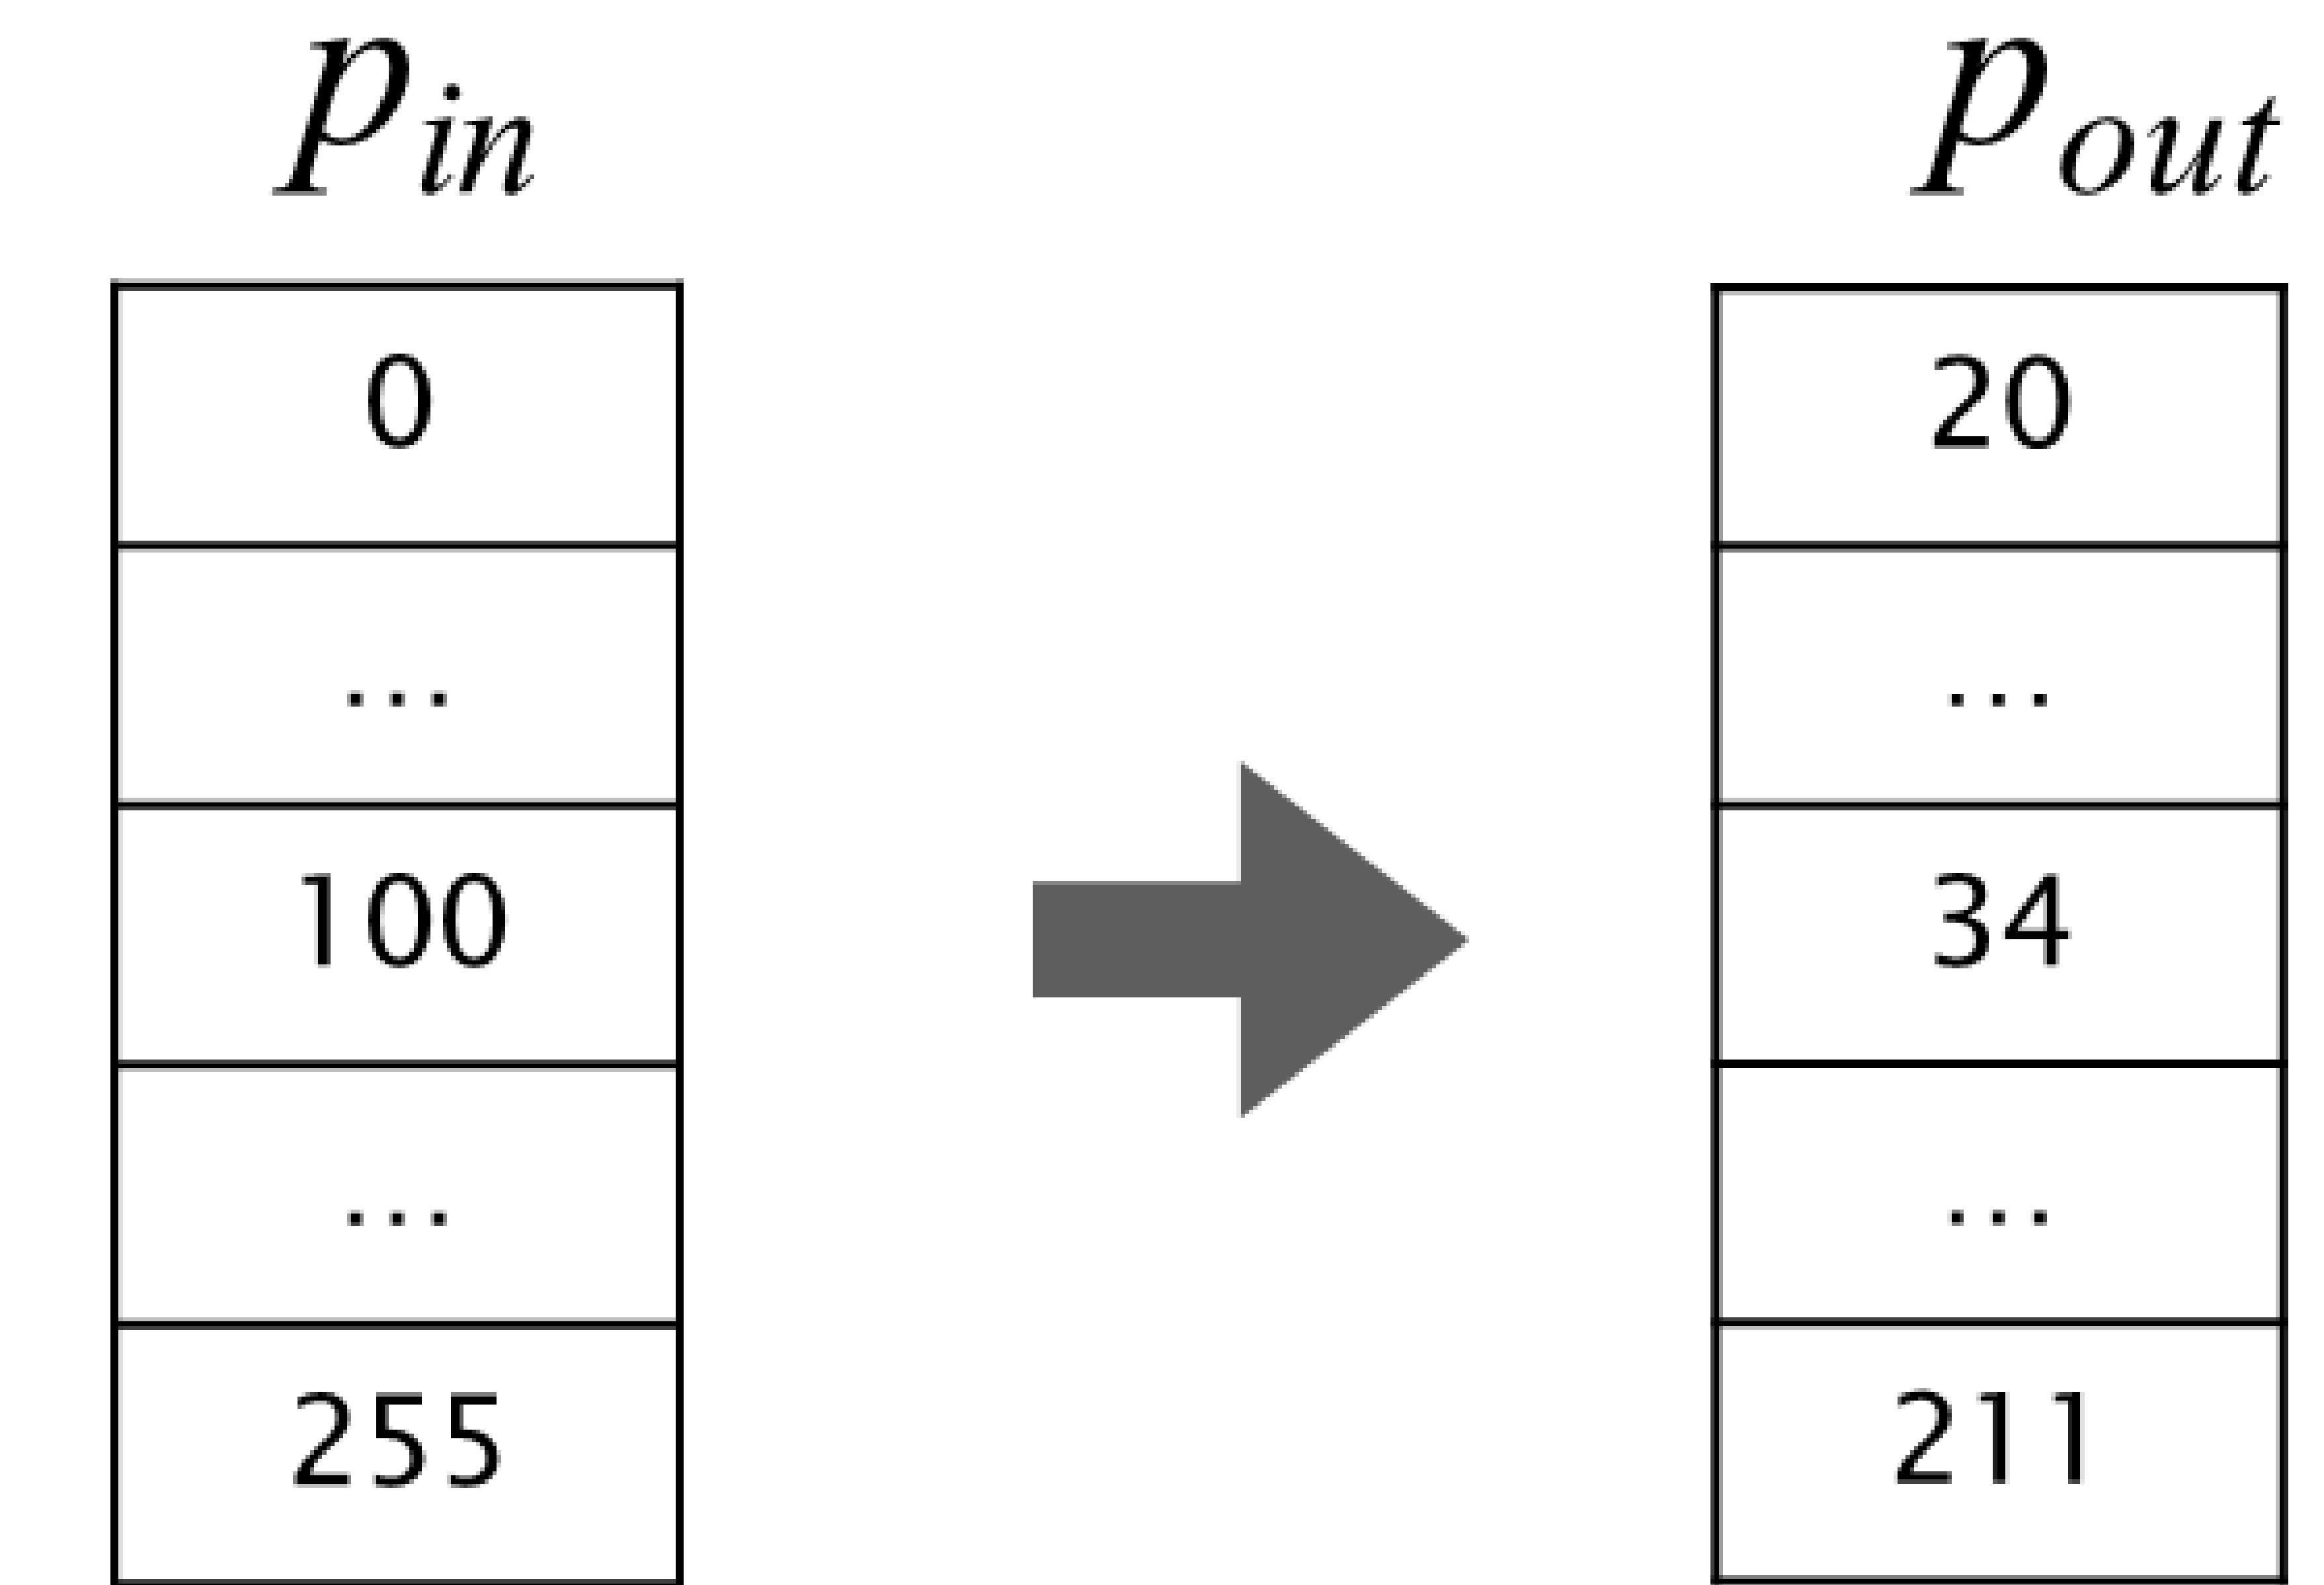


Nonlinear transfer functions

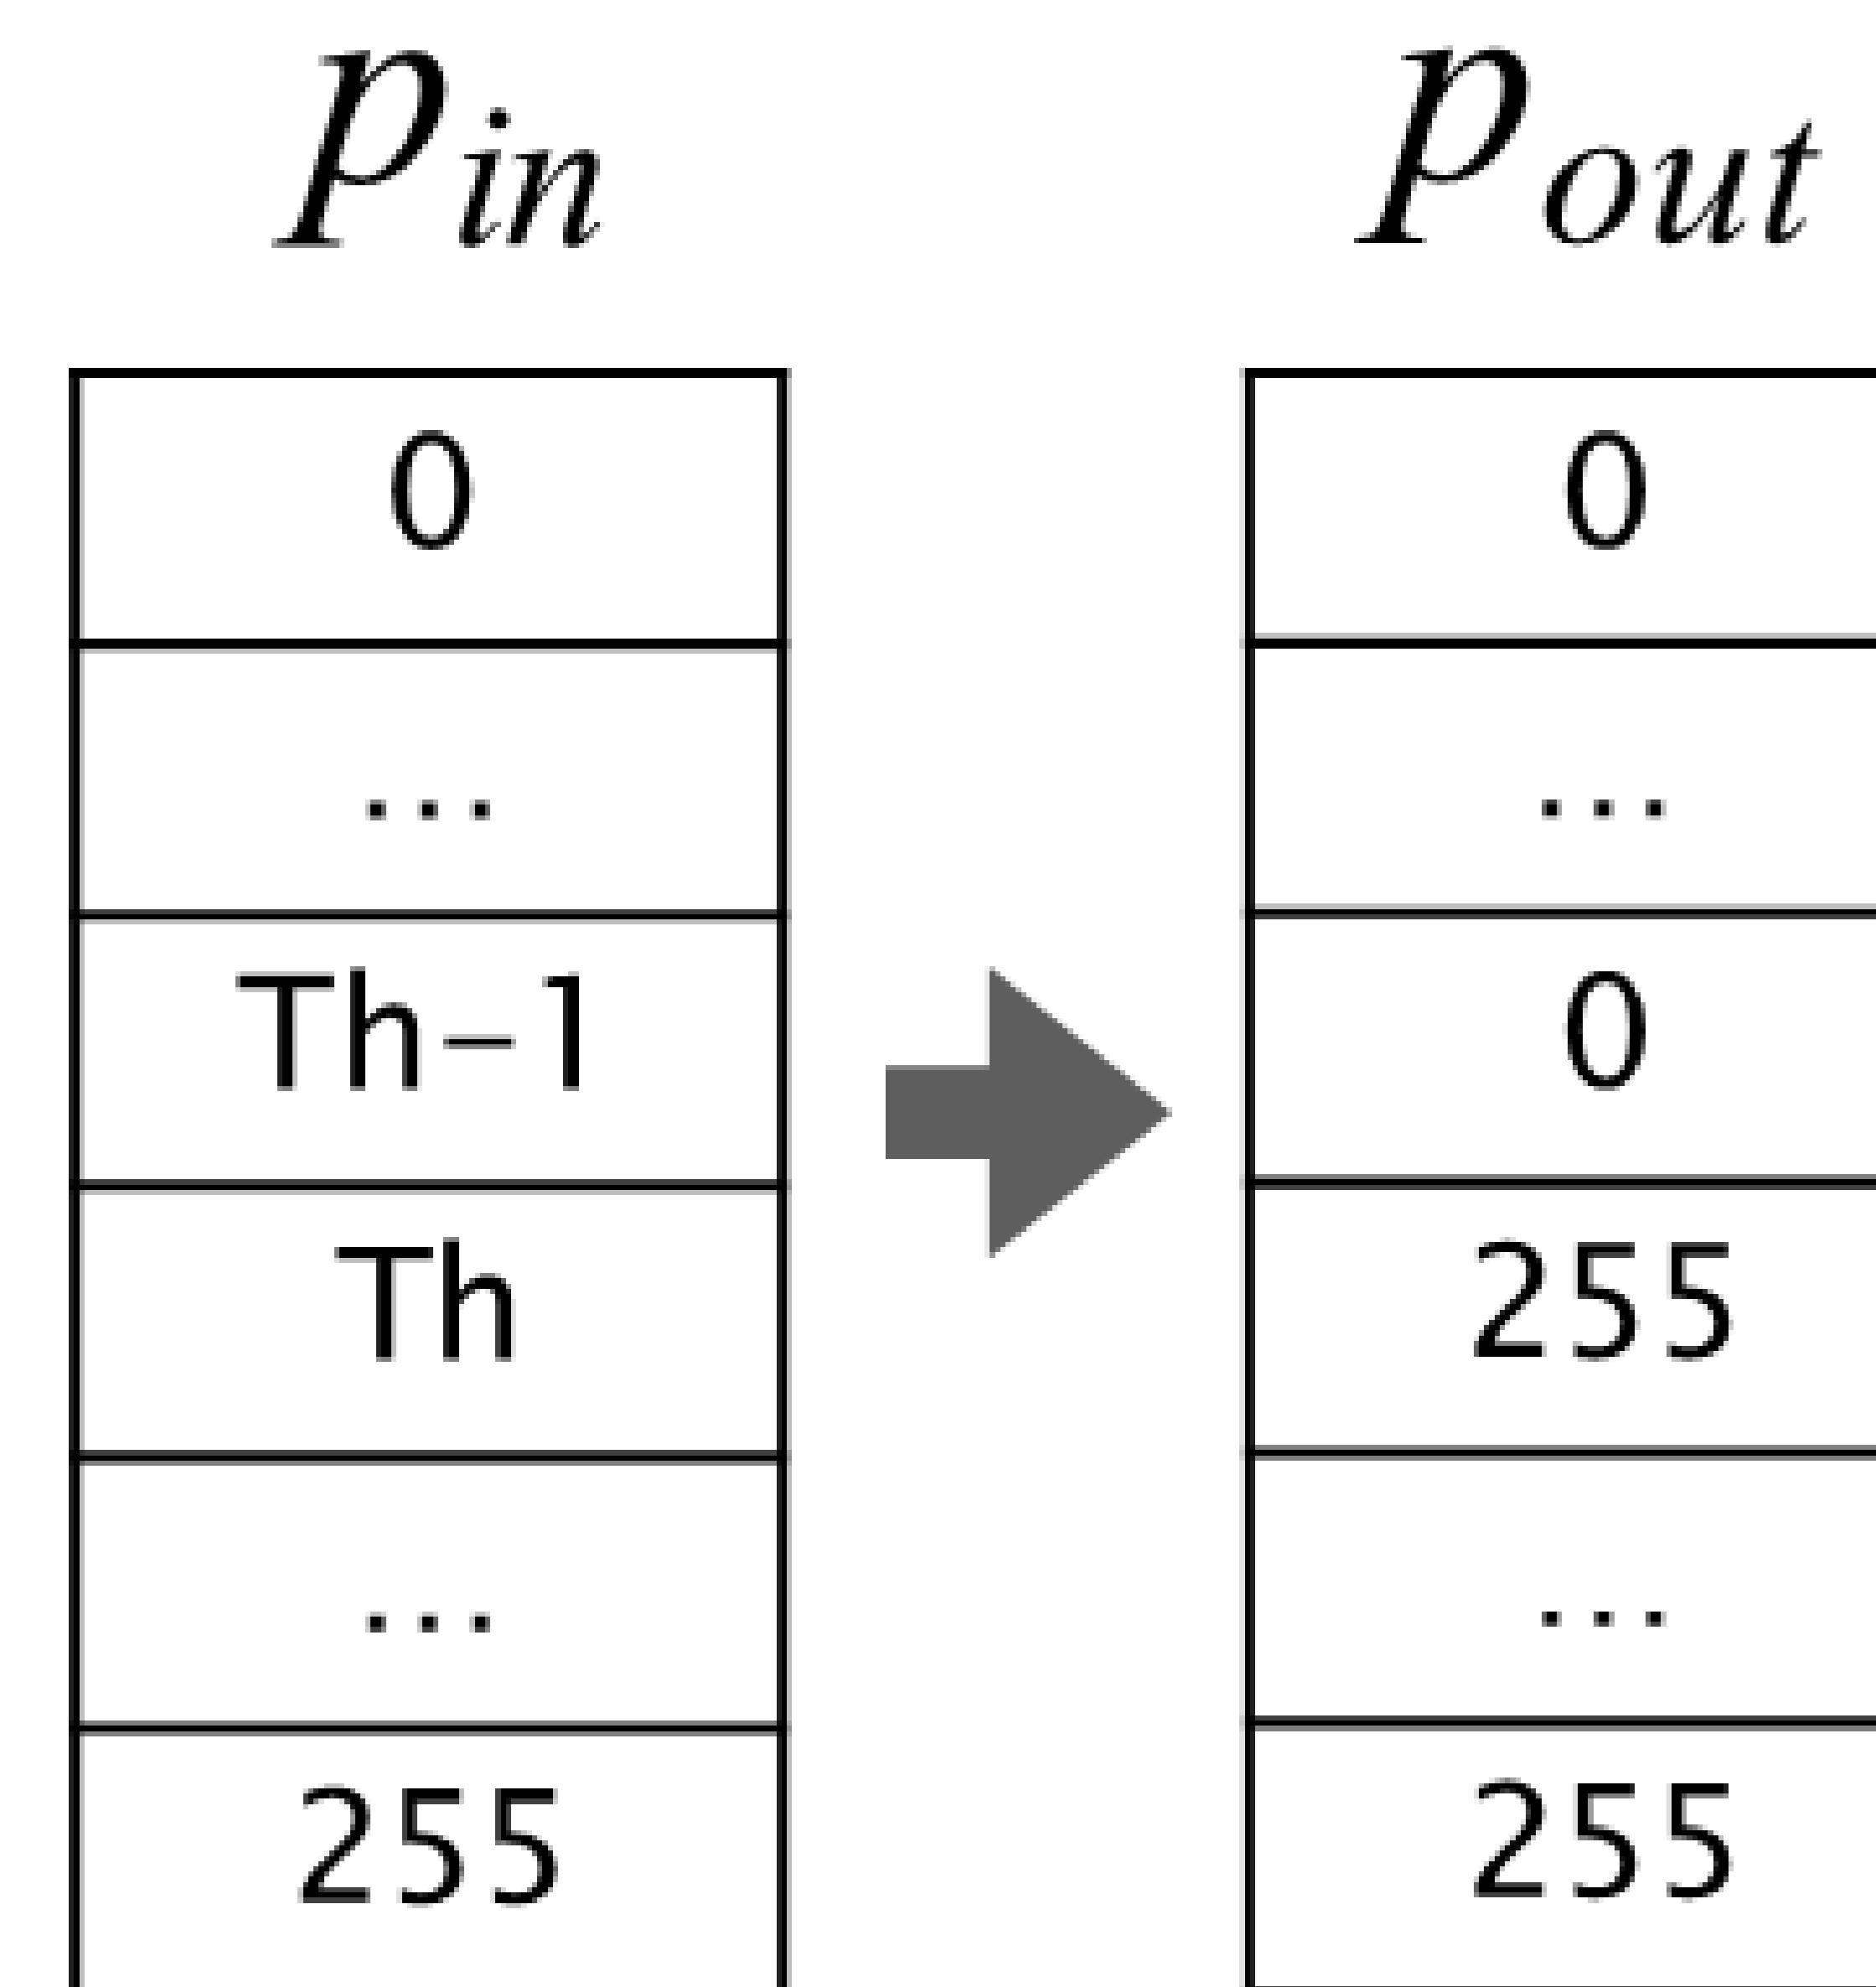
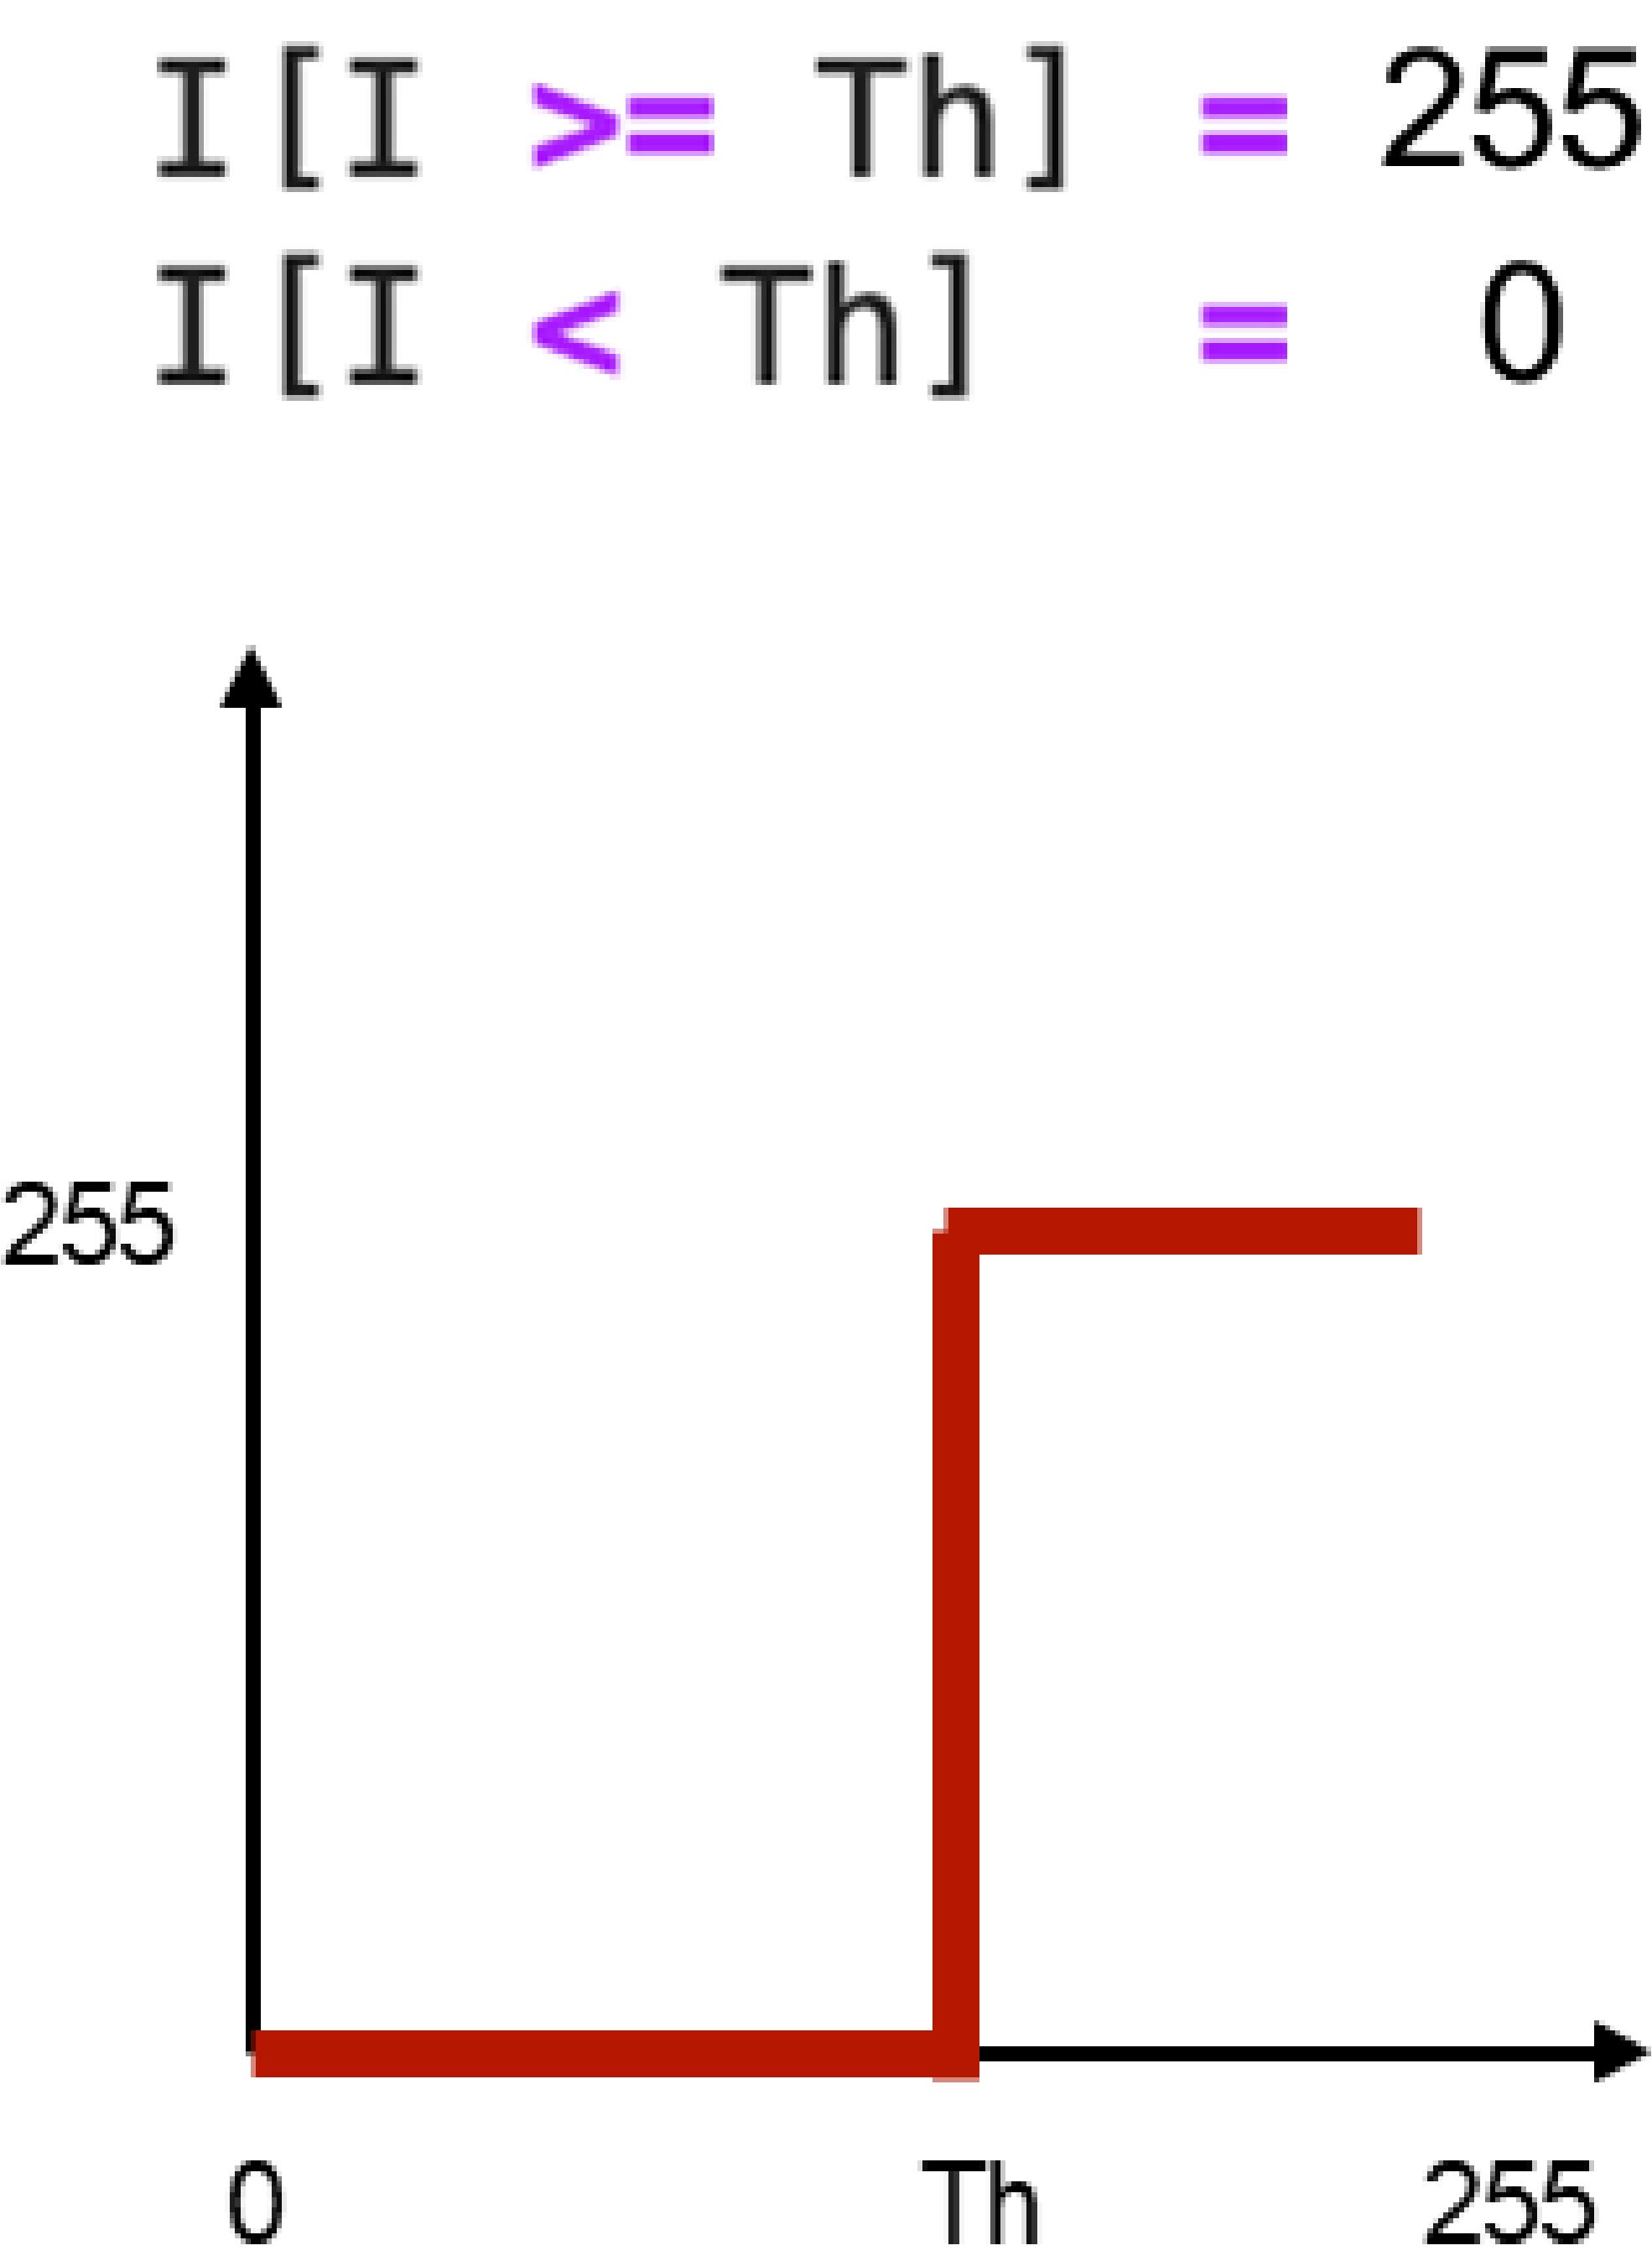
Visualization



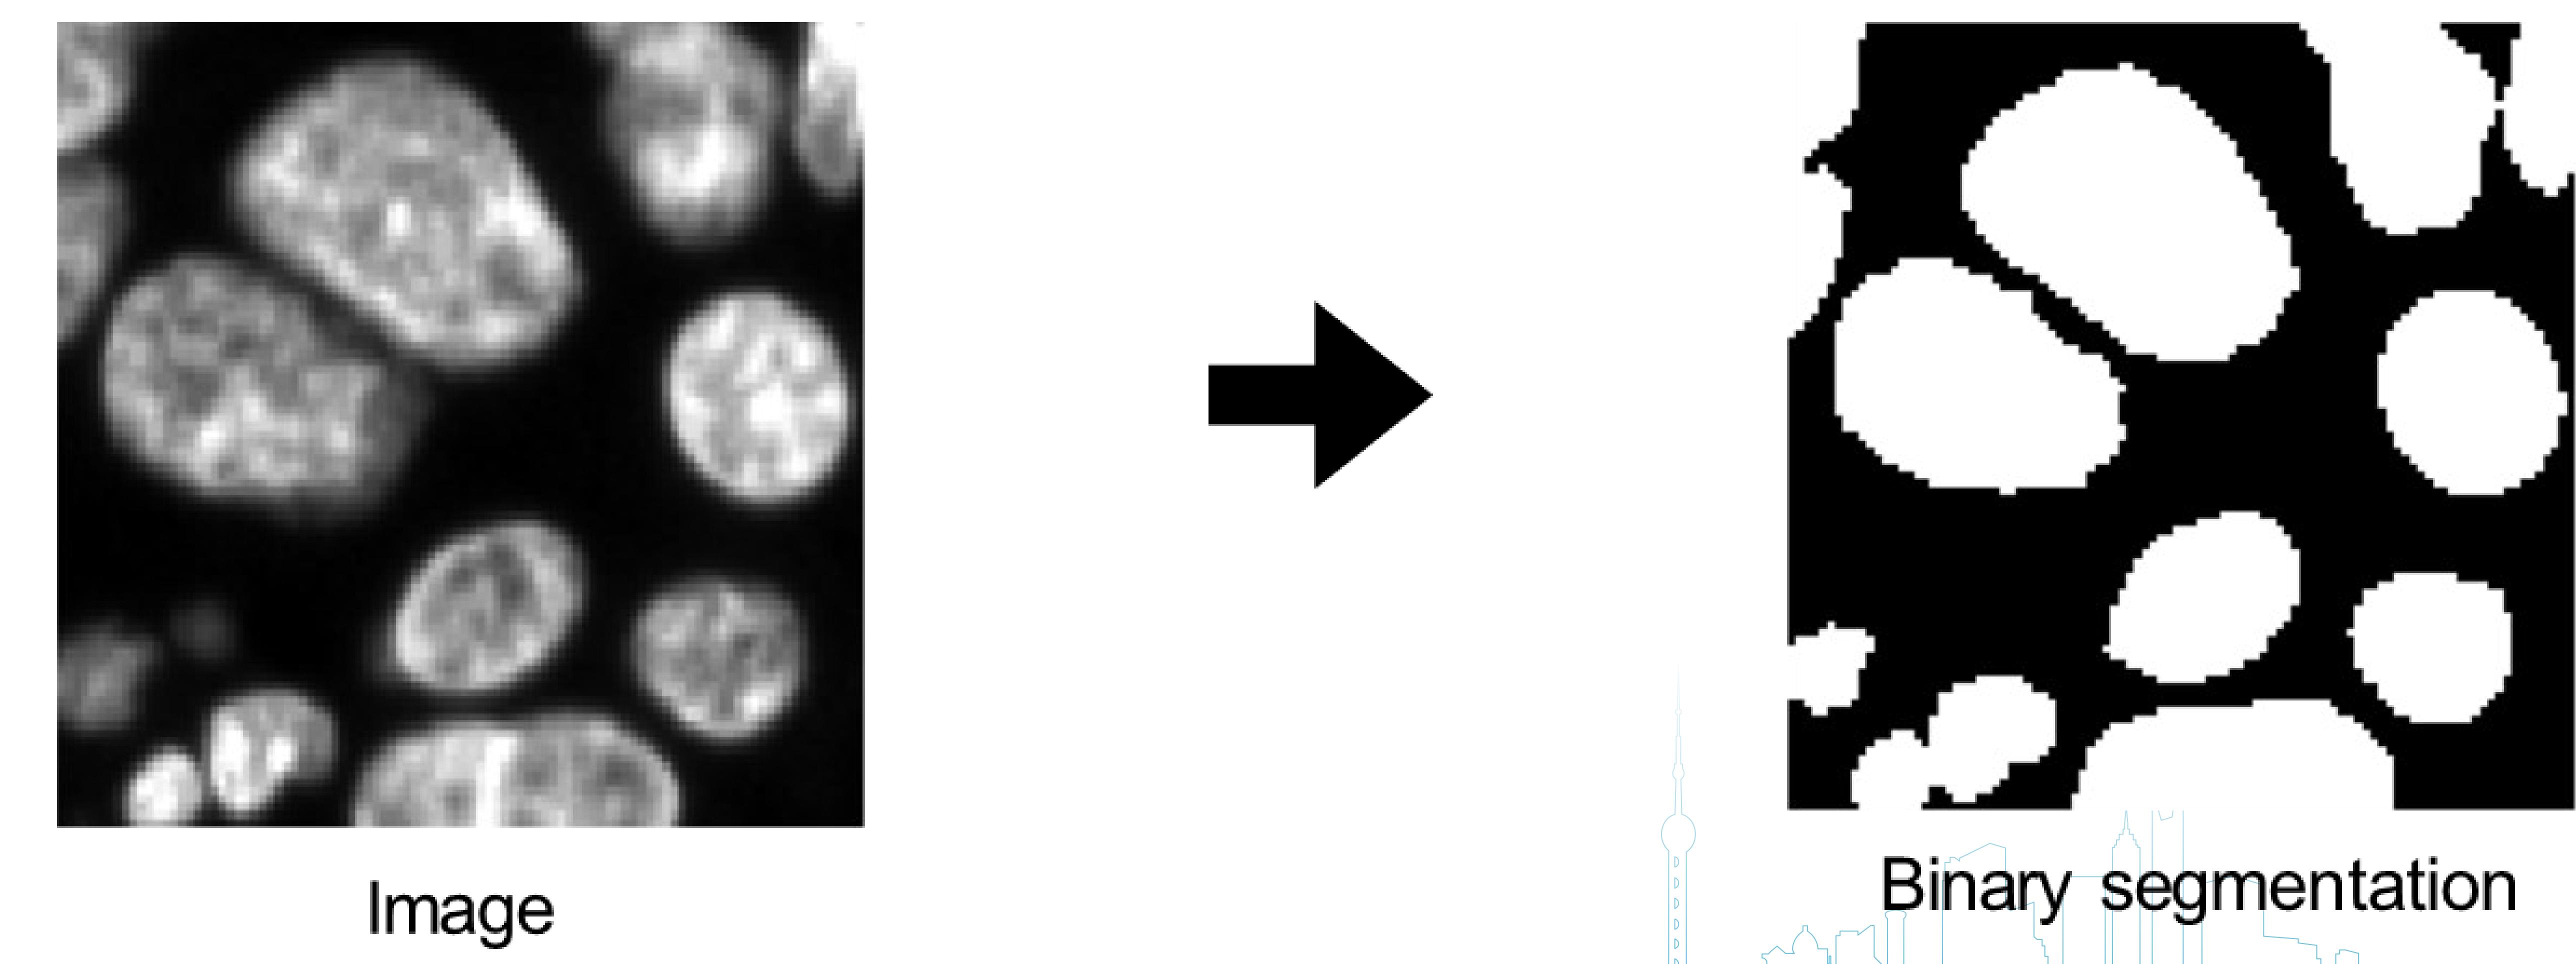
Example: an arbitrary function



Application: image thresholding



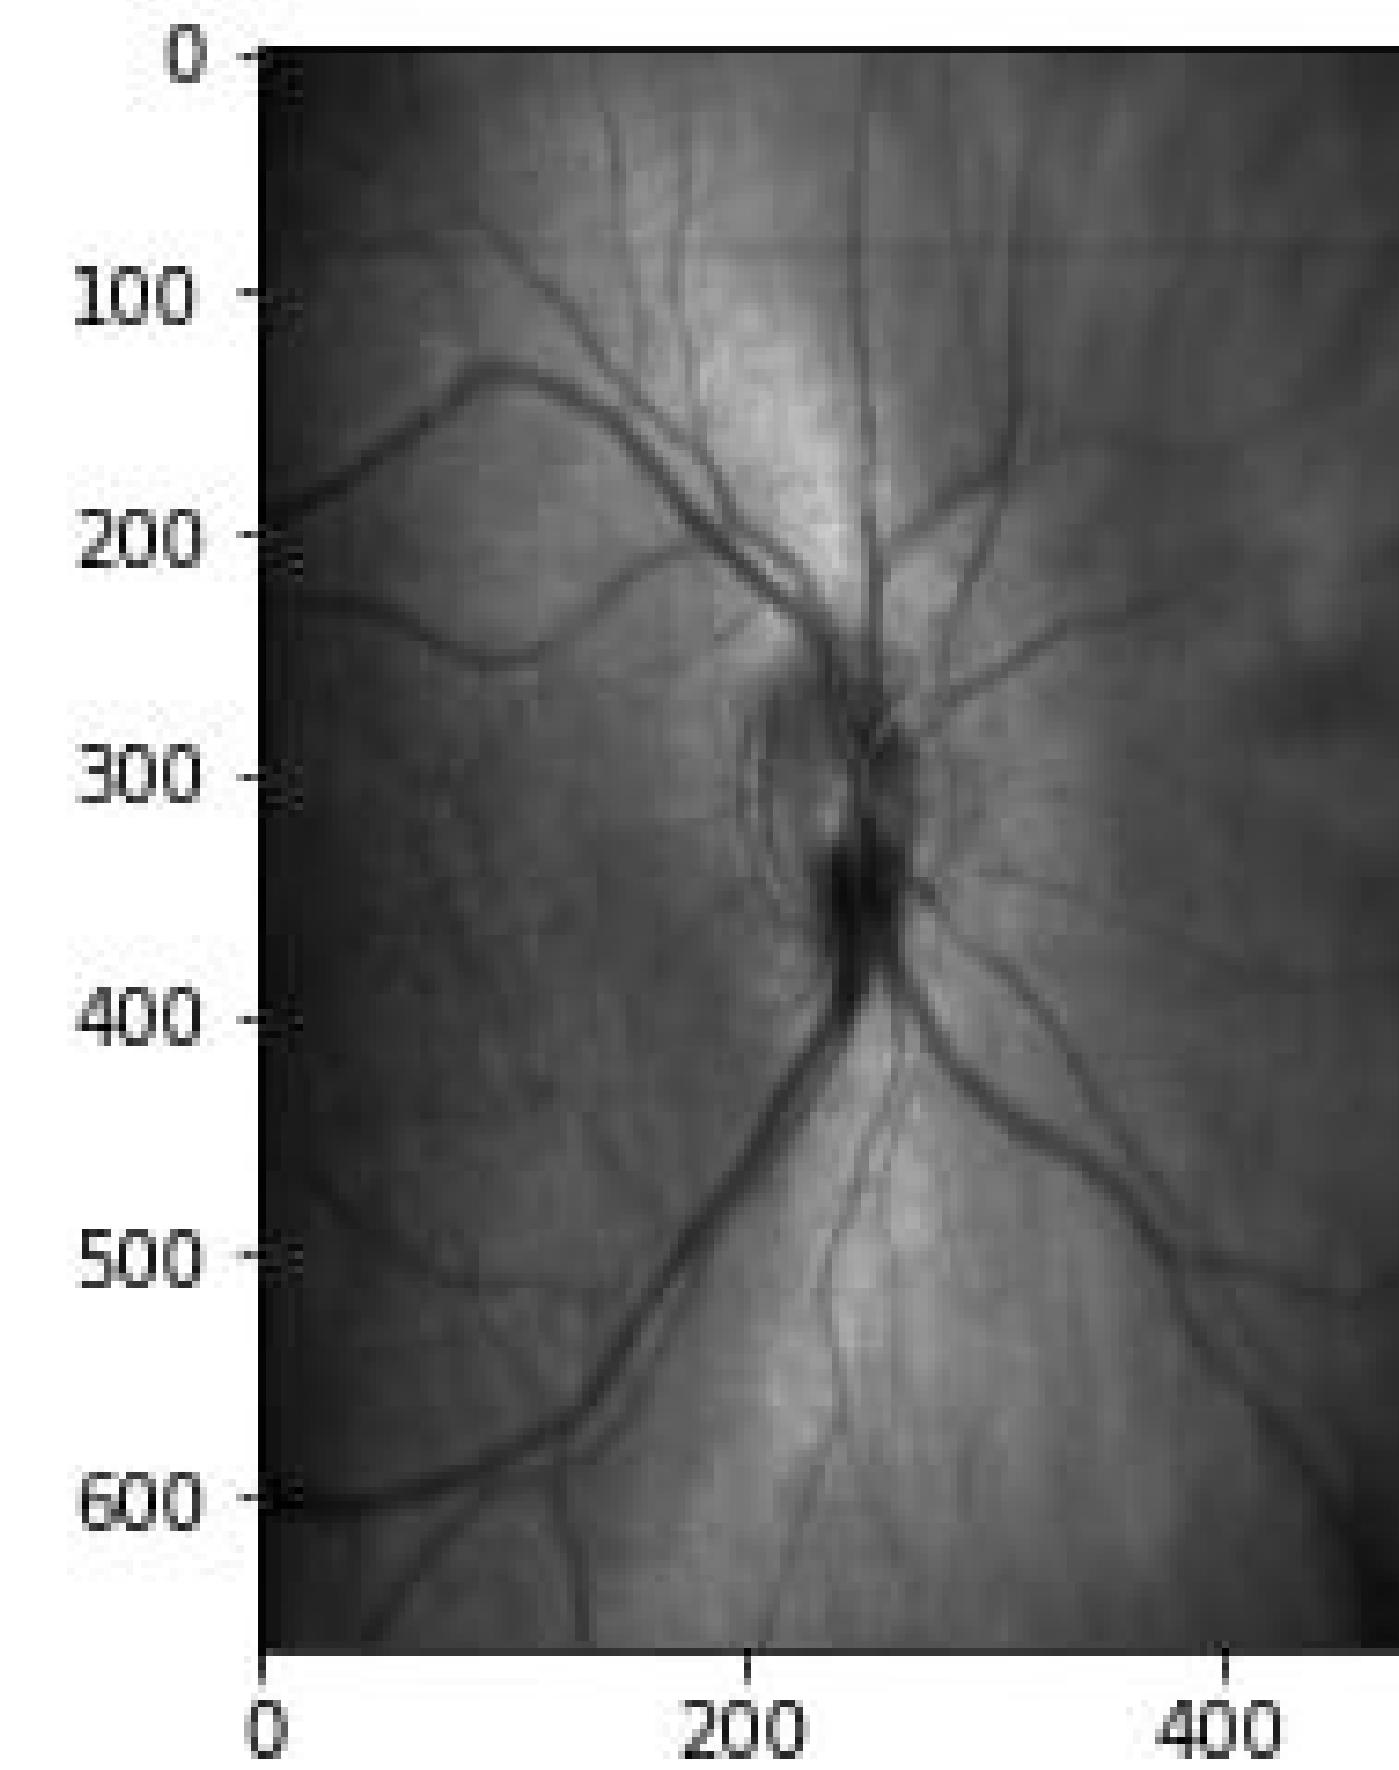
Binary segmentation:
Find the threshold



Application: image enhancement

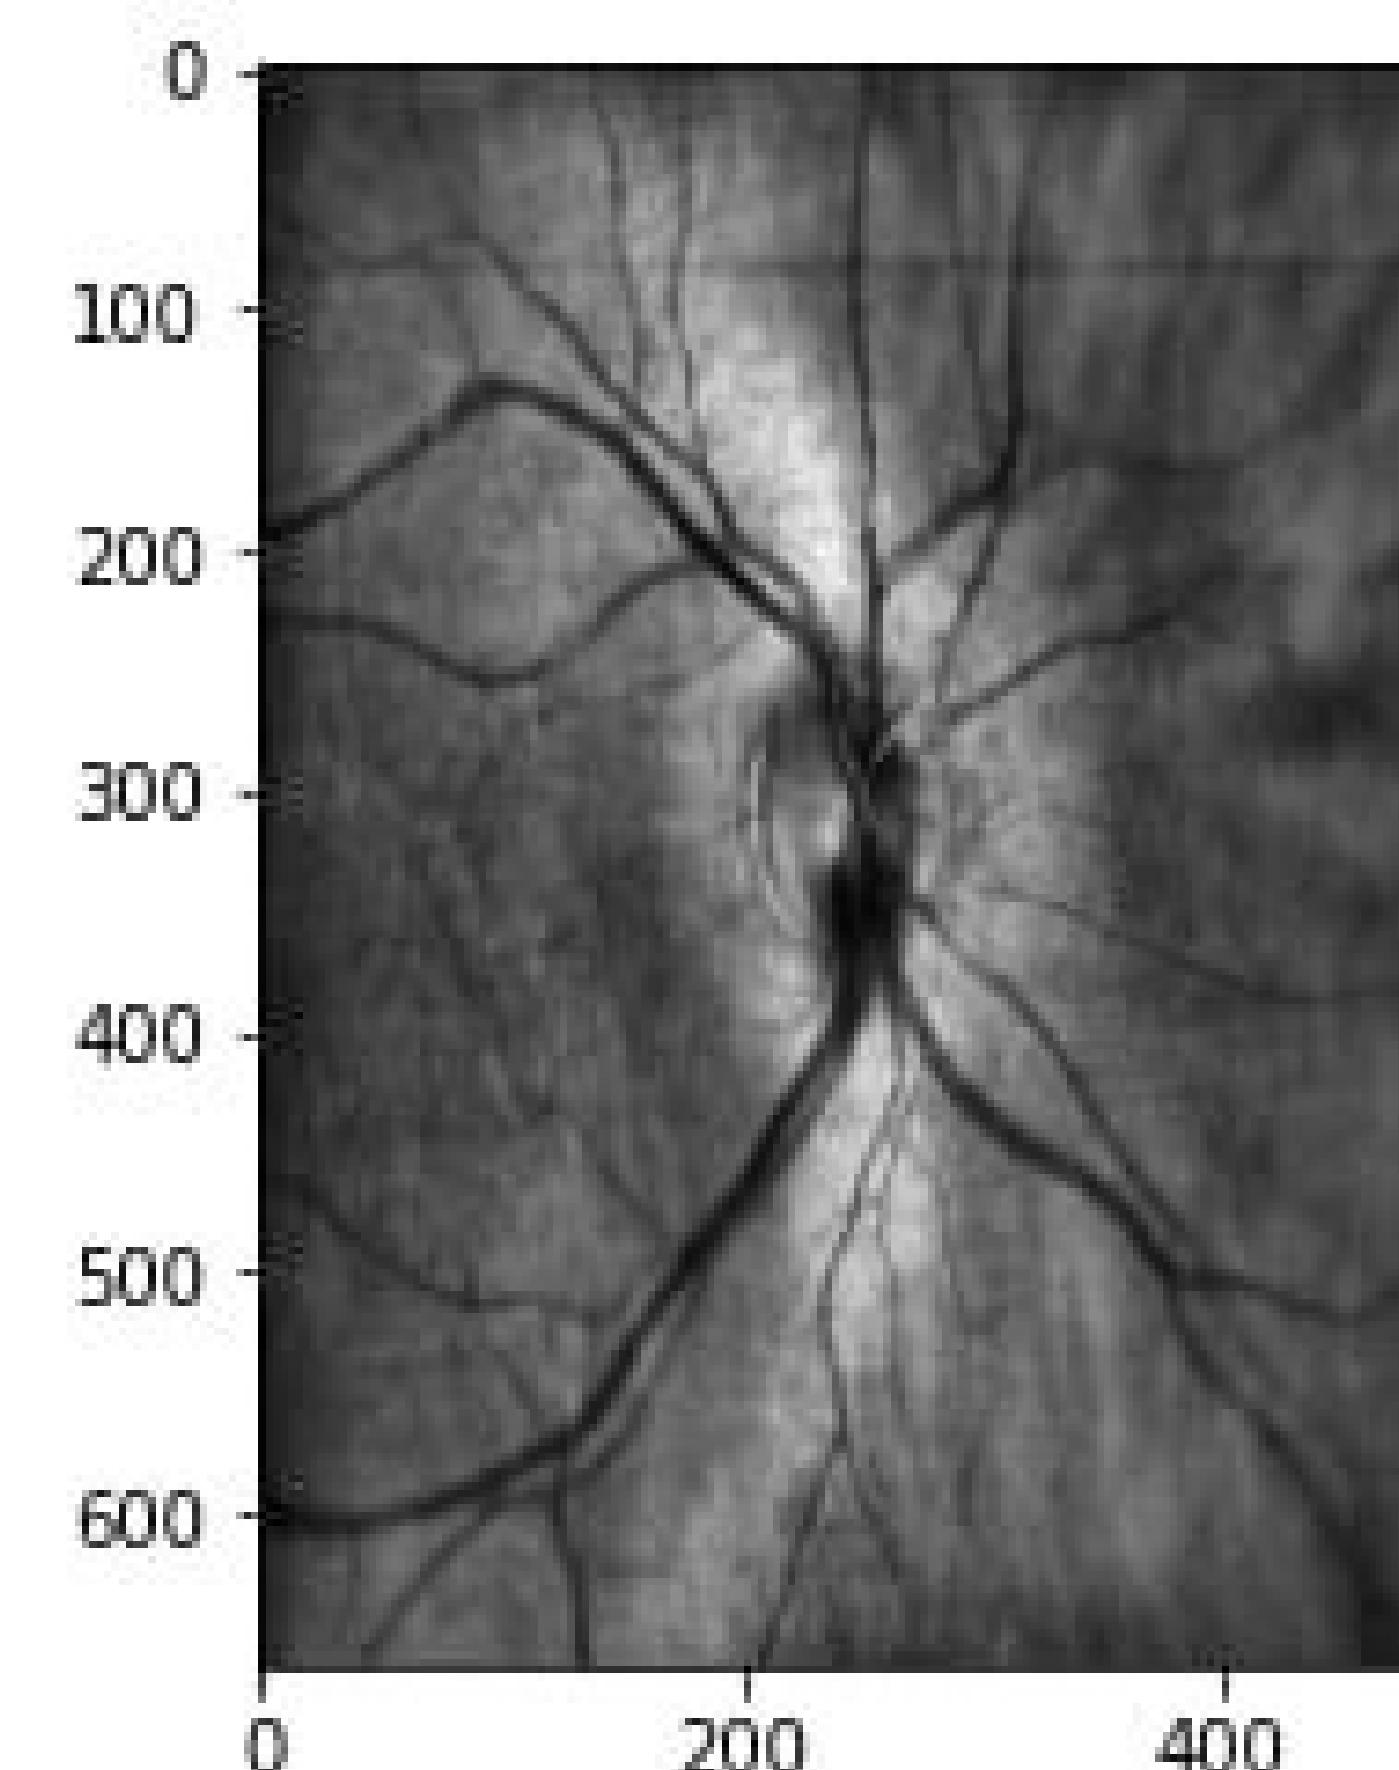
```
In [107]: 1 plt.imshow(dataf['slo_fundus'], cmap='gray')
```

```
Out[107]: <matplotlib.image.AxesImage at 0x7f64260349a0>
```



```
In [108]: 1 clahe = cv2.createCLAHE(clipLimit=4.0, tileGridSize=(8,8))
2 fundus_slo_enhance = clahe.apply(dataf['slo_fundus'])
3
4 plt.imshow(fundus_slo_enhance, cmap='gray')
```

```
Out[108]: <matplotlib.image.AxesImage at 0x7f6425798fd0>
```



Patch-level processing

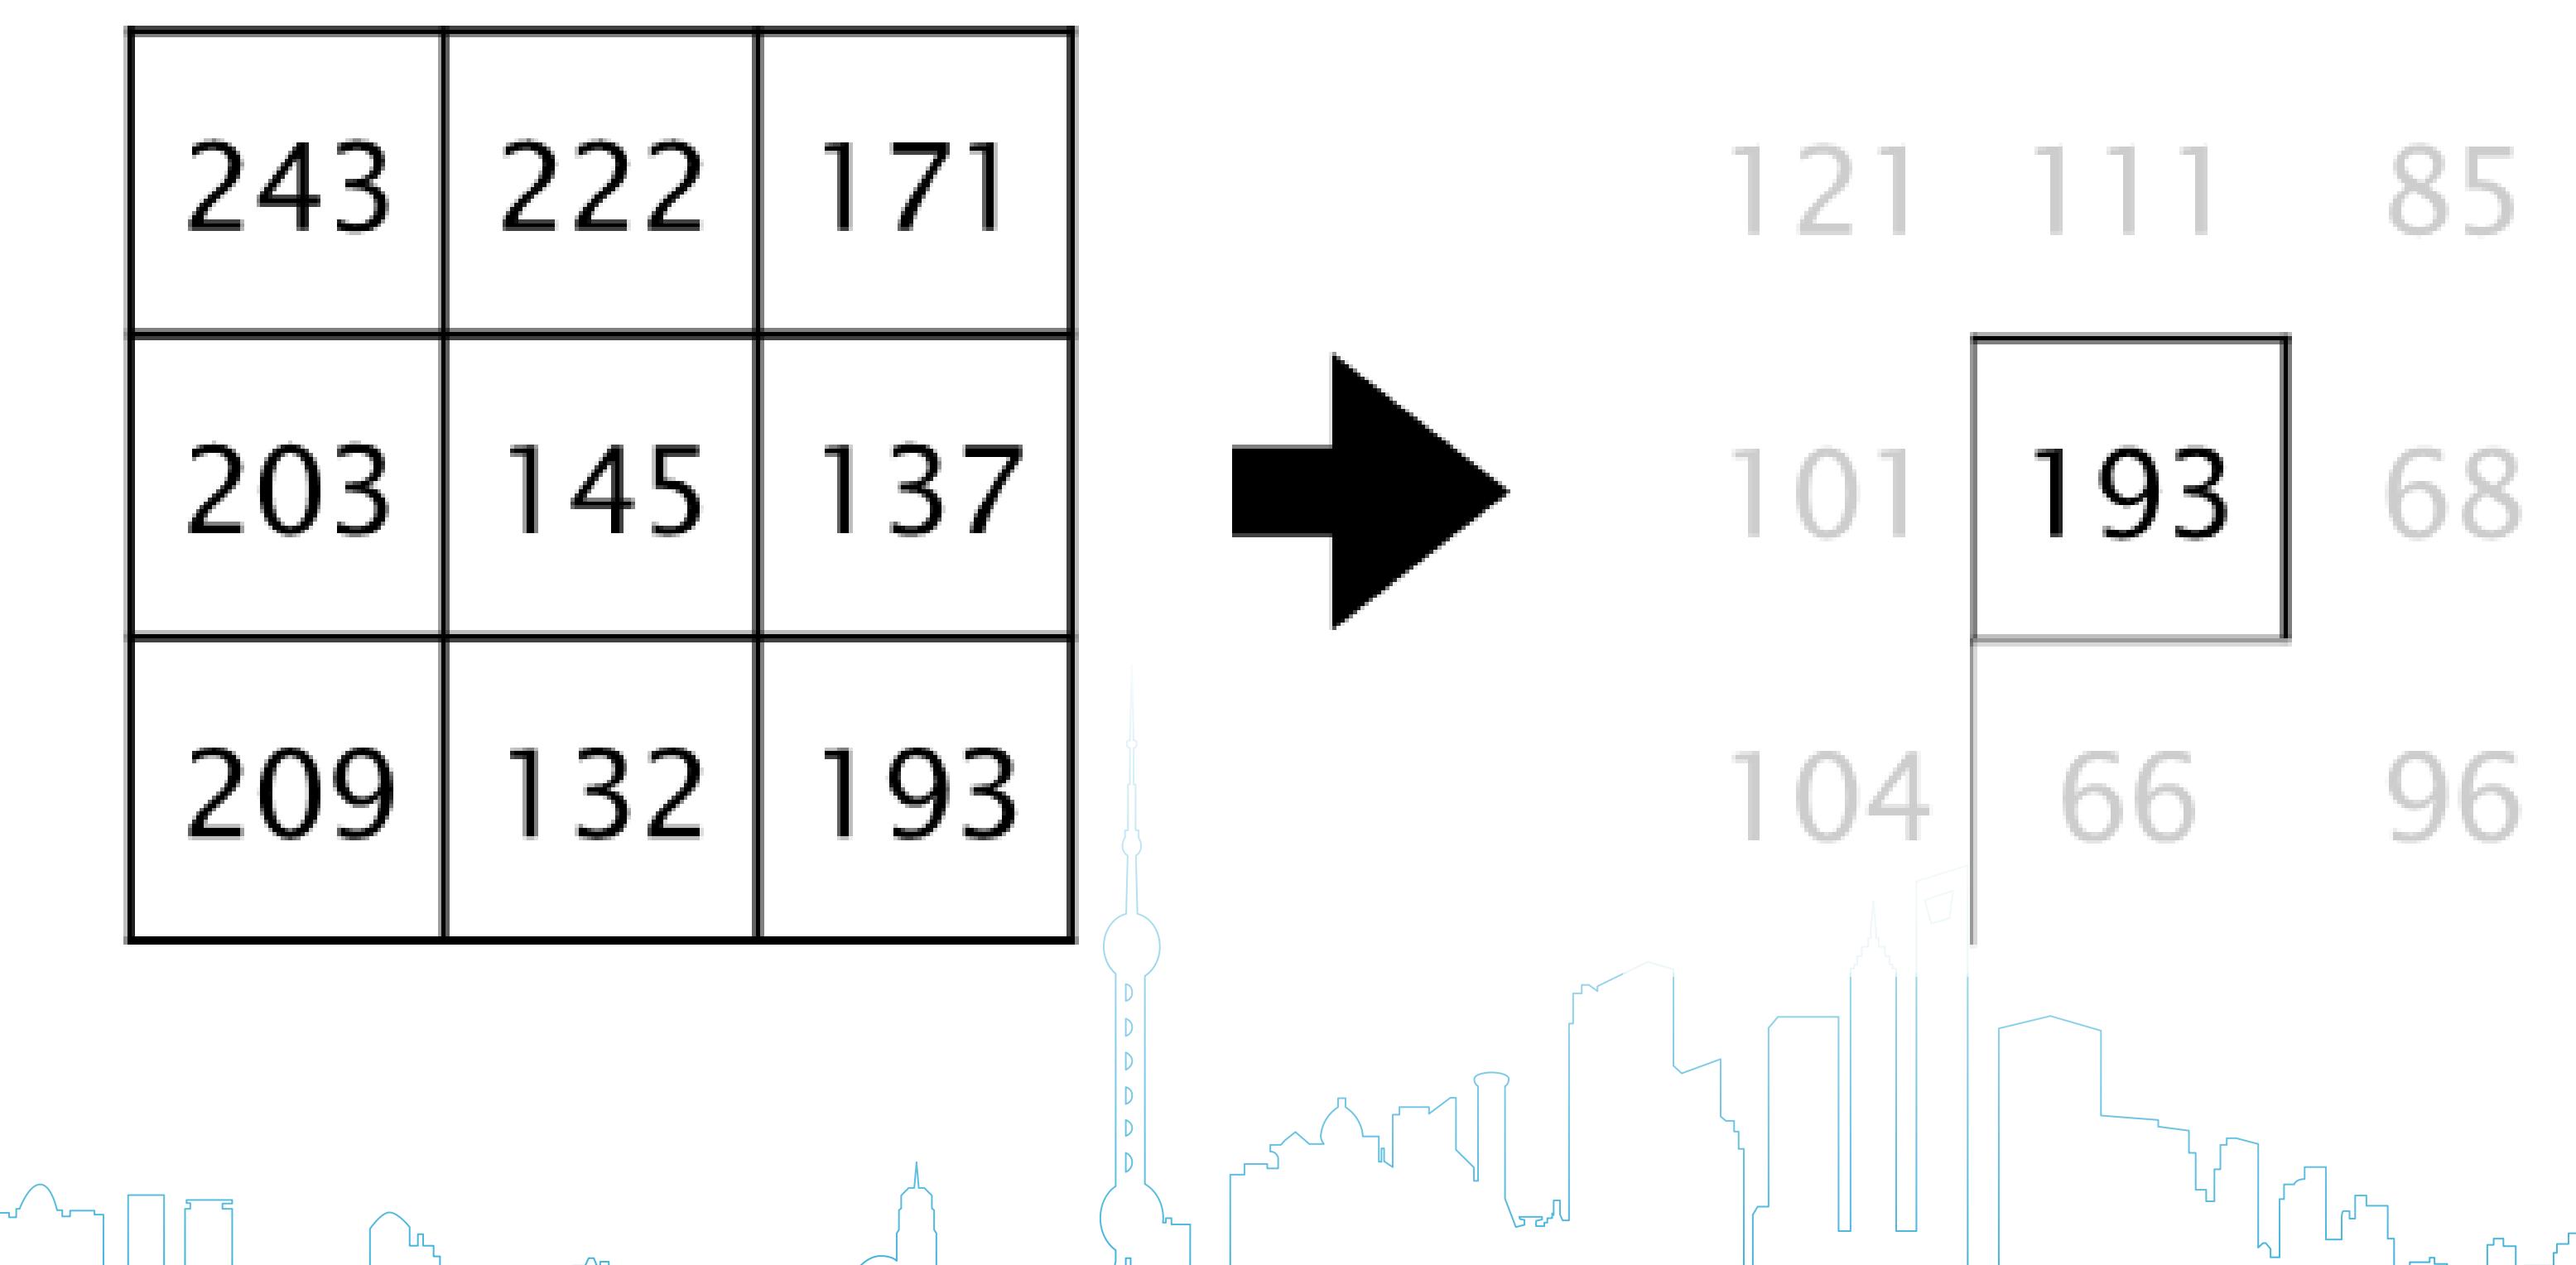
Pixel-level
(Mon)

90	3	241	65	213
242	143	122	71	8
203	103	45	167	240
29	239	2	93	218
243	242	183	54	123

Patch-level
(Today)

90	3	241	65	213
242	143	122	71	8
203	103	45	167	240
29	239	2	93	218
243	242	183	54	123

Intuition: Need bigger input domain (patch)



Patch-level processing: image noise



Clean image



Noisy image
(Salt&pepper noise)

$$f: p \rightarrow p'$$

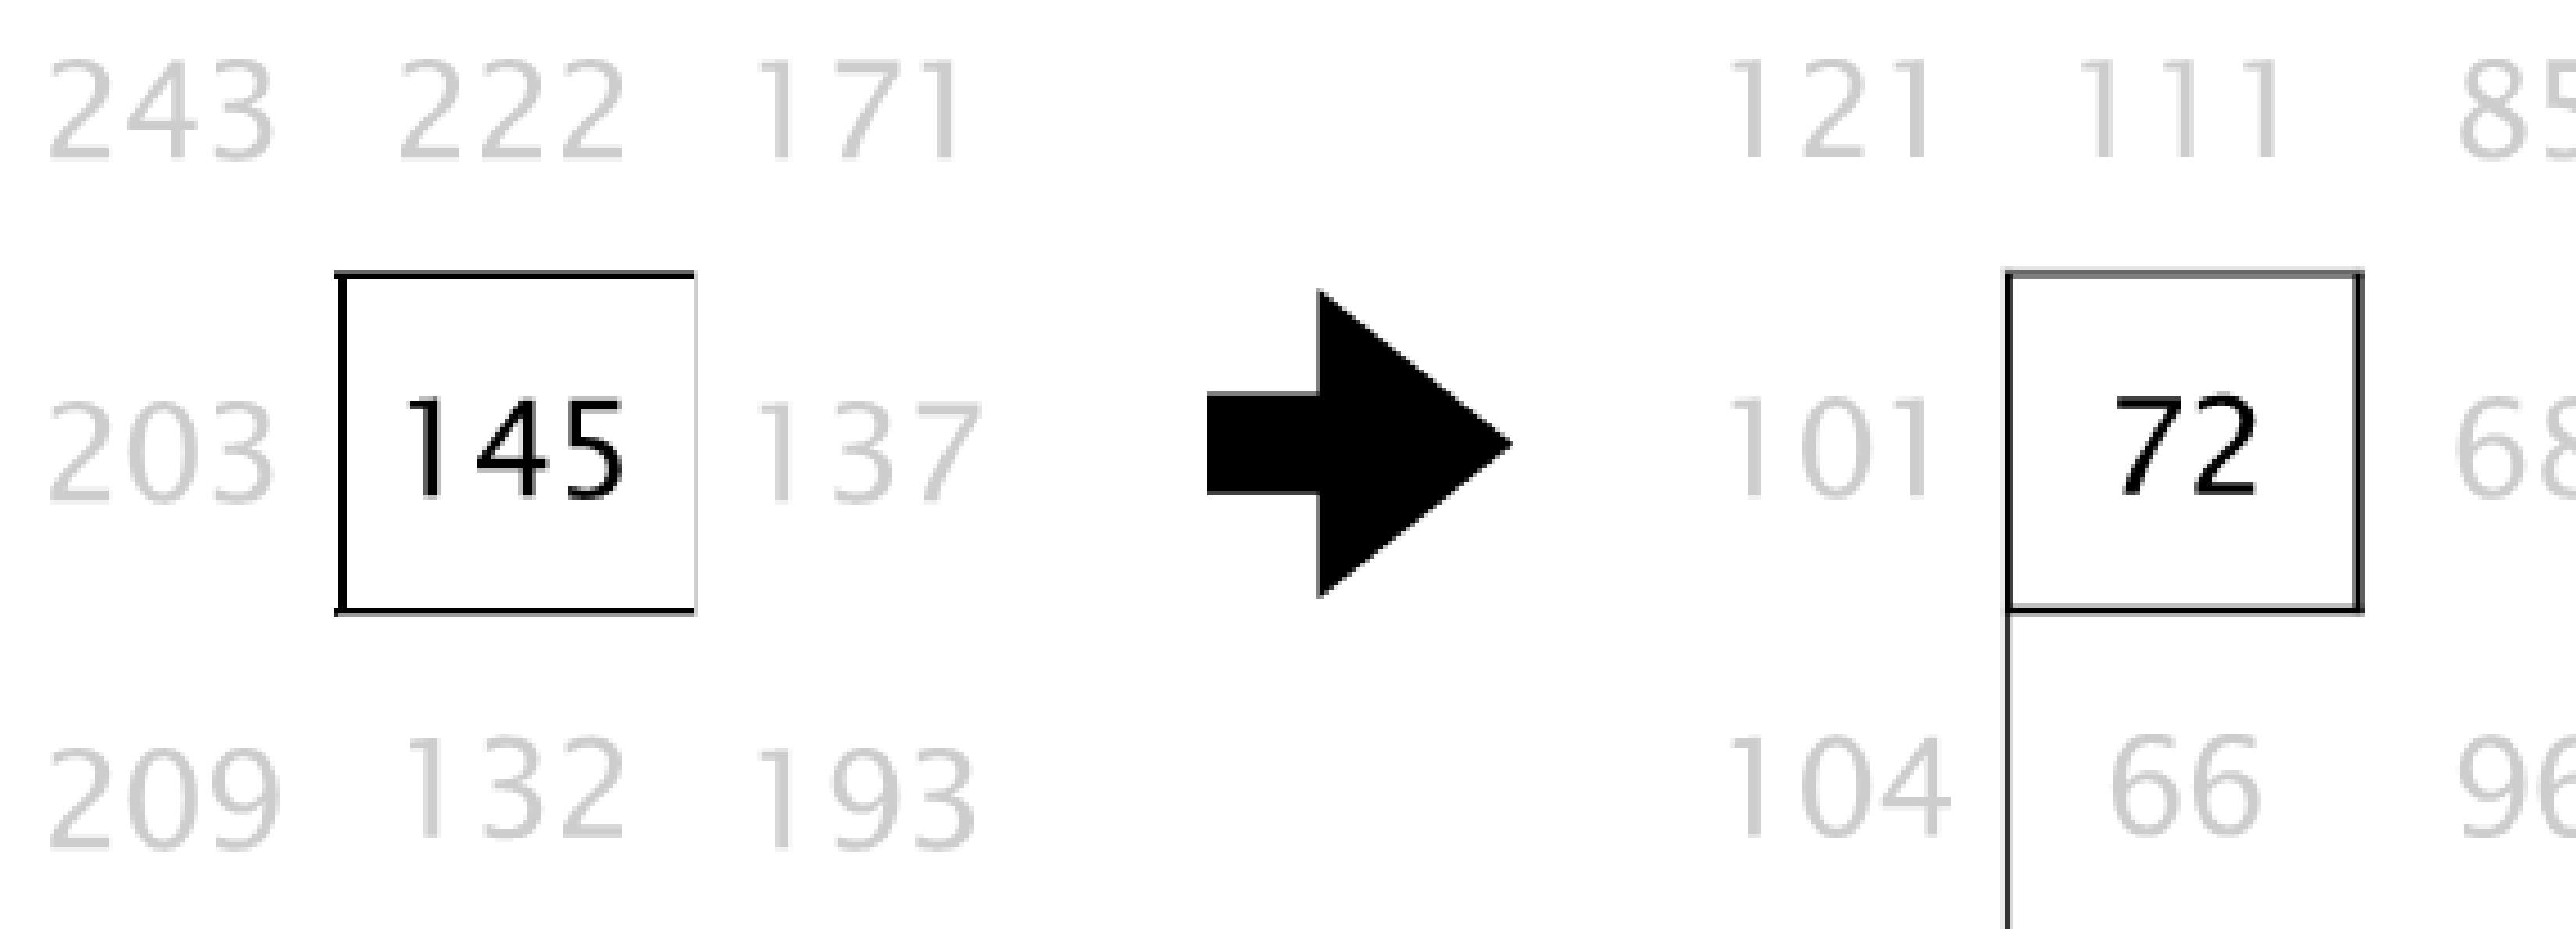
Input pixel Output pixel

0	0
...	...
100	50

Transfer function doesn't help
- Same pixel value has different
“ideal” values

Patch-level processing

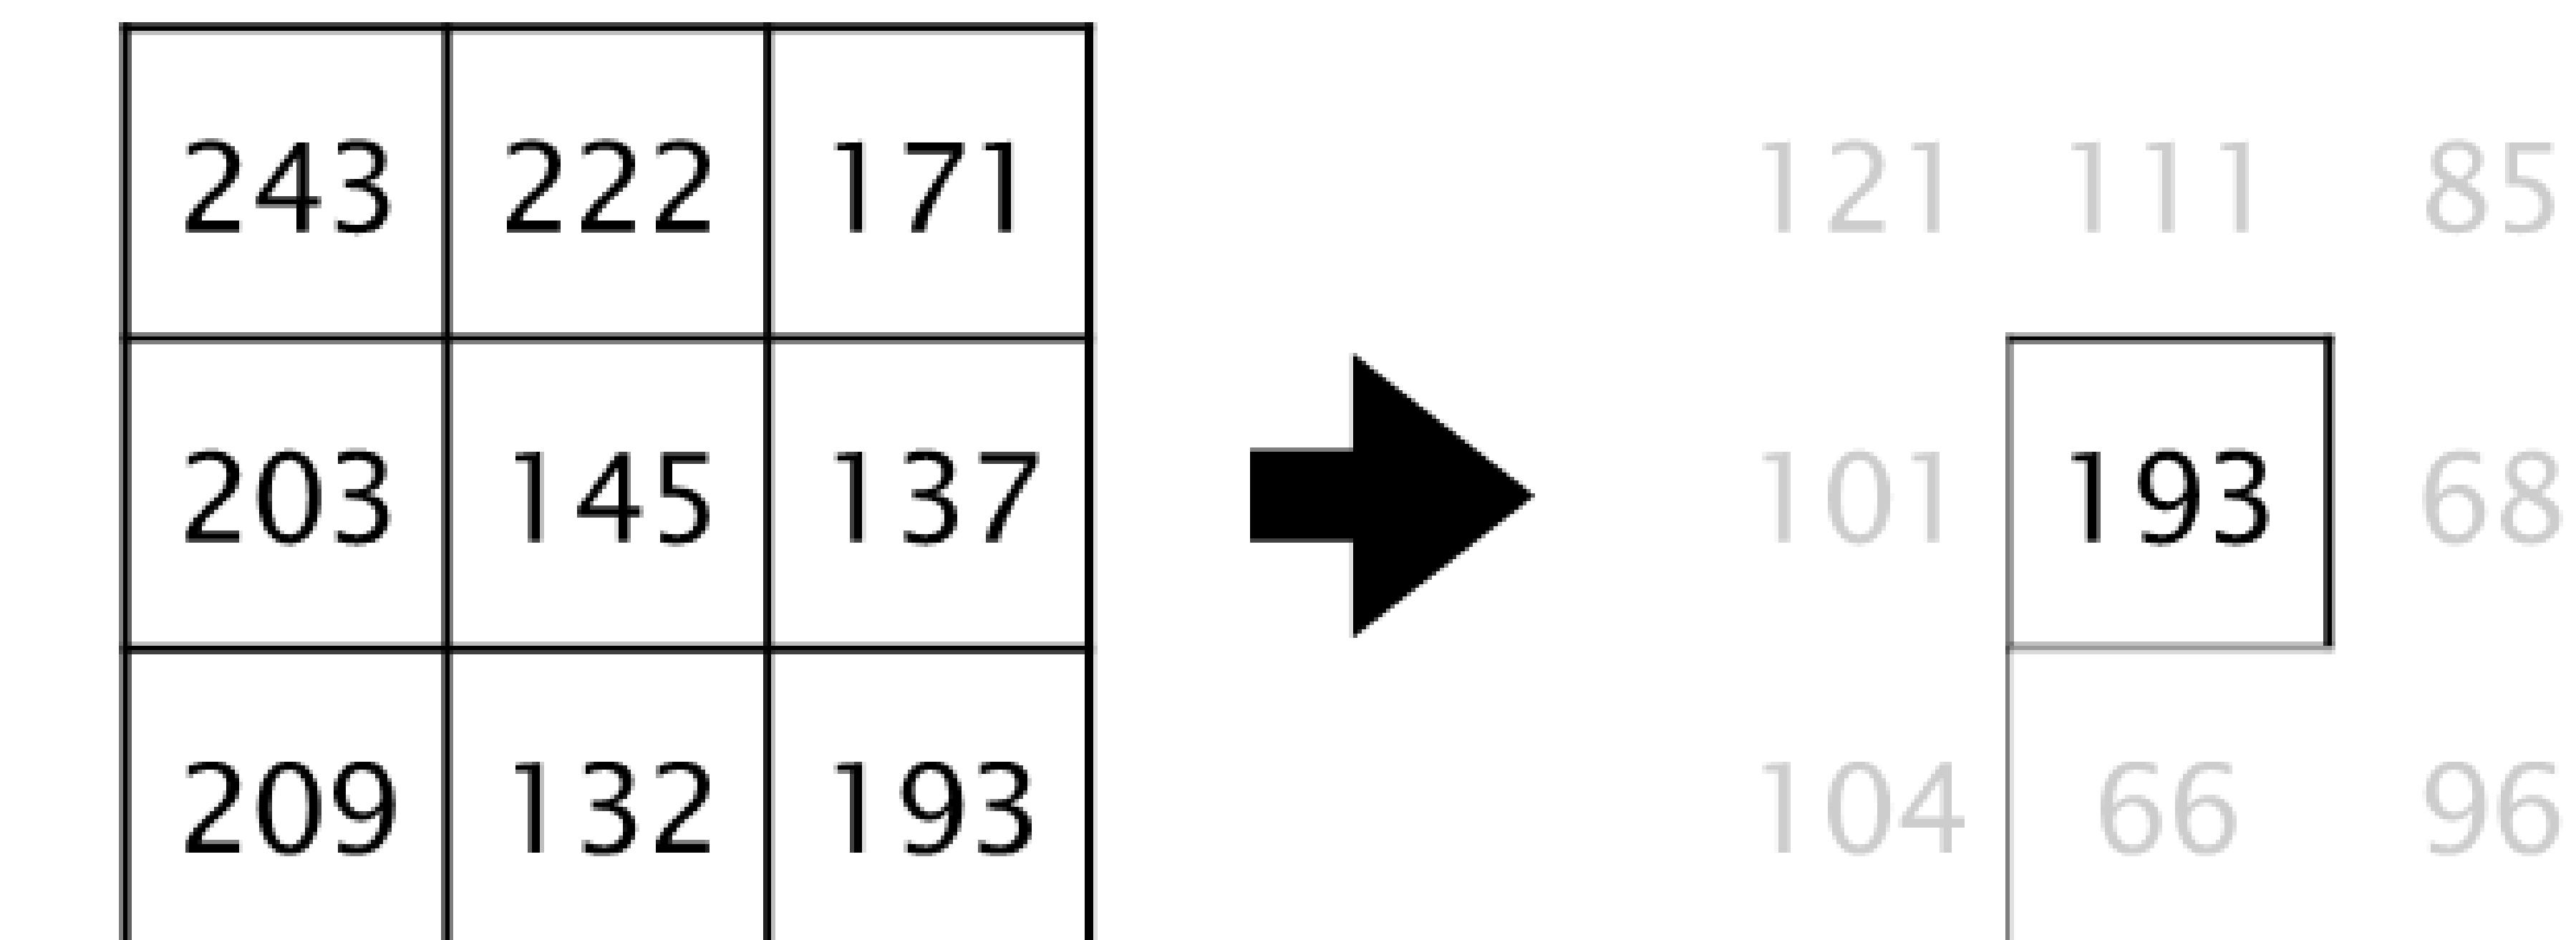
Maths: make a patch function



$$f: p \rightarrow p'$$

Input pixel

Output pixel



$$f: pa \rightarrow p'$$

Input **patch**

Output pixel

Patch-level processing

Example: Median filter

243	222	171	
203	145	137	→
209	132	193	

121 111 85
101 72 68
104 66 96

243	222	171	
203	145	137	→
209	132	193	

121 111 85
101 193 68
104 66 96

$$f: p \rightarrow p'$$

Input pixel Output pixel

0	0
...	...
100	50

$$f: pa \rightarrow p'$$

Input **patch** Output pixel

$$p' = \text{median}(pa)$$

Transfer function (lookup table)

Patch-level processing

Example: Median filter

30	31	32	3	4
0	6	99	50	90
99	35	33	32	98
0	90	90	36	31
32	31	0	90	90

Input

				?

Output

Patch-level processing

Example: Median filter

30	31	32	3	4
0	6	99	50	90
99	35	33	32	98
0	90	90	36	31
32	31	0	90	90

Input

				32

Output

Patch-level processing

Loop through each column

30	31	32	3	4
0	6	99	30	90
99	35	33	32	98
0	90	90	36	31
32	31	0	90	90

Input

			32	?

Output

Patch-level processing

Loop through each column

30	31	32	3	4
0	6	99	30	90
99	35	33	32	98
0	90	90	36	31
32	31	0	90	90

Input

		32		?

Output

Patch-level processing

Loop through each column

30	31	32	3	4
0	6	99	30	90
99	35	33	32	98
0	90	90	36	31
32	31	0	90	90

Input

		32		32
32			32	

Output

Patch-level processing

Loop through each row

30	31	32	3	4
0	6	99	30	90
99	35	33	32	99
0	90	90	36	31
32	31	0	90	90

Input

	32	32	32	
	?			

Output

Patch-level processing

Loop through each row

30	31	32	3	4
0	6	99	30	90
99	35	33	32	99
0	90	90	36	31
32	31	0	90	90

Input

	32	32	32	
	35			

Output

Patch-level processing

Loop through each row

30	31	32	3	4
0	6	99	30	90
99	35	33	32	98
0	90	90	36	31
32	31	0	90	90

Input

	32	32	32	
	35	35	36	
			36	

Output

Patch-level processing

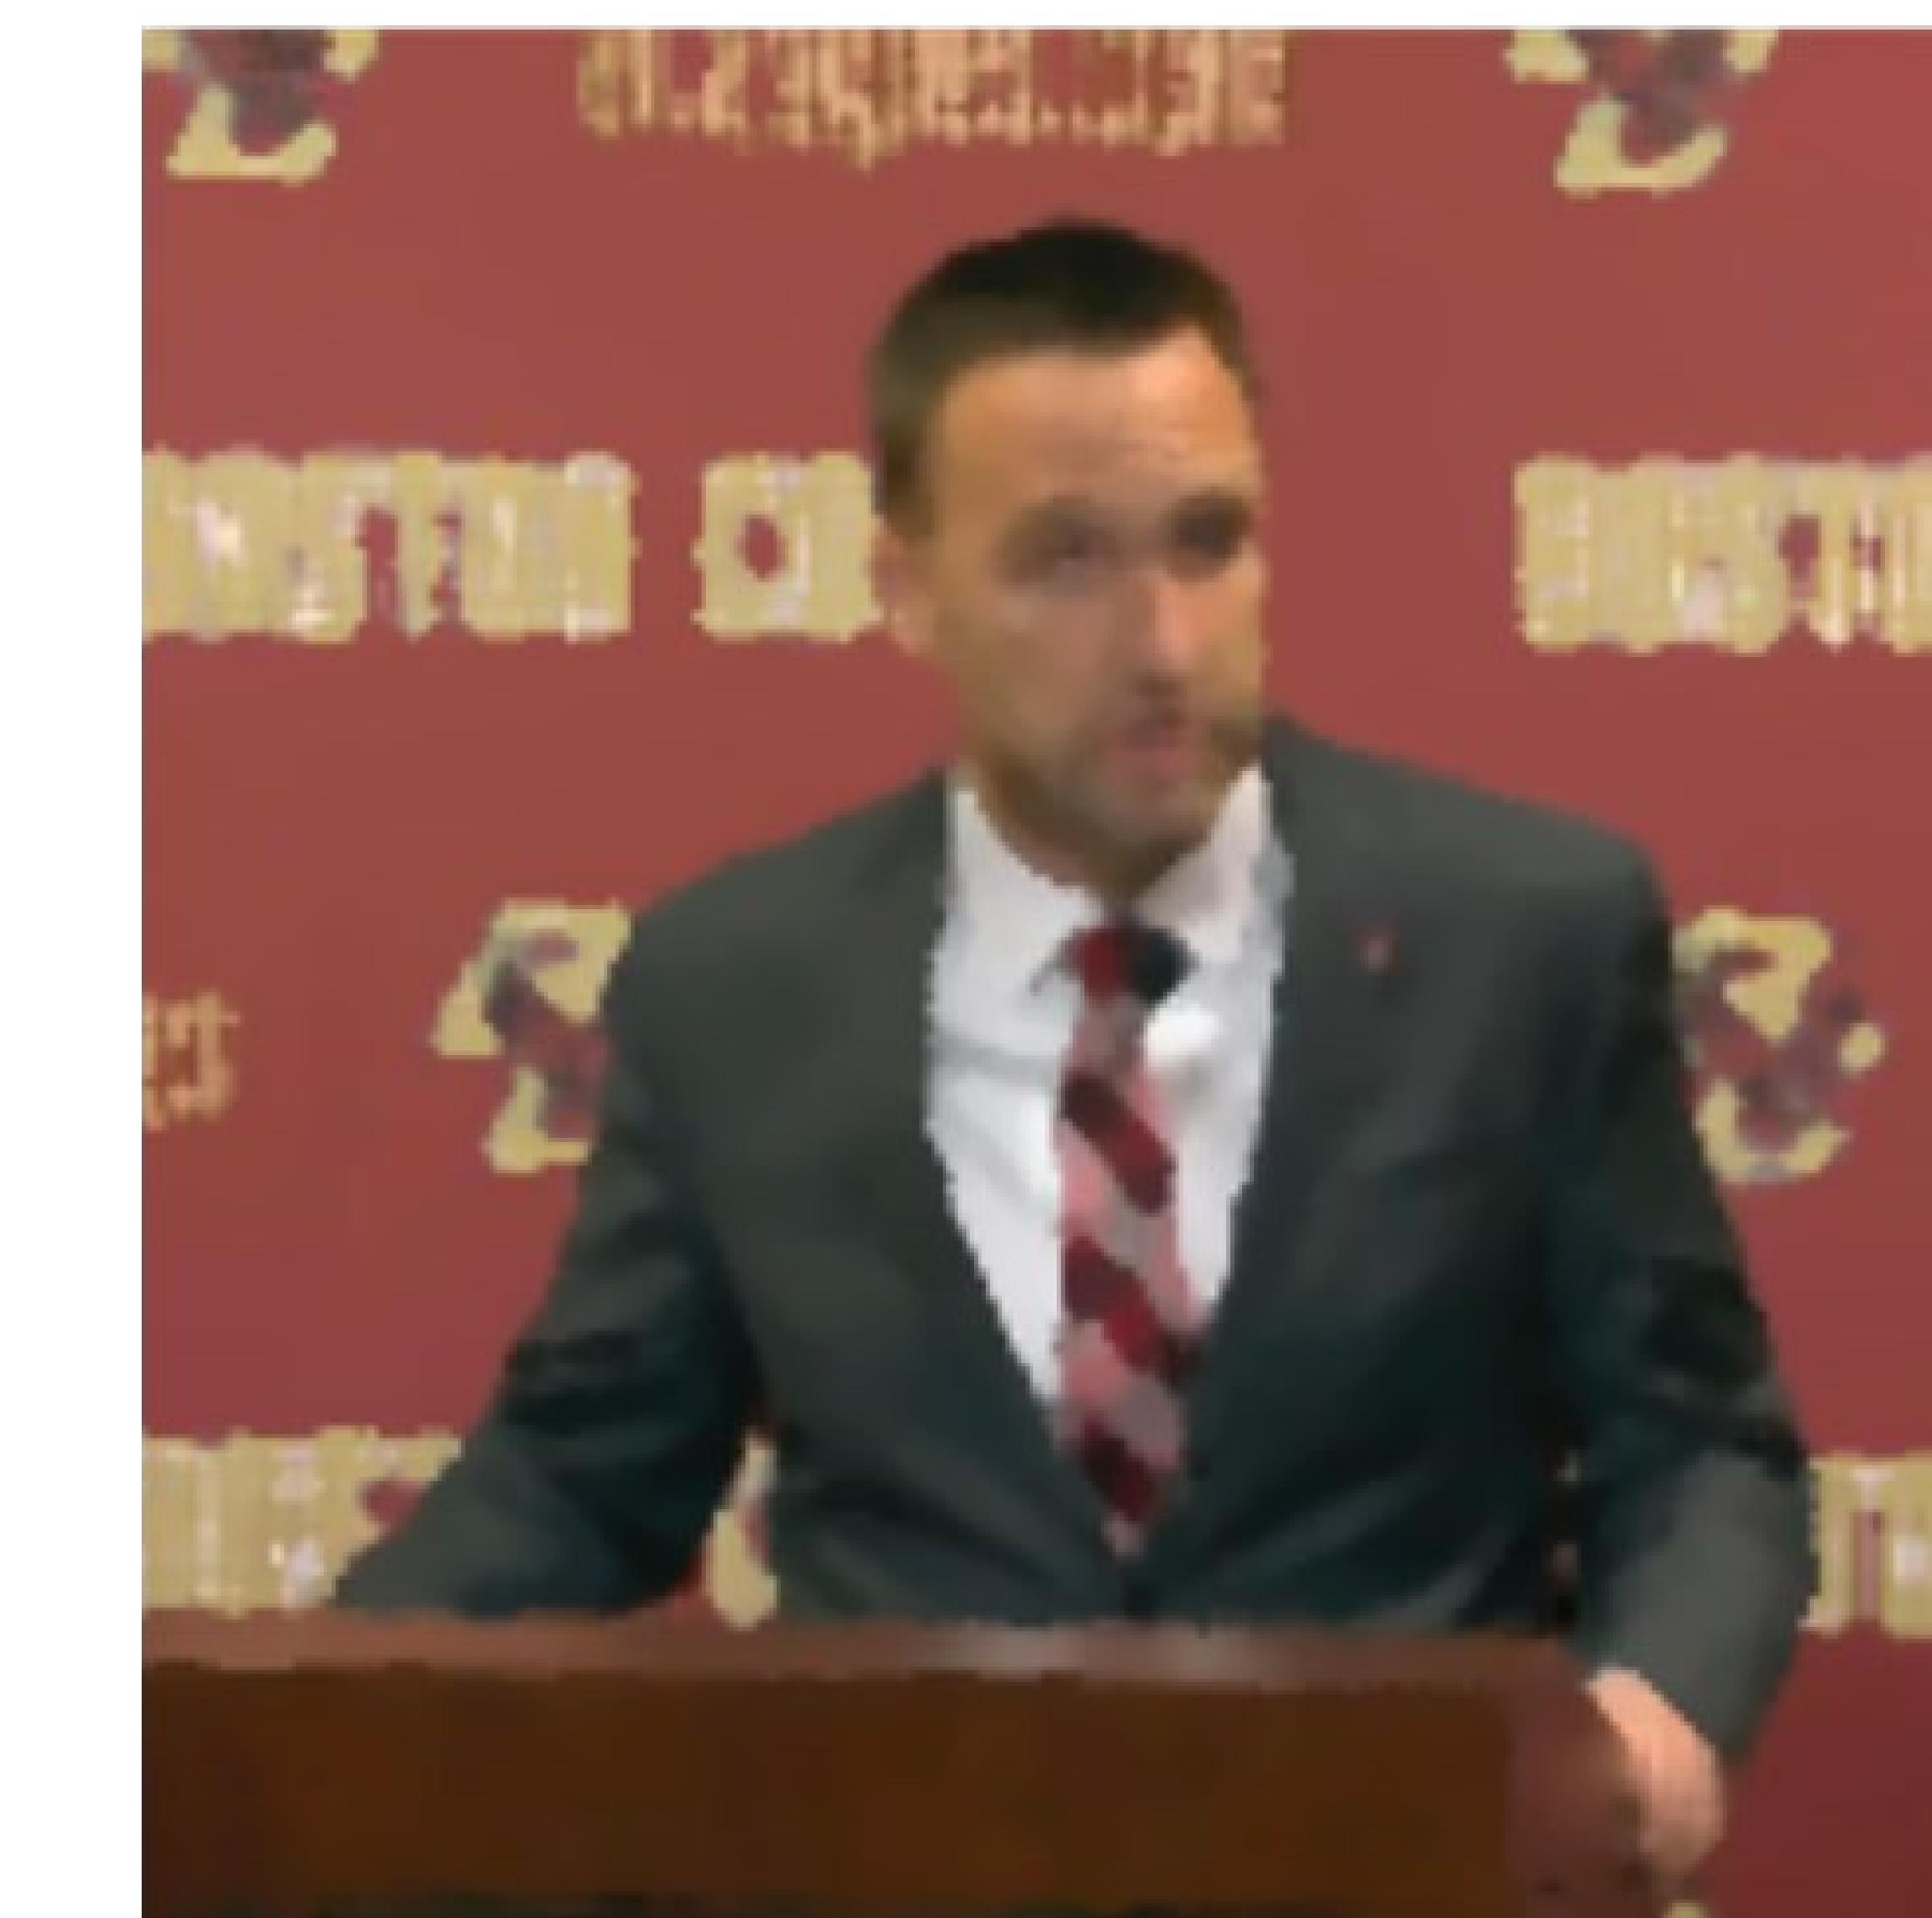
Median filter: image denoising



Ideal image

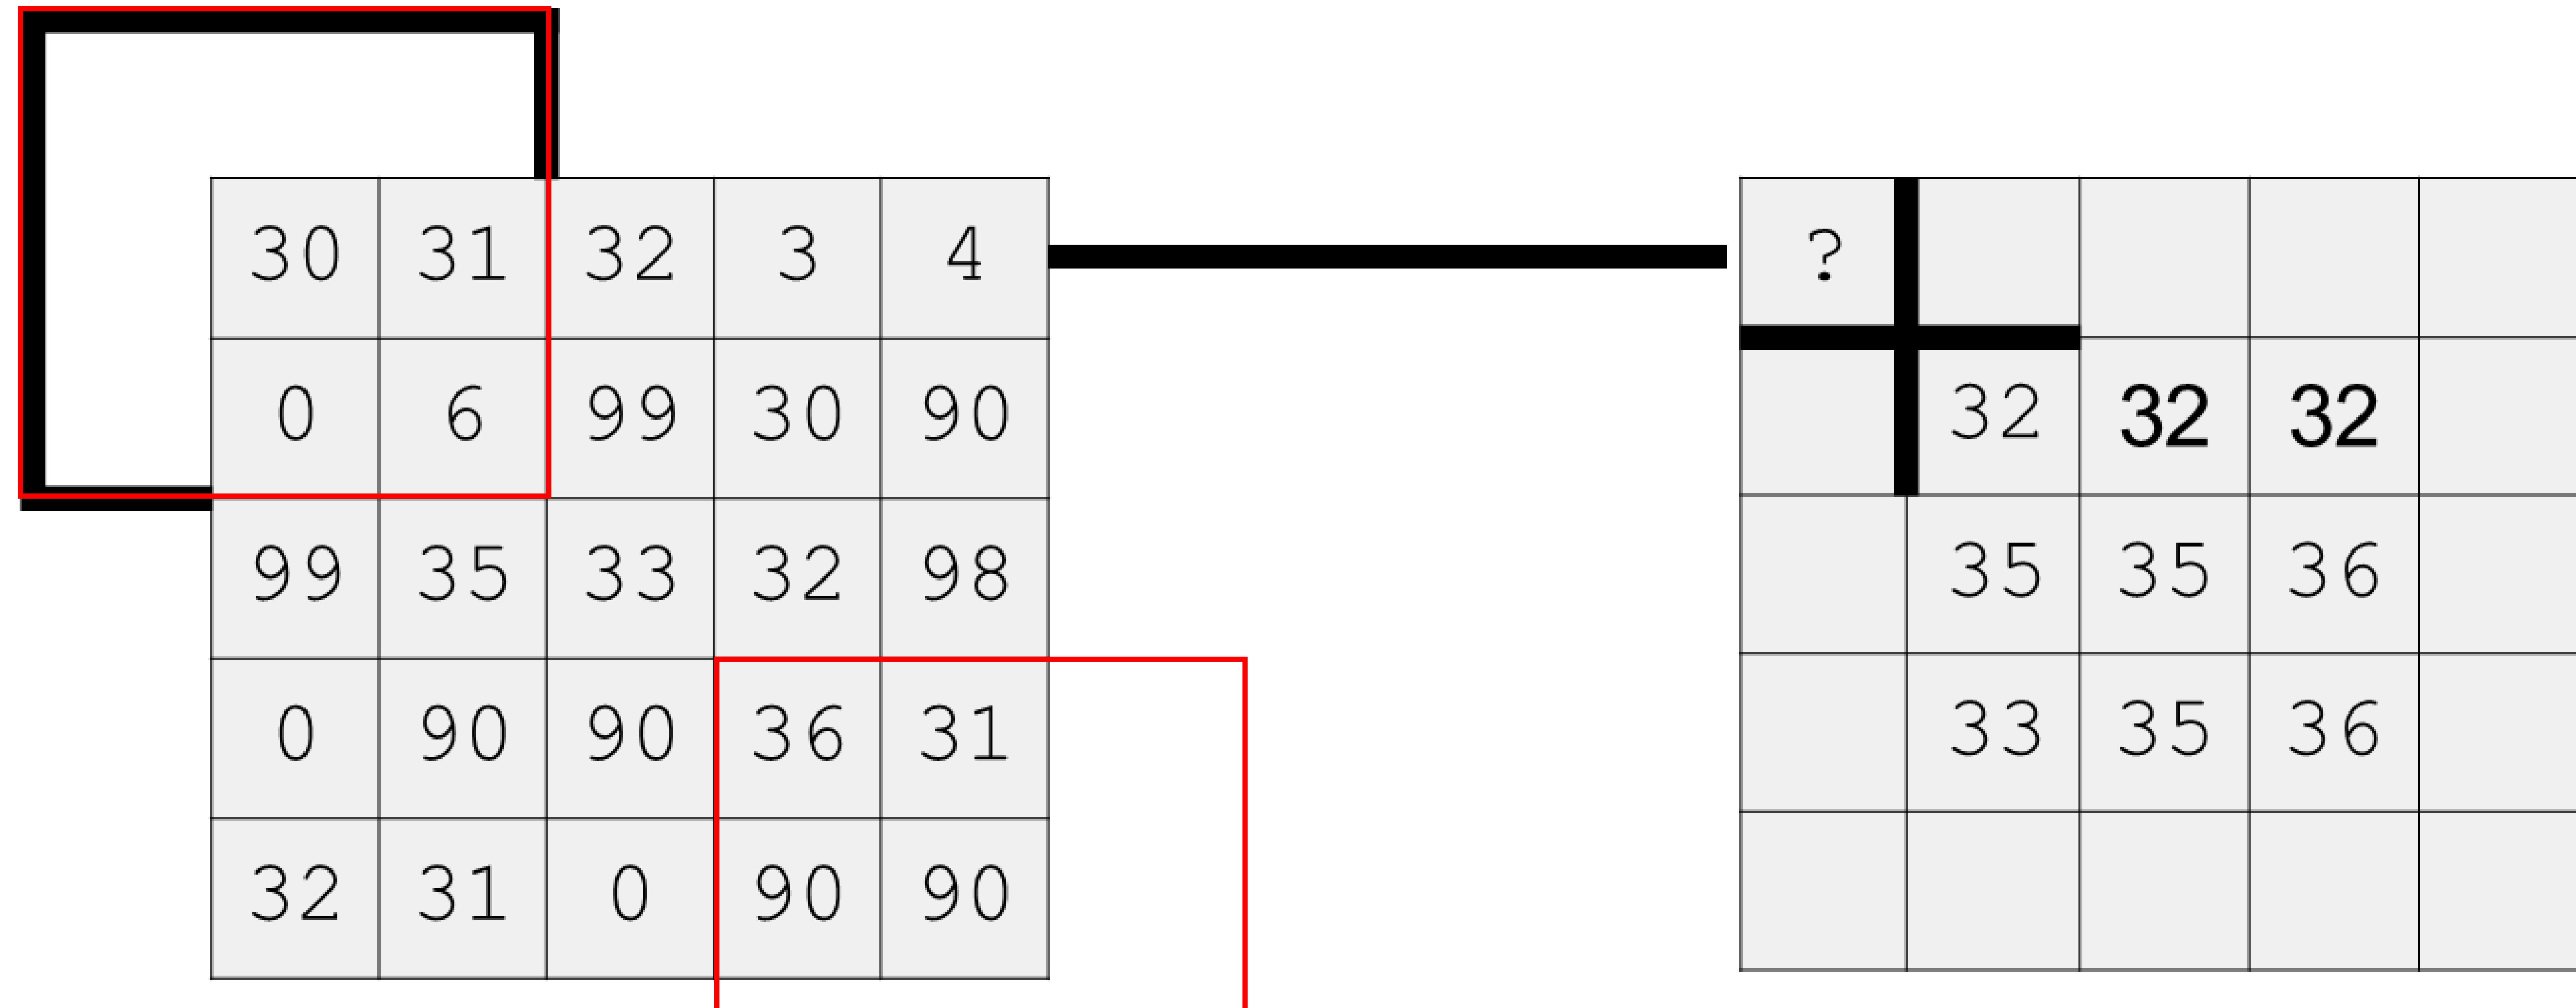


Real-world image



Median filter result

Patch-level processing



Input

Output

Patch-level processing

Solution: Pad the input image

0	0	0	0	0	0	0
0	30	31	32	3	4	0
0	0	6	99	30	90	0
0	99	35	33	32	98	0
0	0	90	90	36	31	0
0	32	31	0	90	90	0
0	0	0	0	0	0	0

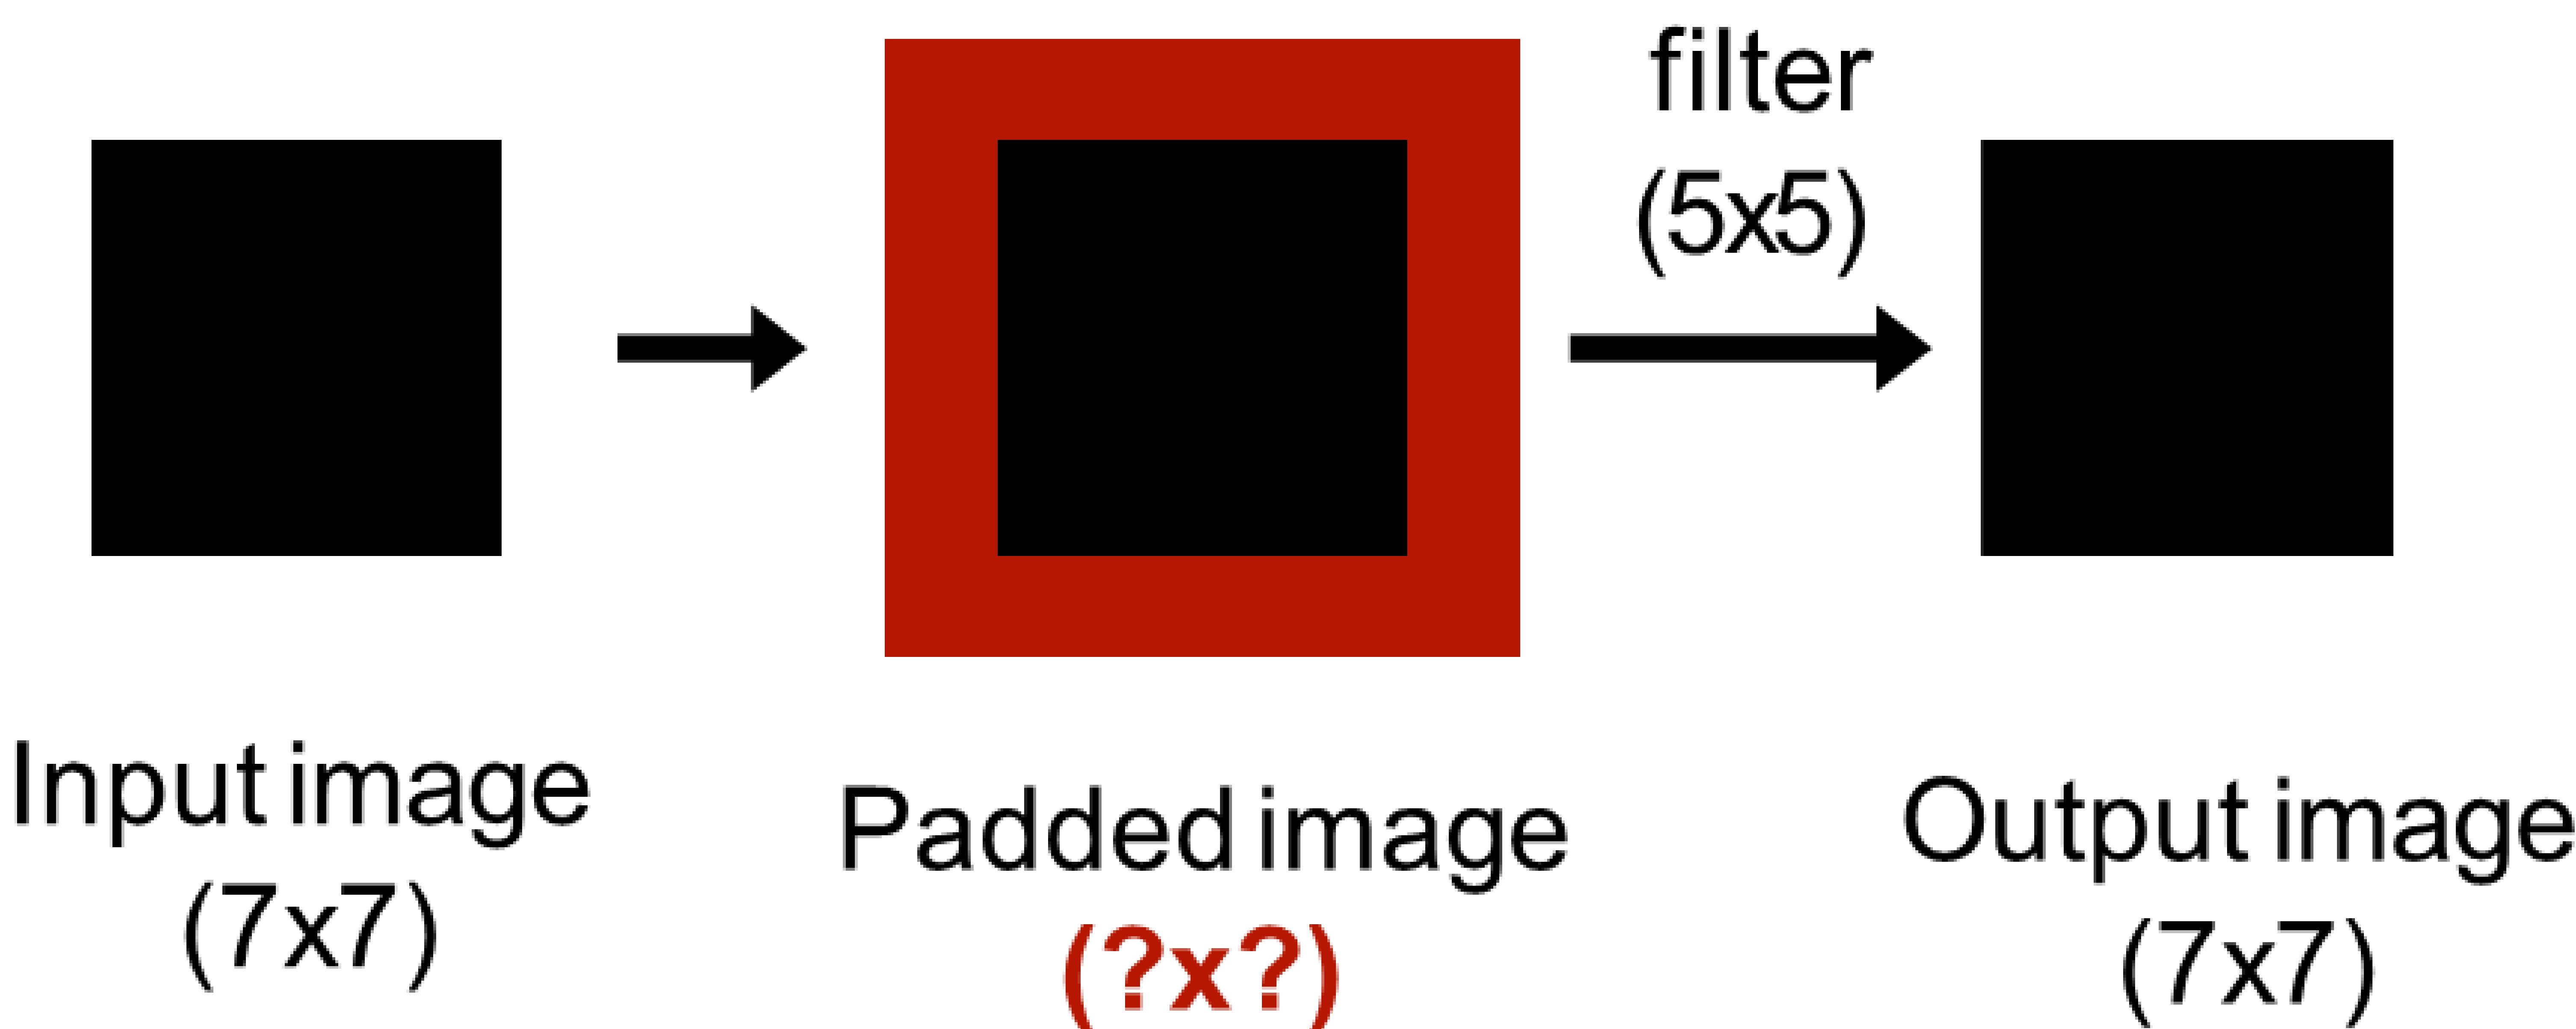
Input

?				
	32	32	32	
	35	35	36	
	33	35	36	

Output

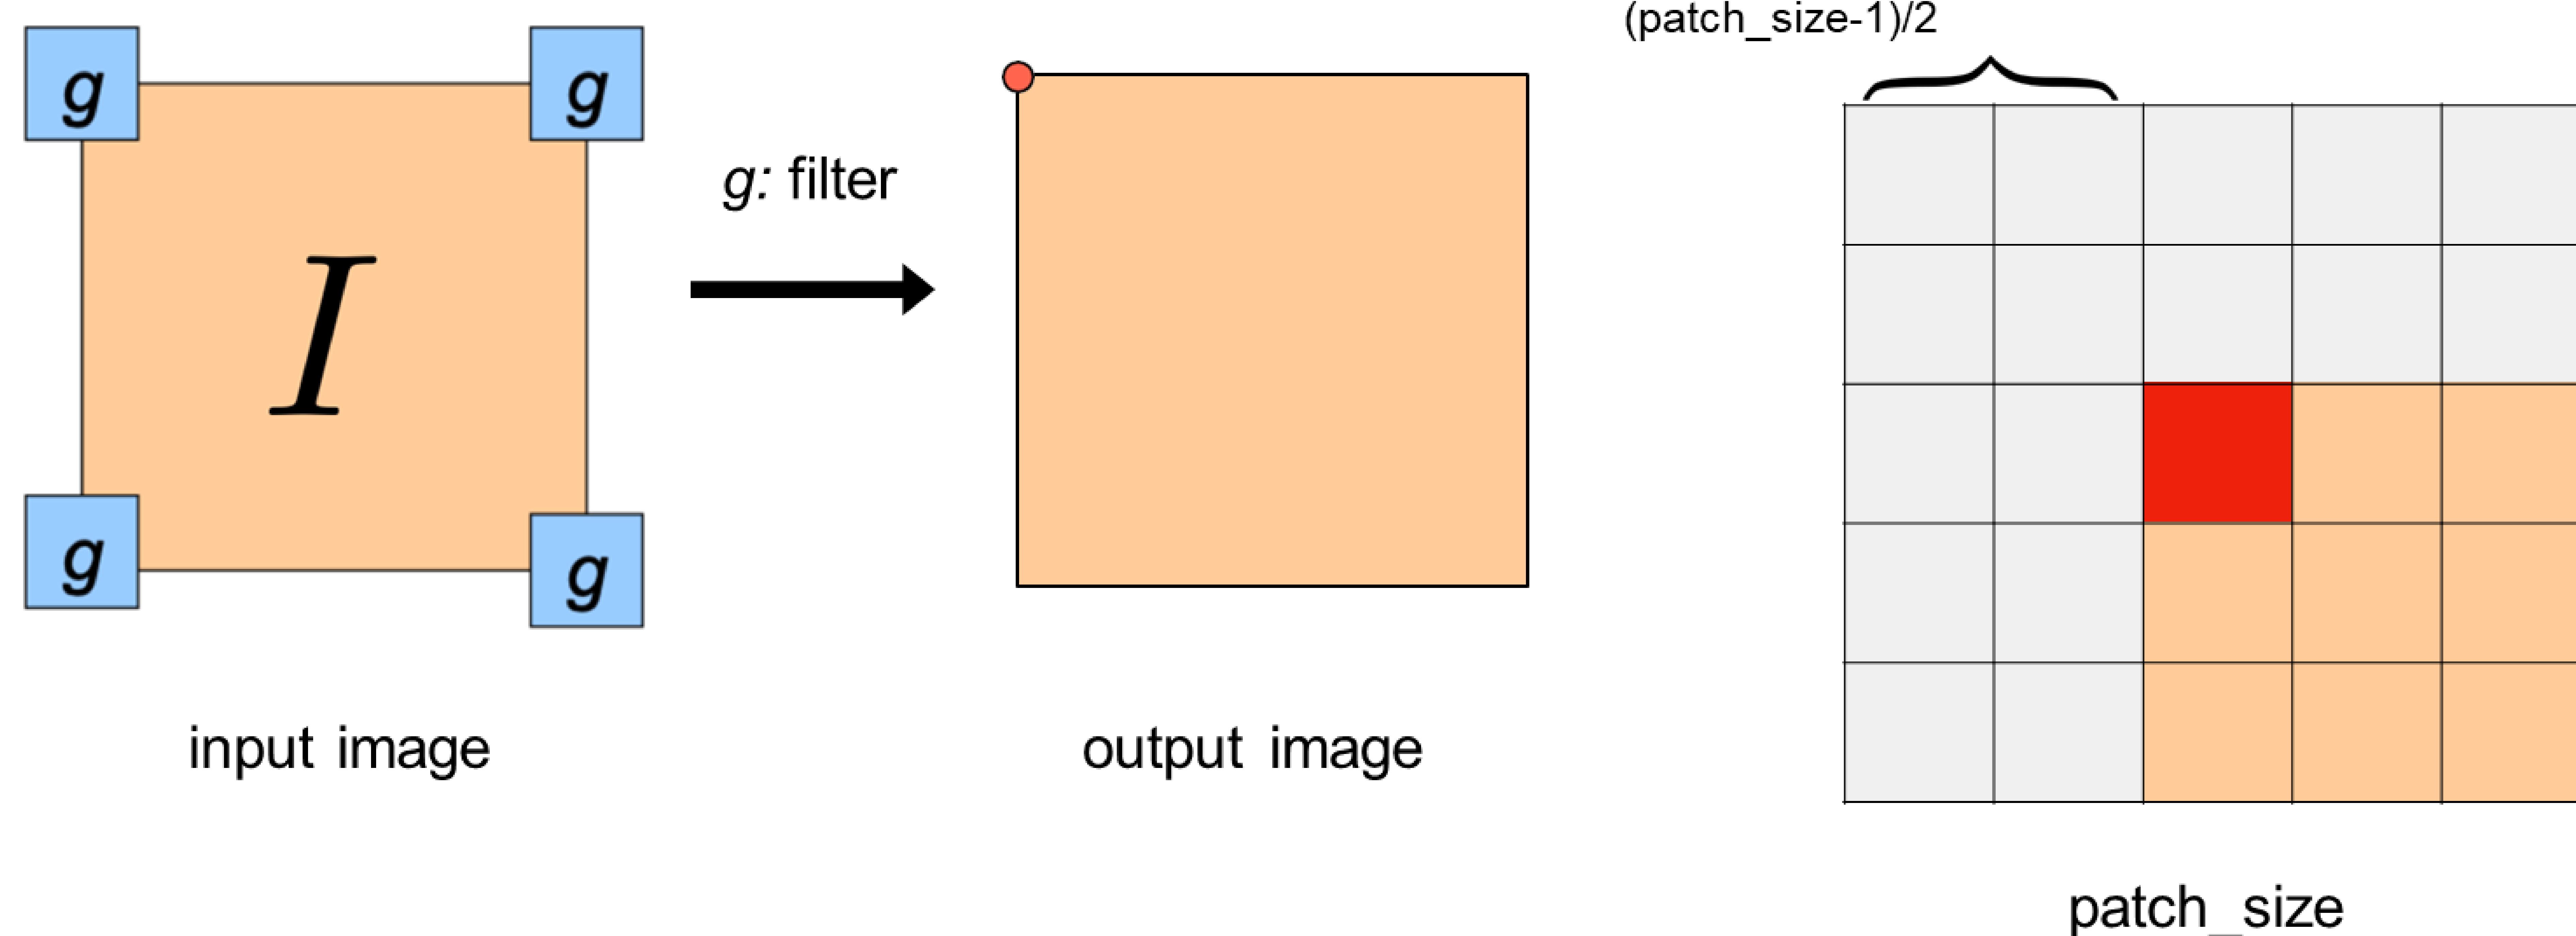
Patch-level processing: padding

How many pixels to pad
on each side?



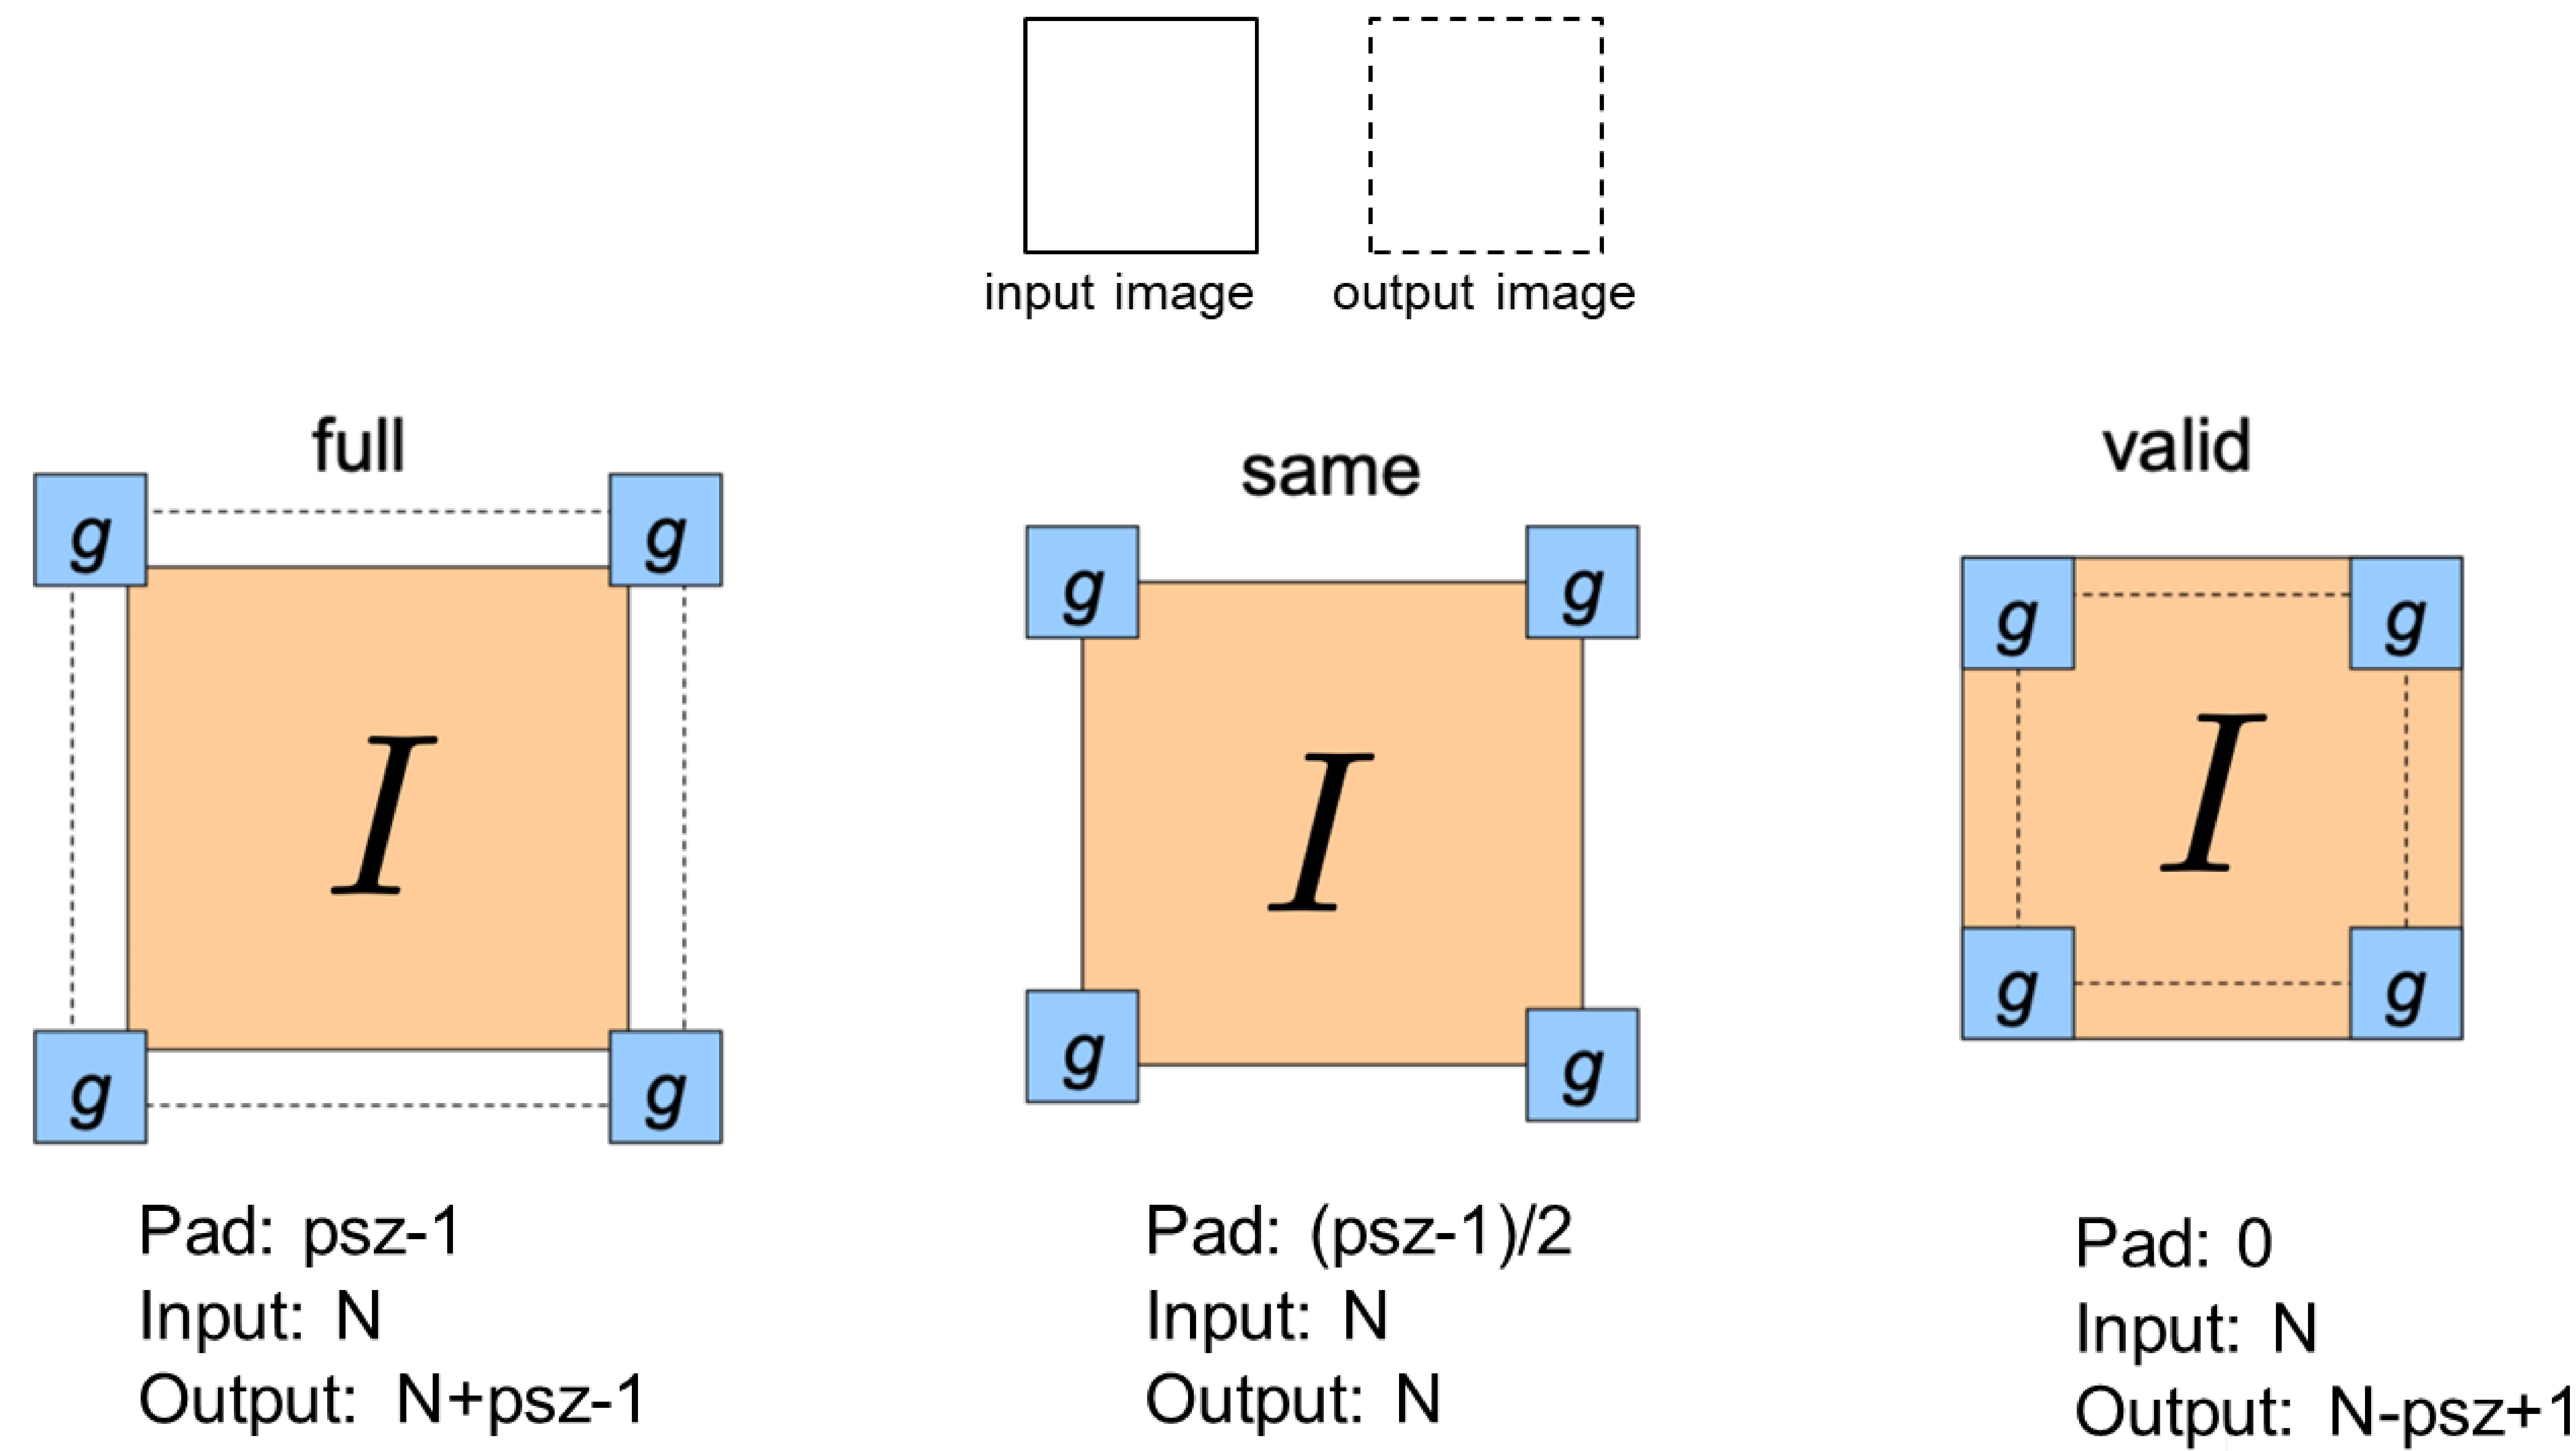
Patch-level processing: padding

Zoom into the corner pixel



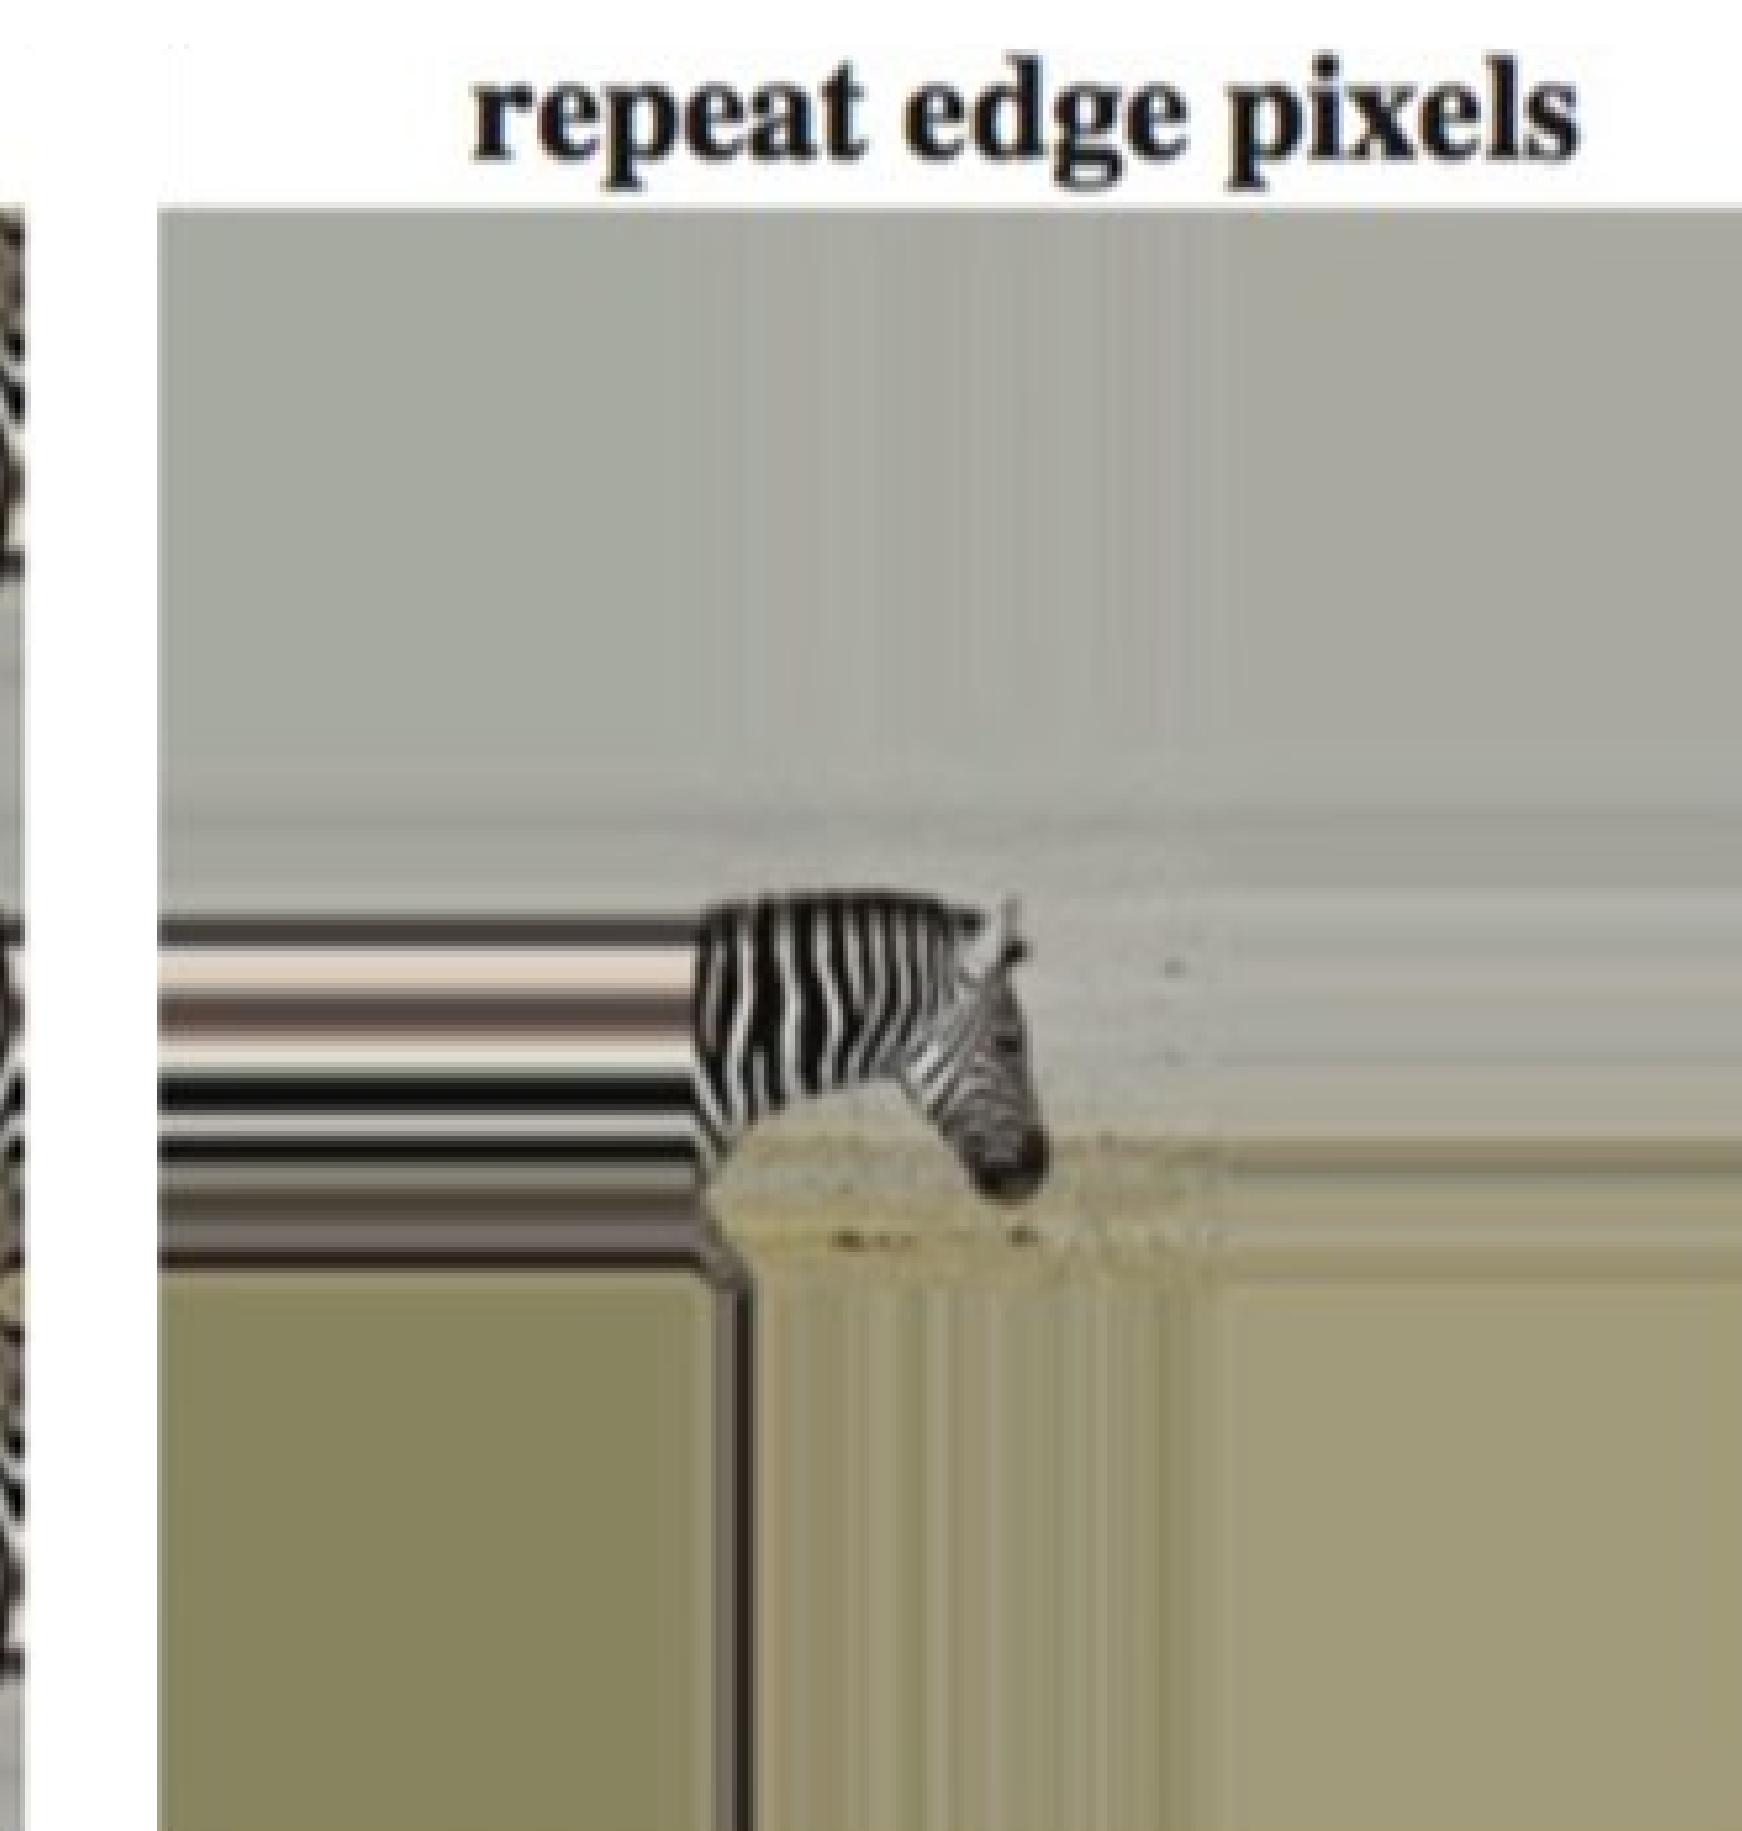
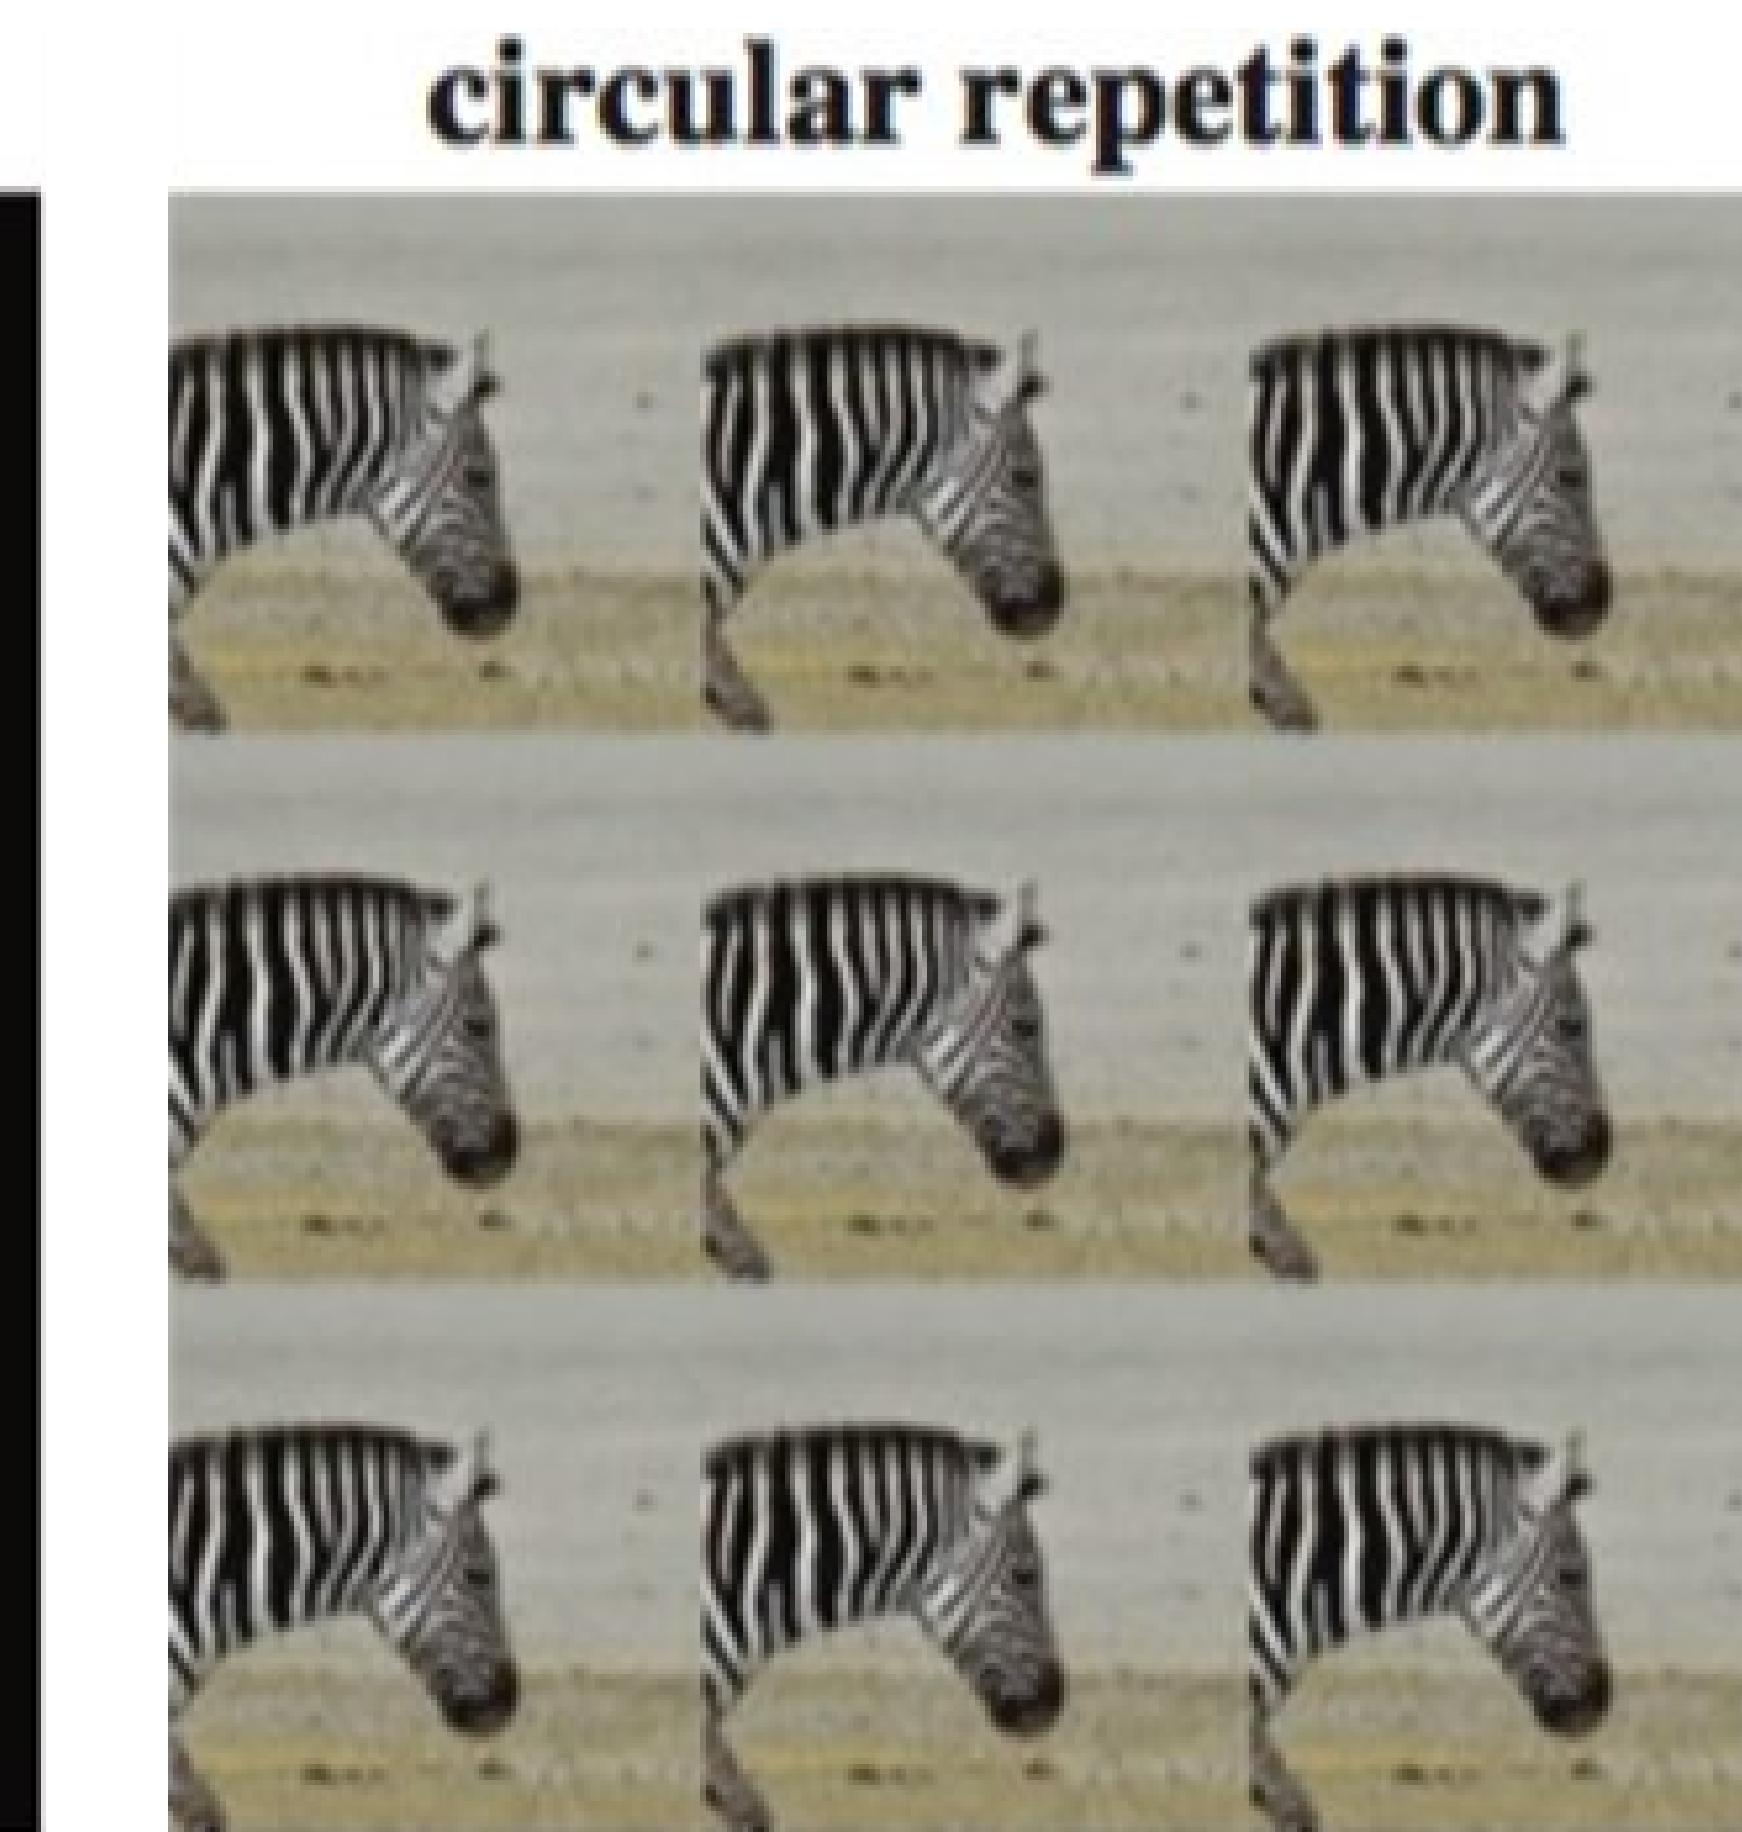
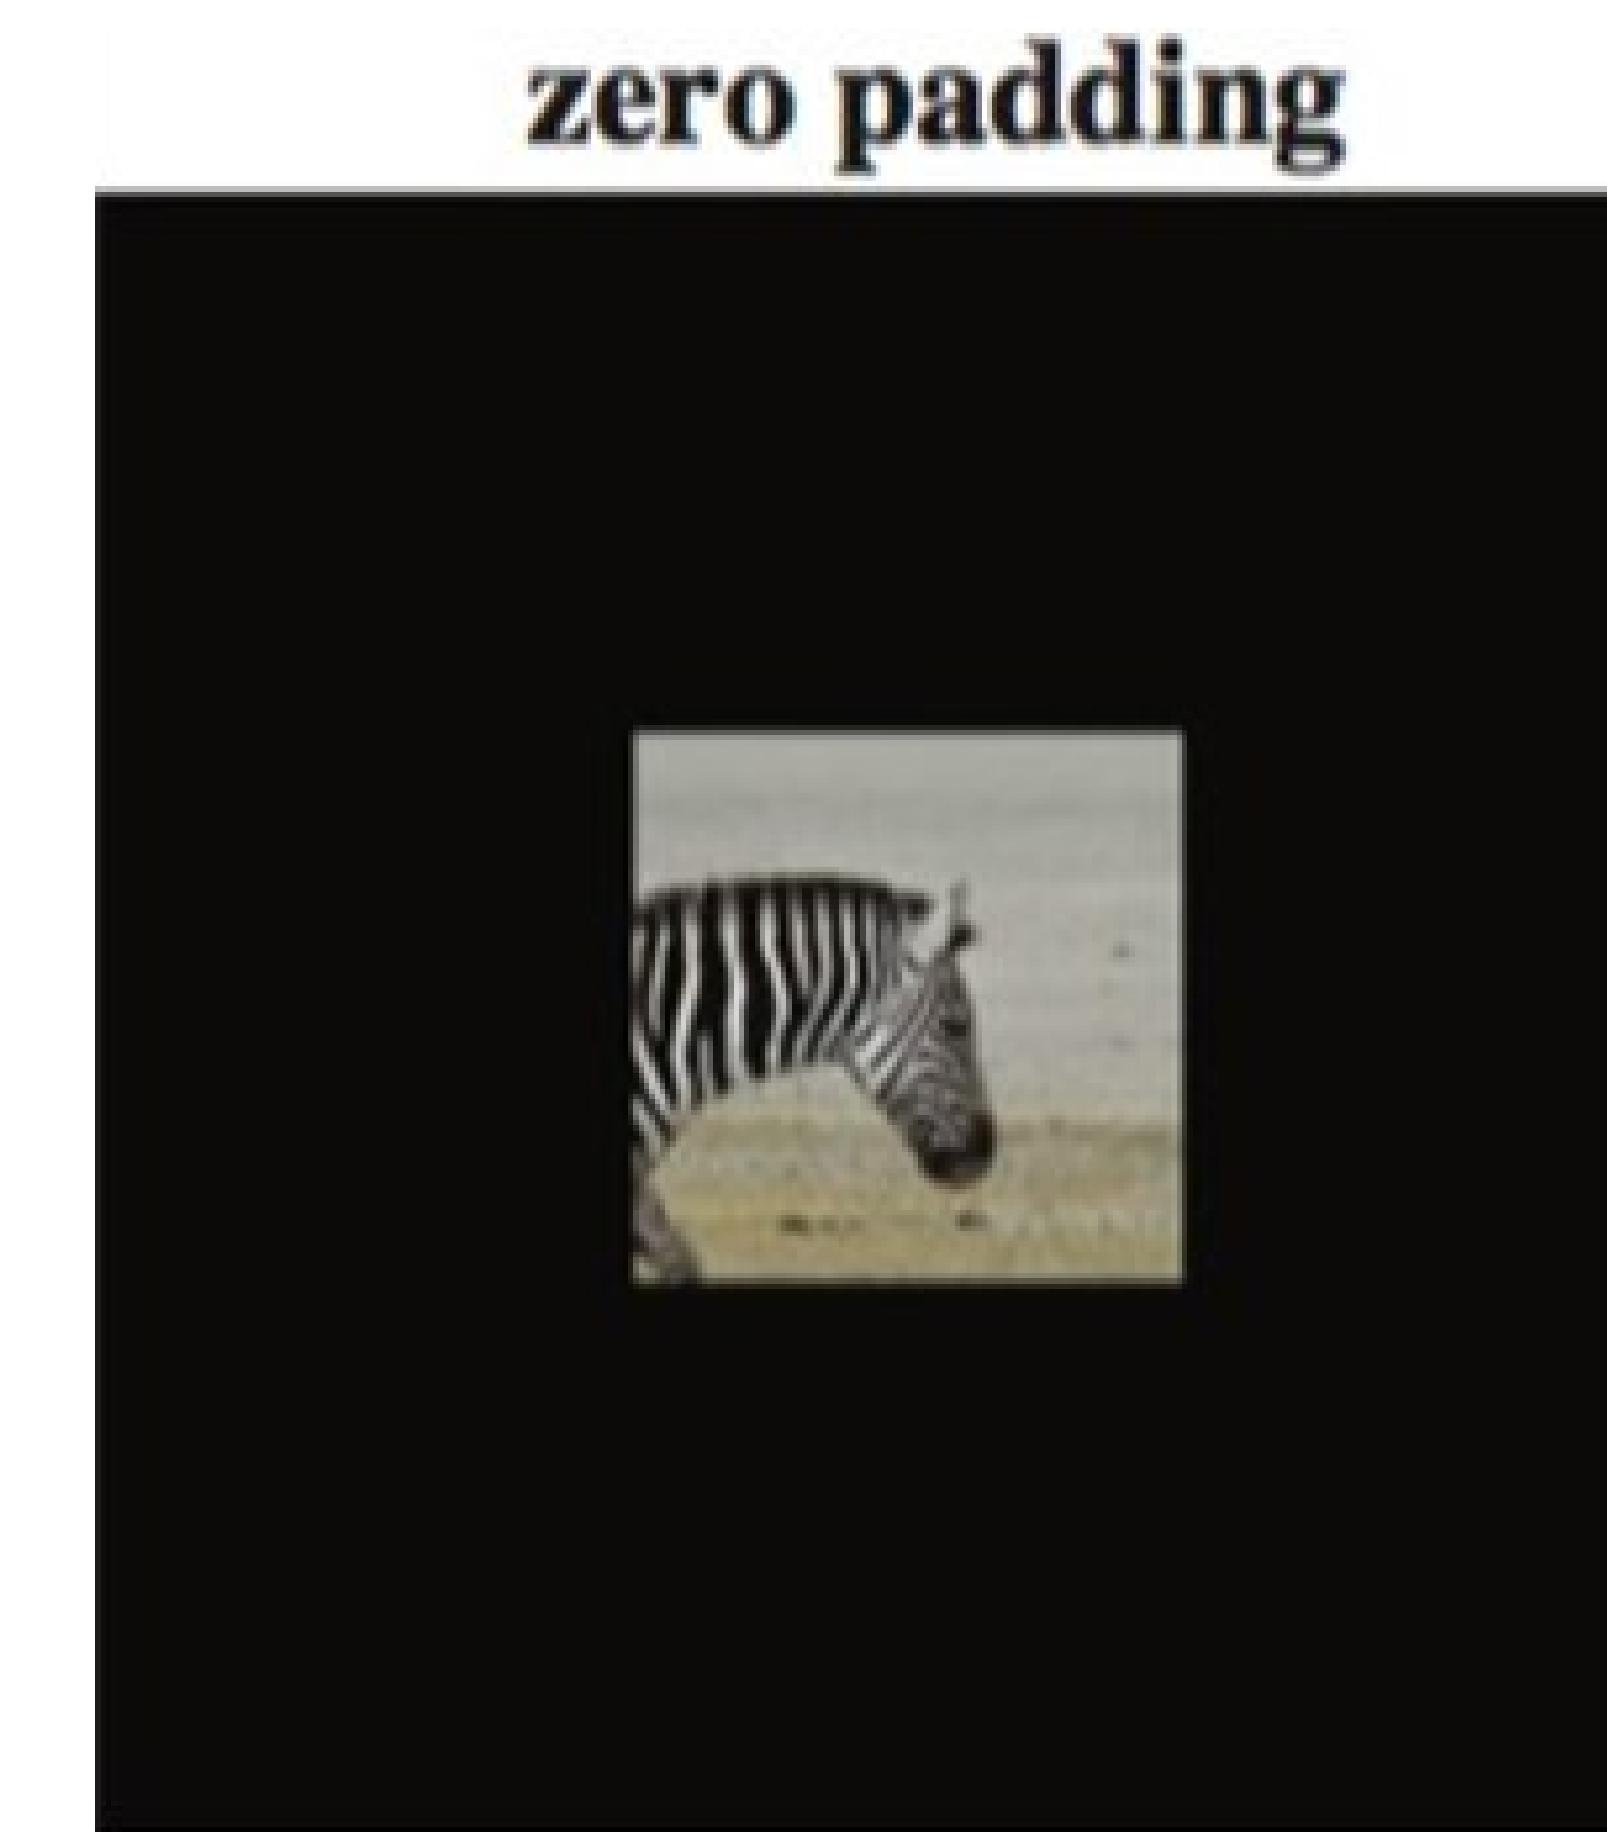
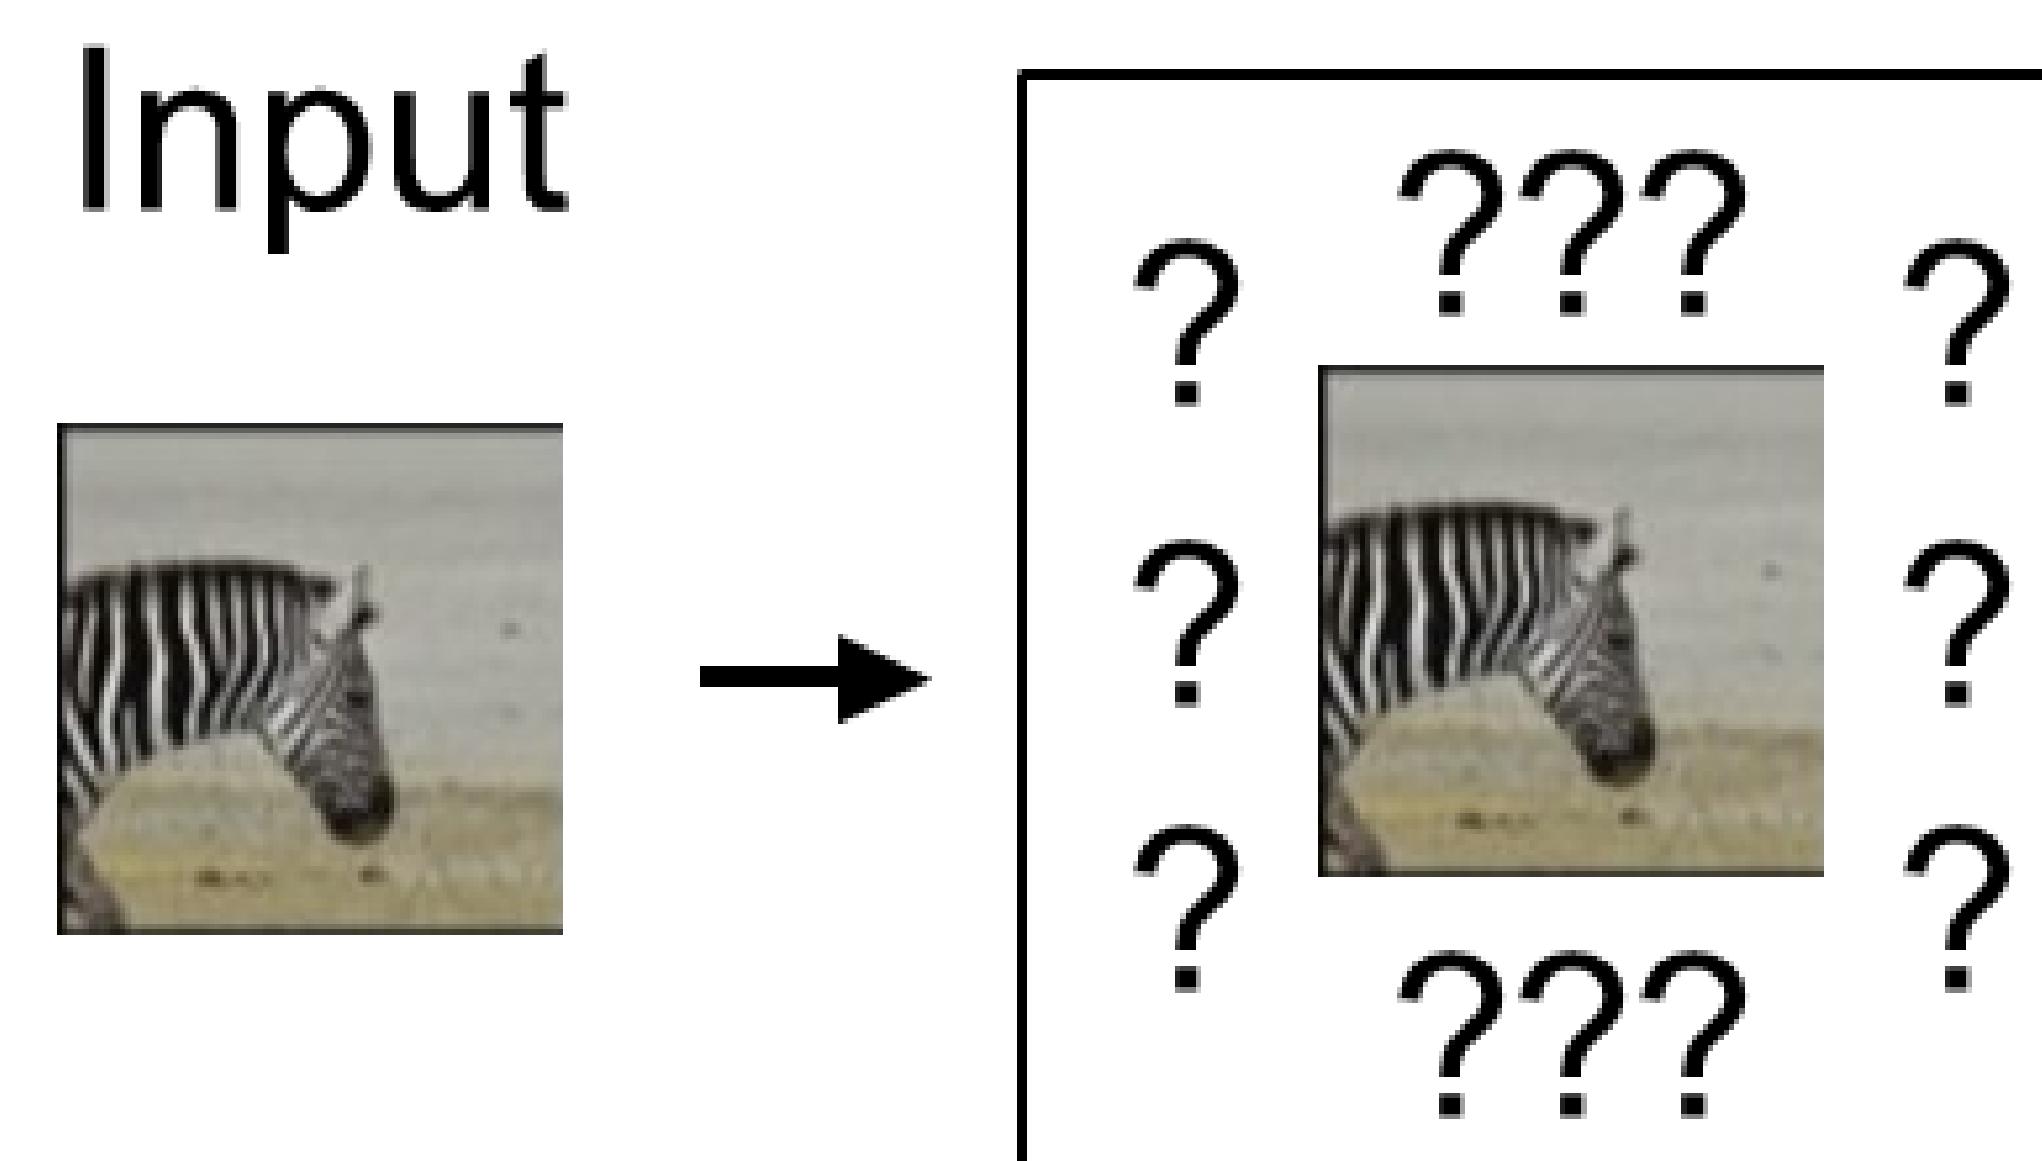
Patch-level processing: padding

Design choice 1: How big to pad



Patch-level processing: padding

Design choice 2: What to pad





Email: min.shi@louisiana.edu